
HackPit

Full Capability Assessment & Agreed Plan

Date: 26 July 2026

Scope: backend · frontend · pipeline · docker · docs · knowledge base

Measured against: OSCP · PNPT · eCPPT · HTB CPTS · CRTP · HTB/CTF · bug bounty · client pentests

HackPit — Capability Assessment & Build Log

0. Verdict

1. Remaining gaps

1.1 The PTY gap — now closed, without losing the transcript

1.2 The human-approval model — settled, not open

2. Feature-by-feature audit

Companion (KB / search / attack paths)

Cockpit (execution)

AD graph + orchestration

Detection footprint

Code scan (SAST)

Frontend

3. Knowledge base assessment

Findings

Recommended ingestion, in priority order — status

4. Source-usage audit — "did I use my sources properly?"

4.1 The headline finding: the pipeline only reads *.md

4.2 Consolidation is lossy for some sources

4.3 The biggest missed idea — PentestGPT's Pentest Task Tree — is now BUILT

4.4 The structural pattern

5. What is genuinely excellent — do not refactor

6. Security posture of HackPit itself

7. What the assessment examined

PART II — WHAT REMAINS

Standing policy (unchanged, still governs the project)

Explicitly deferred (with the reason)

reconFTW — recon gaps mined, and what to skip

PART III — BUILD LOG (Phases 1–5, shipped 2026-07-26)

Decisions taken (D1–D22)

Phase 1 — reach a real target

Phase 2 — make it remember

Phase 3 — make it fast to drive

Phase 4 — content & goal-specific

Phase 5 — the PTY gap, and finishing the KB

Windows execution backend — the WinRM driver (shipped 2026-07-26)

Post-assessment refinements (2026-07-27)

reconFTW review + incorporation (2026-07-27)

Persistence / backdoors (TA0003) enrich (2026-07-27)

AV/EDR evasion + traffic obfuscation (2026-07-27)

Gate-integrity audit (2026-07-27)

Gate integrity — build #5 (2026-07-27)

Live fire, image pinning, and the prevalidated danger re-check — build #7 (2026-07-27)

Live fire: what ran, and what it found

The image is now byte-reproducible — and its smoke test had never passed

The prevalidated path now re-checks danger

The three live-fire defects, fixed — and a stale precedent decided (2026-07-28)

The Sliver implant build: a subcommand that does not exist

Both listener lifecycles: a status that was assigned, never observed

The stale precedent, decided: both listeners are now gated

Two further defects, found the same way

Both fixes verified against a rebuilt image, not just against a diff

What the live fire now shows

Closing the not-runs: iodine demonstrated, dnscat2 diagnosed (2026-07-28)

Build #8 — the reasoning copilot, and a measured substrate (2026-07-28)

The substrate actually runs, and here is the number

The reasoning proposer (2.1–2.8) — deeper proposals, not autonomy

Web-exploitation and privesc depth (Tasks 3 & 4, still human-fires)

What is re-locked, unchanged

KB exploitation-writeup corpus — the case-based material 2.7 keys into (follow-on, now landed)

Verification

Status

Build #9 — the Windows/AD path, driven live against a real domain (2026-07-31)

Build #10 — the opt-in C2 exposure, and an honest NOT-RUN (2026-07-31)

Build #11 — the launcher, and operator identity (2026-07-31)

A dozen surfaces were built and then invisible

The status rail answers a question the UI could not

The report could not identify its own author

A mistake worth recording

What is locked

KB enrichment — the first external corpus, and what it cost to read it (2026-07-31)

KB enrichment — 24 cert/CTF/pentest repos, and a prediction that was wrong (2026-08-01)

KB enrichment — 7 GitBook certification spaces, and a zero that took three probes to earn (2026-08-01)

KB enrichment — the 0xdf fingerprint corpus, and the series close-out (2026-08-01)

The series, whole (all batches)

Batch 3 — PDFs, pages, and CPENT: the close-out (2026-08-01)

Build #12 — the capabilities nothing could reach, and the first CI (2026-08-03)

Ten endpoints that existed and could not be reached

The first CI, and the corpus problem it exposed

The first CI run failed, and that is the result worth recording

The drift job had never run — and could not have told anyone if it failed

A latent finding, recorded not fixed

Housekeeping, and a document that had been wrong for a year

Build #13 — where a callback lands (2026-08-03)

What was built

Three decisions went against the first draft, all from pushing back on it

The hole self-review found, and the one the suite found

Measured, not assumed

What it still cannot do

Build #13 part 2 — the evasion engine can now deploy (2026-08-03)

Two primitives, deliberately separated

The guard that fired found a design mistake, not an import

Smaller things the tests forced

Deferred to a later session, deliberately

Status

Build #13 part 3 — the out-of-band canary (2026-08-03)

HackPit does not provision, and that is a decision

The first internet-facing component, and why it is safe to expose

Four things the design got right only after being questioned

Verified end to end, on loopback, for real

The rest of part 1a — poll client, templates, panel (2026-08-03)

What is not built yet

Build #13 part 4 — the public C2 redirector (2026-08-03)

This one carries real AUP weight, and the safety argument had to be rebuilt

Extending part 1 rather than sitting beside it

One deploy engine, not a second path

The one awkward fact, stated rather than papered over

A test bug worth recording

Verification

Frontend

Build #14 part 1 — ZAP, the first web vulnerability scanner (2026-08-03)

Why not Burp Suite Professional

The whole build is seven touch points and no new architecture

The split, and an inconsistency stated rather than hidden

Two things the existing guards caught that the plan had not anticipated

The image build caught what the whole test suite could not

The trap that would have shipped silently

Verified, and what is not

Timeout, resolved without touching a global

Deferred to part 2

Build #14 part 2 — the recording proxy (2026-08-03)

The daemon part 1 refused, built anyway — because the transport changed

Four defects the guards caught, none of them test bugs

Two things only a live process could find

Secrets: raw in the panel, masked in reports

One invariant with nothing behind it, stated plainly

Verification

Build #14 part 3 — the scanner learns to aim (2026-08-04)

The finding that came first, and still blocks the other half

The gate: an honest surface, not a new one

Part 2's central lock could not be restated, so it was replaced

A second bound, enforced by ZAP rather than by us

The shape trap, caught by measuring instead of assuming

What the proof found that no test could

Verification

Closing the two gaps that were left open — and what the second one found

Full-project audit — the whole surface, driven end to end (2026-08-04)

The finding that mattered most was a combination, not a bug

Every gate fires, and every gate can fail — 27 checks, both directions

The missing gate is the mirror of the never-fired guard

Two hypotheses in the brief were wrong, in the product's favour

Defects fixed in this commit

Verification

Build #15 — browser interception (2026-08-04)

The measurement that unblocked it was a correction, not a discovery

Part 1 — the port is published through exposure.py, not hardcoded

What the design spec did not state, and the implementation had to answer

The key is closed twice, and one of the two cannot regress

Part 2 — the AJAX spider, and the red-confirm that could not fire

Why ZAP's spider and not our own headless browser

An OK is not a result

The isolation argument was rebuilt, not relaxed

Two defects the proof found that nothing else could — and the second was hidden by the first

Two more defects, this time in the proof itself — and one is a lesson arriving backwards

CI caught a test depending on Docker — for the SECOND time, in the same file

Verification

The acceptance test, run live (2026-08-04) — part 1 passes, part 2 does not

Two smaller things the live run surfaced

The audit punch list, worked — build #16 (2026-08-04)

D8 — the planner emitted commands nobody could run, and nothing said so

D3 — nothing throttled anything

Two shipping defects the pacing work found in the proxy flag it was mirroring

D5 — the safety suite rewrote the live KB on every run

D6 — nothing noticed a derived artefact going stale

D7 — every successful AD command looked like a failing one

The bug-bounty submission fields — and one that nothing could set

The known-issue check — flag, never suppress

Authenticated-scan session expiry — a silent wrong answer made loud

The phantom class, and what the type-checker cannot see

D9 — the two surfaces that read as a different product

The browser pass found two things nothing else could

What was not done

The three things the punch list left open, closed

Two defects only the browser found, again

Build #17 — closing out build #15 (2026-08-04)

17 is that close-out, plus a diagnostic that must land before anything is built on top of it.

The rules-of-engagement free-text field — reviewed and ACCEPTED

Item 1 — the record was corrected, not made enforceable

Item 4 — choosing the endpoint IS the safety decision

The pacing added in build #16 does not cover this path — checked, not assumed

The defect that blocked item 4 — the product goes blind to its own daemon, silently

The fix — the key is recovered, and the clash check can finally see

The defect underneath it: a capture was unreadable on Windows, and read as no capture

Item 5 — the browser must talk to the target and nothing else

Item 2, RUN — the discriminator is not the User-Agent, and that is why the honest fix works

Item 4, RUN — the attack half IS blocked, and the first verdict said the opposite

Item 3 — option A, built on the measurement rather than on the suspicion

The defect that hid item 4's own evidence: "recent" traffic was the oldest traffic

The one place this build said "you may not" — removed

Verification of item 3, live — and the runbook is now actually finished

What was prepared and has now run

Verification

Build #18 — reaching the targets that refuse us, and scanning behind a login (2026-08-05)

Item 1 — the bypass header, and the fact that its value is a credential

Item 2 — is it fronted, and what is behind it? Passive, and unknown is a real answer

Item 3 — a scan policy is a list of checks to switch off, each with a reason

Item 4 — the scanner cannot shape payloads, and saying so is the result

Item 5 — curl-impersonate, and why a User-Agent could never have worked

Items 6 and 7 — the hard part already existed; what was missing was telling ZAP what it means

Item 8 — the silent-empty sweep, and the two that mattered

Item 9 — one authored entry, and the KB was checked rather than assumed

What was NOT built, and why

Verification

Build #18, RUN — the image rebuilt, the proofs executed, and eight defects the read-backs caught

The frontend — built, and LOOKED AT

The parse-drop count, closed

The curated exploit overlay (2026-08-05)

The overlay is SOURCE; the mirror is DATA

gte — the index could not say "there is no fix"

What was NOT indexed, and why that is the discipline working

An existing lock broke, and it was right to

Two sources evaluated and declined (2026-08-05) — with the reasons, so nobody re-reads them

tim-barc/ctf_writeups — 209 PDFs, and the commands are pictures

sources/some vul.md — declined for the KB, redirected to the index

An undeclared dependency, found on the way

Build #19 — the interactive half, and eyes for the agent (2026-08-05)

Item 1 — the go/no-go, and it was a GO for a reason nobody would have guessed

*** 5. A drop WITH NOTHING HELD PERMANENTLY WEDGES THE BREAK MANAGER. ***

Item 2 — the repeater cookie jar, and a value that never leaves the module

Item 3 — history filtering, where the counts ARE the feature

Item 4 — interception, which adds no gate because there is nothing to bypass

Item 5 — the intruder, and why one approval may buy thousands of requests

Item 6 — the MCP server: eyes, not hands

Addendum (2026-08-08) — extended for the Q2/Q3 builds, plus an opt-in execution door

Six defects this build found in its own work, and each needed a different kind of check

The screens were LOOKED AT, and here is how, because the usual way was unavailable

What was NOT done, and why

Verification

Build #19, CI — main was already red, and then this build broke it a different way

Source evaluated: sources/insane ctf writeups and vulnerabili.txt (2026-08-05)

KB — 3 authored entries (2744 → 2747)

CVE→exploit index — 2 curated rows (5 → 7)

Verified, not assumed

A finding that evaporated when the tree stopped moving

Build #20 — GraphQL, end to end (2026-08-05)

Item 1 — the go/no-go, and *** IT CHANGED THE PLAN ***

61's request has the request, not the schema.

Seven more findings from the same script, each one measured

The measurement item 5 exists for

*** THE CONTAINER IS NOT THE IMAGE, AGAIN — AND THIS ONE IS A CAVEAT ON EVERYTHING ABOVE ***

Item 2 — detection, by BODY SHAPE and never by path

Item 3 — the repeater's round trip, and the operator's raw body WINS

Items 4 and 5 — recon, and a scan behind the existing four gates

Item 6 — the arsenal, and *** A SUPPLY-CHAIN TRAP AIMED AT AGENTS ***

Item 7 — the screens were LOOKED AT, and it found five things nothing else could

One safety predicate FIXED rather than narrowed

Verification

What was NOT done, and why

Build #21 — pin the engine, then read a schema nobody will give you (2026-08-05)

Item 1 — the pin, and *** THE ASSERTION CAUGHT A BROKEN PIN ON ITS FIRST RUN ***

The runtime half, measured in both directions

The re-measurement — *** THE 56/0 BASELINE REPRODUCED EXACTLY ***

Item 2 — engine fingerprinting, and the unit is the CORE

Item 3 — field-suggestion enumeration, and the measurement that shaped it

Three more cores, 2026-08-05 — and they went three different ways
.NET is two implementations, and they agree on nothing
The long tail, triaged — one question each, no parsers written
FIVE outcomes, and four of them are an empty list from a naive implementation
Item 4 — what the recovered schema is FOR
Item 5 — the screens, and looking at them found two things nothing else could
Verification
What was NOT done, and why
interact.sh — a second, zero-infrastructure OOB backend (2026-08-06)
Credential attack — spray captured creds, crack captured hashes (:credentials) (2026-08-06)
Nuclei template scan — the bug-bounty staple, wired (:nuclei) (2026-08-06)
Documentation & housekeeping (2026-08-06)
Guided recon → ranked attack surface — the front door (:recon) (2026-08-06)
Cloud attack surface & IAM privesc graph — the cloud parallel to BloodHound (:cloud-graph) (2026-08-06)
Web SSRF → cloud credentials — the IMDS seam (:cloud) (2026-08-07)
AD CS ESC1–8 routed in the graph — the AD parallel to DCSync synthesis (:ad-graph) (2026-08-07)
Unconstrained delegation edge + golden/silver ticket forging — the delegation-family closer (:ad-graph) (2026-08-07)
AI code-audit fan-out — open-kritt's decomposition, HackPit-gated (:code-scan AI mode) (2026-08-07)
Finding pipeline — dynamic schema · auto-dedup · pluggable rankers · post-scripts (cross-cutting) (2026-08-07)
Web3 / smart-contract audit — three playbooks on the fan-out (2026-08-07)
Formal engagement governance — RoE / ConOps / Deconfliction / OPPLAN (2026-08-07)
Reusable prompt-workflow builder — author your own playbooks over the fan-out (:workflows) (2026-08-07)
Interactive persistent-session engine — named tmux sessions + automatic prompt-detection (:terminal) (2026-08-07)
Binary-RE / pwn + forensics / CTF arsenal expansion — the HexStrike coverage gap (2026-08-08)
Token workbench — JWT / OAuth / OIDC / SAML analysis, tamper and a gated crack (:tokens) (2026-08-08)
Single-packet race tester — beat a check-then-act window (:race) (2026-08-08)
Request-smuggling / desync detector — timing-differential detect + gated confirm (:smuggle) (2026-08-08)
Web cache poisoning / deception detector — unkeyed-input detect + gated poison-plant (:cache) (2026-08-08)
Cross-domain kill-chain graph — the capstone that stitches web → cloud → on-prem into one route (:killchain) (2026-08-08)
Parameter / content discovery — hidden params + paths as a first-class :recon step (:discover) (2026-08-08)
JS recon → secret / endpoint mining — the S3→bundle→secret chain, live (:jsrecon) (2026-08-08)
Dual-candidate "second opinion" on generated commands — foundation (2026-08-09)

HackPit — Capability Assessment & Build Log

Date: 2026-07-26 **Scope:** every backend module, the frontend, the pipeline, the docker stack, the docs, and the live KB **Targets this is measured against:** OSCP · PNPT · eCPPT · HTB CPTS · CRTP · HTB boxes/CTF · bug bounty · real-world client pentests.

Latest (build #12, 2026-08-03). A frontend-to-backend coverage sweep found **ten endpoints with no caller at all** — D19's "built and then invisible" one layer deeper, including both depth halves of build #8 — and all ten are now reachable. The project also has its **first CI:** the 56-file hermetic suite, the frontend build and an eslint-baseline check run on every push, with the corpus those tests need supplied by a derived, equivalence-proven fixture rather than skipped (D26). See Build #12 at the end of Part III.

Status note (updated). This began as an assessment (Part I) and an agreed plan (Part II). **All four phases of that plan have since been built**, and a fifth followed — the execution substrate (Phase 1), the state model (Phase 2), the drive-speed gaps (Phase 3), the content-and-goal-specific work (Phase 4: KB Route-B ingest, the evasion/OPSEC channel, the HTTP repeater, pivot/tunnel routing, exam-mode report templates), and **Phase 5**, which closed the two things Part I had left standing: a real PTY (added as a *second* surface, so the auditable one survives) and the recommended KB enrichment (PortSwigger Academy, HackTricks Cloud, the HTB slice, and a CVE→exploit index). Part I has been trimmed to what genuinely remains; the completed blockers, decisions and phases live in **Part III — Build Log**.

0. Verdict

The **architecture is genuinely strong** — the safety engineering and the LLM-grounding discipline are better than most commercial tooling.

The original problem this assessment found was *reach*: nothing in HackPit was fake or stubbed, but the container it executed into contained seven programs, so it could not finish a real HTB box, touch a real AD domain, run a bug-bounty recon sweep, or complete an OSCP-style engagement. **That execution substrate has now been rebuilt (Phase 1), the loop reasons over structured state instead of stdout tails (Phase 2), the drive-speed gaps are closed (Phase 3), the content/goal work is done (Phase 4), and the interaction and knowledge gaps Part I left standing are closed (Phase 5).** See Part III.

Since then **Phase 5** closed the PTY gap (a real terminal as a *second* surface, keeping the auditable one) and finished the KB enrichment the assessment recommended, and the **Windows execution backend** (a WinRM driver) has since made the CRTP toolset, the OSCP AD set and real internal pentests *executable* — and made the existing AD attack-path graph run **live** instead of on synthetic data. What remains is a short list of *deliberately deferred* items — a VPS for blind callbacks, and the last few thin KB categories — each deferred for a stated reason, not left undone. The Windows/CRTP execution target is **no longer among them**: it is closed by the WinRM driver against an external VM you run (see the Windows execution backend in Part III and D9). Risk-tiered approval is also gone — **decided against** (see D5). See §1 and Part II's "Explicitly deferred".

1. Remaining gaps

The blocker sections for everything the five phases fixed have been removed (see Part III). What is left is genuinely small.

1.1 The PTY gap — now closed, without losing the transcript

This section previously read "no true PTY — deliberate partial", and framed it as a trade: readable logged transcripts *or* full-screen tooling. **Phase 5 rejected that trade.** The premise was wrong — the two are not alternatives if they are separate surfaces.

`:kali` is unchanged: still a persistent shell delimiting each command with a sentinel over a plain pipe, still producing the clean, escape-free, per-command transcripts reports are built from. Alongside it, `:terminal` now carries **two** surfaces onto the same open sandbox: a **real PTY** (so `vim`, `top`, a curses tool render) and — the 2026-08-07 addition — **named persistent tmux sessions** with per-session cwd and **automatic interactive-prompt detection**, so an interactive `msfconsole` / `sliver-client` / `evil-winrm` becomes first-class (the UI raises "interactive — send input" and the human types the next line; the orchestrator never can). Both are audited; both carry identical containment; input is human-only on all of them. See Part III, Phase 5 and the dated session-engine section.

1.2 The human-approval model — settled, not open

Current state: every command needs individual `approved=true`; anything the danger heuristic flags additionally needs `dangerous_ack`; `:kali` and `:terminal` are human-driven (typing *is* the approval). Phase 3 added **one-keystroke approve** (`Enter` = approve · `S` = skip · `Esc` = stop; a dangerous command never fires on `Enter` alone).

This section previously proposed **risk-tiered approval** — auto-running "passive" commands, batch-approving a plan of read-only ones — as open work. **That is now decided against (D5).** Per-command approval is the project's single load-bearing safety property in engagement mode: Wall A is down, the sandbox reaches the internet, the host and the LAN, and nothing else bounds *where* a command can go. A tier that auto-runs anything would remove the only gate on the one mode that needs it most, and it would do so by trusting a classifier — a new component, on the wrong side of the safety boundary, to decide when the human can be skipped. The convenience it buys is a keystroke that has already been optimised away.

Per-command approval is therefore **standing policy, not a deferred item**, and does not appear in the deferred list. The prerequisite work that would have enabled tiering (`_looks_like_host()`, Phase 3 step 12) was worth doing on its own merits — the scope check is trustworthy now regardless. **Unchanged and non-negotiable: never remove the gate for anything the danger heuristic flags.**

2. Feature-by-feature audit

Companion (KB / search / attack paths)

Component	State	Notes
KB load + serve	Solid	1,601 entries in memory, exclusions dropped at the door and re-filtered defensively at query time.

Component	State	Notes
Hybrid search (<code>pipeline/search.py</code>)	Solid	BM25 (stdlib) + <code>nomic-embed-text</code> cosine, weighted RRF (lex 1.0 / vec 0.5 — correct call for a KB full of exact identifiers), tier boost, whole-query title bonus. Degrades to lexical when Ollama is down instead of 500-ing.
Attack-path composition	Solid, genuinely novel	Writeup-first mode + KB-first mode. Retrieval seeded per-phase so no phase starts empty. Step eligibility filtering (<code>is_step_eligible</code>) is unusually careful — excludes writeups, defensive-hardening guides, tool-install meta-docs, personal logs, grab-bags.
Grounding / validation	Best-in-class	Cited <code>entry_id</code> s must resolve or the step downgrades to <code>ai_suggested</code> ; commands come from the KB entry, never the model; <code>target_adaptation</code> is dropped entirely if it names an FQDN not in the supplied facts.
<code>substitute_target</code>	Excellent	CIDR sentinel to avoid double-rewrite, host-position gating so <code>.env.example</code> and <code><script></code> survive, refuses to rewrite what it can't rewrite confidently.
<code>foreign_refs</code>	Excellent	Honesty marker — reports a foreign host/AD domain rather than guessing a replacement. Right call.
Channel-2 context grounding	Good	Writeup/methodology content as reasoning background, strictly separated from the step pool, with a leak guard over model output.
Engagement sessions	Works	SQLite, per-step checked + pasted results, debounced autosave.
Report generation	Very good, now multi-template (Phase 4)	Evidence built programmatically , not by the LLM — the model was observed mis-transcribing a port number. Exam/format templates: OSCP (per-host walkthrough + a <code>proof.txt</code> table spliced from state), CPTS (exec-summary + findings register), H1/Bugcrowd (impact-first + a CVSS 3.1 score <i>computed</i> , not asserted). <code>proof.txt/local.txt</code> are real per-host state (auto-captured when a command reads a flag file, or pasted). Collision-proof fences; lab vs REAL-TARGET labelled; detection footprint appended grounded-only; optional red-team OPSEC roll-up (D10).
Assistant chat	Works	Session-aware, KB-grounded, deterministic citation extraction (scans the reply for ids that resolve).
Scripts arsenal	Works	1,028 scripts deduped from the KB across 6 groups.

Cockpit (execution)

Component	State	Notes
Gated executor	Excellent design, now properly tooled	Four ordered gates, mode split, argv-only (never a shell), SSE streaming, run records persisted. The image it execs into now ships ~63 of the 73 catalogued tools (Phase 1).
Isolation gate	Real	<code>assert_isolation_proven()</code> structurally inspects every attached network for <code>internal: true</code> . Not a comment — an actual check.
Engagement mode	Correct	Explicit entry required, fail-closed on unknown/exited id, per-command approval re-checked even on the prevalidated path (belt-and-suspenders in <code>iter_run</code>).

Component	State	Notes
Scope model (<code>scope.py</code>)	Good	Hosts, <code>*.wildcards</code> , CIDRs, <code>!exclusions</code> , <code>* opt-out</code> . Fail-closed. No DNS at match time (avoids the "resolve an out-of-scope name to decide it's out of scope" leak). The <code>_looks_like_host()</code> false-positives are fixed (Phase 3).
Recon-driven expansion	Good	Mines run output for hosts, sorts by scope, in-scope join the live allowed set, out-of-scope surfaced read-only. Cannot widen the scope. Capped and never silently truncating.
Orchestrator loop	Now state-grounded	Proposes one command, never executes. Feeds on the structured state model + live task tree (Phase 2) instead of stdout tails.
Engagement state model	New (Phase 2)	hosts/services/endpoints/credentials/findings, upsert-only, fed by output parsers + loot-file ingest; drives the planner and a UI panel. Executes nothing (AST-asserted).
<code>:kali</code>	Now a persistent shell (Phase 3)	One long-lived <code>docker exec -i sh ; cd /env/background</code> jobs persist. Same containment (hardcoded open container, human-only, audited, no isolation gate). Deliberately no PTY — that is what keeps its transcripts clean (§1.1).
<code>:terminal</code>	Real PTY + named tmux sessions (Phase 5 · session engine 2026-08-07)	Two surfaces onto the same open box. PTY: <code>pty.fork()</code> driven over a framed WebSocket to xterm.js; full-screen tools render and resize. Named sessions : parallel persistent tmux sessions with per-session cwd, automatic interactive-prompt detection (<code>msfconsole</code> / <code>sliver</code> / <code>evil-winrm</code> /REPLs → "interactive — send input"), auto-background at 60s with notify-once completion, wedge/pipe-degradation recovery. Both HUMAN-ONLY (source-scan locked), no new gate, audited; ported from Decepticon's <code>tools/bash</code> . Does not replace <code>:kali</code> (§1.1).
<code>:exploits</code>	New (Phase 5)	Version-keyed CVE → exploit lookup over the sandbox's local exploit-db catalogue — 47k exploits, 25k distinct CVEs. Read-only; executes nothing.
Live sessions	Well-designed, now tooled	Start is a gated command; stdin is human-only and source-scan locked.
HTTP repeater	New (Phase 4)	Compose/send/replay/diff. Argv-only curl inside the hardcoded open box (no shell parses a request field; body on stdin), human-only + source-scan locked like <code>:kali</code> , scope-checked against a named engagement, every send run-recorded.
Pivot / tunnels	New (Phase 4)	chisel/ligolo-ng lifecycle (human-only start/stop) + pure route resolution and a visible proxychains rewrite applied <i>before</i> the approval screen. A tunnel's subnet enters scope only via an explicit, audited amendment; recon expansion still cannot widen.
Credential attack (<code>:credentials</code>)	New (2026-08-06)	Spray captured/OSINT creds across a service (netexec / kerbrute / hydra) or crack captured hashes (hashcat, mode auto-detected). ONE approved job per spray/crack with an ungated stop — the intruder's shape on a long process, gated by the same <code>validate_request</code> , no new gate . Secrets go to loot files, never an argv; a hit writes a validated credential + finding into state and marks the AD node owned . Planner (<code>cockpit/credattack.py</code>) executes nothing (AST-asserted); execution is the gated worker (<code>cockpit/credjobs.py</code>).

Component	State	Notes
Nuclei template scan (:nuclei)	New (2026-08-06)	Scoped target(s) → templates → severity-ranked Finding s. ONE approved job with an ungated stop — the <code>ffuf</code> / ZAP-active-scan shape, gated by the same <code>validate_request</code> , no new gate ; the per-mode sandbox is resolved by <code>executor.resolve_mode</code> (isolated lab / open engagement). Default targets seed from the session's in-scope endpoints; results dedupe by (<code>template-id</code> , <code>matched-at</code>) and upsert into the same engagement Finding store the report renders. Pure planner/parser (<code>cockpit/nuclei.py</code> — <code>argv</code> + JSONL parse, AST-asserted no-exec); the worker gates then spawns. Verified live against the lab (Prometheus-metrics medium + tech fingerprints).
Cloud IAM privesc graph (:cloud-graph)	New (2026-08-06)	The cloud parallel to the AD graph. Enumeration (<code>ScoutSuite</code> + <code>Prowler</code> , with <code>pacu</code> / <code>cloudfox</code> added to the arsenal + sandbox image in this build) is ONE approved job — the recon/nuclei shape, gated by the same <code>validate_request</code> , no new gate , engagement-bound. Its JSON is parsed into a typed IAM privilege-escalation graph (<code>cloudgraph/</code> , a near-clone of <code>adgraph/</code>): principals + resources wired by abusable IAM relationships (<code>sts:AssumeRole</code> , <code>iam:PassRole</code> / <code>AttachRolePolicy</code> / <code>CreatePolicyVersion</code> , <code>lambda:UpdateFunctionCode</code> , Azure <code>Owner</code> -on-self / app-cred-add, GCP <code>serviceAccountTokenCreator</code> / <code>actAs</code>). BFS routes to an admin/owner principal; the orchestrator picks an EDGE INDEX, never authors a command (regression-locked by an AST + source scan in <code>test_cloudgraph_safety.py</code>) — the abuse is KB-grounded (534-entry cloud corpus, precise CLI catalog behind it) and runs only through the gated executor, with advance requiring an approved exit-0 run checked server-side. Privesc paths + Prowler misconfigs land as engagement Finding s. Multi-cloud by construction (provider on every node); AWS end-to-end, Azure/GCP node/edge + technique support in place.
Web SSRF → cloud creds (IMDS bridge)	New (2026-08-07)	The seam between the web/cockpit half and <code>cloudgraph/</code> . A captured instance-metadata (IMDS) response — pasted, from a repeater exchange, or an OOB callback body — is parsed into an owned cloud principal seeded into the session's IAM graph, so the privesc walk starts from the SSRF/RCE-stolen identity. <code>cloudgraph/imds.py</code> is a pure parser (executes nothing, imports no network/exec module — AST-asserted in <code>test_cloudgraph_safety.py</code>): AWS (IMDSv1 + the IMDSv2 token-PUT/creds-GET two-step + role listing + instance-identity doc), Azure managed-identity JWTs (<code>oid</code> / <code>appid</code> /tenant decoded), GCP SA tokens (with the Metadata-Flavor requirement flagged); malformed/truncated blind-SSRF bodies degrade to warnings , never a crash. No new gate — there is nothing to gate: the request that touched <code>169.254.169.254</code> already ran through the human-approved repeater/executor; the bridge only parses a string and seeds. The cross-cutting seed route (<code>POST /cockpit/cloud/seed-imds</code>) lives in <code>main.py</code> (so <code>cloudgraph</code> never imports <code>cockpit.loot</code> — the decoupling rule): it add_node s the owned identity (matching an already-enumerated node when the stolen identity coincides, so a route to admin lights up), stores the secret in the engagement vault + loot file , and records a high-severity Finding carrying provider/identity/expiry but never the secret . Parser fixtures for all three clouds + the secret-never-in-the-finding invariant are pinned in <code>test_cloud_imds.py</code> .

Component	State	Notes
Token workbench (:tokens)	New (2026-08-08)	The token parallel to the GraphQL panel — JWT / OAuth-OIDC / SAML made a first-class surface. A pure analysis/tamper core (<code>cockpit/tokens.py</code> — imports nothing that executes, AST-asserted with a planted control in <code>test_tokens_safety.py</code> , and modelled point-for-point on <code>cockpit/graphql.py</code>): recognises a token by shape, not path (three base64url segments — Authorization: Bearer , a cookie, a <code>?param</code> , a body field), decodes header/claims/signature for a pasted token and flags the classic misconfigs (<code>alg=none</code> , missing <code>exp</code> , weak HMAC secret, injectable <code>kid</code> / <code>jku</code> / <code>jwk</code> / <code>x5u</code>). Tamper primitives each produce a token string, and send nothing: <code>alg=none</code> (with case variants), RS256→HS256 confusion signed with the server's public key as the HMAC secret, <code>kid</code> injection (traversal/SQLi/command), <code>jwk</code> / <code>jku</code> / <code>x5u</code> header injection, edit-any-claim-and-re-sign. OAuth: parse an authorization request/callback and build <code>redirect_uri</code> bypasses (the open-redirect table), <code>state</code> drop, PKCE downgrade, forced <code>response_mode</code> , implicit leak. SAML: parse a Response (base64[+deflate]/XML), locate whether the assertion vs response is signed, and emit XSW1-8, signature strip, comment-injection, unsigned-assertion — each mutated Response round-trips to well-formed XML (verified in <code>test_tokens.py</code>). NAMES, NEVER VALUES: a token auto-detected in captured traffic is modelled claim-<i>names</i>-only (never a value or signature — a JWT claim is routinely a secret), the same rule <code>GraphQLArgument</code> follows. The one thing that runs — the weak-secret crack (<code>hashcat -m 16500</code> over a wordlist, <code>cockpit/tokenjobs.py</code>) — is ONE gated job behind the same four gates: the token goes to a loot file (never the argv), a recovered secret goes to loot (never the finding) and lands a high Finding , engagement-bound, stop ungated. No new gate — a mutated token reaches a real endpoint only through the human-approved, scope-checked repeater (the core has no path to <code>.send</code>). Surface at <code>/tokens</code> and mounted in <code>:proxy</code> .
Single-packet race tester (:race)	New (2026-08-08)	The race-condition primitive the intruder's serial loop cannot do — fire one request N times so the copies arrive in the same instant (HTTP/2 single-packet, or HTTP/1.1 last-byte sync) to beat a target's check-then-act window (limit overrun, TOCTOU, coupon reuse, one-time-token replay). Modelled exactly on the intruder: ONE gated job (<code>cockpit/race.py</code>), the same four gates via the same <code>validate_request</code> and no new gate, run inside the hardcoded open sandbox, scope-checked on the wire (an out-of-scope host refuses the whole batch — all N share one URL — and sends nothing), with an ungated stop. The whole request template <i>and</i> the concurrency count N are in the approved surface (the intruder's Critical-2 completeness rule — a body carrying <code>\ sh</code> reaches the danger gate rather than hiding behind a count); the declared command is clean by binary identity (firing benign duplicated requests is not code execution), so a plain race adds no red-confirm, exactly like a plain <code>ffuf</code> . The engine is a small in-repo client baked into the sandbox image (<code>docker/race_singlepacket.py</code> → <code>race-singlepacket</code> , its own venv for the HTTP/2 <code>h2</code> framing library; h1-last-byte is pure stdlib) — Turbo Intruder is a Burp extension and not headless, so this is the vetted minimal client instead. It is invoked argv-only with the whole request delivered on stdin (no shell parses a request byte), added to <code>arsenal/tools.json</code> (<code>_MUST_NOT_FIRE</code> bucket — same class as <code>ffuf</code> / <code>sqlmap</code>), the Dockerfile.sandbox (<code>--selftest</code> smoke-tested), and <code>docker/proof/race_install_proof.sh</code> . The verdict is the race-window signal, not a raw dump (<code>race.race_verdict</code> , pure, fixture-tested): the N responses are clustered and when the rare outcome a serial baseline returns <i>once</i> appears more than once, that is a race — reported as "K of N requests won the race", and a confirmed race upserts a High Finding in engagement state. Locks in <code>test_race.py</code> ; the engine image layer is a rebuild flag for the operator. Surface at <code>/race</code> , next to <code>:intruder</code> .

Component	State	Notes
Request smuggling / desync detector (:smuggle)	New (2026-08-08)	The other front-end/back-end parsing bug, made testable — probe a target for CL.TE / TE.CL / CL.CL / TE.TE / CL.0 and the HTTP/2-downgrade H2.CL / H2.TE / H2.0 variants. Built around the rule the class demands: detection is safe, confirmation is not, so they are split. Detection (default) is timing-differential and self-contained — each mutation's probe makes a <i>back-end</i> wait for body bytes the ambiguous framing never delivers, so a front-end/back-end split shows as a large timing delta over a baseline while a single RFC-consistent server answers fast; no other user's request is touched. A per-mutation verdict matrix (baseline / probe / Δ / susceptible-clear-inconclusive) renders, and a susceptible mutation upserts a High Finding . Confirmation is a SEPARATE approve-each with a co-tenant warning (<code>POST /smuggle/confirm</code> , stage pinned): socket-poisoning smuggles a partial request onto a shared connection then sends a normal one to observe the poisoned response — <i>this can affect the next request on that connection, possibly another user's</i> — so it carries the plain-language <code>CO_TENANT_WARNING</code> , runs one mutation at a time, and a confirmed desync upserts a Critical Finding . Modelled exactly on <code>:race / :nuclei</code> : ONE gated job per stage (<code>cockpit/smuggle.py</code>), the same four gates via the same <code>validate_request</code> and no new gate class, hardcoded open sandbox, scope-checked on the wire (an out-of-scope host refuses the whole sweep — all mutations share one URL — and sends nothing), ungated stop. The whole request template + mutation set are in the approved surface (Critical-2 completeness — a body carrying <code>\ sh</code> reaches the danger gate); mutations ride the surface dot-free (<code>cl-te</code>) so a dotted <code>CL.TE</code> is not misread as a host by the best-effort target-lock. The engine is a small in-repo client baked into the sandbox image (<code>docker/smuggle_probe.py</code> → <code>smuggle-probe</code> , its own venv for the <code>h2</code> framing library; the <code>h1</code> mutations are pure <code>stdlib</code>), invoked argv-only with the whole request delivered on stdin (no shell parses a request byte); the standard external tools <code>smuggler.py</code> (defparam) and <code>h2csmuggler</code> (BishopFox) ship alongside for manual use. All three added to <code>arsenal/tools.json</code> (<code>_MUST_NOT_FIRE</code> bucket — same class as <code>ffuf / sqlmap / race-singlepacket</code>), <code>docker/Dockerfile.sandbox</code> (<code>--selftest</code> smoke-tested), and <code>docker/proof/smuggle_install_proof.sh</code> . The verdict is pure and fixture-tested (<code>smuggle.detection_verdicts / parse_engine_output</code> , threshold pinned at <code>SUSCEPTIBLE_DELTA_MS</code> ; a manual <code>smuggler.py</code> transcript maps into the same table via <code>parse_smuggler_output</code>). Locks in <code>test_smuggle.py</code> ; the engine image layer is a rebuild flag for the operator. Surface at <code>/smuggle</code> , next to <code>:race</code> .

Component	State	Notes
Web cache poisoning / deception detector (:cache)	New (2026-08-08)	The shared-cache parallel to :smuggle , made testable — probe a target for web cache poisoning via unkeyed inputs (X-Forwarded-Host / X-Forwarded-Scheme / X-Forwarded-Proto / X-Host / X-Original-URL / X-Rewrite-URL / Forwarded , a cloaked query param, a fat GET body) and for cache deception (a dynamic page cached under a static-looking /.../foo.css / ;foo.css / %2f... path). Built around the same rule as :smuggle : detection is safe, confirmation is not, so they are split. Detection (default) plants nothing — each candidate input carries a unique marker in ONE request and the engine reports two observations: is the marker reflected, and is the response cacheable (Cache-Control public/max-age, Age , X-Cache / CF-Cache-Status). Reflected + cacheable ⇒ a candidate (the rule lives in one place, cache.candidate_of , so a test pins it); the path-confusion deception probes run alongside. A per-input verdict table renders, and a candidate or a cached-dynamic deception hit upserts a High Finding . Confirmation is a SEPARATE approve-each with a co-user warning (POST /cache/confirm , stage pinned): it plants the poison — primes the cache with the marker request, then fetches it back with a fresh request that never carried the marker — <i>the poisoned entry a fresh request receives is the same one a real co-user would receive until the cache expires</i> — so it carries the plain-language CO_USER_WARNING , runs one candidate at a time, and a confirmed poisoning upserts a Critical Finding . Modelled exactly on :smuggle / :race : ONE gated job per stage (cockpit/cache.py), the same four gates via the same validate_request and no new gate class, hardcoded open sandbox, scope-checked on the wire (an out-of-scope host refuses the whole sweep — all inputs share one URL — and sends nothing), ungated stop. The whole request template + input set are in the approved surface (Critical-2 completeness — a body carrying \ sh reaches the danger gate). The engine is a small stdlib-only in-repo client baked into the sandbox image (docker/cache_probe.py → cache-probe , no venv needed — urllib / ssl / re only), invoked argv-only with the whole request delivered on stdin (no shell parses a request byte); Hackmanit's wcvs (web-cache-vulnerability-scanner, Go) ships alongside for manual use, built in its own temporary-toolchain Docker layer. Both added to arsenal/tools.json (_MUST_NOT_FIRE bucket — same class as ffuf / sqlmap / smuggle-probe), docker/Dockerfile.sandbox (--selftest smoke-tested), and docker/proof/wcvs_install_proof.sh . The verdicts are pure and fixture-tested (cache.detection_verdicts / deception_verdicts / parse_engine_output ; a manual wcvs transcript maps into the same table via parse_wcvs_output). Locks in test_cache.py ; the engine + wcvs image layers are a rebuild flag for the operator. Surface at /cache , next to :smuggle .
Guided recon → ranked surface (:recon)	New (2026-08-06)	The front door: a scoped domain → recon as approved jobs → a ranked attack surface. A passive sweep (default, bug-bounty safe) chains subfinder → dnsx → httpx → gau/waybackurls/katana ; an active sweep (one more approval) chains naabu → nmap - sV . Each sweep is ONE approved job with an ungated stop — the crack-worker shape, gated by the same validate_request , no new gate; engagement-bound (the open sandbox has egress + the scope). Output is parsed to Host / Service / Endpoint (new state/parsers.py parsers: subfinder/dnsx/naabu/url-lister, query-param NAMES mined not values) and upserted. Scope discipline is a correctness property, not a gate: discovered hosts are sorted by the declared scope via engagement.record_discoveries + recon.filter_in_scope — in-scope names join the live allowed set and are the <i>only</i> hosts the probing tools are pointed at; out-of-scope names are surfaced read-only and never scanned or upserted (regression-locked end-to-end in test_recon_safety.py). recon.rank_surface scores each host by likely-exploitable (open services, CVE-worthy stacks via :exploits , parameter-rich endpoints, auth surfaces, findings) — advisory, executes nothing (AST-asserted) — and hands off into :attack-paths / :nuclei . Pure planner/parser (cockpit/recon.py); the worker gates then spawns.

AD graph + orchestration

Component	State	Notes
BloodHound parser	Genuinely good	Handles v4/v5/CE, zip/dir/json/bytes/mapping, reconciles naming drift, synthesizes DCSync from <code>GetChanges</code> + <code>GetChangesAll</code> , emits coverage warnings for missing collection methods. Real work.
AD CS ESC1–8 graph	New (2026-08-07)	<code>certipy find -json</code> folded into the SAME graph: <code>certtemplate</code> / <code>certauthority</code> nodes + synthesized composite ESC1...ESC8 edges , the AD parallel to the DCSync synthesis (a predicate over <code>template-vuln × enroll-right</code> → one edge from a low-priv enrollee to Domain Admins). ESC1/6/8 direct; ESC4/ESC7 emit the two-hop reconfigure-then-abuse shape through the <code>template/CA</code> node; ESC2/3 modeled; ESC9–11 catalog-cited. Each edge grounds to <code>certipy req → auth</code> (+ native <code>Certify.exe</code>), all destructive and oracle-locked to trip the red-confirm on both transports. <code>certipy find</code> runs as a gated scope-locked enum job. No new gate — the orchestrator still picks an edge index, never authors a command. <code>test_adcs_graph.py</code> + extended <code>test_adorch_safety.py</code> ; <code>Certify.exe</code> added to the arsenal. See the dated section below.
Path engine	Correct	BFS shortest path over abusable edges only, abuse-rank tie-break, k-shortest-ish alternatives.
Technique catalog	Good	25 edge kinds, KB-grounded with catalog fallback, target-type specialization (<code>GenericAll</code> on a group → <code>AddMember</code>). Each edge now also carries a native Windows variant (PowerView/Rubeus/Mimikatz) for live WinRM execution alongside the Linux <code>impacket/evil-winrm</code> one.
Orchestrator	Excellent safety design	The model picks an edge index , never a command. Cannot invent a host, cannot author a command, cannot reach an edge outside the collection. A pick outside the list is refused rather than repaired. Now runs live — the approved command executes over WinRM on a selected Windows target (proposes, never auto-fires; regression-locked for the WinRM path too).
advance endpoint	Excellent	Advancement requires a <code>run_id</code> that was approved and exited 0 , verified server-side — transport-agnostic, so the WinRM path advances the walk identically.
Execution tooling	Now present + verified live	<code>impacket</code> , <code>certipy-ad</code> , <code>bloodyad</code> , <code>netexec</code> , <code>evil-winrm</code> , <code>bloodhound.py</code> , <code>kerbrute</code> , <code>responder</code> , <code>mitm6</code> are in the image (Phase 1). The WinRM driver (Windows execution backend) runs the abuse live against a real Windows/AD box you run in VMware — the graph is no longer synthetic-only. Verified against a real domain controller in build #9 (2026-07-31, <code>backend/Livefire.log</code> : 45 PASS / 0 FAIL / 3 NOT-RUN): a live WinRM round trip reached a real <code>corp.local</code> DC and returned <code>corp\administrator</code> , and <code>bloodhound-python</code> collected the real forest into the typed graph (84 nodes / 625 edges), asserted to be the operator's actual domain rather than <code>sample_data</code> . Unit tests remain hermetic (WinRM mocked) — the live run is the separate evidence, and it found four defects the green hermetic suite could not see.

Detection footprint

Read-only, honest, well-sourced. Real SigmaHQ rule UUIDs, real ATT&CK ids (Enterprise v19.1), and a verifier script (`pipeline/detection_sources.py --verify`) that re-checks every id against live upstream. The "describes what blue sees, never how to be seen less" line is enforced in code. **Scale:** ~133 command specs, 19 distinct ATT&CK techniques — a good purple-team *flavour*, not a coverage tool.

C RTP conflict — now reconciled (Phase 4, D10). CRTP includes AMSI/logging evasion, which the panel used to refuse outright. It now has an **additive second channel**: the blue-team detection view is byte-for-byte unchanged and still passes the never-prescribe guard, and a separate, opt-in `opsec` channel adds the offensive half — what

makes a command loud, the quieter tradecraft, and, mandatorily, *what still records it*. The OPSEC channel has its own guard (it may never advise disabling/clearing/tampering with a sensor); the LLM may extend it for uncatalogued commands, marked `ai_suggested`. See Part III, Phase 4.

Code scan (SAST)

Well-built and genuinely safe. Static-only invariant asserted at every `_spawn()` and again by `test_codescan_safety.py`. Deliberately orthogonal — imports nothing from the engagement/executor/scope model. **Limits — now widened (post-assessment):** the offline bundle grew from 19 rules (Python/JS/TS) to **34 across 8 languages** — Java/Go/PHP/Ruby/C# added (`rules/hackpit-languages.yaml` : command injection, SQLi, unsafe deserialization, code-eval, file-inclusion, SSRF), plus a **ruleset picker** (bundled / per-language / a registry pack for the full online catalogue). Still offline-first. See "Post-assessment refinements".

AI code-audit fan-out (New, 2026-08-07). A second mode alongside the rule scan (`codescan/ai_audit.py`), porting open-kritt's context-saving decomposition onto the `reasoning/` substrate: **map the repo's externally-reachable entrypoints and their flows once, then hand each downstream agent exactly one flow to verify against source** — a concrete vuln-with-attacker-path or an honest no-finding stub — then dedup + severity-rank by `IMPACT_LEVELS`. A non-concrete claim is downranked to a stub by the concrete-or-stub gate; specialists are KB-grounded (`reasoning.retrieval`), the verify step is model-tiered (`reasoning.tiering`), and `patched-since` restricts the whole audit to a git diff. **No new gate** — the proposer reads source and calls the LLM layer, it **executes nothing** (AST-locked in `test_ai_audit_safety.py`): it is one approved job (the ZAP/nuclei justification), any PoC a finding offers is a **string to run approve-each** through the existing executor, and the diff provider + engagement-state sink are **injected from `main.py`** so codescan stays orthogonal (no `state` /git import of its own). Degrades to a deterministic heuristic analyst when no LLM is reachable. See the dated section below.

Finding pipeline (New, 2026-08-07). The AI audit is the heaviest producer of findings, but every surface (recon/nuclei/AD/cloud/IMDS/manual) makes them — so open-kritt's finding-processing machinery is ported **cross-cutting** into a pure-data `backend/findings/` package: a **dynamic/structured schema** (`schema.py` — `FIELD_TYPE_MAP`, `normalize_output_format`, `output_schema`, `validate_payload`; base fields + an engagement-defined `extra` map), **automatic de-duplication** (`pipeline.py` — a stable key over location + type collapses two wordings of the same bug into one, worst-severity-wins, idempotent so re-ingest never multiplies, with a "merged N duplicates" note), **pluggable severity rankers** (`rankers.py` — a per-engagement rule set rescores; ships `default` / `bug-bounty-payout` / `compliance`, the last two being *different lenses over the same findings*), and **post-scripts** (`postscripts.py` — a post-finding hook: `validate` / `report` run in-process and execute nothing, `poc` returns an **approve-each** command; `postScriptLocks` refuse a concurrent double-run). The structured fields were wired through all three schema places (the `state.models.Finding` dataclass + migration-safe `store.py` columns + the frontend), and a round-trip test proves no `response_model` strips them. **No new gate** — ranking/dedup/schema are pure data (AST-locked in `test_finding_pipeline_safety.py`, which imports no cockpit/executor/state and names no gate symbol); only a **command** post-script touches the executor, and only approve-each — the coupling (dict→ `Finding`, command post-script→gated executor) lives in `main.py`. See the dated section below.

Engagement governance — RoE / ConOps / Deconfliction / OPPLAN (New, 2026-08-07). The single most on-brand addition: it turns "*human approves each command*" into "*human approves each command inside a written, agreed operating frame*." Before an engagement goes live, the operator drafts (LLM-assisted, **propose-only**) and approves four governance documents — **Rules of Engagement**, **Concept of Operations**, a **Deconfliction Plan**, and an **OPPLAN** (a list of **objectives**, each with a status **state machine** — `pending` → `in-progress` → `completed` / `blocked` / `cancelled`, `completed` / `cancelled` terminal — one or more MITRE ATT&CK technique ids, an OPSEC level and an optional C2 tier). The data model + state machine are ported wholesale from **Decepticon**

(`tools/opplan.py` / `conops.py` / `roe.py` / `killchain.yaml` , Apache-2.0 — attribution in `NOTICE` / `THIRD_PARTY_LICENSES`) and reshaped onto HackPit's upsert-only SQLite (`backend/state/governance.py` + `killchain.py`). Objectives drive the orchestrator's targeting (the proposer aims *toward an active objective*; an approved exit-0 run records itself as the objective's advance evidence, the same `advance` model the graphs use), a **MITRE ATT&CK coverage matrix** renders which tactics/techniques the engagement exercised, and the whole package flows into the report. The RoE **formalises** the scope handrail — it references the same scope model (`cockpit/scope.py`), and an out-of-RoE scope is **flagged in the UI, not machine-blocked** (the advisory check lives in `main.py` , the only place `state/governance.py` lets scope be imported from). **No new gate** — the whole package is authored + human-approved documentation plus a formalised scope frame; it is **advisory to the human**, never a machine veto, and per-command human approval stays the actual bound (matching the standing "target lock is a handrail" decision). Generation is **propose-only** (`backend/governance_draft.py` , the generative layer like `attack_path.py` — degrades to a deterministic scope-derived skeleton with no LLM). AST-locked in `test_governance_safety.py` (governance + killchain + the drafter make no eval/exec/subprocess/socket/HTTP call, import no cockpit/executor/sandbox and name no gate symbol; the drafter persists nothing and advances no objective; the state machine governs objective status only). See the dated section below.

Web3 / smart-contract audit (New, 2026-08-07). Three built-in **playbooks** on the AI code-audit fan-out — `evm-external-flow` (Solidity), `cosmos-abci-halt` (Go/Cosmos-SDK), `anchor-solana` (Rust/Anchor) — ported from open-kritt's `external-flow-analysis` and `Cosmos ABCI Panic Halt Review` . A **Playbook** steers the SAME three stages: it appends a domain framing to each stage prompt (LLM path), scopes the mapped file extensions to one language, and selects a language-specific heuristic sink group (`ai_audit_web3_rules.json` , the no-LLM demo/degradation path). Findings are **chain/contract/function-tagged** and **KB-grounded** in three new authored methodology entries (external-flow analysis, the four Cosmos panic classes, the Anchor account model). A propose-only **tool pass** (`codescan/web3_tools.py`) builds `slither` / `mythril` / `echidna` / `forge` command STRINGS the operator runs **approve-each** in the `:kali` sandbox and parses their JSON/text output back into the same finding shape — it executes nothing (added to the `test_ai_audit_safety.py` AST no-exec lock). Web3 tooling was added to the arsenal + `Dockerfile.sandbox` + `docker/proof/web3_install_proof.sh` (**image rebuild is the operator's step**). **No new gate** — the analysis reads source and proposes, exactly like the web-app audit. See the dated section below.

Frontend

16 routes, real API wiring throughout, SSE streaming, no mocked data layer (`cockpitSample.ts` is the only sample content, explicitly labelled). New in Phases 1–3: the arsenal availability band, per-run time-budget + detach controls, the engagement-state/task-tree panel, the credential-vault "use" action, and the persistent `:kali` shell UI. **Gaps**: no global current-target/engagement context, no engagement export/import, no multi-target view, and the accepted 10-error `react-hooks` lint baseline (documented; `next build` passes exit 0).

3. Knowledge base assessment

2,617 entries (1,601 at the time of the assessment; +66 in Phase 4, +950 in Phase 5) · median body 3,000 chars.

Source distribution:

Source	Entries
hacktricks	715 (45%)

Source	Entries
madstuff	260
some-hacking-resources	233
writeups	59
payloadsallthethings	49
claudе-bug-bounty	46
peh-notes	43
htb-writeups	41
claudе-red	39
oscp-cpts-notes	36
decepticon	33
htb-academy	13
hackpit-authored	12
htb-box-pdfs	9
galaxy-checklist	7
htb-my-resources	5
shodan-dorks	1

Phase 5 added: hacktricks-cloud 578 · portswigger 372.

Tiers: 2,265 tier-3 · 241 tier-2 · **111 tier-1** (your own) — see finding 2. **Ranking rebalanced (post-assessment):** tier is now a *substance-gated* nudge, not a blanket trust prior — a thin tier-1 stub no longer outranks richer, more relevant lower-tier content, and a command-rich page wins close calls regardless of tier. See "Post-assessment refinements".

Findings

1. **PayloadsAllTheThings payload depth — now recovered (Phase 4, D11/D12).** PATT's 66 `.txt` payload lists + the shodan dork list are ingested at **payload-level granularity** (64 `payload-set` + 2 `dork-list` entries, `no_merge` -guarded so consolidation can never collapse them again), each a searchable entry over a full sidecar corpus mounted read-only at `/payloads` ; 56 `oscp_tools` files went to the Scripts Arsenal. KB 1,601 → 1,667. See Part III, Phase 4.
2. **"Grounded in your own notes" is mostly "grounded in HackTricks."** Search boosts tier-1 (`search.py:58`), but with 111 tier-1 entries that lever has little to pull on. **Growing tier-1 is the highest-leverage KB work still available — still open**, and now the *only* open KB item. Phase 5 grew the KB by 950 entries, all tier-3, so the ratio moved the wrong way: the fix is writing, not ingesting.
3. **Empty categories — mostly fixed (Phase 5).** Was: cloud 2 · mobile 2 · iot 2 · forensics 1 · ics 1 · phishing 1 · supply-chain 1. Now **cloud 535** and **supply-chain 47** (HackTricks Cloud), and web/API coverage is deep

(PortSwigger). **Still thin:** mobile · iot · forensics · ics · phishing — none of which is on the target list, so they stay unfilled by choice rather than oversight.

4. **HTB Academy — narrowed, not closed (Phase 5).** The local folder is an 18-module *slice*, not the CPTS syllabus, and HTB content is proprietary so nothing is scraped. Auditing the slice against what had been ingested found two files silently lost — one tracked in git but absent from disk, one lost purely for having a `.txt` extension — both now recovered. Closing this properly needs the course itself.
5. **Missing as retrievable topics — fixed (Phase 5).** API security (GraphQL/REST), OAuth/SSO, race conditions, business logic, request smuggling, prototype pollution, deserialization and SSRF→cloud-metadata chains all now resolve, most to PortSwigger Academy material with the lab that drills each one alongside.
6. **No CVE/exploit index — fixed (Phase 5).** `:exploits` answers service+version → CVE → public exploit over the sandbox's own exploit-db catalogue, version-compared rather than substring-matched.

Item 2 is the only KB item still open. Items 3–6 were the recommended next batches; they were never in scope for the first four phases and were built in Phase 5.

Recommended ingestion, in priority order — status

1. **PortSwigger Web Security Academy** — the single best web-security source, and structured. **Done (Phase 5, +372).**
2. **HTB Academy modules properly (CPTS syllabus).** **Partially — the local slice is fully ingested; the syllabus needs the course.**
3. **HackTricks Cloud.** **Done (Phase 5, +578).**
4. **PayloadsAllTheThings re-ingested at payload-level granularity (see §4).** **Done (Phase 4).**
5. **An exploit-db / CVE mapping keyed on service+version.** **Done (Phase 5) — and deliberately not as KB prose; see Part III.**

4. Source-usage audit — "did I use my sources properly?"

Measured on disk against the built KB, not inferred.

4.1 The headline finding: the pipeline only reads `*.md`

`pipeline/ingest.py:229` → `sorted(source_path.rglob("*.md"))` ; `pipeline/ingest_notes.py:599` likewise.

Every `.txt` , `.json` , `.py` , `.sh` , `.csv` , `.yaml` file in every source tree was invisible to the pipeline.

Source folder	Files	.md	KB entries	Non-md skipped	Of which is <i>real content</i>
hackdic	807	806	715	1	—
madstuff	798	797	260	1	—
some hacking resources	233	233	233	0	—
PayloadsAllTheThings	481	134	49	347	66 <code>.txt</code> payload lists + 28 <code>.py</code> scripts
oscp-cpts-notes	288	85	36	203	174 PNG screenshots (+ git internals)

Source folder	Files	.md	KB entries	Non-md skipped	Of which is <i>real content</i>
oscp_tools	99	1	0	98	28 .exe binaries, ~22 .ps1 / .sh / .py scripts
HTB_academy	48	16	13	32	3 PNGs — effectively nothing
shodan-dorks	32	1	1	31	1 .txt dork list
Hacking-Tools	31	1	0	30	nothing — git internals only
PRACTICAL ETHICAL HACKING NOTES	110	45	43	65	images
more new resources (writeups)	118	72	100	46	images
htb my resources	5	5	5	0	—
htb boxes	10 (pdf)	0	9	—	PDF path handled separately
PentestTools	0	0	0	—	empty folder

Scale: stripping git internals, images and compiled binaries, the genuinely lost knowledge was roughly **95 files** — chiefly **PATT's 66 .txt payload lists** (the pipeline ingested the *explanations* and discarded the *payloads* — backwards for bug bounty), the **shodan-dorks .txt**, and **~22 oscp_tools scripts**. **Fixed (Phase 4, D11/D12):** an additive Route-B ingester (`pipeline/ingest_corpora.py`) folded all 66 payload lists + both dork lists into the KB as `payload-set` / `dork-list` entries with full sidecar corpora, and 56 `oscp_tools` files into the Scripts Arsenal — without rewriting a single existing entry (verified byte-identical) and idempotently. Two env traps hit and handled: `rglob` silently omitted 2 of 66 files (OneDrive dehydration → union with `git ls-files`) and 22 files were AV-locked (→ git-blob recovery).

4.2 Consolidation is lossy for some sources

`madstuff` 797 md → 260 (33%). `oscp-cpts-notes` 85 md → 36 (42%). `PayloadsAllTheThings` 134 md → 49 (37%). Some is correct deduplication into HackTricks; a 3:1 collapse on your OSCP/CPTS notes is worth an eyeball — those are exactly the entries you'd want surviving as tier-1/tier-2.

Audited (post-assessment) — healthy dedup, no loss. A read-only provenance audit of the OSCP set: **36 survived as their own entries (33 with real commands), 12 folded into stronger canonical pages** (`PayloadsAllTheThings/peh-notes/some-hacking-resources` — preserved as `also_covered_in` + variants, not deleted), and the remaining ~37 were **same-technique OS/lab splits merged together** (the GitBook ships Linux/Windows/lab variants as separate files). **Zero** OSCP entries were dropped as low-value (only 11 were dropped KB-wide). Also corrected: `oscp-cpts-notes` is a **tier-3 public GitBook repo**, not your own writing — your genuine tier-1 comes from `writeups` / `peh-notes` / `htb-my-resources`. **Decision: no re-ingest** — re-adding with `no_merge` would just recreate duplicate technique pages that hurt search. See "Post-assessment refinements".

4.3 The biggest missed idea — PentestGPT's Pentest Task Tree — is now BUILT

The assessment's single highest-value reasoning-layer recommendation was PentestGPT's task tree: a *persistent, hierarchical, status-tracked engagement state the model reads from and writes back to every turn*. This shipped in **Phase 2** (`backend/state/tasks.py`) — layered `1 / 1.1 / 1.1.1` tasks with `todo | done | n/a`, updated as results arrive via model-proposed operations that code validates. It is retained here as the rationale for what was built. (Implementation note: HackPit's variant has the model return *validated operations* rather than rewriting the whole tree, so one bad response cannot silently rewrite the plan.)

4.4 The structural pattern

The pipeline is excellent at *techniques* and lossy on *lists, workflows and payload corpora*. A checklist is a *sequence*, a dork list is a *set*, a payload collection is a *corpus* — all three were forced through a technique-shaped, markdown-only schema and collapsed. **Both fixes shipped (Phase 4)**: the Route-B ingester accepts non-markdown inputs, and the second entry shapes (`meta.kind = payload-set / dork-list`, `no_merge` -guarded) survive consolidation intact. The `checklist` shape is defined and available; no checklist source was in the D12 scope, so none were ingested yet.

5. What is genuinely excellent — do not refactor

- **The safety architecture is real.** Source-scan regression locks on `:kali` and session stdin, four ordered gates, clean lab/engagement mode split, `assert_isolation_proven()` as a structural runtime check. Better than most commercial tooling. (The Phase 1–3 work preserved every one of these invariants; the test suite grew with it.)
- **The grounding discipline is the best idea in the project.** Grounded-vs- `ai_suggested` labelling; `entry_ids` must resolve or the step is downgraded; commands come from the KB, never the model; **evidence spliced programmatically into reports rather than written by the LLM**; `foreign_refs` as an honesty marker; the Channel-2 leak guard.
- **The AD orchestrator picks an EDGE, not a command.** A safety design that also improves output quality. The Phase-2 task tree follows the same principle (validated ops, not free-form rewrites).
- **`advance` requires an approved, exit-0 run** — evidence-based state transition, verified server-side.
- **The consolidation pipeline** — alias-key + cosine match, structural merge, idempotent, provenance preserved.
- **The detection footprint's framing and sourcing** — real Sigma UUIDs with a drift verifier.
- **`substitute_target`** — CIDR sentinel, host-position gating, refusal to rewrite what it can't rewrite confidently.
- **The BloodHound parser** — v4/v5/CE reconciliation with honest coverage warnings.

6. Security posture of HackPit itself

- **No authentication anywhere.** No `Depends`, no key, no session. CORS restricted to `localhost:3000`. Honestly documented in code.
- On a laptop this is fine. **On a VPS attack box it is an unauthenticated RCE endpoint** — `POST /cockpit/kali` (and the new persistent-shell routes) run arbitrary `sh -c` in a container with full network reach to your host and LAN. Auth is a hard blocker before any non-localhost deployment.

- Secrets handling is correct: `llm_config.json` gitignored, API key written to disk but never returned, collector preview redacts the password/hash. The engagement-state DB (which now holds **real captured credentials**, by decision) is gitignored like the rest.
 - The repo is code-only; KB, sources, sessions DB and `.env` are all gitignored. A prior full-history audit confirmed no secrets were ever committed.
 - **Known:** a self-scan flagged SSRF and CORS misconfiguration in HackPit's own backend (memory obs #1496). Worth revisiting before any exposure.
-

7. What the assessment examined

Read in full: every non-test backend module (executor, allowlist, config, sandbox, engagement, scope, kali, session, runstore, router, models; attack_path, main, orchestrator, llm, report, chat, context_channel, sessions; all of adgraph, arsenal, detection, codescan), `pipeline/search.py`, `docker/Dockerfile.sandbox`, `docker-compose.yml`, the QA report, and the live KB.

Structurally mapped: the frontend (routes, API surface, component→endpoint wiring, demo/mock scan — none found beyond the labelled `cockpitSample.ts`); the test files; `pipeline/consolidate.py`.

Source trees inspected on disk: `hexstrike-ai` (151 `@mcp.tool` defs), `mantishack`, `Decepticon`, `PentestGPT` (task-tree prompt read directly), `claude-red`, `claude-bug-bounty`, `Galaxy-Bugbounty-Checklist`, `hackdic`, `htb boxes`, the writeup trees, some hacking resources. File counts and markdown ratios measured per tree.

Verified live (assessment): container tooling (37-tool probe), SYN-scan capability, nuclei templates, container capabilities, KB statistics, ingester file-type globs.

PART II — WHAT REMAINS

Decided with Zaid, 2026-07-26. All five phases and every decision that drove them are in Part III. The build is complete; this section is only what is intentionally left — some of it deferred, some of it rejected outright.

Standing policy (unchanged, still governs the project)

These are not open work — they are the rules the project runs by, and building is finished, so they no longer need a decision table. Kept here as the active policy:

- **D2 — Windows + Docker Desktop; a VPS is added later, only for bug-bounty callbacks.** Docker Desktop is a real Linux VM (WSL2); the one thing it can't give is a publicly reachable listener for blind SSRF/XXE/RCE. The repeater (Phase 4) sends and reads the *direct* response; the VPS piece is still deferred until bounty work needs it.
- **D5 — Per-command approval stays. Risk-tiering is rejected, not deferred.** One-keystroke approve shipped (Phase 3). Tiering — auto-running "passive" commands, batch-approving read-only plans — was reconsidered and **decided against** (§1.2): in engagement mode per-command approval is the *only* thing bounding where a command may go, and a classifier deciding when to skip the human is a new component on the wrong side of that boundary. Every execution surface since preserves this — repeater, tunnels (rewrite *visible before* approval), and Phase 5's `:terminal` (human-only, no agent path).
- **D8 — HexStrike is a reference, never an execution backend.** Its tools bypass all four gates, the target-lock, the scope model and the audit trail. Rejected, not deferred.
- **D9 — CRTP Windows execution: CLOSED via a WinRM driver against an external VM (was "accept the gap").** The original decision accepted that CRTP's PowerShell/.NET tooling can't run on the Linux container. That gap is now closed the honest way: **no Windows VM lives inside the project** — HackPit *drives* one you run in VMware over WinRM (Model A), a new execution transport behind the same gates. Windows-only tools now report runnable when a Windows target is selected (reconciled per active target); on a Linux run they are still N/A and marked. The AD graph executes live against that target. See the Windows execution backend in Part III.

Explicitly deferred (with the reason)

- **A VPS for blind out-of-band callbacks (D2)** — a NAT'd laptop has no public listener; deferred until bounty work needs it.
- **A Windows execution target for CRTP (D9) — DEMONSTRATED LIVE** (build #9, 2026-07-31): the WinRM driver executes CRTP/AD work live on an external VMware VM you run. No longer "built but only unit-tested" — a real PowerShell round-trip lands on the profile host over WinRM, the whole-script danger classifier fires on real input, the credential leaks into nothing, and output ingests into state (Task 1, 14/14). See build #9 in Part III. The one caveat still open is a full Windows-target C2 callback (Task 4), which needs a routable listener.
- **The AD graph run off real collection, not synthetic data — DEMONSTRATED LIVE** (build #9): the gated, scope-locked `bloodhound-python` collected the real `corp.local` (84 nodes / 625 edges), it parsed into the typed graph rather than `sample_data`, and the route to Domain Admin was computed on the real graph (Task 2, 10/10). The live abuse walk then executed a real DCSync over the domain — four NTLM hashes including `krbtgt` — through the danger red-confirm, with real loot ingesting into state and the walk advancing only on

the approved exit-0 run (Task 3). Four defects that only a real domain could surface were found and fixed (credential-in-record leak, impacket parser-name mismatch, KB grounder arming a free edge, collector FQDN classifier), each regression-locked. The native-tooling abuse variants (Rubeus/mimikatz on the DC) stay not-run: a bare Server 2022 has no offensive tooling and staging it was out of scope.

- **Growing tier-1** (§3, finding 2) — the only KB item still open, and the one that cannot be solved by ingesting: it needs Zaid's own writing.
- **The HTB Academy syllabus proper** (§3, finding 4) — the local slice is fully ingested; the rest is proprietary and needs the course.
- **Thin categories with no target on the list** — mobile · iot · forensics · ics · phishing (§3, finding 3). Unfilled by choice.
- **C2-framework install + the persistence evasion/OPSEC half** — **DONE** (build #4, no longer deferred): Sliver is installed and wired to a gated panel, and the persistence OPSEC notes were backfilled. Empire and weeveily remain catalogued reference-only by choice.
- **Live-fire verification of the C2 and tunnel surfaces (build #4) — DEMONSTRATED IN LAB, with exact caveats.** A controlled lab exists (`docker/proof/live_fire_proof.sh` + `live-fire-lab.yml` , four containers on `internal: true` networks) and drives the surfaces through their shipped gated entry points: **66 passed, 0 failed, 3 not-run**. What is demonstrated live: an implant that actually builds and **a beacon that calls back**; a **proxychained request routed through the pivot** into a subnet the operator box has no route to, using the `argv wrap_command` produced; all three lifecycle gates (Sliver server, both pivot kinds, DNS listener) refusing on all three limbs with nothing spawned; every listener coming up, being confirmed bound by `ss` inside the container, and staying bound for the whole phase; the pivot subnet entering scope only via the explicit amendment, with the deep target confirmed unreachable beforehand; the DNS key absent from the pasteable one-liner; and **containment holding live, measured from inside the implant host** — no internet, no host, C2 only. What is **not** demonstrated, and why, kept deliberately separate from the above:
 - **The three defects the first live-fire run found are FIXED (2026-07-28), and the fixes are what the numbers above measure.** The Sliver implant build now runs through the console `sliver-server` hosts (Sliver 1.5 has no `generate` subcommand on either binary) and decides `generated` vs `failed` by reading the artifact back rather than trusting a console exit code; both listener lifecycles now hold a console binary's stdin open and report a status they **observed**. Fixing them surfaced two more of the same family, also fixed: a stop that killed the `docker exec` client while leaving the server running inside the container (next start: `EADDRINUSE`), and a `proxychains` config shipped pointing at Tor rather than the SOCKS5 port HackPit's own rewrite targets.
 - **iodine's IP-over-DNS tunnel is now demonstrated** carrying traffic through the gated surface, confirmed DNS-encapsulated on the wire (2026-07-28). **dnscat2's handshake is decisively an upstream defect** in the pinned build — `tcpdump` shows the server itself answering `NXDOMAIN` to correct queries, reproduced over loopback in one container — sitting above HackPit's proven-bound listener. What stays operator-side: a **delegated DNS zone** (a domain with an NS record pointed at a publicly reachable listener, the D2 VPS gap), which neither tunnel can exercise without it.
 - **ligolo's interface routing** needs `session` then `start` typed into the proxy's console. HackPit holds that console's stdin open and deliberately never types into it, so completing a ligolo session stays the operator's step; the routed-traffic claim is carried by chisel, whose server needs no console. **Detonating an artifact on an instrumented Windows host** remains untouched (below), and pairs with the D9 VM. The original list, for the record:
 - **Catch a live Sliver beacon.** The server lifecycle starts and stops and is audited, but no implant has ever called back. Needs a listener reachable from a victim host — i.e. the same VPS gap as D2, or a lab VM on a

routable segment. Until then the implant argv is only as correct as the catalog's documented syntax; a wrong flag surfaces as `status='failed'`, not as a containment failure.

- **Stand up a real DNS tunnel.** `dnscat2-server` and `iodined` start inside the engage sandbox and the client one-liner is handed back, but no delegated zone exists to test against. Needs an authoritative zone the operator controls with an NS record pointed at the listener. `iodined` also needs a TUN device in the engage sandbox — plausible given its `NET_ADMIN`, but unproven.
- **Run a generated artifact on a Windows box and watch the telemetry.** The evasion engine emits artifacts and the paired footprints claim specific Event IDs and Sysmon codes. Those claims are transcribed from ATT&CK v19.1 and SigmaHQ, not observed here. The honest version of this feature is only fully honest once someone detonates an artifact on an instrumented host and confirms the footprint matches what actually landed in the log. This is the natural pairing with the D9 Windows VM. None of this blocked the build shipping — the containment properties are what the test suite locks, and those are provable without live infrastructure. Build #7 converted the containment half from "locked by tests" to "holds live", and the efficacy half from "documented" to "three specific defects, named and reproducible" and then to "fixed, with the callback and the routed traffic actually observed". What remains before build #4 could be called field-ready on these surfaces is the DNS session and a delegated zone.
- **An independent review of build #4's Tasks 8–13 — DONE** (no longer outstanding). Run in a fresh session (the safety classifier had begun refusing subagent dispatch partway through the build, and it refuses on accumulated context). It found **no containment hole** — no agent path, no shell, the container never request-influenced, no delivery primitive — and five substantive defects, all fixed: a multi-technique request that built one artifact and described a different one; the T1690/T1685 mis-mapping above; two stub headers describing memory patching when the stubs are managed-reflection bypasses; a no-agent-path test whose positive control counted the wrong variable (99 files reported, 5 actually scanned); and a `_gate_request` comment asserting a scope-gate behaviour that does not hold. Plus seven minor items — an image smoke test that could not fail, unpinned installs, dead fields, weak assertions. Every fix carries a test that fails without it.
- **Gate-integrity Criticals 2 and 3 — DONE (build #5, 2026-07-27).** Both closed; full detail in Part III. **Critical 2** (the WinRM first-token classification — the sharpest known hole in the tool, because it defeated the red-confirm by moving a cmdlet one token to the right on the one path that reaches a real domain-joined host) is fixed by classifying the whole joined PowerShell script, deliberately without parsing it, with the gate and the transport now deriving that string from one shared function. **Critical 3** (source-scan locks covering a fraction of the tree) is fixed by one shared scanner: whole-tree `rglob`, repo-relative allow-lists, an AST pass for indirection, and per-lock planted-violation controls — coverage per lock goes 30 of 69 modules to 67. The plan's third task, the one that matters longest, is written down in `backend/AGENTS.md`. **I2, the ungated tunnel listener start, was outside build #5's plan but has since been closed** — `start_tunnel` now runs the real `validate_request` before spawning, `TunnelStartRequest` carries `approved / dangerous_ack` defaulting false, and the route 403s a safety refusal; detail in Part III.
- **Pinning the build-#4 install layer** (audit M2) — **DONE** (build #7, no longer deferred). All three float-installs are pinned to real resolved values: `dnscat2` to commit `42f8d783...`, the gems to `trollop:2.9.10 salsa20:0.1.3 sha3:2.2.4 ecdsa:1.2.0`, `donut-shellcode==1.1`. `test_evasion.py` now asserts full pinning (three installs, three `ARG` pins, a real 40-char SHA, no floating form left behind) rather than tolerating a declared gap. The fresh build also surfaced a smoke test that **had never passed** — `sliver-server version | grep -qi sliver`, against a binary that prints only `v1.5.42 - <sha>` — so that layer had not built since the check was added; both sliver checks now match `v${SLIVER_VERSION}`, which is strictly stronger. `hackpit-kali:build7` builds clean with every pinned tool resolving and smoke-testing by its real invoked name.
- **The prevalidated path's danger re-check — DONE** (build #7). `iter_run(prevalidated=True)` re-checked approval but not danger; nothing was exposed (the single such caller validates first) but the asymmetry was a

trap for the second caller. It now re-checks danger in all three modes using the same per-mode classifier `validate_request` uses, regression-locked in `test_prevalidated_gates.py`.

- **Kali VM as an execution target; HexStrike as an execution backend; risk-tiered / batch approval — rejected, not deferred** (D8, D5, §1.2).
- **Authentication before any non-localhost deployment** (§6) — no decision needed until a VPS enters the picture, but a hard blocker the moment it does. `:terminal` makes this sharper: an unauthenticated *interactive terminal* onto a box that reaches the host and LAN. **Still outstanding, and now quantified (2026-07-31)**. A Cloudflare Tunnel + Access + app-level-auth build was scoped in full and then deliberately not taken. The scoping established the exact state of the gap: `backend/main.py` has **no authentication of any kind on any route** — no `Depends`, no HTTPBasic, no API key — and the CORS allowlist is a *browser* policy that stops nothing which is not a browser. The sharpest edge is the single WebSocket route `/cockpit/terminal/ws`, a live PTY into the Kali container: gating HTTP alone would leave a free shell. Localhost-only remains the default and the mitigation.
- Screenshot capture, recon diffing/monitoring, checklist-driven runs.
- **A README refresh** (build #12) — three items, held for a session of their own: state plainly that `data/kb` (the ~2,743-entry technique/workflow dataset) is **gitignored and not in the repo**, so a fresh clone's search is empty until the KB is obtained from the author; mention the CI now that the suite runs on every push; and correct the stale counters (README says 2,621 entries / 110 tools / 42 test files — measured: **2,743 / 115 / 56**).
- **Continuous integration — DONE** (build #12, no longer outstanding). `.github/workflows/ci.yml` runs the 56-file hermetic suite, `next build` and an eslint-baseline check on every push and PR, with the live ATT&CK/Sigma drift verifier on a weekly schedule rather than the merge path. The KB the fingerprint locks need is supplied by a derived, equivalence-proven fixture (D26) rather than skipped. All three jobs have since been **run on GitHub and passed**, the drift job by manual dispatch after it was found never to have executed at all.

reconFTW — recon gaps mined, and what to skip

reconFTW (six2dez, MIT) was reviewed for what HackPit's *recon* was missing — its thinnest area (the incorporation narrative is in Part III). Respecting the constraint that governs everything here — HackPit is human-gated, so "incorporate" means adopt reconFTW's **knowledge**, never its auto-runner — the concrete gaps, where each landed, and rough effort:

- **URL-triage pipeline** — `gf` + `qsreplace` + `unfurl` (arsenal `web`) + a KB methodology entry. The single biggest bug-bounty gap: HackPit harvested URLs but had no triage step. *Low — done*.
- **Subdomain permutations** — `gotator`, `regulator`, `subwiz` (arsenal `recon`). An entire enum technique HackPit had none of. *Low-med — done*.
- **Mass resolution** — `puredns` + `massdns` + `dnsvalidator`, with a baked resolver list (arsenal `recon` + `image`). HackPit had only `dnsx`. *Med — done*.
- **JS mining** — `jsluice`, `subjs`, `getjs` (arsenal `web`): endpoint/secret extraction, not just crawling. *Low-med — done*.
- **ASN → CIDR expansion** — `asnmap`, `mapcidr` (arsenal `recon`) for wide-target scope work. *Low — done*.
- **External OSINT** — `trufflehog`, `github / gitlab-subdomains`, `dorks_hunter`, `gitdorks_go`, `msftrecon`, in a new `osint` category (HackPit had no OSINT home). *Low — done*.
- **Bucket/cloud enum** — `s3scanner`, `cloud_enum` (arsenal `cloud`): external asset discovery, distinct from the posture scanners. *Low-med — done*.

- **Injection completers** — `commix`, `SSTImap`, `nomore403`, `crlfuzz` (arsenal `web`): fill the cmd-injection / SSTI / 403-bypass / CRLF tool gaps (skills existed; tools didn't). *Low* — *done*.
- **Recon ordering** — the subdomain-enum sequence and the URL→gf-triage→fuzz pipeline, as two KB `checklist` entries (`recon-methodology-*`) + an advisory note in the composer's real-target prompt. The higher-leverage half — tools without the ordering is the cheap half. *Low-med* — *done*.

Skip — conflicts with the human-gated model (rejected, not deferred): the `reconftw.sh` runner and `reconftw.cfg` engine (autonomous chaining is the exact thing HackPit rejects), `interlace` (parallel command execution), `notify` (autonomous alerts), monitor/incremental/diff auto-rescan, the axiom/Ax fleet, `reconftw_ai` and `faraday`. `inscope` is redundant — `scope.py` is stronger (fail-closed, no DNS at match time). `interactsh` /OOB stays deferred on the VPS (D2, unchanged). `waymore` was considered and left out: `gau` + `waybackurls` + `katana` already cover URL harvesting.

PART III — BUILD LOG (Phases 1–5, shipped 2026-07-26)

Branch `sandbox-kali-image` . Full hermetic safety suite green throughout (36 test files); both Docker proofs 4/4 (lab still egress-less; engage fully open); browser-verified with Ollama (`qwen3:8b`).

Decisions taken (D1–D22)

Every decision that drove the build. D1/D3/D4/D6/D7 were the Phase-1 build; D9 was a standing accepted gap that is **now built** (the Windows execution backend); D2/D5/D8 are standing policy (Part II); D10/D11/D12 were Phase-4 work, now built; D13 is this document; D14/D15 shaped Phase 5; **D16/D17 are build #4** — D16 amends D10 and is the one policy reversal in the project's history.

- **D1 — Fix the sandbox image; HackPit is the attack box.** A Kali VM's advantages (root, raw sockets, VPN, `/etc/hosts`) are all reachable in Docker via capabilities + `/dev/net/tun` . *Built (Phase 1).*
- **D3 — Capabilities + root on the ENGAGE sandbox only.** Lab keeps `cap_drop: ALL` + non-root. *Built.*
- **D4 — All three sandboxes get the new toolset; only the lab's *network* isolation stays.** *Built.*
- **D6 — Catalog describes reality, not aspiration.** *Built.*
- **D7 — Startup reconciliation check.** *Built.*
- **D9 — CRTP Windows execution.** Was "accept the gap; no Windows VM." Now **closed by the WinRM driver** (Windows execution backend): HackPit drives an external VMware VM you run — no Windows VM inside the project. Windows-only tools run when a Windows target is selected; still N/A + marked on a Linux run. *Built (the AD graph executes live).*
- **D10 — Allow evasion/OPSEC content as an additive second channel, keeping the blue view.** *Built (Phase 4).*
- **D11 — Fix the KB gap via Route B (additive merge), not a rebuild.** *Built (Phase 4).*
- **D12 — Ingest PATT payloads + shodan dorks into the KB; `oscp_tools` into the Scripts Arsenal.** *Built (Phase 4).*
- **D13 — Plan lives in this document.** *Done.*
- **D14 — The PTY is a SECOND surface, never a replacement.** The old framing treated auditable transcripts and full-screen tooling as alternatives. They are only alternatives on one surface. `:kali` keeps its sentinel; `:terminal` gets the pty; both are audited and identically contained, and a test asserts `kali.py` never grows a pty. *Built (Phase 5).*
- **D15 — The CVE index is a lookup table, not a KB batch.** "Find the exploit for this version" wants an exact, deterministic table; embedding it as prose would make the one query it exists to answer fuzzy. Its own data shape, search path and surface — and it executes nothing. *Built (Phase 5).*
- **D16 — Lift the OPSEC sensor-tamper ban; the honesty marker becomes the sole invariant.** Amends D10. The OPSEC channel previously refused any sensor-blinding phrasing, which became incoherent once build #4 shipped an engine that emits AMSI-patch and ETW-blind artifacts: the tool would have been producing artifacts it was forbidden to describe. The ban is lifted **entirely**, and two invariants carry the whole weight

instead — every note must name what **still records** the activity, and the blue-view footprint is always produced alongside and can never be suppressed. The blue-side describe-never-prescribe guard and all four execution gates are **unchanged**. A posture shift on a public repo, taken deliberately. *Built (build #4).*

- **D17 — Split the gating for C2 rather than gate it uniformly.** Server and listener **lifecycle** is human-only with no red-confirm (operator infrastructure on the operator's own sandbox — no target is touched, so clicking start is the approval); **artifact generation** is a fully gated command (approval + scope + red-confirm) because it produces something that will run on someone else's machine. Uniform gating would have meant either a meaningless confirm on every server start or no confirm on a payload build. *Built (build #4).*
- **D18 — The loop may REASON, but it still never executes, and approval stays per-command.** Build #8 makes the proposer genuinely deeper — working memory (a tried/failed ledger), hypothesis-first proposals, a scored candidate frontier, failure diagnosis, a refute-first critic, domain-specialist routing, fingerprint-keyed retrieval, and a model-tier config lever. Every one of those produces **proposals and rationale only**. The decision recorded here is the boundary they were built against and re-locked to: `orchestrator.py` and the whole new `reasoning/` package have **no path to any execution surface**, the frontier holds untried leads rather than a queue that fires, and there is **no auto-approve / batch-approve / run-the-chain** anywhere — a human still approves every single command. Deeper proposer, identical autonomy. *Built (build #8), regression-locked by an extension of the loop's source-scan onto the whole reasoning package.*
- **D19 — The launcher lives ON the home page, not on its own route.** Roughly a dozen surfaces were built and then invisible — `:arsenal`, `:c2`, `:tunnels`, `:windows`, `:repeater`, `:scripts`, `:code-scan` and the AD graph were reachable only by typing the URL, and the operator's own notes had flagged it twice before it was fixed. Both placements were built as mocks and compared. The deciding argument was that *a page you have to navigate to fixes discoverability only if you remember to navigate to it* — so the index belongs on the screen you already land on. Cost measured before committing: +1,024px on a page already ~1,936px, i.e. ~2 screens to ~3. The hero, the stat counters and the intro are untouched. *Built (build #11).*
- **D20 — Operator identity is configuration, never a constant.** This repo is public, and a real name, email or OSID written into source is in the public git history permanently. Identity therefore lives in a gitignored `backend/operator.json` with env overrides, read through two deliberately separate accessors: `public_profile()` (name + handle) for the browser, `report_identity()` (adds OSID / handle / contact) for a report handed to an examiner. A page has no use for an OSID, so it never receives one. The report block is *spliced by code*, never prompted for — handing the model a real name or an OSID invites it to transcribe or invent one. *Built (build #11).*
- **D21 — A source is DISTILLED into the KB, never parroted into it, and a source that teaches something wrongly is rewritten or dropped.** Set while folding in the first external corpus of live-hunting transcripts. Three rules came out of it and now govern every enrichment batch. (a) **Check before writing:** every candidate technique is grepped against the whole KB first, and anything already covered is skipped and reported as skipped — that batch added 13 entries while deliberately re-ingesting nothing of the 24 existing Java-deserialization, 19 cache-poisoning, 11 open-redirect/OAuth or 3 mass-assignment rows. (b) **Zero is a valid result:** a source that yields nothing is reported as yielding nothing, because a batch padded to look productive is worse than a small one. (c) **Correctness outranks fidelity to the source:** the corpus taught "price tampering" while conflating it with bounty manipulation and discussing selling exploits instead of reporting them; the concept was kept and rewritten from scratch, and that framing was not carried across. *Applied (KB enrichment, 2026-07-31).*
- **D22 — A command is gated on what it RUNS, not on what it is named.** Set while cataloguing the pivot tools. The danger heuristic classified `argv[0]` and stopped there, so `weeveily` demanded the red-confirm and `proxychains -q weeveily ...` demanded nothing — while `cockpit/tunnels.py` builds exactly that argv for the operator to approve. The consequence was precisely inverted: **routing a command through a tunnel made its**

gate weaker, so the further into a network you reached, the less the confirm applied. The rule now is that a wrapper is peeled off and the command that actually executes is what gets classified, with the wrapper named in the reason so the warning still matches the argv on screen. The corollary is the part worth keeping: a catalogued entry is classified by **what it does**, which is why `proxychains` is clean (it routes; it carries nothing) while `sshuttle` sits beside `chisel` and `ligolo` (it needs no agent and no listener, which makes it the quietest way to put a subnet inside reach — not the least dangerous one). *Built (KB repo batch, 2026-08-01), regression-locked through `validate_request` rather than through the predicate, with both halves asserted: the wrapped dangerous command must fire, and a wrapped benign one must stay silent.*

- **D23 — 0xdf is drawn from under the project-wide distil-not-parrot rule; the standing ban on it was reviewed and reversed.** When the fingerprint corpus was first scoped, `ingest_exploitation_writeups.py` carried a hard sourcing line naming 0xdf, lppsec and individual HTB walkthroughs as sources *never to draw from* — and the batch declined the ingest on that line. The operator reviewed that position and reversed it: a third-party writeup is treated exactly like every other source this KB absorbs (D21) — **you learn the technique from it and write an original entry, crediting the source by URL in references ; you never reproduce its prose, structure, screenshots or phrasing.** The reversal is the decision; the honest record is that 0xdf was declined first and that position did not survive contact with the rule already governing every other source. Two things were changed and one was deliberately left alone. The docstring that forbade what the pipeline now does was rewritten to state distil-not-parrot and the attribution requirement, because a comment contradicting practice is worse than either policy. The link-index skip at `consolidate.py:2324` — which drops a 0xdf *link-list* file as "no technique" — was **left untouched**, because a list of links carries no technique regardless of sourcing policy, the same reason `awesome-oscp` yields nothing. *Applied (0xdf fingerprint batch, 2026-08-01).*
- **D24 — a shared version predicate names its boundary convention EXPLICITLY, because two subsystems need opposite ones.** `_version_verdict` in the CVE→exploit index is the single comparison both the exploit lookup and the fingerprint corpus (2.7 retrieval) key on. They disagree on what the stored endpoint *means*: the CVE index stores `versions[-1]` as the **fix** version (first patched → exclusive `<`), while the fingerprint corpus stores the **last vulnerable** version (→ inclusive `<=`, which is what the `lte` kind name literally says). One predicate, an unstated convention, and nothing asserting the obvious — so **35 of the 38 versioned fingerprints silently failed to match the very version they were written about**, the most precise hit there is. The decision, and the reason it is a D-entry: the fix is the predicate, not either caller (per `backend/AGENTS.md`). `_version_verdict` took an explicit `inclusive` argument, and **every** caller now states its convention — `search_service` and `critic` pass `inclusive=False`, the fingerprint matcher passes `inclusive=True` — so neither relies on a default. Option (b) (rewrite the 35 corpus entries to store fix versions) was rejected: it is smaller but leaves the convention implicit, and the next entry author gets it wrong again. The CVE→exploit index is provably untouched — all **110,695 verdicts across 47,108 entries byte-identical** before/after, `test_exploits.py` green. A second defect rode alongside: `fingerprint()`'s first-token heuristic collapsed `Apache Tomcat` and `Apache httpd` onto the same `apache/<ver>` key, so two unrelated products were told apart only by version — a *confidently wrong* hit, not a missing one; it now resolves the product (`tomcat`) not the vendor. Both are regression-locked by tests that iterate the real corpus and carry a positive control (`test_fingerprint_versions.py` , `test_fingerprint_norm.py`), lifting the covered-banner hit rate 70%→93% at 96% precision with near-miss false-fire held at 0%. **This is the THIRD shared-predicate defect in the project's history** — build #5's WinRM `argv[0]` classification and D22's `proxychains` red-confirm laundering were the first two — and the pattern is now worth watching for by name: a guard or verdict shared by two callers with an unstated convention, sitting behind no test that pins the obvious. *Built (fingerprint fix, 2026-08-01).*

- D25 — `fingerprint_match` is reserved for a STRUCTURED match; a bare product-name hit is a distinct, weaker signal.** The last measured gap in 2.7 retrieval was a **20% false-fire on services the corpus does not cover** (`pure` → `Pure-FTPd`, `node.js` → `Node.js`, `minio` → `MinIO`), stable across the corpus's growth. Its cause was a *second* code path from the structured matcher: the version-less substring fallback in `rerank()` set `fingerprint_match=True` whenever the scanned product merely appeared as a substring anywhere in a plain entry — so an entry that had nothing to do with the service was surfaced with the same confidence as a real fingerprint, and the grounding line claimed "this exact stack was solved by X" unearned. The decision was made from measurement, not taste: the eval was first instrumented to split covered fires into structured vs fallback, and it showed the fallback contributes **0 of the 28 covered fires** while being the **whole** of the false-fire — so this was never a stricter-matcher-vs-hit-rate trade-off, and the fallback could be demoted at no cost. `fingerprint_match` now means a structured `meta.fingerprint` match only; an unstructured product-name hit becomes a labelled `fallback_match` that still outranks a pure token match (a version-less scan keeps some signal, which is legitimate) but never claims the exact stack, and the substring was tightened to a word boundary. The fallback was **not deleted** — the instruction and the measurement agreed to keep and label it rather than drive a number to zero by removing real behaviour. Result: UNCOVERED false-fire **20%→0%** with covered hit rate, precision, near-miss and self-match all unchanged. Regression-locked by `test_fingerprint_fallback.py` (real corpus + 15 real uncovered services + positive control), the last of the retrieval-path correctness gaps to close. *Built (residual fix, 2026-08-01).*
- D26 — a corpus CI needs is DERIVED from the real one and proven equivalent; it is never authored.** Set while wiring the first CI. Three fingerprint locks iterate the live KB deliberately (`backend/AGENTS.md` §1: draw from the real population, so a fingerprint added tomorrow is covered without anyone editing a test) — but `/data/` is gitignored, so on a clean checkout they did not skip, they **crashed**. The obvious fix, a small hand-written sample KB, is precisely what §1 forbids and for a documented reason: the gate audit's worst finding was a test asserting on a synthetic value the real system never produces, which stayed green while eight dangerous tools passed the danger gate. The decision is the third option. One measurement made it available: `retrieval._entry_blob` reads `title`, `text`, `body`, `summary`, `tags`, `product`, and the live KB carries **no text, body or product field at all** (it stores `body_md` and `steps`, which that function never consults) — so the fingerprint path only ever sees title + summary + tags. The committed fixture is therefore the **complete** corpus — all 2,743 entries, all 105 structured fingerprints — projected onto the fields the matcher provably reads: 22 MB becomes 1.1 MB with **zero** information loss for these tests. It is not a sample and answers to nobody's judgement about what was "representative". The claim is checked rather than asserted: `test_kb_fixture.py` re-derives the fixture from the live KB and requires byte-identity (staleness), compares **20 probes × 2,743 entries** verdict-for-verdict between the two corpora (equivalence), reads the matcher's field tuple out of `retrieval.py` by AST so teaching `_entry_blob` a new field fails the test instead of silently blinding the fixture (faithfulness), and carries a positive control that a truncated fixture is caught. The two checks that need the live KB report **NOT-RUN**, loudly, in both the test output and the CI job summary — the `test_proof_honesty` discipline applied to CI itself, because a green badge that quietly means "and four checks were skipped" is the same defect as an unfilled proof slot scoring as a pass. *Built (build #12, 2026-08-03).*

Phase 1 — reach a real target

Commits `e4e8de5`, `80a18d5`, `2d0928f`.

- New sandbox image** (`docker/Dockerfile.sandbox`) on `kalilinux/kali-rolling` + `kali-linux-headless`, with `SecLists`, the Kali `wordlists` set (rockyou unzipped) and **~13,400 nuclei templates** baked in (the lab has no egress, so nothing can be fetched at runtime). Thematic tool layers for the specific tools the assessment

found missing — `impacket`, `netexec`, `evil-winrm`, `bloodhound.py`, `certipy-ad`, `bloodyad`, `responder`, `mitm6`, `smbmap`, `enum4linux-ng`, `gobuster/nikto/feroxbuster/subfinder/httpx/amass`, `metasploit`, `exploitdb`, `hashcat/john`, `chisel/ligolo-ng/proxychains/socat/ncat`, `openvpn` — plus `katana/kerbrute/waybackurls/dalfox/gau/rustscan/jwt_tool` as pinned release binaries (no Kali package). **Result:** ~63 of 73 catalogued tools present. Kali's `setcap'd nmap / fping` are stripped at build so `no-new-privileges` doesn't refuse to exec them.

- **Privilege split** (`docker-compose.yml`) — the **engage** sandbox runs `user: root` with Docker's default capability set + `NET_ADMIN` and `/dev/net/tun`, so `-sS / -sU / -O`, `masscan`, `tcpdump`, `privileged-port` binds (`responder/ntlmrelayx`) and VPN all work. **Lab and :kali are untouched** — `uid 1000`, `cap_drop: ALL`, no devices. Verified: `engage` does SYN scans and binds `:445`; `lab/ :kali` refuse `-sS`.
- **Per-request timeouts** (default 180s, ceiling 3600s) on `/cockpit/exec` and `:kali`; **background/detached jobs** returning a `run_id` and a reconnectable replay stream (`/cockpit/runs/{id}/stream`), with a reaper. `:kali`'s old hardcoded 60s is gone.
- **Loot volume** — `backend/data/engagements` → `/loot` on the engage and `:kali` sandboxes (deliberately **not** the isolated lab); each engagement works in `/loot/<id>` via `docker exec -w`, so `nmap -oA` output survives `docker compose down`.
- **Startup reconciliation** (`GET /tools`) — probes the running sandbox with one `command -v` sweep over every catalogued name+alias, filters the planner's prompt to installed tools, and surfaces the gaps in the arsenal UI. Windows-only tools (`rubeus/powerview/mimikatz/winpeas`) are marked `platform: windows`, kept for planning/write-ups, and never proposed.
- **Arsenal catalog** reconciled against the shipped image; the orchestrator prompt no longer hardcodes `gobuster / nikto`.

Phase 2 — make it remember

Commits `c606f6d` (backend), `a8f40ae` (UI).

- **backend/state/ package** — the structured engagement state the assessment identified as the deepest gap: `hosts` → `services`, `endpoints`, `credentials`, `findings` as dataclasses over a shared SQLite store where **every write is an upsert** (re-running a scan corrects, never duplicates). Output **parsers** (`nmap XML`, `httpx/ffuf/nuclei JSON`, `secretsdump`) as pure functions, and a post-run **ingest** that reads both a run's stdout and any file it wrote into its loot directory (that is how `nmap -oA` populates state). The package **executes nothing** — AST-asserted, because ingest runs automatically after every command.
- **Pentest Task Tree** (`state/tasks.py`) — PentestGPT's central idea (\$4.3): layered `1 / 1.1 / 1.1.1` tasks with `todo | done | n/a`, persistent per session, updated as results arrive. The model returns **validated operations** (`add_subtask / mark_done / mark_na / ...`) that code applies against the stored tree; one bad op is rejected on its own and the rest still apply.
- **Orchestrator grounding** — the proposer prompt now leads with the accumulated state + live task tree, with raw output demoted to a short tail. Measured: **460 chars of state vs 6,600 of stdout tails** for 12 runs of one `nmap`, and each fact appears once instead of once per run. A tail window forgets; state does not.
- **Credentials in the prompt: real values** (Zaid's explicit decision) so the planner writes fully-formed commands — one constant (`render.INCLUDE_CREDENTIAL_SECRETS`) with a test proving the flip works, so the choice stays reversible.
- **UI** — the `CockpitState` panel: counters, the task tree (with "seed from plan"), `hosts+services` grouped, `endpoints` with status colours, `credentials` masked with per-row reveal + `validated/untested`, `findings` by

severity.

- **Bonus fix found by the browser pass** — the arsenal `/arsenal` response model was silently dropping `platform / runs_here`, so every tool badged "windows only". Fixed + regression-tested.

Phase 3 — make it fast to drive

Commit `b316b74`.

- **Step 11 — one-keystroke approve.** The guided loop takes `Enter` = approve, `S` = skip, `Esc` = stop, hints shown on the buttons. A dangerous proposal never fires on `Enter` (the explicit danger confirm is still required); shortcuts are ignored while a text field is focused.
- **Step 12 — `_looks_like_host()` fixed.** Flipped from "assume host unless the last segment is a known file extension" to a **positive test**: a dotted token is a host only if its last label is a plausible alphabetic TLD and not a file extension. Stops the false rejections (`directory-list-2.3-medium`, `Mozilla/5.0`, `-oA scan.1.2`, version strings) while still catching real hosts; a genuine off-target host is still refused. This was the prerequisite for any future auto-run tier.
- **Step 13 — persistent `:kali` shell.** One long-lived `docker exec -i sh` instead of a fresh exec per command; each command is delimited over the pipe with a per-command sentinel so output and exit code read back cleanly. `cd /env/background jobs persist` (verified in-browser: `cd /tmp → export F00=... → echo cwd=$(pwd) foo=$F00` returned `cwd=/tmp foo=...` across three separate API calls). Every `:kali` containment invariant intact. No PTY, by design (§1.1). (Caught the Windows `\n → \r\n` stdin-translation bug in the process — the Linux shell was receiving `pwd\r`; fixed with bare-LF stdin.)
- **Step 14 — credential vault.** Captured credentials fill the `<user> / <password> / <ntlm-hash> / <domain>` placeholders the AD templates carry, one click instead of retyping a hash. The placeholder→field mapping lives server-side (`state/credvault.py`) so front and back can't drift; a password fills `<password>` not `<hash>` and vice-versa; only credential placeholders are touched. Verified in-browser: "use `svc_sql`" turned `nmap -u <user> -p <password> ...` into `nmap -u svc_sql -p S3cr3t! ...`.

Phase 4 — content & goal-specific

Commits `90fe260`, `885d776`, `bf46695`, `8486c56`, `937c420`.

- **Item 1 — KB Route B additive ingest (D11/D12).** `pipeline/ingest_corpora.py` — a targeted ingester over only the missing files, run *additively* on the built KB (never through the real pipeline, which would revert downstream enrichment and rewrite a 15 MB gitignored, Defender-sensitive file). Two regression-locked guarantees: existing entries pass through as **bytes** (no key/escape drift on the 1,601 it never looked at), and every line it owns carries `meta.corpus_ingest` so a re-run is byte-identical. Shapes (§4.4): **64** `payload-set` + **2** `dork-list` entries, all `no_merge` so consolidation can't collapse a corpus into a technique page; each holds a capped excerpt while the full list is a **sidecar** under `data/kb/payloads/`, mounted read-only at `/payloads` in all three sandboxes so ffuf/wfuzz point straight at it. **56** `oscp_tools` files went to the Scripts Arsenal as file-backed rows (a path to copy, not 20,000 lines of PowerView), 38 Windows-only ones kept and marked `runs_here=false` (D9). Two env traps handled: rglob's OneDrive omission (→ union with `git ls-files`) and AV-locked files (→ git-blob recovery). KB 1,601 → 1,667; arsenal 1,028 → 1,137. The Defender exclusion on `HackPit\data` was added first.
- **Item 2 — evasion/OPSEC as an additive second channel (D10).** The blue-team detection view is **byte-for-byte unchanged** and still passes `assert_describes_not_prescribes` (a test proves every blue field is identical with

and without the offensive half attached, and that the default `footprint()` carries no `opsec` key — so tagging/reports/run-annotation are untouched). `detection/catalog.py` gains a curated `OPSEC` table (10 notes keyed by spec key: what makes it loud, the quieter tradecraft, the tradeoff, and — mandatory — `still_recorded`, the honesty marker). `resolver.py` gains an opt-in `include_opsec` channel with its own guard (`assert_opsec_is_separate`): a note that fabricates the honesty marker or advises disabling/clearing/tampering with a sensor is rejected. `report.py` gains an off-by-default red-team OPSEC roll-up. The panel renders it as a distinct amber section. LLM may extend to uncatalogued commands, `ai_suggested`.

- **Item 3 — HTTP repeater.** `cockpit/repeater.py` mirrors `:kali` containment (hardcoded open container, human-only + source-scan locked, audited) and adds a scope check `:kali` lacks. **Argv-only curl**, never a shell — method/headers/URL are discrete tokens, the body rides on stdin, so a header of `a; rm -rf /` is one literal `-H`. A human clicking Send is the approval (no per-send prompt); the orchestrator/agent have zero path to `send`. When the send names an active engagement the URL host is scope-checked (out-of-scope → 403, nothing sent). Frontend `/repeater`: compose/send, response view, per-session history with edit-to-resend and a line-level body diff. The VPS-for-callbacks piece stays its own decision (D2).
- **Item 4 — pivot/tunnel routing.** `cockpit/tunnels.py` — chisel/ligolo-ng lifecycle (start the listener in the engage sandbox, hand back the agent one-liner to paste on the compromised host; human-only start/stop, source-scan locked). `route_for` / `wrap_command` are **pure** — they compute a **visible** `proxchains -q` prefix (chisel/SOCKS) or leave the command unwrapped with a route note (ligolo/interface), applied *before* the approval screen so the human approves the exact argv (silent routing was rejected for exactly this reason). A tunnel's subnet enters scope only via `engagement.add_pivot_subnet` — the one deliberate, audited widening path; recon expansion still cannot widen. Frontend `/tunnels`: start form, live tunnels with the copy-ready one-liner + per-subnet "add to scope (by hand)", and a pure route preview.
- **Item 5 — exam-mode report templates + proof.txt as per-host state.** `state` `Host` gains `local_txt` / `proof_txt` + `ownership()`; captured automatically when a command reads a flag file (`proof.txt` / `root.txt` → `proof`; `local.txt` / `user.txt` → `local` — a stray 32-hex hash is never mistaken for a flag) or pasted via `POST /sessions/{id}/state/proof`; the state panel badges foothold vs owned. `report.py` gains four templates — standard, **OSCP** (per-host walkthrough + a `proof.txt` table spliced from state at `{{PROOF_TABLE}}`, computed not model-written), **CPTS** (exec-summary + findings register), **H1/Bugcrowd** (impact-first + a **CVSS 3.1** base score computed by `cvss31_base`, matching the official calculator). Every template keeps the shared grounding rules + the `{{EVIDENCE}}` splice. Frontend: a template picker on the report screen.

Phase 5 — the PTY gap, and finishing the KB

Commits `e8eeef`, `8f0b91b` (terminal), `84e31bf` (CVE index), `f6fb6a4` (KB batches).

- **Item 1 — a real PTY, as a SECOND surface.** §1.1 used to call this a trade: readable transcripts or full-screen tooling. It is not, if they are separate surfaces. `:kali` is untouched — still sentinel-delimited over a plain pipe, still the clean per-command transcripts reports are built from — and `cockpit/terminal.py` adds `:terminal` beside it. `docker exec -t` refuses without a client TTY and a WebSocket handler has none, so the pty is allocated **inside** the container by a constant driver that `pty.fork()` s bash and multiplexes it over the pipes. Its stdin is **framed** (`0x00` = keystrokes, `0x01` = `COLSxROWS` → `ioctl TIOCSWINSZ`), which is what makes resize possible without an in-band escape a keystroke could forge; its stdout is the raw pty stream, straight to `xterm.js`. Containment mirrors `:kali` point for point: hardcoded container (the request carries `session_id` / `cols` / `rows` and nothing else), driver a raw-string constant passed as one argv element, no

isolation gate, **human-only and source-scan locked** like `run_kali`, and the raw stream audited to the run store — capped, and checkpointed every 30s so a crash still leaves a transcript. The WS route pins `Origin` by hand, because CORS middleware does not cover WebSockets. `test_terminal.py` (16 checks) locks all of it *plus the additive property itself*: `kali.py` must still carry its sentinel and must never grow a pty, so the clean transcript cannot quietly die. Verified live — `tty` reports `/dev/pts/0`, `top` and `vim` render, resize propagates $30 \times 100 \rightarrow 43 \times 132 \rightarrow 60 \times 200$.

- **Item 2 — CVE → exploit index (§3, finding 6).** The OSCP inner loop, built as a **keyed lookup rather than another KB batch**: "find the exploit for *this* version" wants an exact table, not prose retrieval, so it gets its own data shape, search path and surface. `pipeline/ingest_exploitdb.py` reads `/usr/share/exploitdb/files_exploits.csv` out of the sandbox image — 47,108 exploits, 27,384 with a CVE, 25,041 distinct — nothing fetched from the internet. exploit-db titles follow `<Product> <Version> - <Vuln>` on ~100% of rows, so parsing the left side turns a flat catalogue into `(product, version) → exploits` with the constraint typed (exact / lte / range / wildcard / none); 32,807 rows yield a version. `backend/exploits/` then does what `searchsploit`'s substring match cannot: `2.4.49` satisfies `< 2.4.50`, sits inside a range, and matches the `2.4.x` line — each with a *different stated verdict*, and the verdict is the ranking's primary key, with token similarity only breaking ties inside a tier. (Blending them let `Apache 2.4.x` outrank the entry naming `2.4.49` exactly — caught on real data, now regression-locked.) Surfaces: `/exploits`, plus a per-service `exploits → link` in the state panel, so a fingerprinted service reaches its exploits in one click. Every hit names the file **already inside the sandbox**. It executes nothing, and `test_exploits.py` asserts the package contains no subprocess path and never reaches the execution layer — finding is not running.
- **Item 3 — PortSwigger Web Security Academy (+372).** `pipeline/fetch_portswigger.py` turns the site into the same shape every other source has, so nothing downstream is special-cased. URLs come from the publisher's own `sitemap.xml` rather than from crawling — it is complete (the all-topics page renders client-side, so scraping it yields 19 of ~140) and nothing is ever guessed; `robots.txt` restricts only `/bappstore/bapps/download/`. 398 pages, 273 of them labs, whose *Solution* steps survive. Two rules written by what dry runs actually did: **labs and sub-topic pages are no_merge** (the first run folded ten Academy pages — Reflected/Stored/DOM XSS, contexts, CSP, dangling markup — into one pre-existing grab-bag called "xss resource", so "Reflected XSS" stopped being retrievable as itself; only a topic *root* may consolidate, which the Academy expresses as URL depth, taking 90 merges down to 26); and the two cheat sheets use `<section>` not `<main>`, so without a fallback they were 2 of 18 pages silently extracting zero bytes.
- **Item 4 — HackTricks Cloud (+578), and the cloud gap closed.** Same GitBook layout and adapter as the HackTricks book already ingested. Two corrections, both from dry runs: it is **not class-grouped** (grouping collapsed 44 distinct CI/CD pages — Jenkins RCE, GH Actions cache poisoning, Okta, Cloudflare — into one candidate), and `CANON` gains only **technique-level** cloud classes. "kubernetes", "ci-cd" and "supply-chain" name a *platform*, not a technique, and `CANON` is the authority on what consolidates; including them collapsed every "Kubernetes X" page into one entry. Their absence is now documented in `CANON` so they are not re-added. **cloud 2 → 535, supply-chain 1 → 47.**
- **Item 5 — the HTB Academy slice, audited (§3, finding 4).** Auditing what the local folder holds against what had been ingested found two files silently lost, both to one root cause — bare `rglob` over `*.md`. `File_Inclusion/README.md` (4.5 KB of LFI module content) is tracked in git but **gone from disk**, the dehydration/quarantine pattern this repo has hit before — and that module is full of web-shell examples, which is exactly what trips the signature. `File_Upload_Attacks/1.txt` (8 KB of module answers and webshell walkthroughs) was lost purely for its extension: §4.1's headline finding in miniature. Now `_all_md` (`rglob U git ls-files`, with git recovery) plus `.txt`.

KB 1,667 → 2,617, 0 malformed rows, embeddings rebuilt incrementally, scripts index rebuilt, `entries.jsonl` verified present and counted after every pass (the Defender-quarantine trap).

Windows execution backend — the WinRM driver (shipped 2026-07-26)

Commits `72b0959` (transport + profiles + `windows` mode), `cf64b8f` (profile CRUD/picker + per-target `/tools`), `a8854dc` (native Windows abuse variants + AD live), `1ce53ce` (frontend), the VM guide, and this doc. Closes D9 the honest way. See `docs/WINDOWS-EXECUTION.md` + `docs/WINDOWS-TARGET-SETUP.md`.

- **A new execution transport, behind the same gates.** `docker exec` is swapped for a WinRM call; nothing about the safety model changes. A third mode, `windows` (alongside `lab` and `engagement`), is selected by `ExecRequest.windows_profile_id`. The target is the profile's **host** — **hardcoded server-side, never a request field** (the same containment shape `:kali` gets from its hardcoded container: a command physically cannot reach a box you did not pick). Gates: `windows` (profile exists) → `target` (host in the engagement scope, if one is also named) → **never-auto-run** approval → danger red-confirm. No isolation gate — it is a real external box, like engagement.
- **Model A — HackPit drives, it does not own.** No Windows VM lives inside the project. The operator runs a Windows/AD VM in VMware Workstation; HackPit opens a WinRM session and runs one PowerShell command string on the box (Rubeus/PowerView/Mimikatz/.NET all run *there*). `pywinrm` is **lazy-imported** so the hermetic suite needs no dependency and no network; the AD-live path is unit-tested with a **mocked** transport.
- **Saved connection profiles** (`cockpit/winprofiles.py`) — a "Windows targets" store in the gitignored `sessions.db`: name, host, transport (WinRM; SSH a documented later seam), port, username, **password or ntlm-hash** (pass-the-hash, presented `LM:NT`), domain. The secret is **write-only** — masked to `has_secret` in every view, read only by the transport, never in a response, record or command line. A **captured vault credential can fill a profile**, resolved server-side so the secret never round-trips.
- **The AD graph executes live.** Every abusable edge gained a **native Windows variant** (PowerView/Rubeus/Mimikatz) beside the Linux one; the walk/orchestrator's proposed command runs over WinRM when a Windows target is picked, and `advance` still requires an approved + `exit-0` run. The danger heuristic learned the native destructive cmdlets (`Set-DomainUserPassword`, `Add-DomainObjectAcl`, `Set-DomainObjectOwner`, `Add-DomainGroupMember`, `Set-DomainRBCD`, `Invoke-Mimikatz`, ...), so a native DCSync/password-reset trips the **same** red confirm as its Linux cousin — the oracle test now checks both transports.
- **Per-target tool reconciliation.** Windows-only tools (Rubeus/PowerView/Mimikatz/winPEAS) report runnable when a Windows target is selected and N/A on a Linux run — reconciled per active target (`/tools?windows_profile_id=...`), never a global flip. Served from `main.py` so the cockpit stays arsenal-blind.
- **Frontend.** A **Windows targets** page (`/windows`) — profile CRUD, masked secrets, connectivity test — and a "run on" picker in the AD walk (Linux sandbox vs a WinRM target); the confirm/flags/command shown follow whichever command will actually run.
- **Safety regression-lock** (`test_winrm_safety.py`): host-locked to the profile / no gate bypass / secrets never leak / **the orchestrator cannot auto-run WinRM** (transport reachable only from the gated executor + the human router probe, source-scanned). Functional path (`test_winrm.py`) uses a mocked transport.

Deferred to a live VM: the live/browser verification against a real Windows box, until the operator's VM is up (build + unit tests are hermetic) — the same way tunnels and AD-live execution were deferred.

Post-assessment refinements (2026-07-27)

Three items from a read-through of this assessment, each done the low-risk way. All query-time / config / ranking changes — **no KB re-ingest, no 15 MB rewrite, no Defender risk**.

- **KB consolidation audited (\$4.2) — no re-ingest needed.** A read-only provenance audit confirmed the OSCP "3:1 collapse" is healthy dedup, not loss: 36 command-rich survivors, 12 folded into stronger canonical pages (still attributed), ~37 same-technique OS/lab splits merged, **0 dropped as low-value**. Also corrected the framing — `oscp-cpts-notes` is a tier-3 public GitBook repo, not the author's own writing. Decision: leave the KB untouched (re-adding with `no_merge` would recreate duplicate technique pages that hurt search).
- **Relevance-first KB ranking (\$3, tiers).** The composer was already relevance-first (BM25 + cosine RRF; tier a small additive nudge), but a *thin* tier-1 stub still got a flat boost that could edge out richer, more relevant lower-tier content. Fixed in `pipeline/search.py`: the tier-1 boost is now **gated on substance** (a stub with no commands + little body gets nothing), and a **completeness nudge** (command-rich, any tier, saturating/capped) lets content decide close calls. Trust breaks ties among substantive entries; it no longer manufactures relevance. Locked by `test_search_ranking.py`.
- **SAST coverage widened (\$ Code scan).** The offline Semgrep bundle grew from 19 rules (Python/JS/TS) to **34 across 8 languages** — Java/Go/PHP/Ruby/C# added (`rules/hackpit-languages.yaml`) — plus a **ruleset picker** (bundled / per-language, or a registry pack for the full online catalogue). Default scan loads the whole offline directory. Locked by `test_codescan_rules.py` (rules well-formed + 8-language coverage + `semgrep -validate` , resolved from the venv the runner uses so it validates rather than skips). The new rules were confirmed **valid** (semgrep `--validate` : 0 errors, 34 rules) and to **fire** on real vulnerable samples (Java/Go/Ruby/C#). One robustness fix fell out of that check: on some Windows builds semgrep itself *crashes* on a `.php` file, and the scan treated semgrep as fatal — so a crash now **degrades to a warning** (the scan still returns with whatever else ran, like bandit), instead of 502-ing the whole scan. (*Engagement export/import was considered and dropped — the existing report generator already covers a shareable dump.*)
- **UI cleanup (Frontend).** Small polish from a working session: the two shell tiles were merged to **one :kali tile that opens the real PTY** (`/terminal`) — full-screen tools render from the everyday shell, while the sentinel "transcript shell" (clean per-command records) is preserved off-nav at `/kali` ; both backends and their containment tests stay intact. The **phishing** category is now surfaced on the front-page grid (its KB entries already existed; forensics/ics/supply-chain stay hidden). The **top nav** dropped the redundant `:library` tile (home is the wordmark), the search field was widened to a one-line field, the unused accent-swatch picker was removed (amber is the default and only accent), and the nav is **optically centered** (a small left nudge balances it against the heavier bordered search box, so it reads centered rather than geometrically dead-centre). All display-only; no backend touched.

reconFTW review + incorporation (2026-07-27)

A reconnaissance engine — reconFTW (six2dez, MIT) — was reviewed end to end for anything worth mining. The short version, and what was folded in:

What it is. An autonomous external-recon / bug-bounty runner: one Bash orchestrator plus eight modules driven by a config file, chaining ~100 tools end to end (optionally across a cloud fleet) and emitting a consolidated report. Its pipeline is OSINT → subdomain enumeration (passive → certificate transparency → ASN/CIDR → permutations → mass-resolve + wildcard-filter → NOERROR → web-metadata scraping → recursive → reverse-IP) → resolution/scope → host/port scan → web probe + screenshots → URL discovery + `gf` triage + JS mining + param discovery + fuzzing → vuln checks → report.

Bottom line: mine the knowledge, never the runner. reconFTW's mechanism — autonomous, fire-and-forget chaining — is the exact opposite of HackPit's one-gated-executor, one-approval-per-command model, so none of the runner, config engine, axiom fleet or monitor/notify machinery is adopted (see Part II's skip list). Its *value* is a

battle-tested set of recon tools, proven flag patterns, and a recon ordering — and that lands squarely on HackPit's thinnest area, external attack-surface discovery. So the worthwhile half was folded in as **data and grounding, not a new execution path**:

- **Arsenal** (`backend/arsenal/tools.json`) — +28 tools, 73 → 101, in a new `osint` category (6) plus `recon` (10), `web` (10) and `cloud` (2): the URL-triage set, permutation engines, mass-resolution, JS mining, ASN expansion, external OSINT, bucket enum and the injection completers, each with `<target>`-based invocation templates. `executes_nothing` is unchanged, and the arsenal safety suite (schema + inertness + no template hardcodes a host) passes at **101 tools / 257 templates** — a rendered invocation is still a string until a human approves it through the same executor.
- **KB methodology** — two `checklist` entries (`recon-methodology-*`): the subdomain-enum ordering and the URL→gf-triage→fuzz pipeline, folded in by a new additive ingester (`pipeline/ingest_recon_methodology.py`) that mirrors the corpora discipline — byte-preserving pass-through, its own idempotency marker, re-run byte-identical, and it segregates the corpus block to the tail so `ingest_corpora` 's byte-identity invariant still holds (verified — `test_corpora` green). KB 2,617 → 2,619.
- **Composer grounding** — an advisory recon-ordering note in the real-target proposer prompt: the sequence as guidance, still one gated command at a time, never a chain.
- **Sandbox image** (`docker/Dockerfile.sandbox`) — the 28 tools installed the way each ships (`go install` for the Go set, `venv` + `PATH` wrapper for the Python set, `apt` for `massdns` / `commix`), with `gf` pattern packs and a static resolver list baked for the no-egress lab. `subwiz` is catalogued but **not baked** — it fetches an ML model at runtime, so it runs only in the egress-enabled sandboxes (`engagement / :kali`), never the airgapped lab; `gotator` / `regulator` cover permutations offline. The ~8–10 GB image was rebuilt and smoke-tested — **all 26 baked tools resolve** in the fresh `hackpit/kali-sandbox:m1` (`trufflehog` moved to its release-binary installer after a recent release began requiring `Go ≥ 1.25`; `subwiz` intentionally absent), with the `gf` pattern packs (37) and a 12,956-line resolver list baked in.

The methodology, the arsenal entries, the composer note and the image rebuild are all complete and verified — arsenal suites green at **101 tools / 257 templates**, `test_corpora` byte-identity intact, the full hermetic safety suite green, and the fresh image resolving every baked recon tool.

Persistence / backdoors (TA0003) enrich (2026-07-27)

The post-exploitation **persistence** layer was reviewed the same way — audit first, add only what is genuinely missing — because HackPit already *executes* persistence on scoped targets through the one gated executor (`engagement` + `WinRM` + `:kali`) and `msfconsole` is installed. This is an enrich: **knowledge, catalog and describe layers only — no new execution capability, no persistence engine or autorunner**.

What the audit found. The earlier " ~60 persistence entries " figure conflated **cloud** persistence with host persistence: of the ~72 persistence-relevant KB entries, **52 are cloud** (IAM / AAD / GCP backdoors) and only 2 were dedicated host-persistence entries. The TA0003 *mechanisms* were present, but only at the command level, scattered through the OSCP/HackTricks corpus (scheduled-task/cron 64 hits, services 59, backdoor accounts 70, web shells 46) — what was missing was the **organised, per-mechanism methodology** tying them together, exactly the reconFTW shape (the raw material existed; the map did not). On the detection side, `attck.py` carried almost none of the persistence technique rows.

- **KB methodology** — two `persistence` `checklist` entries (`persistence-methodology-windows` , `persistence-methodology-linux`): a TA0003 mechanism map per OS — Registry Run keys, Startup folder, scheduled tasks, services, WMI subscriptions, accessibility/IFEO, backdoor accounts and web shells on Windows; cron, systemd

units/timers, SSH `authorized_keys`, shell-init profiles and backdoor accounts on Linux — each mechanism a single gated command plus the footprint it leaves, cross-referenced to the detection panel. Folded in by a new additive ingester (`pipeline/ingest_persistence_methodology.py`) mirroring the recon-methodology discipline (own idempotency marker, byte-preserving pass-through, corpus block segregated to the tail, re-run byte-identical). KB 2,619 → 2,621; `test_corpora` byte-identity intact. Rootkits and bootkits are named as knowledge only, never tooling.

- **Arsenal** (`backend/arsenal/tools.json`) — +4 tools, 101 → 105, in a new `persistence` category: **SharPersist** (`platform:windows` , add/remove templates, mirroring `rubeus`) plus **Sliver**, **Empire** and **weeveily** as **reference-only** — catalogued but not installed, so reconcile reports them not-present and the planner will not propose them (the same treatment as `subwiz`). Their install and C2 wiring are **deferred to build #4**. `executes_nothing` unchanged; arsenal suites green at **105 tools / 262 templates**.
- **Detection footprint** (`backend/detection/`) — 13 `FootprintSpec` s + 11 ATT&CK rows. `attck.py` gained the persistence technique rows (T1547.001, T1543.003/.002, T1053.003/.006, T1546.003/.004/.008, T1136.001, T1505.003, T1098.004) — additive, transcribed from ATT&CK v19.1 and trimmed to the Windows/Linux/network log-source subset, verified against live upstream. `catalog.py` gained a describe-side footprint per mechanism (Event IDs / Sysmon / Autoruns / auditd, loudness, and a `why_rating` that carries the honesty marker — *quiet = a defender coverage gap, not an operator advantage*). This is the **blue/describe half only**: no OPSEC/evasion notes were added (that channel is build #4), so the `assert_opsec_is_separate` and never-prescribe guards are untouched. Aliases were added for unambiguous binaries only (`schtasks` , `crontab` , `useradd` , `sharpersist` , `weeveily`); multi-purpose `reg / sc / net / systemctl` stay on the `argv` path so a `reg query` is never mislabelled persistence.

All three landed additively and verified — the full hermetic safety suite green (`test_detection` , `test_detection_safety` , `test_arsenal` , `test_arsenal_safety` , `test_corpora`), and `pipeline/detection_sources.py --verify` clean against live ATT&CK v19.1 + SigmaHQ. No image rebuild (OS built-ins and the installed `msfconsole` cover the mechanisms), and no new execution path — the gated executor already ran these commands, one human approval at a time.

AV/EDR evasion + traffic obfuscation (2026-07-27)

Build #4 closes the C2/evasion channel the persistence enrich deferred. Unlike that enrich, this one **does add execution capability** — a Sliver C2 surface, DNS-tunnel listeners, and a generate-only evasion engine — so the whole design is about making the new surfaces inherit the existing containment rather than sit beside it. It also makes a **deliberate policy change**, recorded as D16 below.

D16 — the OPSEC channel may now prescribe evasion. Until this build, `assert_opsec_is_separate` refused any sensor-blinding phrasing outright: the channel could say "rate-limit and add jitter" but not "patch AMSI". That ban is **lifted entirely and deliberately**. The reason is coherence: this build ships an engine that *emits* AMSI-patch and ETW-blind artifacts, and a tool that produces an artifact it is forbidden to describe is worse than one that describes it honestly. **Two invariants survive and are now the whole contract**: every OPSEC note must carry a substantive `still_recorded` naming what catches the technique anyway, and the blue-view footprint is always produced alongside and is never suppressible. What did **not** change: the blue-side describe-never-prescribe guard (`_evasion_prescription`) is untouched, so the defender's copy is byte-identical to before; and the four execution gates — human approval per command, scope lock, red-confirm, audit — are untouched. This is a **posture shift on a public repo** and is stated plainly rather than buried: HackPit now documents in-process sensor tradecraft, always paired with its detections.

- **Sliver C2 (backend/cockpit/sliver.py) — a split-gated surface.** Server lifecycle (start/stop/list) is **human-only** with no red-confirm, mirroring `tunnels.py` : it is operator infrastructure on the operator's own sandbox, so clicking start is the approval. Implant **generation** is a **gated command**, mirroring `session.py` : it builds an `ExecRequest` and runs the real `executor.validate_request` — approval, scope, red-confirm — before anything is produced. It only generates; there is no delivery or execution route, and live beacon catch is deferred. `<listener>` is the operator's callback address and passes through verbatim, never target-substituted. A defect found while building it: `allowlist._FRAMEWORKS` held `"sliver"` , which never matches `sliver-client` , so an implant build would have passed the danger gate with **no red-confirm at all**; the real binary names were added.
- **DNS-tunnel obfuscation (backend/cockpit/obfuscation.py) — dnscat2 / iodine listeners**, entirely human-only. The far-side client half is returned as a **string for the operator to carry across by hand**; the module has no delivery primitive and must never gain one. The tunnel's pre-shared key is masked **at source** — `client_command` is built from a masked copy, so the real key never occupies a field that crosses the HTTP boundary. That mattered: the obvious fix (`model_dump(exclude={"secret"})`) still leaks, because the key is embedded verbatim inside the one-liner.
- **Bespoke evasion engine (backend/evasion/) — generate-only, with forced honesty.** A top-level package (peer to `detection/` and `state/` , per the cockpit/arsenal decoupling rule). It runs donut/ScareCrow as argv-only `docker exec` into a container resolved from the execution mode — never a request field — and **never runs or deploys what it builds**: the produced artifact is never `argv[0]` , and no deploy primitive exists in the package. **The honest half is computed before anything is built**: if the engine cannot resolve the blue-view footprint for a technique it refuses outright rather than emit a footprint-less artifact. Both the footprint and the `still_recorded` note are required fields on the result model, so no route can return one without the other, and the two PowerShell stubs carry the same honesty header *inside the artifact* so it survives being copied out of HackPit.
- **Detection (backend/detection/) — 7 new footprint specs and 20 new OPSEC notes.** TA0011 gained `c2_dns_tunnel` , `c2_malleable_profile` , `c2_jitter_beacon` and `c2_domain_fronting` (describe-only; domain fronting is documented as largely dead); TA0005/TA0112 gained `evasion_packed_loader` , `evasion_amsi_patch` and `evasion_etw_blind` , which are what the engine maps onto. The 13 persistence specs from build #3 were backfilled with per-mechanism OPSEC notes — 7 new notes alongside the new specs plus those 13 is where the 20 comes from. **SPECS 46 → 53, OPSEC notes 10 → 30** (both re-counted against the live catalog; an earlier draft of this bullet said "4 OPSEC notes", which contradicted its own 10 → 30). Two accuracy corrections were made under review: T1029 had been filed under TA0011 when *Scheduled Transfer* is TA0010 Exfiltration (it was dropped rather than retagged, since an Exfiltration label on a beaconing footprint reads as wrong), and the telemetry strings on five new rows had been written rather than transcribed — replaced with real ATT&CK v19.1 values, `--verify` clean against live upstream. No new ATT&CK rows were needed for the evasion specs: under v19 naming T1562.001 and T1562.006 (Indicator Blocking) *both* revoke into **T1685**, which was already present. An earlier version of this claim paired T1562.006 with T1690 and mapped the ETW footprint onto it — T1690 is the successor of T1562.003 (*Impair Command History Logging*) and covers shell history, not sensors. The Tasks 8–13 review caught it; the spec now cites T1685. `--verify` did **not** catch it and could not have: it checks that a cited id's name, tactic and log sources match upstream, not that the technique is the right one for the activity.
- **Arsenal — 105 → 110 tools**, new `c2` and `evasion` categories; Sliver moved out of `persistence` into `c2` . `<tunnel-zone>` and `<tunnel-net>` are new placeholders rather than reuses of `<domain>` , because the composer's target substitution rewrites any dotted token — spelling them `<domain>` would have pointed an operator's tunnel at the system under test. A latent data bug was found and fixed here: `placeholders`

declared "`<payload>`" **twice**, so every `json.load` silently discarded the first description; a duplicate-key check now guards it.

- **Image (`docker/Dockerfile.sandbox`) — Sliver, dnscat2, donut and ScareCrow installed; iodine was already present.** Three of the plan's assumptions did not survive contact with the base image and were caught by the layer's own smoke tests: `dnscat2` and `ScareCrow` are **not packaged** in kali-rolling, and there is **no Go toolchain** in the image, so release binaries and a from-source `dnscat2` build were the only options. Two further faults surfaced only because the smoke test *invokes* each tool rather than running `command -v`: `dnscat2`'s server dies on `require 'ecdsa'` (four gems missing), and `donut-shellcode` is a C extension with no console script, so `python -m donut` cannot work and a real CLI had to be written.

Containment, restated. No orchestrator, agent or loop module can reach any of the three surfaces — asserted by whole-tree source scans, not by convention. Nothing here is autonomous, no gate was weakened, and the only capability the agent gained is none: every new action is human-initiated and audited.

Gate-integrity audit (2026-07-27)

After build #4 shipped, the whole project's **guards** were audited with one question: *would this actually fire?* Not "is a guard present" — the build had already shown that a guard can be present, look correct, pass its tests, and never trigger. **24 guards were probed by constructing a deliberate violation for each. Seven silently failed to fire.** Full detail in `docs/GATE-AUDIT-FINDINGS.md`; the three criticals were:

1. **Eight catalogued tools passed the danger gate with no red-confirm** — demonstrated end to end through `executor.validate_request`. `weeveily` (which both generates a webshell *and* drives it), `dnscat2`'s client and server, `iodine / iodined`, `commix`, `SSTImap`, `SharPersist` and `Invoke-Obfuscation` all returned zero reasons, as did the argument-driven `sqlmap --os-shell`, `commix --os-shell` and `netexec -x`. **FIXED and pushed (4492fec)**.
2. **The WinRM transport classifies only the first token of what is a whole PowerShell script.** `executor.py` joins `command+args` and `run_ps()` executes the lot, so `;` and `|` are live separators: `Write-Host go ; Invoke-Mimikatz` is silent, while the same cmdlet as `argv[0]` fires. **FIXED — build #5 (see "Gate integrity — build #5" below)**. This was the most consequential finding of the whole review cycle, because it executes on a real domain-joined host under real credentials.
3. **The `:kali` human-only lock never opens 39 of 69 backend modules** — its scan globs only `backend/*.py` and `cockpit/*.py`. A planted `from cockpit.kali import run_kali` in `adgraph/orchestrator.py` passes. The same narrow glob was copied into the tunnels, repeater, terminal and WinRM locks, and the `cockpit→arsenal` lock covers 5 of 22 modules. **FIXED — build #5 (see "Gate integrity — build #5" below)**.

What Critical 1's fix actually changed, because the shape matters more than the list. The `.exe / .py / .ps1` normalisation that the AD sets already did was **collapsed into one shared normaliser used by every set** — that asymmetry was the root cause, not a symptom, and it was why `powershell.exe`, `nc.exe` and even `sliver-client.exe` produced no reason at all on the transport where `.exe` is the ordinary spelling. Two new sets were added rather than padding `_FRAMEWORKS`, so the reason the operator reads is true: `_TUNNEL_TOOLS` (a tunnel is a C2 path and an exfil path, not a payload generator) and `_RCE_TOOLS`. `_FRAMEWORKS` ' `ligolo` entry was corrected to `ligolo-proxy` — the binary the repo actually runs — an entry that could never match, added *after* the `sliver / sliver-client` bug was fixed, and worse than no entry because it reads as coverage.

The finding underneath all three. `test_arsenal_safety` claimed to verify that a catalogued dangerous invocation demands the red-confirm — while testing `python3`, **a command that is not in the catalog at all**. That single choice is why eight tools shipped ungated. It now pins all **176** catalogued invocation names (every name,

alias and template `argv[0]`) into exactly one bucket, so an unclassified tool fails the suite; five planted regressions were run against it and all five fail correctly. Every critical in this audit, and two of the five findings in the Tasks 8–13 review, reduce to the same root cause: **a test that could not fail, or that tested a value the real system never produces**. That is what build #5's third task turns into a written repo convention.

Gate integrity — build #5 (2026-07-27)

The audit's two remaining criticals, and the testing pattern that hid them. Plan: `docs/superpowers/plans/2026-07-27-build5-gate-integrity.md`.

Critical 2 — the WinRM danger gate read the first token of a whole PowerShell script. `cockpit/executor.py` joins `command` + `args` into one string and the Windows transport hands that string to PowerShell, where `;`, `|` and newlines are live statement separators. The heuristic classified `basename(argv[0])` and scanned only the `args` for markers. So the gate was defeated by moving a cmdlet one token to the right, on the one path that executes against a real domain-joined host under real credentials:

Script	Before
<code>Write-Host go ; Invoke-Mimikatz -DumpCreds</code>	no reason, allowed
<code>Get-DomainUser \ Set-DomainUserPassword</code>	no reason, allowed
<code>IEX (New-Object Net.WebClient).DownloadString(...)</code>	no reason, allowed
<code>Write-Host go + newline + Invoke-Mimikatz</code>	no reason, allowed
<code>powershell -enc <base64></code>	fired by accident — <code>-enc</code> splits into <code>-e</code> , <code>-n</code> , <code>-c</code>

All five became tests first, verbatim, and were confirmed failing before anything was written.

The fix shape, and why it is not a parser. Four options were weighed; the choice was to **scan the whole joined string for every marker, wherever it appears**, and explicitly *not* attempt statement splitting. Correct PowerShell tokenising has to handle quoting, `$ ( )` subexpressions, the `&` call operator, backticks and line continuations — and any parser written here becomes a new bypass surface, which is precisely the bug being fixed: **the classifier's model of the input was narrower than the input**. Whole-string scanning has no "position" to exploit, and that is the property the first-token design lacked. The cost asymmetry settles it: this gate raises a **red-confirm, not a block**, so a false positive costs one click while a false negative runs Mimikatz on a domain controller.

Two things a name-based scan cannot see were closed alongside it. **Download cradles** — `IEX (New-Object Net.WebClient).DownloadString(...)` never spells the payload's name anywhere, so the *cradle* is the marker. **Encoded commands** — `-enc` defeats every text scan by construction, so the payload is decoded and re-scanned, matching every PowerShell prefix of the flag rather than only the spellings a human types; a blob that will not decode is itself reported rather than passed.

The root-cause lock. The bug was never only "the heuristic is too narrow". The string the gate classified and the string the transport executed were **built in two places**, and scanning a different string than the one that runs would have reproduced the bug somewhere new. Both now derive from `executor.join_ps_command()`, and a test asserts on the source — transitively, with a negative control — that they still do. The AD orchestrator's advisory pre-check was moved onto the same union, so the panel can no longer promise a confirm the gate would not require, or stay quiet where it would.

The false-positive cost, measured rather than promised (plan Task 1 Step 6). Run across every AD/Windows template in `tools.json` and every command in the AD graph's edge definitions: **78 of 145 invocations in the full population (53%), and 27 of 43 in the WinRM-reachable subset (62%), now demand a confirm.** The second number is the one that matters — the catalog's Linux impacket invocations go through the docker path, which is untouched. Of those 27, **none is a false positive**: every one is a genuine destructive abuse (Set-DomainUserPassword, Add-DomainObjectAcl, Add-DomainGroupMember, Invoke-Mimikatz, SharPersist, Invoke-Obfuscation, Rubeus ptt, Set-DomainObjectOwner). Read-only enumeration stays silent — Rubeus kerberoast and asreproast, `Get-DomainUser -SPN`, `Find-InterestingDomainAcl`, `Find-LocalAdminAccess`, winPEAS, `Get-ADServiceAccount`, `Get-DomainObject` — and that half is pinned as a control, because a banner on ordinary enumeration is how an operator learns to click through the one that matters. The measurement also found **two real false negatives** the sets had missed: `Invoke-SQLOSCmd` (OS command execution on a SQL host) and `Enter-PSSession / New-PSSession` (a further hop out of an existing session). Both were catalogued; both were silent.

Because whole-string scanning over-flags by design, **every reason names the marker it matched and the offset it matched at** — `powershell script: 'invoke-mimikatz' at offset 14 - dumps/ replicates domain credentials`. A reason the operator can evaluate is a gate; a banner they cannot is decoration, and trading a silent bypass for a decorative one is not a fix.

Critical 3 — eleven source-scan locks, nine of them narrower than their own docstrings. The `:kali` human-only lock is the entire reason the one deliberately-unbounded arbitrary-shell surface is considered safe. Its docstring said *scan the whole (non-venv) source tree*; it globbed `backend/*.py` plus `backend/cockpit/*.py`, which is **30 of 69 modules**, so all of `adgraph/`, `arsenal/`, `codescan/`, `detection/`, `evasion/`, `exploits/` and `state/` were never opened and a planted `from cockpit.kali import run_kali` in `adgraph/orchestrator.py` — a literal orchestrator module — shipped green. Three distinct defects, now fixed structurally in one shared `backend/test_support/scans.py` rather than by asking each suite to remember:

1. **Narrow file selection** — an `rglob` over the whole backend. Coverage per lock goes 30 to 67 modules; there is no per-suite glob left to get wrong.
2. **Baseline allow-lists** — `{"router.py"}` matched against `f.name` exempted `adgraph/router.py`, `detection/router.py`, `arsenal/router.py` and every other `router.py` in the tree by accident. Allow-lists are now keyed on repo-relative POSIX paths only.
3. **Four literal substrings** — an AST pass alongside them catches what the audit planted and the old predicate missed: aliased imports, imports opened inside a function body, `import_module("cockpit." + "kali")`, `getattr(m, "run_" + "kali")`, and f-strings with constant parts.

Prose is stripped so a module that *documents* a rule does not violate it — this is not hypothetical, documenting the Critical 2 fix tripped the WinRM lock — but ordinary string literals are deliberately **kept**, because a scanner that blanked every string would go blind to `import_module("cockpit.kali")`, which is the indirection the lock exists to catch.

Two more locks were closed with it: the **cockpit→arsenal** lock checked a hardcoded list of five filenames while printing "the cockpit package has zero references to the arsenal" (the package has 22 modules, 17 of which post-date the invariant), and the **evasion agent-path** scan still filtered by *filename*, so `from evasion import engine` in `cockpit/executor.py`, `chat.py` or `adgraph/techniques.py` was never looked at. A module named something nobody predicted is exactly what a name filter cannot cover.

Widening the scans produced two false positives, and both were fixed in the predicate. `cockpit/reconcile.py` "referenced the arsenal" — as a *parameter name*; it takes the loaded catalog as an opaque injected object precisely so it has no import-time dependency. `detection/{catalog,resolver,router}.py` and `report.py` "reached the evasion engine" — they are the *anti-evasion* guard, the code that refuses prescriptive evasion copy. Both locks

now assert the claim the invariant actually makes (no import, in any form the AST can see). **Narrowing the file set to silence a false positive is how these guards got broken in the first place**, so that is written into the convention as a rule.

Task 3 — the convention, which is the part that outlasts both fixes. Every critical in this audit, and two of build #4's five review findings, reduce to one root cause: a test that could not fail, or that tested a value the real system never produces. `test_arsenal_safety` claimed to verify that a catalogued dangerous invocation demands the red-confirm while testing `python3`, a command absent from the catalog it was guarding — one choice that hid eight tools. `backend/AGENTS.md` now carries the rule, with the five failures that produced it in a table: **a safety test must iterate real data, assert on what it actually checked, and prove it can fail**. Concretely — draw inputs from the real source of truth so a tool or module added tomorrow is covered by nobody remembering; assert on the *filtered* count, never files opened; and carry a planted-violation positive control in the same test. `test_scans.py` runs **first** in the suite for that reason: ten locks rest on the shared scanner, and it must demonstrate it catches a planted violation before any of them means anything.

Sweeping for the *pattern* rather than the files the plan named found three more, none of which was in the task list. The **live-session stdin lock** — the load-bearing one, since anything able to type into an already-approved running session is executing un-gated commands — said "the whole source tree" and globbed three directories. `test_sliver_safety` and `test_obfuscation_safety`, the two suites the audit called correct, incremented their scanned counter *above* the `test_skip`, so their `>= 40` controls counted files opened rather than files judged. And re-running the audit's own probe list surfaced a real heuristic gap: `sekurlsa::pth ... /run:cmd.exe as argv[0]` fired only because `cmd.exe` happens to sit in its args (change it to `/run:powershell.exe` and it went silent), while `kerberos::list /export` — which writes every ticket in the session to disk — matched nothing at all.

A follow-on fix, after the plan's three tasks: the tunnel listener start is now gated (I2). It was outside build #5's plan and was recorded here as open, then closed immediately after. `cockpit/tunnels.start_tunnel` reached `subprocess.Popen` directly with no `validate_request` call, and `POST /cockpit/tunnels` carried no `approved / dangerous_ack` field — so a pivot listener, which is a C2 path and an exfil path in one, was raised on a plain POST with human-only as its *sole* bound, strictly less than every other execution surface gets. The docstring's claim that "the rewritten pivoted command still runs through the gated executor" was true and beside the point: the listener start *is* the tunnel primitive. Fixed test-first (the unapproved-start test was confirmed failing on the old code): `start_tunnel` now runs the real `executor.validate_request` before anything spawns, `TunnelStartRequest` carries `approved + dangerous_ack` defaulting **false** (an omitting client is refused, never silently granted), and the route maps a safety refusal to 403 naming the gate while an availability problem stays 409. The danger verdict comes from the server binary — `chisel` and `ligolo-proxy` are both in `_TUNNEL_TOOLS`, pinned by a positive control so the leg cannot pass vacuously — and a start now requires an engagement, because without one it would resolve to lab mode and fire the lab's isolation gate about a container the listener does not run in. The gated argv and the spawned argv derive from one `server_argv_for`, the same drift lock Critical 2 uses. This is deliberately **not** the D17 shape: for Sliver, C2 lifecycle is human-only while artifact generation is gated, because a listener nothing has connected to is inert; a pivot listener's whole purpose is to become a route into an unseen network, so it is gated like an execution.

What still remains open. `iter_run`'s prevalidated path re-checks approval but not danger (nothing is exposed today, because the single `prevalidated=True` caller validates first, but the asymmetry is the kind of thing a future caller trips over), and there is still no route-level authorisation on any of the routes, which remains the known localhost-only posture rather than a new finding.

Live fire, image pinning, and the prevalidated danger re-check — build #7 (2026-07-27)

Three items from Part II's deferred list, none of which adds offensive capability. One makes the image byte-reproducible, one closes a latent asymmetry in a belt-and-suspenders gate, and the largest — live-fire verification — went looking for evidence that the C2 and tunnel surfaces work and came back with three defects and a lot of confirmed containment.

Live fire: what ran, and what it found

`docker/proof/live_fire_proof.sh` builds a lab that is ours end to end — four containers on four `internal: true` networks (`docker/proof/live-fire-lab.yml`), no third-party target, no public callback — and drives the surfaces through their shipped entry points via `docker/proof/live_fire_driver.py` . Every gated step **proves the refusal first**: the same request with the acknowledgement missing must be refused with nothing spawned, because a proof that only shows the happy path cannot tell a working gate from an absent one. Results are reported as PASS / FAIL / **NOTRUN**, and a step that could not run is never folded into a pass. The run stands at **26 passed, 3 failed, 5 not-run**.

The gates and the containment held, live and without exception:

- The **I2 pivot gate fires on all three limbs** against a real `ligolo-proxy` start — no engagement refuses at `gate=engagement` , unapproved at `gate=approval` , un-red-confirmed at `gate=danger` — and the approved start is accepted. This is the first time that gate has been exercised outside a stub.
- The **Sliver implant gates fire** — `gate=approval` unapproved, `gate=danger` without the red-confirm.
- The **pivot subnet enters scope only through the explicit amendment**, verified against a before/after snapshot of the engagement record, with the control taken first: the deep target is confirmed *unreachable* from the operator box before any tunnel exists, so "we reached it through the tunnel" would have meant something.
- **Containment holds live, measured from inside the contained host**. The implant host cannot reach the internet (`1.1.1.1`) or the host (`host.docker.internal`), and *can* reach the C2 — so the beacon has nowhere else it could go, which is what makes demonstrating a callback safe in the first place.
- The **DNS tunnel key never enters the pasteable one-liner**.
- A real Sliver server runs. `sliver-server` daemon starts through `cockpit/sliver.py` , unpacks its embedded assets and holds `:31337` .

Three things **failed**, and they are defects in the surfaces rather than in the gates. All three are the same shape: a claim that unit tests asserted structurally and nothing had ever executed.

1. The **Sliver implant build cannot succeed as shipped**. `_implant_argv()` builds `sliver-client generate --os ... --mtls <listener>` , and `sliver-client` has no `generate` subcommand — it has `import` , `version` , `completion` , `help` . In Sliver 1.5 `generate` is an *interactive console* command. The existing tests assert the `argv`'s **shape** (that `<listener>` is emitted verbatim, that the target is never emitted) and never that Sliver accepts it, so a build that could not run looked correct for as long as nobody ran one. The console route was verified working during this build with the identical flags over stdin — `echo "generate --os linux --arch amd64 --format exe --mtls H:P --save ..." | sliver-client` produced a 14 MB implant in 27 s — so the fix is known and small. **It was deliberately not applied here**: making a non-functional offensive path functional is a capability change, and this build adds none. It is the operator's call.
2. `tunnels.start_tunnel` reports `status="listening"` for a process that has already exited. `ligolo-proxy` is an interactive console binary; spawned with no usable stdin it binds the port, prints `Listening on`

0.0.0.0:11601 , reads EOF and exits 0. The status is assigned at `Popen` time and never observed, so the panel shows a live pivot over a dead process. Confirmed by `proc.poll()` returning 0 four seconds after the model claimed it was up, with nothing bound in the container.

3. `obfuscation.start_listener` has the same defect for `dnscat2-server` , for the same reason.

Because the implant never built and the listeners never survived, the three end-to-end claims those steps feed are recorded as **not-run, not as passed**: no beacon callback was caught, no proxychained traffic reached the deep subnet, and no dnscat2 session was established. The harness now asserts process liveness directly, so these three become regressions the moment they are fixed.

Two further items are **not-run because they need operator infrastructure**, and no amount of local wiring substitutes for either:

- **A delegated DNS zone.** Only the direct-to-server dnscat2 path can be exercised locally. A genuine delegated tunnel needs a domain the operator controls with an NS record pointed at the listener. "Demonstrated in lab" and "needs a real delegated zone" are different claims and are kept apart.
- **iodine**, which needs a TUN device at *both* ends plus that same zone; the engage sandbox has `/dev/net/tun` , the lab client container does not.

A side finding, recorded rather than acted on: `obfuscation.start_listener` justifies being human-only-and-ungated by citing "the identical reasoning the pivot-listener lifecycle uses" — and the pivot listener was tightened in build #5 (I2) to require engagement plus approval plus the red-confirm. The cited precedent no longer holds. Whether the DNS listener and the Sliver server start should follow it is a design decision, not a defect, and it is left open.

One environmental fact surfaced immediately and is worth writing down: the running stack was still on `hackpit/kali-sandbox:m1` built *before* build #4, so none of the C2 tooling existed in the sandbox the gated code execs into. The compose file rebuilds that tag from the current Dockerfile; nobody had re-run it with `--build` . The stack was realigned and both network proofs re-verified afterwards — `isolation_proof.sh` 4/4, `engage_open_proof.sh` 4/4.

The image is now byte-reproducible — and its smoke test had never passed

`docker/pin-build4-versions.sh` resolved the real upstream values and rewrote the layer: dnscat2 to commit `42f8d783...` (it has no release tags, so a SHA is the only thing there is to pin to), the four Ruby gems to `trollop:2.9.10 salsa20:0.1.3 sha3:2.2.4 ecdsa:1.2.0` , and `donut-shellcode==1.1` . Two bugs in that script had to be fixed before it could run: it required `python3` on `PATH` (only the backend venv exists here) and its verification step used `grep ... && { exit 1; }` , which under `set -e` exits 1 on the *success* path. `test_evasion.py` now asserts the stronger claim — all three installs pinned, three `ARG` version pins, the SHA a real 40-char object id, no floating form left behind — instead of tolerating a documented gap.

The fresh build then failed, and not because of the pins. The consolidated smoke test asserted `sliver-server version 2>&1 | grep -qi sliver` , and `sliver-server version` prints exactly `v1.5.42 - <sha>` — no substring "sliver". **That assertion could never pass, so this layer had never built since the check was added**; the image in the daemon predates it. The static guard that watches this block (`test_evasion.py`) reads the *shape* of these lines — `|| true , | head` — and by construction cannot know whether a grep pattern matches anything, so only a real build could surface it. Both sliver checks now match `v${SLIVER_VERSION}` , which is strictly stronger: it fails when the binary is missing, when it dies, and when it is not the version pinned above. `hackpit-kali:build7` builds clean, and all three pinned tools resolve and smoke-test by their real invoked names, with the shipped versions confirmed inside the image against the pins.

The prevalidated path now re-checks danger

`iter_run(prevalidated=True)` skips `validate_request` because the router already ran it to decide the HTTP status. It has re-checked **approval** for engagement and windows since build #4 — never-auto-run being the sole floor on a real target — but had no equivalent **danger** re-check. Nothing was exposed: the one `prevalidated=True` caller validates first. But "no caller does this today" is a coincidence, not an invariant, and the asymmetry was a trap set for the second caller.

Test-first per `backend/AGENTS.md`: the failing test showed a dangerous, un-acked, prevalidated request reaching `start` and spawning. The re-check mirrors the approval one and uses the **same per-mode classifier** `validate_request` uses — `windows_danger_reasons` on the joined script for windows, `dangerous_command_heuristic` on the argv for lab and engagement. The windows probe is deliberately a command only the script classifier catches (`Write-Host go ; Invoke-Mimikatz`, the Critical-2 shape), so a re-check wired to the generic heuristic fails the suite rather than passing quietly.

The refactor that made this work is worth recording because the project's own guard caught it. Routing all three validators through one shared `danger_reasons_for_mode()` dispatcher looked strictly better — divergence becomes structurally impossible — but `test_winrm_safety` asserts over the **static call graph** that `_validate_lab` can never reach `join_ps_command`, and that assertion is the positive control proving the reachability check can fail at all. The dispatcher reaches the PowerShell derivation on its windows branch, so wiring the docker-path validators through it emptied the control while every test still passed. The validators call their classifier directly and deliberately; only `iter_run`, which knows a `mode` string rather than a branch, uses the dispatcher. The classifier each side calls is identical either way — only the call graph differs, and the call graph is what the control reads.

One existing test needed updating rather than the guard being loosened: `test_winrm_safety`'s host-lock case drove `Invoke-Command` prevalidated with no ack. `Invoke-Command` is flagged, so the router would have demanded the red-confirm before ever setting `prevalidated=True`; the test had been relying on the gap it was sitting next to. It now carries the ack, which is what the real path sends.

The three live-fire defects, fixed — and a stale precedent decided (2026-07-28)

Build #7 found three surfaces that did not work as shipped and deliberately left them alone, on the grounds that making a non-functional offensive path functional is a capability change and that call belongs to the operator. The call was made: fix all three, and decide the design question the run had left open. This section is what that took, including two further defects the fixing surfaced — both of the same family, and both found the same way.

The Sliver implant build: a subcommand that does not exist

`_implant_argv()` built `sliver-client generate ...`, and neither Sliver binary has a `generate` subcommand — `sliver-client --help` offers `completion`, `help`, `import`, `version`, and nothing else. In Sliver 1.5 `generate` belongs to the interactive **console**, which `sliver-server` hosts when run with no subcommand. A build now runs `docker exec -i <container> sliver-server` — still an argv list — and writes one `generate ...` line to its stdin, followed by `exit`. Verified against the pinned 1.5.42 image: 16 s, a 14 MB implant on disk.

Two consequences follow, and both are handled rather than hoped about.

The console line is **text** reaching Sliver's own flag parser, so it is a parser this code did not previously feed. Everything interpolated into it is bounded before it arrives: `os / arch / format / transport` already came from fixed sets, `listener` is now pinned by a pattern that admits a host or IP with an optional port (or an angle-bracket placeholder) and rejects whitespace and a leading `-`, and the server-chosen artifact path is checked

whitespace-free rather than trusted — it is derived from a caller-supplied engagement id, and a space in it would split into extra console tokens. Nothing here reaches a shell: the argv is still a list, so `/bin/sh` is not in the picture at all.

And the console exits 0 whether or not the build inside it worked, so an exit code cannot decide success. `generated` versus `failed` is now decided by reading the artifact back out of the container with `stat`. That is a stricter test than the one it replaces, in both directions: a console that exits 0 having written nothing is `failed`, and a console that complains while producing a real artifact is `generated`. The gate surface moved with the binary — `_gate_request` classifies `argv[0]` itself rather than a constant, so the audit row names the process that actually runs.

Both listener lifecycles: a status that was assigned, never observed

`ligolo-proxy` and `dnscat2-server` are interactive consoles. `docker exec without -i` hands a container process a closed stdin no matter what the host process got, so both printed their banner, read EOF and exited — while `start_tunnel` and `start_listener` returned `status="listening"`, assigned at `Popen` time and never looked at again. The operator was then holding an agent one-liner for a port with nothing behind it.

Both halves are fixed in one place, `backend/cockpit/lifecycle.py`, now shared by all three lifecycle surfaces:

- **A console binary gets a stdin that does not end.** `-i` is per-binary and off by default — `ligolo-proxy` and `dnscat2-server` need it; `chisel server`, `iodined -f` and `sliver-server daemon` are real daemons and are spawned without one, verified rather than assumed.
- **The write end is a raw file descriptor, deliberately.** Passing an int fd rather than `subprocess.PIPE` leaves `proc.stdin` as `None`, so there is no writer object for any caller to reach for. Holding a listener's stdin open is a liveness requirement; being able to *type into* a live C2 console is a capability, and this grants the first without the second. A whole-tree source scan with a planted-violation control keeps the handle inside the four modules that own a listener.
- **The status is observed.** After a settle window: a process that exited is a **refusal** carrying the output it died on; a live process whose port `ss` confirms bound inside the container is `listening`; a live process whose bind could not be confirmed is `starting`. The probe is tri-state on purpose — "could not probe" is not "not bound", and neither is ever rendered as a confirmed listener. Every model now carries a `liveness` field holding the evidence, not just the verdict.
- **Output is drained** into a bounded ring buffer, so a chatty server cannot fill the 64 KB pipe buffer and wedge mid-engagement looking exactly like a hung tunnel.

The stale precedent, decided: both listeners are now gated

`obfuscation.start_listener` justified being ungated by citing "the identical reasoning the pivot-listener lifecycle uses" — and build #5's `I2` finding had already tightened that pivot listener to require engagement plus approval plus the red-confirm. The citation was stale. The larger problem was that `I2`'s actual argument applies to a DNS tunnel with full force: *a listener whose purpose is to become a route into a network the scope gate has never seen is an execution, not a lifecycle no-op*. A DNS-tunnel listener is the server end of a covert exfil channel — carrying arbitrary traffic out of a network through that network's own resolvers is the thing egress controls exist to stop. "No target" was true and beside the point, exactly as it was for the pivot listener.

The DNS listener is now gated the same way: engagement, approval, red-confirm, through the real `executor.validate_request`. `dnscat2-server` and `iodined` were already in the covert-tunnel danger set, so the heuristic fires on the binary and the red-confirm is always required.

The Sliver server start is gated too, on a weaker argument that is worth stating honestly. A Sliver daemon is genuinely less exposed than either listener: starting it opens no channel toward anything under test, it binds a multiplayer port on the operator's own sandbox for the operator's own console, and the steps that *do* reach a target — creating a listener, generating an implant, delivering it — are separate and separately gated. That argument is real, and it is not quite sufficient, because Sliver's server config can **persist listener jobs** that come back up with the daemon. "Starting it is inert" is therefore a property of a config file, not a property of this code, and a gate should not rest on something the module cannot see.

Human-only remains the load-bearing property for all three surfaces and is untouched: the whole-tree source scans still prove no orchestrator, agent or loop path reaches any of them. The gates are belt to those suspenders. What still differs between the two halves is the **target**: a C2 server and a DNS listener have none, so neither carries a target-lock, while implant generation is scope-checked as before — and `SliverServerRequest` is asserted to have no `target` field, because one would mean a scope gate firing on the operator's own box.

Two further defects, found the same way

Re-running the live fire against the fixed surfaces surfaced two more, both invisible to every unit test and both the same shape as the originals: a state reported rather than observed.

A stop did not stop. Killing a `docker exec` client does not kill the process it started inside the container. After stops that reported `status="down"`, `ss` inside the sandbox still showed `iodined` holding UDP/53 and two `chisel server` processes listening — and the next start failed with `EADDRINUSE`, a refusal about a listener the operator had been told was already stopped. A stop now closes stdin (which is enough for a console binary), kills the client, and then reaps the server inside the container by **its own argv**: `pkill -f` against the full command line, which carries the listener's port, so it stops that listener rather than every process of the same kind. A blanket `pkill -f chisel` would have stopped somebody else's tunnel.

proxychains pointed at Tor. Debian ships `/etc/proxychains4.conf` configured for `socks4 127.0.0.1 9050`, while HackPit's chisel plan brings up a reverse SOCKS5 at `127.0.0.1:1080` and `wrap_command` rewrites a command to `proxychains -q ...` on the strength of it. With the stock config the rewritten command the operator approves would have quietly tried Tor and timed out. The Dockerfile now replaces that line rather than appending to it — `proxychains` reads the last `[ProxyList]` section, and leaving 9050 in place would make the chain depend on ordering. This is exactly the class of gap a proof harness exists to find: every unit test passed, and the shipped path could not work.

Both fixes verified against a rebuilt image, not just against a diff

Two of the five fixes carried a claim that had not itself been executed, which is the exact failure mode this build exists to eliminate, so both were closed properly rather than left as edits.

The `proxychains` line had never been built. The fix was written into `Dockerfile.sandbox` while the running image predated it, and the harness applied the same one-line edit at runtime — correct for making the proof valid, but it left the Dockerfile layer itself unproven. `hackpit-kali:build7b` now builds from the current Dockerfile with exit 0, and the config inside the image reads `socks5 127.0.0.1 1080` with **zero** `socks4` lines remaining. The stack was rebuilt and recreated from that Dockerfile, and the live fire re-run against it takes the *"already points at the chisel SOCKS5 — baked into the image"* branch rather than the runtime-patch branch, with the proxychained request still routing. The runtime patch stays in the harness on purpose, so the proof remains valid on an older image, and it says which branch it took.

The stop-reap was unit-tested but never demonstrated against the failure that motivated it. Run live: a `chisel server` started through the lifecycle and confirmed bound; killing only the `docker exec` client left the process running and the port still bound — the defect reproducing exactly as described; the reap then cleared both. And it is targeted, which is the part that matters: reaping one tunnel left a second `chisel server` on a different port still bound, where a blanket `kill -f chisel` would have stopped it too.

That second run also surfaced a sixth defect of the same family, now fixed. On a freshly recreated container `ligolo-proxy` took longer than the settle window to bind, so it was honestly reported `starting` — and `route_for` correctly refuses to route through an unconfirmed bind. But the status was **observed once and never re-observed upward**: the list functions only ever demoted to `down`, so a listener that bound a second after its window closed stayed `starting` for its whole life, and therefore permanently unroutable. The refresh now moves in both directions and only on evidence — a dead process becomes `down`, a `starting` listener whose port has since appeared becomes `listening`, and a probe that says "not bound" or "could not run" promotes nothing. Only `starting` listeners are re-probed, so a confirmed bind does not cost a `docker exec` on every poll of the panel. The negative controls are in the same test, because a refresh that promoted on an unknown would be laundering exactly the uncertainty the settle-window observation exists to preserve.

What the live fire now shows

66 passed, 0 failed, 3 not-run (after the not-run-closing pass below). The defects are fixed and their assertions pass; the gates and the containment continue to hold, and iodine now demonstrates the DNS-tunnel class end to end.

- **A real implant builds** through the gated path — `status='generated'` with a 14 MB artifact read back out of the container by size — and **the beacon calls back**: the implant, run on a lab host whose only network is internal, registers a session with the operator's server. Containment is measured from inside that host: no internet, no Docker host, C2 only.
- **A proxychained request routes through the pivot** to a host on a subnet the operator box has no route to, returning the deep target's known body — and the argv that ran is the one `wrap_command` produced, so what was demonstrated is what a human would have approved. The same tunnel cannot reach `1.1.1.1`.
- **All three lifecycle gates refuse on all three limbs** — no engagement, unapproved, un-red-confirmed — with nothing spawned, for the Sliver server, both pivot kinds and the DNS listener. Each surface's status is separately asserted to have been *observed*, against the process handle and against the `liveness` evidence.
- **Every listener stays up** and its port stays bound for the whole phase.

Four things are **not-run, and they are not passes**:

- **ligolo's interface routing**. ligolo routes at the interface, and creating that route means typing `session` then `start` into the proxy's console. HackPit holds that console's stdin open and deliberately never types into it — an unapproved command on a live C2 console is precisely what the gates exist to prevent. Completing a ligolo session stays the operator's own step. The routed-traffic claim is carried by chisel, whose server needs no console at all, and that is why both kinds are exercised.
- **The dnscat2 client handshake** — now diagnosed decisively (see the next subsection): `tcpdump` shows the `dnscat2-server` process itself answering `NXDOMAIN` to correctly-formatted queries, reproduced with both ends in one container over loopback, which isolates it to the pinned dnscat2 build rather than to HackPit or the lab. The listener HackPit starts is gated, bound and stable throughout.
- **A delegated DNS zone**. Unchanged and unfakeable: a genuine delegated tunnel needs a domain the operator controls with an NS record pointed at the listener. "Demonstrated in lab" and "needs a real delegated zone" stay different claims.

- **A delegated DNS zone** — the public-hierarchy hop in front of either tunnel. iodine's IP-over-DNS channel itself is now **demonstrated** (below), carrying traffic and confirmed DNS-encapsulated on the wire; what stays not-run is only the delegation in front of it, which needs a domain and a public listener the operator does not have.

One harness lesson is worth recording, because the harness was briefly lying about the product. A console listener's lifetime is tied to the process holding its stdin — in the backend that is the long-lived server process, which is correct. But each driver invocation was its own short-lived Python process, so a listener came up, was confirmed bound by `ss`, and died the moment the driver exited, after which the shell's follow-up check reported "did not reach LISTEN" for something that had demonstrably been listening seconds earlier. Client connect-backs now happen inside the same process that holds the listener, so every assertion is measured while the thing it is about is alive. Three further false findings had the same character and are documented in place: a `pskill` sweep that killed its own wrapper shell on the first pattern and silently abandoned the rest, a `sed` whose absolute path Git Bash rewrote into a Windows path so the config edit never applied, and a `dnscat2` client invocation refused for passing both a `--dns` argument and a bare zone. A proof that reports its own bugs as product defects is worse than no proof.

Closing the not-runs: iodine demonstrated, dnscat2 diagnosed (2026-07-28)

Four things were left not-run after the surfaces were fixed. Two are now resolved — one by demonstration, one by a decisive diagnosis — and the remaining two are stated for what they actually are rather than carried as vague gaps.

iodine's IP-over-DNS tunnel is now demonstrated carrying traffic, through the gated surface. A dedicated lab fixture (`iodine-client` in `live-fire-lab.yml`) carries the one thing iodine genuinely needs and the `dnscat2` host must not have — a `/dev/net/tun` device, `NET_ADMIN`, and root, because iodine brings up a whole tunnel *interface* and a `cap_add` is not effective for a non-root process. The listener still comes up through HackPit's gated `start_listener(kind="iodine")`, refusing on all three limbs first; iodine is the other tool on the same gated path, not a bypass. The client is then forced onto the DNS-encapsulated path with `-r` — load-bearing, because on a flat lab bridge iodine will happily fall back to raw UDP and report a "tunnel" that never touched DNS, which would be a dishonest pass. The proof confirms the real thing on the wire: four ICMP echoes reach the server's tun endpoint through the tunnel with 0% loss, and a `tcpdump` on the server's UDP/53 shows the traffic arriving as DNS-encoded queries under the tunnel zone — eight of them for that ping. That is a genuine interface-level DNS tunnel, gated, carrying traffic, confirmed DNS-mode. (One harness bug surfaced and was fixed in the same pass, and it is the honesty guard earning its place: `tcpdump` line-buffers, so a first draft grepped the capture file before the packets had flushed and reported "no DNS packets" for a tunnel that was demonstrably DNS-encapsulated. The check now waits out the capture window.)

dnscat2's failure is now diagnosed decisively, and it is not HackPit's. A `tcpdump` on the server shows the client's encoded queries arriving correctly formatted — `<data>.lab.hackpit.internal`, alternating MX/TXT/CNAME, which is `dnscat2`'s own protocol — and the `dnscat2-server` process itself answering `NXDOMAIN` to every one. Put both ends in the *same container* talking over `127.0.0.1` — no network, no resolver, no HackPit code, the server's DNS driver confirmed started (`New window created: dns1`) — and it still `NXDOMAIN`s, in every client shape. That isolates it fully to the pinned `dnscat2` build (commit `42f8d783...`): a defect in the tool's own client/server handshake, sitting entirely above the listener HackPit starts, which is proven gated, bound and stable throughout. It is not chased further — fixing upstream `dnscat2` is out of scope, and iodine now demonstrates the DNS-tunnel class end to end regardless. The proof records this as a precise finding rather than a bare "did not work".

The two that remain are genuinely operator-side, and cannot be faked into a pass. A *delegated* DNS zone — the public-hierarchy hop in front of either tunnel — needs a domain the operator controls with an NS record pointed at the listener, plus a publicly reachable listener, i.e. the same VPS gap as D2. The operator has neither, so it stays recorded as a prerequisite. And ligolo's *interface* route is a deliberate **won't-do**, not a not-yet: completing a ligolo session means typing `session` then `start` into a live C2 console, and the only way HackPit could do that is a surface that writes to a listener's stdin — precisely the capability `test_lifecycle_safety` exists to forbid (the held stdin is a raw fd specifically so no writer object exists). chisel already carries the routed-traffic claim end to end, so nothing is lost by leaving ligolo's console step to the operator.

Build #8 — the reasoning copilot, and a measured substrate (2026-07-28)

Two things this build set out to establish, one about the agent and one about the ground it stands on: make the orchestrator a genuinely *deeper* proposer without giving it one inch more autonomy, and turn "110 tools" from a claim into a measured number.

The substrate actually runs, and here is the number

`reconcile.py` has long answered *installed?* with a `command -v` sweep. That is a weaker claim than it sounds: a binary can resolve on PATH and then die the moment the dropped capability is needed. So `backend/substrate_probe.py` asks the harder question directly — it invokes every catalogued tool in the **live sandbox image, under its real profile** (`no-new-privileges + CapDrop: ALL`) with a trivial call (`--version / --help`) and records three tiers: *catalogued* → *installed* → *actually-runs*.

97 of 104 catalogued Linux tools actually execute (93.3%), measured, not asserted. The breakdown is reported honestly, because a tool that does not run is never counted as covered:

- **2 installed-but-do-not-run** — `amass` and Empire (`powershell-empire`). Both are the documented landmine class, though not the one expected: their packaged wrappers shell out to `sudo`, which `no-new-privileges` refuses, so the binary resolves and then fails. (nmap, the *expected* casualty, runs fine — the image was hardened since that trap was recorded.)
- **5 not-installed** — `proowler`, `scoutsuite`, `kube-hunter` (cloud-audit tooling), `ghidra`, `subwiz`. Genuinely absent from the image; reported as gaps, not covered.
- **6 windows-only** — Rubeus/PowerView/Mimikatz/winPEAS-class entries, not applicable to a Linux sandbox by construction (D9).

The probe earned its place immediately: it surfaced two real catalog defects of exactly the kind Task 1 warns about — **a package that installs under a different binary name than the catalog uses**. Sliver resolves as `sliver-server` / `sliver-client` and Empire as `powershell-empire`; the catalog listed neither alias, so both read as "not installed" until the aliases were added — and adding them then forced those binaries through the arsenal danger gate under their *real* invoked names (`sliver-server`, `sliver-client`, `powershell-empire` now each demand the red-confirm, the same `sliver` -vs- `sliver-client` bug the gate audit fixed once already, caught again by the same test). The full per-tool report is committed at `docs/substrate-coverage.md` (+ `.json`); the static half (catalogued names vs. Dockerfile install lines) is unit-tested and needs no stack.

The reasoning proposer (2.1–2.8) — deeper proposals, not autonomy

The orchestrator used to pattern-match: it saw a port and reached for the tool it always reached for. `backend/reasoning/` makes it reason, and every component is additive, independently tested, and produces **proposals + rationale only**:

- **2.1 working memory — a tried/failed ledger.** The single biggest anti-looping lever. The proposer is fed the full state plus a ledger of what has been tried and how it turned out (read from the exit code *and* the output — a tool can exit 0 having failed), and a lead ruled out by the critic is persisted as a dead end. A failed lead does not get re-proposed. Validated live: with a `curl` to a closed port recorded failed, the model proposed the SQLi step against the real finding instead of repeating the dead lead.
- **2.2 hypothesis-first proposals.** The schema now leads with `hypothesis` + `expected_signal`, then the command — enumerate-before-exploit as a shape, and a wrong guess made visible. A proposal missing hypothesis, expected_signal, or a citation **fails a validation gate** (invariant 3, expressed as code with a control that shows it can fail).
- **2.3 a candidate frontier, not a greedy single step.** The proposer surfaces N leads, each scored by *evidence strength × expected payoff*; the top is pursued and the rest are **persisted as untried leads**, so a dead-end recovers to the next-best instead of looping. This generalizes the AD graph's edge-frontier to the whole engagement. It is a set of leads, **never a run queue** — nothing in it fires.
- **2.4 failure diagnosis.** A connection-refused output yields a reachability check *before* another attempt, not a blind repeat.
- **2.5 a skeptic/critic pass.** A second look tries to *refute* the proposal against the evidence, reusing the CVE→exploit index's rule that **the version verdict outranks token similarity**: a proposal citing a CVE whose exploit targets Apache 2.4.49 is caught and downranked against an observed 2.4.58, no matter how well the words match. This is the check that kills the top hard-box failure mode — confident hallucination. The critic is **advisory only — it downranks and flags, it never suppresses**: a refuted proposal is surfaced with its concerns and a lowered confidence so the human sees it is shaky, but the proposer stays free to raise it again, and nothing is ever blocked from being run by hand. (The first live run against Ollama found a real bug in this pass — it read an IP octet as a port — which was fixed and regression-locked; the honesty guard earning its place.)
- **2.6 domain-specialist routing.** A web finding routes to a web lens, an AD state to an AD lens, a foothold to a privesc lens — instead of one generalist voice.
- **2.7 fingerprint-keyed retrieval.** An exact service+version fingerprint (case-based: "this stack was solved by X") ranks ahead of generic token matches. The plumbing is built now; **growing the exploitation-write-up corpus it keys into is a follow-on** — noted, not faked: on today's KB the ranker degrades to the base order.
- **2.8 model-tier routing.** A config lever, default OFF and inert: with no `reasoning_tiers` configured, every step uses the base config byte-for-byte; configured, a hard step can be routed to a more capable model.

The honest ceiling. Reasoning depth is mostly the base model. A small local model caps it no matter how good the scaffolding is — the same live run that produced a fully-cited, correctly-routed proposal produced, on a re-run, one with no hypothesis at all (which the schema gate then flagged, correctly). This build moves the loop *up the curve* on easy and known-pattern-hard boxes; it does not turn a 9B local model into something that "solves anything hard." What it reliably changes is the loop's behaviour: it stops re-proposing dead leads, it states what it is testing, and it refuses to cite a CVE that does not fit the version in front of it.

Web-exploitation and privesc depth (Tasks 3 & 4, still human-fires)

Two drafting surfaces, both propose/ground/generate only:

- **Web (`reasoning/webexploit.py` , `POST /sessions/{id}/webexploit/draft`).** Given a web finding in state, HackPit drafts the actual exploit — the sqlmap invocation, the SSRF metadata probe, the IDOR replay, the parameter-pollution request, the LFI traversal — with an explanation grounded in the bug-bounty KB and citations back to the finding, handed to the human to fire through the **existing** repeater/executor. Nothing in the path sends without the human.

- **Privesc** (`reasoning/privesc.py` , `POST /sessions/{id}/privesc/ingest`). Paste linpeas/winpeas output; it identifies the vectors (NOPASSWD sudo, SUID GTFOBins, PwnKit, capabilities, SelImpersonate, AlwaysInstallElevated), drafts the escalation for the strongest one grounded in KB + state, and hands it over. Execution stays human-approved.

What is re-locked, unchanged

The whole point is that the proposer got smarter and nothing else moved. `test_loop.py` 's source-scan lock is **extended onto the shared scanner across the entire `reasoning/` package**: `orchestrator.py` plus all eleven reasoning modules are proven to have no execution path — no subprocess, no executor run method, no `:kali /sandbox` — by AST call-analysis (so a *drafted payload string* that contains `os.system(...)` as exploit text is correctly **not** a violation, while a real call is), with positive controls. A companion scan proves there is **no auto-approve / batch-approve / run-the-chain** anywhere in the proposer path or the frontier. The four execution gates (scope, approval, danger red-confirm, isolation) are untouched and still fire.

KB exploitation-writeup corpus — the case-based material 2.7 keys into (follow-on, now landed)

Build #8 shipped 2.7's fingerprint retrieval as plumbing and flagged the corpus it queries into as a follow-on; that corpus now exists. `pipeline/ingest_exploitation_writeups.py` seeds **78 distilled exploitation writeups** keyed by service+version fingerprint — the initial 16 (vsftpd 2.3.4, ProFTPD mod_copy, Samba usermap, SMBv1/EternalBlue, Apache 2.4.49/.50 traversal→RCE, Shellshock, Jenkins, Tomcat, Drupalgeddon2, GitLab ExifTool, Log4Shell, unauth Redis, BlueKeep, Kerberoasting, AD CS ESC1, Zerologon) plus a second batch of 19 (Confluence OGNL, Struts2, Citrix, F5 BIG-IP, PrintNightmare, ProxyShell, ProxyLogon, Spring4Shell, FortiOS, PHP-CGI, Grafana LFI, WebLogic, phpMyAdmin, MSSQL xp_cmdshell, SNMP, PostgreSQL, Webmin, Joomla, Elasticsearch), and a third of 43 across four themes — **AD depth** (AS-REP roasting, DCSync, noPac, unconstrained/RBCD delegation, PetitPotam+ESC8, GPP cPassword, LAPS, pass-the-hash, golden ticket), **unauth network services** (MongoDB, Memcached, CouchDB, MySQL UDF, JMX/RMI, IPMI, VNC, SSH user-enum, anon FTP/LDAP, SMTP enum), **Linux privesc** (DirtyPipe, Dirty COW, Baron Samedit, Docker/LXD group, NFS no_root_squash, PATH hijack, LD_PRELOAD, wildcard injection) and **modern enterprise CVEs** (Spring Boot Actuator, Laravel Ignition, Text4Shell, ColdFusion, PaperCut, TeamCity, OFBiz, vCenter, Cacti, ThinkPHP, WordPress, GraphQL, Zimbra) — each carrying a structured `meta.fingerprint ( service + version_kind + versions + cve + solved_via + a one-line detection footprint)`, whose `version_kind / versions` shape **mirrors the CVE→exploit index exactly** so the version verdict stays authoritative.

Three properties matter and each is verified:

- **Distilled, not verbatim.** Every entry is a general technique note — the approach, the command pattern, what it leaves in logs — distilled from the offensive- / *bug-bounty methodology and public CVE facts*. *Not one line is a copyrighted third-party walkthrough. Operator-supplied writeup notes are handled like the phishing lures: imported at runtime with `--operator` from a gitignored engagement directory, marked `operator` , and never committed (entries.jsonl is itself gitignored, so the seed ships only as `ingester code*` that regenerates it).*
- **Additive, idempotent, consolidation-safe.** The ingester mirrors the recon-methodology template: existing lines pass through as raw bytes, its own lines carry `meta.exploitation_writeup` , and a re-run is **byte-identical** (verified by hash). Each entry is `category="writeup" + meta.no_merge` , which `consolidate.py` treats as standalone — so no duplicate is ever created, without re-running the full pipeline (which would risk the downstream-enrichment revert and a Defender quarantine). The file was confirmed present and countable after every write. KB 2,621 → 2,699.

- **Reachable by 2.7, with the version verdict intact.** A fingerprint query for a discovered service returns the seeded writeup **ranked ahead of generic token matches** — verified end-to-end through the real KB search (vsftpd 2.3.4, log4j 2.14.0, gitlab 13.9.0, drupal 7.50, apache 2.4.49 all range-match their writeup at rank 1) — and an **out-of-range** version does not wrongly match (log4j 2.17.0, apache 2.4.58, gitlab 14.0.0 correctly do not). `retrieve()` gained an optional entry-hydrator so the structured `meta.fingerprint` (which the lightweight search index does not carry) is available to the range match; the exact-version substring path remains the fallback.
- **Now wired into the live loop.** 2.7 is no longer only plumbing: `main.py` injects the KB retriever into the orchestrator at startup, and the proposer prompt carries a **FINGERPRINT-MATCHED WRITEUPS** block for the exact service+version fingerprints in state — "this stack is commonly solved via X", with the `solved_via` lead and a citable entry id. Verified live: a `vsftpd 2.3.4` service in state pulls its backdoor writeup into the prompt. Additive and injected: with no retriever wired (the hermetic tests), the block is empty and the prompt is unchanged; it retrieves, and like the rest of the proposer path it runs nothing. The remaining growth is content — more distilled fingerprints — not code.

Verification

- **Hermetic safety suite** (`sh backend/run_safety_tests.sh`) — green after every phase, expanded across all five with `test_phase1_runtime`, `test_state`, `test_scope_hostcheck`, `test_credvault`, `test_corpora`, `test_detection` / `test_detection_safety` (OPSEC channel + blue-view-unchanged), `test_repeater`, `test_tunnels`, `test_report_templates`, persistent-shell containment tests in `test_kali`, Phase 5's `test_terminal` (PTY containment + the sentinel shell provably untouched) and `test_exploits` (version comparison, tiered ranking, executes-nothing), and the Windows backend's `test_winrm` + `test_winrm_safety` (host-locked / no gate bypass / secret never leaks / orchestrator can't auto-run WinRM) with the AD oracle extended to the native Windows variants, plus the post-assessment `test_search_ranking` (substance-gated tier boost + completeness nudge) and `test_codescan_rules` (8-language rule bundle + resolver), plus build #4's six new suites — `test_sliver` / `test_sliver_safety`, `test_obfuscation` / `test_obfuscation_safety` and `test_evasion` / `test_evasion_safety` (no agent path, gated-vs-human-only split preserved, `<listener>` never substituted, the tunnel key never crossing the HTTP boundary, the artifact never executed, and a negative control proving that stripping the OPSEC note makes the engine **refuse** rather than degrade), plus build #5's `test_scans` — the shared source scanner's own regression file, which runs **first** in the suite and plants a violation of each shape (unscanned package, basename-exempted path, aliased import, in-function import, `import_module` string concatenation, f-string) into a temp tree to prove the scanner catches it before ten locks rest on it, and the `I2` follow-on's four new `test_tunnels` cases (a start needs approval + the red-confirm, both server binaries actually trip the heuristic so the danger leg is live, the gated argv is the spawned argv, and the request carries the gate fields defaulting false), plus build #7's `test_prevalidated_gates` — the belt-and-suspenders re-checks on the prevalidated path, proving danger is re-checked in **all three** modes with a windows probe (`Write-Host go ; Invoke-Mimikatz`) that only the whole-script classifier catches, so a re-check wired to the generic heuristic fails rather than passing quietly, carrying positive controls that the same request **with** the ack still runs in all three modes and that a benign request is untouched. plus build #7's continuation — `test_lifecycle_safety`, the lock on the shared listener lifecycle (a watched listener exposes **no stdin writer**, because the child is handed a raw pipe fd rather than `subprocess.PIPE` so `proc.stdin` is `None`; a whole-tree scan with a planted-violation control keeps the handle inside the four modules that own a listener; every branch of the status derivation is exercised including the one that must never claim a listener, an unprobed port; a dead process is reported dead **with the output it died on**; and a stop is asserted to reap the container-side server targeted at its own argv rather than

the tool in general), the new per-surface cases for the two gated lifecycles (engagement + approval + red-confirm each refusing with nothing spawned, both server binaries in each pair actually tripping the heuristic so the danger leg is live, `-i` present for the console binary and absent for the daemon, and a dead listener refusing rather than reporting up), and the Sliver build's new contract (`generated` vs `failed` decided by the artifact read-back, not the console's exit code, asserted in **both** directions), plus build #8's two new suites — `test_reasoning` (the reasoning copilot's plumbing: the ledger records a failed lead and blocks its re-proposal; the hypothesis/expected_signal/citation schema **fails** an uncited proposal; the frontier persists untried leads and recovers a dead-end without resurrecting a killed one; a connection-refused output yields a reachability diagnostic; the **positive control** that a version-mismatched CVE is caught and downranked while the matching version passes; web/AD/privesc specialist routing; fingerprint retrieval outranking a token match; the model-tier lever inert when unset; and the web-exploit + privesc drafts grounded and cited) and `test_substrate` (the Task-1 verdict function separating a landmine from a version banner from an absent binary with a control, the whole pipeline classified over the **real** catalog with an injected container, and the static Dockerfile coverage) — and the extension of `test_loop` itself onto the whole `reasoning/` package (no execution path anywhere in orchestrator + eleven reasoning modules, by AST call-analysis with a control that a payload *string* is not a false positive, and no auto/batch-approve in the proposer path or frontier), and the exploitation-writeup corpus's retrieval-alignment case (`test_reasoning`): a structured-fingerprint writeup range-matches an in-range version and floats above a token match, while an **out-of-range** version does not match — the version verdict holding through the corpus. Build #9 then added three regression files/cases for the defects its live fire surfaced — `test_secretargs` (credential values redacted out of the persisted record, per-tool, with `nmap-ports` and `password-with-@` controls), the `impacket-` prefix parser-name case in `test_state`, the inherited-rights-never-armed case in `test_adorch_safety`, and the `-dc -needs-FQDN` ordering case in `test_adgraph_collector`. **48 test files, 529 checks, 0 failures.**

- **Live-fire proof** (`sh docker/proof/live_fire_proof.sh`) — **66 passed, 0 failed, 3 not-run** against the rebuilt image, after the defects the earlier runs found were fixed and iodine was demonstrated end to end. Passing: an implant that builds through the gated path and a **beacon that calls back**; a **proxychained request routed through the pivot** into an otherwise unreachable subnet, using the argv `wrap_command` produced, with the same tunnel unable to reach `1.1.1.1`; all three lifecycle gates refusing on all three limbs with nothing spawned; every listener observed bound by `ss` inside the container and still bound at the end of its phase; the pivot subnet entering scope only via the explicit amendment, with the deep target confirmed unreachable first; and containment measured **from inside** the implant host — no internet, no host, C2 only. Also passing: iodine's IP-over-DNS tunnel, through the same gated surface, carrying ICMP with 0% loss and confirmed DNS-encapsulated by a server-side `tcpdump`. Not-run and reported as such: ligolo's interface route (a deliberate won't-do — it needs commands typed into a live C2 console, which HackPit will not do), the dnscat2 client handshake (decisively an upstream defect in the pinned build, above HackPit's proven-bound listener), and a delegated DNS zone (a domain and public listener the operator does not have).
- **The audit's "probed and holds" list, re-run after build #5** — all 41 items still hold, including the `.exe` spellings from `I1`, the four gate orders, the no-silent-downgrade-to-lab path, and Critical 1's eight catalogued tools. The five demonstrated Critical 2 bypasses are all refused. The planted `from cockpit.kali import run_kali` in `adgraph/orchestrator.py` now **fails** the scan (verified against a temp mirror; the repo was never modified). Read-only enumeration — `nmap`, `sqlmap --dbs`, `netexec --shares`, `ldapsearch -x`, `hashcat`, `nuclei` — stays clean on both paths.
- **Docker proofs** — `isolation_proof.sh` 4/4 (lab still cannot reach internet or host), `engage_open_proof.sh` 4/4 (engage has full reach). Both re-verified in build #7 **after** the sandbox stack was rebuilt and recreated from the current Dockerfile — twice, the second time after the `proxychains` fix was added to that Dockerfile — which is a real change to the running containers rather than a no-op re-run.

- **Image build** — `hackpit-kali:build7b` builds from the current `Dockerfile.sandbox` with exit 0 and carries the `proxychains` fix baked in (`socks5 127.0.0.1 1080` , zero `socks4` lines); the stack was rebuilt and recreated from it, and the live fire re-run against that image reports the config as coming from the image rather than from its own runtime patch. `hackpit-kali:build7` builds from the pinned `Dockerfile.sandbox` with exit 0, and every pinned tool resolves and smoke-tests by its real invoked name (`dnscat2-client` , `dnscat2-server` via its four gems, `donut`), with the shipped `dnscat2 HEAD` , gem versions and `donut-shellcode` version confirmed **inside** the image to equal the pins.
- **Browser** — every UI surface exercised against a live Ollama backend per the testing rule: the Phase-4 surfaces (payload-set arsenal rows, the OPSEC red-team channel, repeater send/replay/diff, tunnels route preview, exam report templates with the proof table) and the Phase-5 ones — `:terminal` running `top` and `vim` with live resize, `:exploits` resolving `vsftpd 2.3.4` to the backdoor exploit and its CVE, the state panel's per-service jump into it, and `/category/cloud` at 535 entries. **Windows targets** (`/windows`): profile create/list/test/delete and the AD-walk "run on" picker verified against the backend. As of build #9 the connectivity **test** and a live WinRM round-trip are no longer deferred — they are **demonstrated live** against a real DC through the same HTTP API the UI drives (see the build #9 live-fire below); the browser render of `/windows` and `/cockpit` was confirmed serving 200 while that run executed.
- **Live fire against a real Windows/AD domain** (`backend/livefire_windows.py` → `backend/livefire.log` , build #9) — **45 passed, 0 failed, 3 not-run** against the operator's VMware Server-2022 `corp.local` DC, driving the live backend over its real HTTP API. Demonstrated live: a WinRM round-trip landing on the profile host with the credential absent from every record, the whole-script danger classifier firing on real input, real `bloodhound-python` collection parsing into the typed graph (84 nodes / 625 edges) with the DA path computed on it, a real DCSync returning `krbtgt` through the red-confirm with loot ingesting into state, and the walk advancing only on the approved exit-0 run. It also found and fixed four defects invisible to the hermetic suite (credential-in-record leak, impacket parser-name mismatch, KB grounder arming a free edge, collector FQDN classifier), and truth-checked the detection catalog against the DC's own Event Log — confirming DCSync's 4662/DS-Replication signal (12 of 12) and **correcting two catalog claims the real telemetry contradicted** (secretsdump's mode-specific artefacts, and `ad_collect` raising no 4662). Not-run and reported as such: a Windows-target C2 callback (needs a routable listener), the native `Invoke-Mimikatz` DCSync variant (no offensive tooling on a bare DC), and a model-spontaneous DCSync pick.
- **Frontend** — `tsc` clean, lint at the pre-existing baseline (11 errors + 1 warning, unchanged — verified by stashing the changes and re-running), `next build` exit 0 (routes include `/terminal` , `/exploits` , `/windows`).
- **The launcher, verified against real state rather than a mock** (build #11) — the status rail was exercised in both directions on the same page with no code change: with Docker down it reported `sandbox stack down` with **7 tiles dimmed and 7 stack down markers**; after `docker compose up -d` it reported `up · engage up` with **0 dimmed and 7 needs the stack** . The freshly created networks were re-checked against the isolation floor — `assert_isolation_proven()` passes and `hackpit-isolated` is `internal=true` while `hackpit-open` and `hackpit-engage` are deliberately not. 4 bands, 15 surface tiles and 30 category cards render live.
- **Operator identity, checked against git itself** (build #11) — `test_operator.py` asserts `backend/operator.json` is ignored by asking `git check-ignore` rather than by reading the ignore file and trusting the pattern, and carries a control proving the predicate can answer "no" for a tracked file. The staged diff was grepped for the real name and identifiers before committing; only synthetic fixtures remained.
- **KB enrichment, measured rather than asserted** (2026-07-31) — the first external-corpus batch is verified four ways instead of by row count: **per-source counts diffed before and after** (2,699 → 2,712, every other source byte-for-byte unchanged), the KB file confirmed to still exist afterwards (Defender has deleted it before), **6 of 7 natural-language retrieval probes returning the new entry at rank 1**, and the full safety suite at 52 files. A first suite run aborted at `test_corpora.py` and passed on re-run and standalone; the cause is environmental and

is recorded rather than hidden — `pipeline/ingest_corpora.py:139` shells to `git show HEAD:<file>` expressly "to recover AV-locked / dehydrated files", and rewriting the 22 MB `entries.jsonl` triggers a Defender sweep that briefly locks corpus sidecars.

- **KB enrichment batch 2 — 24 cert/CTF/pentest repos** (2026-08-01) — verified the same four ways, and the numbers are small on purpose. **Per-source counts diffed before and after:** 2,712 → 2,714, all 22 sources still present and every one of the other 21 unchanged to the row; `pivoting` 6 → 8. The KB file, `embeddings.npy` and `ids.json` all confirmed present afterwards (22,487,315 bytes; Defender has deleted the KB before, so it is backed up first and re-checked after). **Retrieval: 6 of 6 natural-language probes surfaced a new entry in the top 5, four of them at rank 1** — including the ones a command reference cannot answer ("the second pivot host cannot reach my attacking machine, how do I get a shell back", "why does nmap find nothing through proxychains"). Full hermetic suite green at **52 files, every one exit 0**; `test_corpora.py` aborted the first run and passed standalone and on re-run, the documented Defender/sidecar flake, recorded rather than hidden. Saturation was measured rather than asserted: 308 terms appear in ≥3 repo files and never in the KB, and all 308 are URL slugs, hostnames, playlist IDs, filenames or typos. Arsenal: 110 → 115 tools, all 188 catalogued invocation names carrying a pinned danger verdict, all 286 templates across 115 tools rendering target-faithfully with no foreign host. All five new binaries were confirmed **present and runnable inside the live sandbox image** before being catalogued, rather than assumed from the Dockerfile. All 1.6 GB of clones deleted afterwards, with `git status` confirming the manifest is the only thing in `sources/` that git can see.
- **KB enrichment batch 4 — the 0xdf fingerprint corpus** (2026-08-01) — unlike batches 1–3 this one **added** rows, so the verification is that it added exactly what it claimed and touched nothing else. Per-source counts diffed before/after: **only hackpit-distilled moved, 78 → 97 (+19); no other source lost a single row**, and a total alone was not accepted as proof. `data/kb/entries.jsonl` still exists (22.5 MB, 2,733 rows) after both the ingest and the embed — the Defender-quarantine check that has bitten before. `embed.py` reports **19 new, 2,714 cached**, so the vector space grew by exactly the new entries. Retrieval was run through `search.search(entries, query, top)` (positional, no `k=`) on realistic scan strings — `OpenSMTPD smtpd 6.6.1 25/tcp`, `Apache Tomcat 9.0.27 8080`, `Apache Solr 8.3.1 wt velocity`, `jetdirect printer 9100 snmp`, `IPMI 623 BMC RAKP`, `gitlab 11.4.7 redis ssrf`, `netdata ndsudo setuid`, `xwiki solr search groovy`, `pgadmin query tool 9.1`, `mariadb wsrep 10.3.25` — and each new fingerprint ranked first or top-three for its own banner while the controls `vsftpd 2.3.4` and `Apache 2.4.49` still won theirs. `sh backend/run_safety_tests.sh` → **52 test files, every one exited 0**, `test_corpora` included and with no flake this run.
- **KB enrichment batch 3 — 7 GitBook certification spaces** (2026-08-01) — the verification here is of a **zero**, which is a different claim and needs different evidence: not "the ingest did no damage" but "no entry was warranted." The KB is byte-unchanged at **2,714 rows**, and no ingester was run. Saturation was established at three independent levels. **Index level, before any page was requested:** the 545 page URLs the six reachable spaces publish were tokenised against all 2,714 rows, and of 109/88/59/37/56/166 slug words per space, the words absent from the KB number **6/5/3/1/2/2** — every one an author name, a certification acronym, or a typo (`foundamentals`, `accross`, `uncostrained`, `gnereral`). **Page level**, for the one space fetched in full: 175 distinct leading commands across its 214 code blocks, of which 10 appear nowhere in the KB and all 10 resolve to covered material. **Concept level**, for the two survivors that a token diff called new — and this is the leg that mattered, because `seshutdownprivilege` (11 occurrences, 0 in the KB) is a real technique that the KB **already carries under a different spelling**, `ex-windows-checking-services` step 4: *"reboot if you hold SeShutdown."* A token miss is not a coverage miss. Full hermetic suite green at **52 test files, every one exit 0**, run after the batch; `test_corpora.py` did not flake this time, which is consistent with its documented cause — no 22 MB KB rewrite happened to trigger the Defender sweep. Politeness is verifiable rather than claimed: `robots.txt` is parsed and honoured **before the first request**, and it refused a space during this batch rather than in theory.

- Fingerprint retrieval — measured, then fixed, then re-measured** (2026-08-01, docs/FINGERPRINT-EVAL.md) — a measurement session first drove the real 2.7 path (`orchestrator._fingerprint_reference` → `reasoning.retrieval.retrieve`) against a test set that was **not** drawn from the corpus alone: 30 covered service+version strings in real `nmap -sV` formatting, 15 near-miss versions outside the stated range, 15 uncovered services. It found the corpus firing at only **70%** on covered banners with two structural defects behind the misses (D24), and — the number that mattered — **0% false-fire on near-miss**, the safe direction. The fix session then lifted covered hit rate to **93%** at **96%** precision, took the exact-boundary case from **0/4 to 4/4** and the corpus-wide self-match from **62/97 to 97/97**, **held near-miss false-fire at 0%** (a version one step above the boundary still does not fire, 0/3 — the change is inclusive `<=` at the boundary and nothing looser), and left the CVE→exploit index **provably byte-identical** (110,695/110,695 verdicts unchanged, `test_exploits.py` green). Two residuals were left deliberately and named rather than folded in: a 20%-unchanged uncovered false-fire from the version-less substring fallback (a different code path), and 2 covered misses that are a base-retriever top-5 recall limit, not a fingerprint one. Locked by two new suites that iterate the real corpus with positive controls — `test_fingerprint_versions.py` (every fingerprint matches its own stored version) and `test_fingerprint_norm.py` (`Apache Tomcat` ≠ `Apache httpd` , no collision). **54 test files, every one exit 0.**
- Fingerprint retrieval — the last residual closed** (2026-08-01, D25) — the one gap the fix session named but left, a **20% false-fire on UNCOVERED services**, is now fixed and re-measured on the same three groups. It was the version-less substring fallback in `rerank()` setting `fingerprint_match=True` on a bare product-name substring; the eval was instrumented before the change to prove the fallback contributed **0 of the 28 covered fires** and was the **whole** of the false-fire, so demoting it (reserve `fingerprint_match` for structured hits; a product-name hit becomes a labelled, lower-ranked `fallback_match` ; word-boundary match) cost nothing on hit rate. Measured before/after on all six: **UNCOVERED false-fire 20% → 0%**, **COVERED hit rate 93% (unchanged)**, precision **96% (unchanged)**, **NEAR-MISS 0% (held)**, corpus self-match **105/105 (held)**, fire-above-boundary **0/45 (held)** — a clean removal, not a stricter-matcher trade. Locked by `test_fingerprint_fallback.py` (live corpus + 15 real uncovered services + positive control, can-fail proven by monkeypatch). **55 test files, every one exit 0.** KB untouched — the fix is entirely in `reasoning/retrieval.py` .
- KB enrichment batch 3 — PDFs, pages, and the CPENT repo: the series close-out** (2026-08-01) — five sources, **2 entries added**, verified the four required ways. **Per-source counts diffed before/after**: only `hackpit-authored` moved, 27 → 29; every other source unchanged to the row, and the category deltas are exactly `ics` 1 → 2 and `recon` 69 → 70 . `data/kb/entries.jsonl` confirmed present after both ingest and embed (the Defender-quarantine check; the KB was backed up first, then the backup deleted once integrity was confirmed). `embed.py` reports **2 new, 2,741 cached** — the vector space grew by exactly the two entries. **Retrieval through `search.search(entries, query, top)` (positional, no `k=`)**: both new entries rank **first** on all four probes a command list cannot answer — "nmap scan an OT network with fragile embedded PLCs safely", "scan an ICS SCADA network without crashing controllers", "map firewall rules and enumerate an ACL with nmap", "bypass a packet filter with fragmentation decoys and spoofed source port". Saturation of the four zero-yield sources was measured, not assumed: every candidate technique (printer pass-back, IKE aggressive mode, open mail relay, ligolo/chisel/sshuttle pivoting, employee OSINT, o365 spray) grepped as already covered, several better than the source covers them — the two entries that landed came from the CPENT repo's two thinnest modules naming a scan-safety discipline the KB had never written down. Full hermetic suite green at **55 test files, every one exit 0**; `test_corpora.py` aborted the first run on the documented Defender/sidecar flake (the 22 MB KB rewrite) and passed on re-run. The 185-page PNPT space, both ebook PDFs and the parzival blog post all yielded **zero**; raw trees deleted, `sources/pdfs-pages-manifest.md` the only residue.

Status

Build #4 (AV/EDR evasion + traffic obfuscation) is complete and verified: image `hackpit-kali:build4` builds with all eight new binaries smoke-tested by invocation, the full hermetic suite was green at 42 files / 446 checks at that point, and the frontend builds clean. Tasks 8–13 have now had their independent review (see Part II) — no containment hole, five substantive defects fixed, each with a regression test. **Two caveats are carried forward as open items in Part II rather than papered over**, and build #4 should not be called field-ready until they close: the C2 and tunnel surfaces have never been run against a live beacon, a real delegated DNS zone, or an instrumented Windows host, so their *efficacy* claims are documented rather than demonstrated (the *containment* claims are what the suite locks, and those hold without live infrastructure); and three installs in the image layer are still unpinned, so that layer is not yet byte-reproducible.

A note the review sharpened, because it cuts against the instinct to trust the tooling:

`pipeline/detection_sources.py --verify` reported **0 problems** throughout, including while `evasion_etw_blind` carried a technique id about shell history. It verifies that every id, name, tactic, log source and Sigma reference matches live upstream — it cannot verify that the technique chosen is the right one for the activity being described. A clean `--verify` is a check on transcription, not on judgement.

Build #5 (gate integrity) is complete, and the audit's one remaining important item is closed with it. All three gate-audit criticals are now closed: Critical 1 in build #4's fix wave (`4492fec`), and Criticals 2 and 3 here. `I2` — the ungated tunnel listener start, outside the plan's scope — was recorded open, then fixed as a follow-on in the same build. The suite stands at **43 files / 477 checks / 0 failures**, the frontend builds clean, and the audit's full "probed and holds" list was re-run with nothing regressed. What changed is not only three guards but the conditions under which a guard is allowed to be believed: `backend/AGENTS.md` now requires a safety test to iterate real data, assert on what it actually checked, and carry a demonstration that it can fail — and `test_scans.py` runs first in the suite to hold the shared scanner to that standard before ten locks depend on it. The remaining work is the deferred list in Part II — chiefly live-fire verification of the C2 and tunnel surfaces, and the three unpinned installs in the image layer.

Build #7 (live fire, image pinning, prevalidated danger re-check) is complete, and it is the first build whose headline result was a set of defects rather than a set of features. Two of its three tasks closed cleanly: the image layer is byte-reproducible and `hackpit-kali:build7` builds green, and the prevalidated path now re-checks danger in all three modes. The third — live-fire verification, Part II's largest outstanding item — did what it was supposed to do. It proved the gates fire against real binaries and proved containment holds live from inside the contained host. It also found that **three of the surfaces did not work as shipped**: the Sliver implant argv named a subcommand that does not exist, and both listener lifecycles reported `status="listening"` for a process that had already exited. None was a containment failure — nothing escaped and no gate was skipped — but they meant the efficacy claims those surfaces carried were never true. Two of the three were found only because a real build and a real process replaced a structural assertion; the pinning task independently surfaced a smoke test that had never passed for the same reason.

All three are now fixed, and the design question the run left open is decided (2026-07-28). The implant builds through the console `sliver-server` hosts and its success is decided by reading the artifact back rather than by a console exit code; both listeners hold a console binary's stdin open — as a raw fd with no writer object, so the surface gains liveness without gaining the ability to type into a live C2 console — and report a status they observed, with a dead process now a refusal rather than a fiction. The live fire stands at **66 passed, 0 failed, 3 not-run**: a beacon calls back, a proxychained request routes into an otherwise unreachable subnet, iodine's IP-over-DNS tunnel carries traffic confirmed DNS-encapsulated on the wire, and every gate still refuses on every limb. The stale precedent is resolved by tightening rather than by deleting the citation: the DNS-tunnel listener is gated exactly as `I2` gated the pivot listener, because a covert channel is an exfil route; the Sliver server is gated too, on

the narrower ground that its config can persist listener jobs that come up with the daemon, so "starting it is inert" is a property of a config file rather than of this code. Human-only is untouched and remains the load-bearing property; the gates are belt to those suspenders.

Fixing them surfaced two more defects of the same family, both also fixed: a stop that killed the `docker exec` client while the server kept running and holding its port inside the container, and a `proxychains` config shipped pointing at Tor rather than at the SOCKS5 port HackPit's own rewrite targets — a shipped path that every unit test passed and that could not have worked. A second pass then verified both of those against a rebuilt image rather than against a diff — the `proxychains` layer now builds and the config is baked in, and the stop-reap was demonstrated live against the failure that motivated it, including that reaping one tunnel leaves a neighbouring one bound. That pass surfaced a sixth defect of the same family: a status observed once and never re-observed upward, which left a listener that bound just after its settle window permanently `starting` and therefore permanently unroutable. It now refreshes in both directions, and only on evidence. That is the durable lesson of this build, and it is worth stating plainly: **six defects in these surfaces were invisible to a green test suite, and every one of them was a claim asserted structurally that nothing had ever executed.** The suite now stands at **45 files, 502 checks, 0 failures**, with `test_lifecycle_safety` locking the new shared lifecycle — including the positive controls that a dead process is reported dead and that an unprobed port is never called listening. Two of the four not-runs then closed: iodine's DNS tunnel is demonstrated end to end, and dnscat2's failure is pinned decisively to the upstream build rather than to anything HackPit owns. What remains genuinely undemonstrated is stated as such and not folded into a pass: a delegated DNS zone (an operator prerequisite), ligolo's console-driven interface route (a deliberate won't-do), and detonation on an instrumented Windows host.

Build #8 (the reasoning copilot + substrate proof) is complete. It did the two things it set out to do without moving the safety boundary an inch. The substrate claim is now a *measured* one — **97 of 104 catalogued Linux tools actually execute** in the live image under its real profile (93.3%), with the seven that do not reported honestly as two `sudo`-wrapper landmines and five genuine absences, and the probe surfaced two real catalog defects (Sliver/Empire installing under a different binary name) that are now fixed and forced through the danger gate under their real invoked names. The orchestrator is a genuinely deeper proposer — working memory that stops it re-proposing dead leads, hypothesis-first proposals with an enforced citation gate, a scored candidate frontier with dead-end recovery, failure diagnosis, a refute-first critic that kills version-mismatched CVE hallucination, specialist routing, fingerprint retrieval, and a model-tier config lever — and **the honest ceiling is stated plainly: reasoning depth is mostly the base model, so this moves the loop up the curve on easy and known-pattern-hard, not to "solves anything."** The one line that did not move is the one that matters: **the loop reasons but never executes, and approval stays per-command** (D18), re-locked by extending the source-scan onto the whole `reasoning/` package and by a scan proving no auto/batch-approve exists in the proposer path or the frontier. The suite stands at **47 files, 519 checks, 0 failures**; the loop was validated live against Ollama (proposals cited, reacting to prior output, not re-looping, every command still stopping for human approval).

The follow-on the build flagged is now **landed**: the **KB exploitation-writeup corpus** that 2.7's fingerprint retrieval keys into. `pipeline/ingest_exploitation_writeups.py` seeds 78 distilled, non-verbatim technique writeups keyed by service+version fingerprint (the CVE-index `version_kind / versions` shape, so the version verdict stays authoritative), additively and idempotently (byte-identical re-run, `category="writeup" / no_merge` so consolidate never duplicates them, file confirmed un-quarantined; KB 2,621 → 2,699). Operator writeups import at runtime into gitignored engagement data and never reach the repo. Retrieval alignment is verified end-to-end through the real KB search: an in-range version returns the seeded writeup ranked ahead of token matches, and an out-of-range version does not wrongly match. The corpus is intentionally a seed — growing it further stays an ongoing content task, not a code one.

Build #9 (the Windows/AD path, driven live against a real domain) is complete. It is the build that moved the Windows/AD path from "built and locked hermetically" to "demonstrated live," and — like build #7 before it — its lasting value was as much in what a real target exposed as in what it confirmed. It confirmed a great deal: a WinRM round-trip locked to the profile host, the whole-script danger classifier firing on real input, real BloodHound collection parsing into the typed graph, and a real DCSync returning `krbtgt` through the red-confirm with loot ingesting into state — the gates holding at every step, on real input, 45 passed / 0 failed / 3 not-run. It also found **four defects the green hermetic suite could not see, every one a place where two halves of the codebase agreed in a test and disagreed against a real domain:** a domain credential persisted into the run record (and thence into the LLM proposer context), impacket loot ingesting as nothing because the catalog and the parser spelled the tool differently, the KB grounder arming an inherited-rights edge with a destructive command, and a failure classifier that misdiagnosed the collector's FQDN error. All four are fixed and regression-locked; the suite stands at **48 files / 529 checks / 0 failures**. Task 5 was the first time the detection describe-side was checked against real telemetry rather than documentation — it confirmed DCSync's high-fidelity 4662 signal and corrected two catalog claims the DC's own Event Log contradicted. What stays not-run is stated plainly and not folded into the pass: a full Windows-target C2 callback (a routable listener the sandbox does not expose by choice), and the native-tooling abuse variants (offensive tooling on a bare DC, out of scope). The durable lesson repeats build #7's in a new domain: **a live target is the only thing that finds the bugs that live between two components that were each tested alone.**

Build #10 (the opt-in C2 exposure override, and an honest NOT-RUN) — complete. Build #9's Task 4 recorded a Windows-target C2 callback as NOT-RUN for one concrete reason: the Sliver/tunnel listener lives in `hackpit-engage-sandbox`, which by standing invariant publishes no ports, so the lab DC had no route to it. Build #10 supplies exactly that missing route **without widening the default posture:** `docker/proof/c2-lab.yml`, an opt-in compose override that publishes a single port — `192.168.13.1:53/udp`, the iodine tunnel server's DNS port bound to the one host interface (VMnet8) the lab DC can reach — and nothing else, never a wildcard. It is never applied by the documented bring-up; `backend/test_exposure_safety.py` locks three invariants (default compose publishes nothing, every override port names a literal host IP, nothing composes it in implicitly) and plants a violation to prove the scan can fail. The override is not only built but **demonstrated live** — the exposure came up on `192.168.13.1:53/udp` and tore down cleanly. On it sits a gated C2/AD proof harness (`docker/proof/c2_0{1..4}_*.sh` + `c2_lab_proof.sh` orchestrator + in-process driver) built on build #7's discipline: offensive command strings are never inlined — they are operator-supplied paste values read by file path, and an empty value reports **NOT-RUN**, never a fake pass. The four offensive proofs it drives — staging the tunnel client on the DC, the iodine DNS tunnel, a Sliver beacon over it, and a native DCSync behind a scoped Defender **exclusion** (not a real-time-protection toggle) removed in a `trap` guard — are **NOT-RUN**: harness-ready but requiring operator-supplied commands that were deliberately not filled, recorded as such rather than folded into a pass. A final hardening pass closed the build's engineering items and corrected one of its own claims: the suspected `run_safety_tests.sh` gating hole did **not** exist (`set -e` has been in the committed runner all along, and a planted failure was verified to abort it), while the failing `nc` danger-gate test turned out to be an environment leak rather than an over-block — the one test in `test_session.py` that called the gates outside the hermetic fixture, so it reached a live Docker probe. Both are now fixed and the runner additionally checks every test's exit code explicitly. The build also added a read-only DC prerequisite probe and a harness-honesty regression test, and that test immediately caught a real defect: three of the four proof scripts' `[[PASTE]]` slots were not actually empty, so an unfilled proof was one WinRM round-trip away from being scored as a genuine attempt. Suite: **538 checks across 50 files, green**. The four offensive proofs remain **NOT-RUN**.

KB enrichment batch 2 (24 cert/CTF/pentest repos) — complete, and its honest result is two entries. 24 repositories cloned, **22 yielded nothing**, one produced both entries. Under D21 that is the process working: the KB was already saturated in what these repos cover, and 308 candidate "new" terms turned out to be URL slugs and filenames rather than techniques. The batch's scoping prediction was **wrong in a way worth recording** — CPENT

was expected to carry IoT and ICS/SCADA and carries neither; across all 24 repos exactly one file mentions ICS vocabulary and it is a CTF challenge named "Scada" that is really Jinja2 SSTI. `iot` (2) and `phishing` (1) remain the thinnest KB categories and now have a *reason* attached: exam notes are the wrong source class for them. The two entries fill the one genuine gap the batch did expose — every existing `pivoting` row is a command reference, and none teaches multi-hop chain discipline. The unplanned find was worth more than the planned one: the arsenal had **zero** pivoting tools despite `cockpit/tunnels.py` driving `chisel`, `ligolo` and `proxychains` since Phase 4, and cataloguing `proxychains` surfaced a **pre-existing gate hole** — the danger heuristic classified `argv[0]` only, so wrapping a dangerous command in a tunnel wrapper stripped its red-confirm, making the gate weaker the deeper you reached. Fixed in the predicate and regression-locked through `validate_request` (D22). Suite: **52 files, all green**; KB 2,712 → 2,714; arsenal 110 → 115. All 1.6 GB of clones deleted; `sources/repos-manifest.md` is the reproducibility record and required a deliberate, single-file gitignore exception to survive.

KB enrichment batch 4 (the 0xdf fingerprint corpus) — complete, and written at the time as the series close-out (superseded: batch 3 was restored as the actual closer and has since run — see the batch-3 records above, which are authoritative for the series' completion). This was the first non-cert-note source, and the first batch to add rows: **19 entries from 12 posts, hackpit-distilled 78 → 97, KB 2,714 → 2,733**, all in the thin categories a fingerprint corpus is for (services, pivoting, credentials, persistence, reversing/privesc). The standing ban on 0xdf was reviewed and reversed to the project-wide distil-not-parrot rule (D23), the ingester docstring was corrected to match, and the `consolidate.py` link-index skip was deliberately left alone. Phase 1 cost two requests and mapped 614 posts / 276 CVE tags, **175 of them never before in the KB** — the measurement that told us this source was unlike the cert notes before a page was read. 55 posts were selected by KB gap and fetched (Defender deleted 5 mid-triage; triage tolerated it and re-fetched), every candidate was concept-grepped before writing, and that grep killed ~43 shortlisted posts as already-covered and caught two pure slug collisions (`schallenge` → `XSSChallengeWiki`, `wsrep` → `wsrepl`) that a token count would have miscalled — the same `SeShutdown` trap batch 2 found. **The series, whole: five batches, of which the PDF/pages batch has NOT run** (it is written but never executed, and is not counted as complete). The four that ran added 34 entries (13 + 2 + 0 + 19). **The headline is the cert-note yield — 31 sources produced 2 entries** — a real measurement of the KB's maturity on exam-syllabus material, set against 19 from one narrative source that supplied the scan→root shape nothing else did. Carried forward: the token diff **nominates, never confirms** (grep the concept); `themastermindnotes` **declined** as a commercial product; `mqt.gitbook.io` **refused itself** by `robots.txt`; **0xdf declined-then-reversed**. KB now at 2,733; suite 52 files all green.

0xdf pass 2 (2026-08-01) — the one amendment to that close-out. Batch 4 read as the series end, but the retrieval eval it prompted found the corpus half-firing and the fix (`6c3ba42`) made new entries fire, so a final novel-CVE pass ran: **8 fingerprints from 51 posts** (HFS 2.3, IIS 6.0 WebDAV, Icinga Web 2, OpenTSDB, SQLPad, ES File Explorer, Next.js middleware bypass, Strapi), `hackpit-distilled` 97 → 105, KB 2,733 → 2,741, 0xdf total **27**. The yield is a declining tail (19/55 → 8/51) into a residue of web-CMS boxes the KB already covers, so **0xdf is now recommended closed as a fingerprint source**. Measured, not assumed: all 8 landed in `category="writeup"` (the ingester forces it), **not** the thin `services` / `pivoting` categories — correcting pass-1's claim about where its rows went; and the eval's UNCOVERED false-fire residual **held at 20% (3/15), zero delta**. Suite **54 files, all green** (the two fix-session locks now validate all 105 fingerprints).

KB enrichment batch 3 (7 GitBook certification spaces) — complete, and its honest result is zero entries. Batch 2 ended by predicting that four of these seven spaces were near-certain duplicates; the prediction held, and the two it named as genuine unknowns resolved against it too. **Five of the seven were never fetched at all** — four because an index-level gate settled them from their published page lists alone, and one because its operator has opted out of machine collection. Of the two fetched, one was taken only as its 15-page delta over a space already mined, and that delta turned out to be an Active Directory chapter that has been outlined but not written (1,058 words, **zero code blocks**, 24-word pages). The remaining space was fetched in full, deliberately, because its hypothesis was

methodology structure rather than technique novelty and no index gate can test shape — 20,248 words and 214 code blocks, and it yielded nothing either: the KB already carries **118 checklist-* entries** in exactly that shape. The build's carry-forward is a method, not a technique. The batch's one plausible find, `seshutdownprivilege`, was absent from the KB as a *string* and present as a *technique* under the spelling `SeShutdown` — so the token diff that batch 2 established as proof-of-saturation is now bounded by an explicit rule: **it can only ever nominate candidates, never confirm a gap; the concept has to be grepped before a zero or an entry is claimed**. Under D21 this is the process working: 21,306 words were read and nothing was written, because writing anything would have meant duplicating rows already present. KB unchanged at 2,714; suite **52 files, all green**; `sources/gitbooks-manifest.md` is the reproducibility record.

Fingerprint retrieval (2.7) is measured, fixed, and locked (2026-08-01). The eval that closed the 0xdf series asked the obvious unasked question — does the fingerprint corpus actually fire on real scanner output — and the answer was "half the time," behind two defects (D24): a shared version predicate with an unstated boundary convention that made **35 of 38 versioned fingerprints miss the exact version they were written about**, and a normaliser that collapsed `Apache Tomcat` and `Apache httpd` onto one key. Both are now fixed at the **predicate**, not the caller: `_version_verdict` names its boundary (`inclusive=`) so the CVE index (exclusive, fix-version) and the fingerprint corpus (inclusive, last-vulnerable) each state their own, and the CVE→exploit index is provably untouched (**110,695/110,695 verdicts identical**). Covered-banner hit rate went **70% → 93%** at **96%** precision; the exact-boundary case **0/4 → 4/4**; corpus-wide self-match **62/97 → 97/97**; and the safety number held — **near-miss false-fire stayed 0%**, with a version one notch above the boundary still refusing to fire. This is the **third shared-predicate defect** in the project (WinRM `argv[0]`, D22's proxychains laundering, now this), and it is called out as a pattern to watch. Two new regression suites iterate the real corpus with positive controls; **54 test files, all green**. Two residuals are named not hidden — a version-less substring-fallback false-fire and a base-retriever top-5 recall limit — neither blocking. No KB entry was touched: this was a matcher fix, and a second 0xdf pass is now worth running at the corrected rate.

The last retrieval residual is now closed too (D25, 2026-08-01). A follow-on session fixed the version-less substring fallback that was the whole of the 20% UNCOVERED false-fire: `fingerprint_match` is now reserved for a structured `meta.fingerprint match`, and an unstructured product-name hit is a distinct, lower-ranked `fallback_match` that never claims the exact stack (word-boundary matched). The decision was measurement-led — the fallback was shown to contribute **0 of the 28 covered fires** before it was demoted, so **UNCOVERED false-fire went 20% → 0% with covered hit rate, precision, near-miss and self-match all unchanged**. `test_fingerprint_fallback.py` locks it; the suite is now **55 files, all green**. That closes the last known correctness gap in the 2.7 path. Separately, the enrichment series' last batch (batch 3 — PDF/pages/CPENT) has since **run and closed the series** — see the batch-3 Status paragraph below and the *Batch 3 — the close-out* record.

KB enrichment batch 3 (PDFs, pages, CPENT) — complete, and it closes the series. The last batch ran its five sources — two redistributed cert-notes PDFs (OSCP, eCPPT), a 185-page PNPT GitBook space (mis-scoped in its own prompt as a "single web page"; its sitemap resolved to 185 pages, so it went through `fetch_gitbook.py`), an OSCP-notes blog post, and the CPENT cheat-sheet repo — and added **2 entries**, both distilled (D21) from the CPENT repo's two thinnest modules: `authored-perimeter-filter-mapping` [recon] (reading a packet filter as an object of enumeration — ACK/window scans, firewalking, scan-time evasion with the decoy caveat) and `authored-ot-safe-scanning` [ics] (scanning an ICS/OT segment without faulting a controller — never `-sV` / `-A`, `-sT` over `-sS`, `--max-parallelism 1`, abort path in the ROE). Both were **genuine gaps confirmed by retrieval**, not by a token count: `--spoof-mac` / `--data-length` / `firewalk` / `--mtu` / `nmap -f` and `--max-parallelism` were all 0 hits pre-ingest, and the `ics` category held one entry that was an operator persona. The four cert-*notes* sources — the two PDFs, the blog, the 185-page PNPT space — yielded **zero**, every candidate already covered (printer pass-back, IKE aggressive mode, open mail relay, pivoting, spraying, employee OSINT), several better than the source. **The prediction held**: batch 3 was five more cert-note sources, predicted ~0–1 entries before the run; it yielded 2, close

enough that the lesson stands. The final scoreboard: **44 entries across the whole series; cert notes 36 sources → 4 entries vs narrative writeups 106 posts → 27** — the ~20× per-source difference that is the series' reusable lesson.

The thin categories the series set out to fill still barely moved (ics 1→2 , recon 69→70 , pivoting 6→8 ; services / credentials / persistence / phishing / iot unchanged), and this record says so. KB 2,741 → 2,743; suite **55 files, all green** (test_corpora flaked on the 22 MB rewrite and passed on re-run, as documented). Records carried forward: the token diff **nominates but never confirms** (its top-two batch-3 nominations were both already covered); a prompt's description of a source's **scope is a claim to verify** (PNPT was 185 pages, not one); and **Defender deleted a file mid-run a fourth time** (bypassing-amsi.md , errno 22 while still on disk) — now a standing repo hazard, not an anecdote. sources/pdfs-pages-manifest.md is the reproducibility record; the raw trees are deleted.

Build #9 — the Windows/AD path, driven live against a real domain (2026-07-31)

Everything in HackPit's Windows/AD path was built and locked **hermetically**: test_winrm / test_winrm_safety monkeypatch the WinRM transport, and the AD graph, orchestrator and technique catalog were all exercised against adgraph/sample_data.py — a GOAD-shaped synthetic domain. Nothing had ever touched a real domain.

Build #9 stands up the operator's own lab (VMware Windows Server 2022, promoted to a corp.local forest, WinRM over HTTP:5985, Administrator@CORP) and drives the **live** backend over its real HTTP API against it. It adds no new capability and no autonomy: every live command clears the existing gates and the human approves each one; the orchestrator proposes, it never fires. The evidence is backend/livefire.log , produced by backend/livefire_windows.py — a by-hand harness that is deliberately **not** part of the hermetic suite (which stays runnable with no VM, no network and no pywinrm). It records every check as PASS / FAIL / NOT-RUN, and a check that could not run is NOT-RUN, never a silent pass — the whole point being that "demonstrated live" and "still needs X" are different claims. **45 passed, 0 failed, 3 not-run.**

Task 1 — the WinRM round-trip and the gates on real input (14/14). A real PowerShell command runs on the DC over WinRM and comes back with the box's own identity (WIN-990RALNGERV , corp\administrator , corp.local) — the proof the host-lock held, since the request carries no host field and the destination is resolved server-side from the profile id. An unapproved command is refused live at the approval gate; a foreign host named in the args cannot redirect the run (it is echoed by the DC, which is where the run still landed); and build #5's whole-script danger classifier fires **on real input** against Write-Host go ; Invoke-Mimikatz — the exact argv[0]-vs-whole-script case that motivated it — then clears once the human acks. The credential appears in none of the run record, the event stream, or the recorded command line; the run is audited and retrievable by id; and real WinRM stdout ingests into state (an OSCP-style proof flag attributed to the profile host). This is the live half of the WinRM deferral, closed.

Task 2 — the AD graph off real collection (10/10), not sample_data . The gated, approved, scope-locked bloodhound-python collects the **real** corp.local over the engagement executor — 84 nodes, 625 edges, the live DC's own computer object present — parses into the typed graph, and the route to Domain Admin is computed on that real graph. The collector preview redacts the credential; an unapproved collection is refused live at the approval gate; and an **off-scope DC is refused live at the target gate**. This closes both "AD off synthetic data" and "touch a real AD domain."

Task 3 — the live abuse walk (12/12, 2 not-run). The orchestrator proposes edges off the real graph; the human approves; the executor runs them. A real DCSync (impacket-secretsdump -just-dc) executes against the domain and returns real credential material — four NTLM hash lines including krbtgt — after the danger red-confirm is demanded and supplied; the dumped credentials ingest into state (administrator , krbtgt , win-990ralngerv\$); and the walk advances **only** on that approved, exit-0 run (advancing a commanded edge with no run is refused). A

native Windows AD action (`Get-ADGroupMember 'Domain Admins'`) runs on the DC over WinRM. Two edges are recorded NOT-RUN, honestly: the model happened to pick `GenericWrite` over `DCSync` on the run, so the destructive step was operator-directed rather than model-selected; and the native `Invoke-Mimikatz` `DCSync` variant cannot run because a freshly promoted Server 2022 carries no offensive tooling, and staging some onto the DC was out of scope — the Linux/impacket variant is what executed.

Four defects the live run found that the hermetic suite could not — all fixed, each now regression-locked. Every one was a claim that two halves of the codebase agreed on in a test but disagreed on against a real domain:

- **The domain credential leaked into the persisted record.** The WinRM path never puts a secret on a command line — the profile resolves it server-side — but every credentialed Linux tool (`bloodhound-python -p ...` , `impacket-secretsdump DOM/user:pass@host`) carries it as an argv token, which was stored verbatim in the run record, and from there into rendered reports **and the LLM proposer context**. So collecting a domain shipped the DA password to the model. `cockpit/secretargs.py` now redacts credential *values* out of the argv at the record boundary (per-tool, never a blanket flag list — `-p` is a password to netexec and a *port list* to nmap), wired into both `RunRecord` constructions. A first cut split the impacket positional on the *first* `@` and leaked a password fragment out the "host" side; the fix anchors the host to the last `@`. `test_secretargs.py` locks both the positive cases and the negative controls (nmap ports survive untouched), including a password-with-`@` fragment check.
- **Real DCSync loot ingested nothing.** The AD technique catalog proposes Kali's `impacket-secretsdump`; the state parser registry was keyed only on upstream's `secretsdump / secretsdump.py`. Two halves disagreeing about one tool's name meant four dumped hashes ingested as zero credentials. `state/ingest.py::program_name` now strips the `impacket-` prefix; `test_state.py` locks that every spelling reaches the same parser (and that `impacket-wmiexec` does *not* become a `secretsdump` parser).
- **The KB grounder armed a "free" edge with a destructive command.** `MemberOf` and `HasSIDHistory` are inherited rights — nothing to run — but their KB seeds still matched a real ACL-abuse entry, whose example commands were adopted wholesale. On the real graph the orchestrator therefore proposed a runnable `net rpc password` (a destructive reset) for `ADMINISTRATOR -MemberOf-> DOMAIN ADMINs`, an edge with no action, rendered `resolution="ready"`. The gates all held — nothing fired — so this is a correctness lock, not a containment one, but a panel that asks a human to approve a domain change for a free step is how a hurried operator makes an unintended one. A `no_command` flag on the spec keeps the KB *citation* while refusing to let it manufacture an *action*; `test_adorch_safety.py` locks it with an adversarial grounder that always offers a destructive command, plus a positive control that an actionable edge is still groundable.
- **The collector's failure classifier had nothing to say about the FQDN error.** `bloodhound-python` refuses an IP for `-dc` ("looks like an IP address, but requires a hostname (FQDN)"), and that same message also mentions a "DNS server IP", so the generic DNS signature matched it and told the operator to fix their nameserver when `-dc` was the real problem. A signature for it now runs **first**; `test_adgraph_collector.py` locks the ordering.

Task 4 — Windows-target C2 (feasibility, measured; not run). Build #7 demonstrated Sliver + iodine against Linux. Completing that against this Windows target needs a callback path from the DC to the listener, and the harness **measures** rather than assumes one: the DC reaches the host's VMnet8 gateway and has DNS egress, but the Sliver listener lives in `hackpit-engage-sandbox`, which publishes no ports and sits on a Docker bridge the DC has no route to. Recorded NOT-RUN with the exact unblock (publish the listener port from the sandbox to the host, point the implant at the gateway) and the reason it was not done here: publishing a port changes the sandbox's network posture, and this build's standing invariant is that it adds no new exposure.

Task 5 — the detection footprint, truth-checked against real telemetry (8/8) — the purple-team validation the describe-side had never had. The catalog's footprint claims were written from documentation; this reads the DC's own Event Log back after techniques that really ran. The audit policy is read first (so "not observed" is

distinguished from "never audited"), and events are counted as **before/after deltas keyed on EventRecordID** — the VM's clock drifts and w32time hauls it back, so a timestamp window silently lands in the future and reads as "the box logged nothing" (an earlier iteration of the harness "disproved" claims that were in fact correct, exactly this way; the log also flushes asynchronously, so the query settles on the DC first). Confirmed against real telemetry: DCSync raises Security **4662 with the DS-Replication GUIDs** — 12 of 12 events carried them — under the DC's **default** audit policy, so "loud" is the right rating; and T1003.006 is the right id. **Two catalog claims the real box contradicted, now corrected (atomic with attck.py, per the detection trap)**: (1) the `secretsdump` telemetry list lumped SAM/LSA-mode artefacts (5145 admin-share access, 7045 short-lived service, hive saves) together with the `-just-dc` path — a live `-just-dc` run added **zero** of them, because it replicates over DRSUAPI and touches no share, service or hive; the entry now says which mode each artefact belongs to. (2) the `ad_collect` entry promised "Security 4662 ... if object auditing is enabled" — but on a DC that auditing is on by default, and a full `-c ALL` collection added **zero** 4662 (4662 needs a SACL on the objects read; LDAP enumeration produces network logons, not directory-access events); the entry now points at LDAP diagnostics (1644) as the dependable host signal, while the "loud" rating stands on query volume.

`test_detection / test_detection_safety` stay green after both edits.

The honest bottom line: the WinRM round-trip, real AD collection, the real DA path and the executed DCSync abuse are **demonstrated live**, with the gates holding on real input; a full Windows-target C2 callback and the native-tooling abuse variants still need more lab infrastructure (a routable listener, offensive tooling staged on the DC), and are recorded as not-run, not as passed.

Build #10 — the opt-in C2 exposure, and an honest NOT-RUN (2026-07-31)

Build #9 closed the Windows/AD path live but left one item explicitly NOT-RUN: a full Windows-target C2 callback. The reason was precise and worth not papering over — the Sliver/tunnel listener runs inside `hackpit-engage-sandbox`, which by standing invariant publishes no ports, so the lab DC (on VMware's NAT subnet) had no route to it. Closing that needs a change to the lab's exposure posture, and build #9's rule was "no new exposure by default." Build #10 supplies the missing route as an **opt-in** and does the surrounding engineering; it does **not** fire the offensive proofs, and says so plainly.

The exposure override — built, locked, and demonstrated live. `docker/proof/c2-lab.yml` is the one place HackPit publishes a host port, and it is never applied by the documented bring-up — reaching it takes a second, explicit `-f docker/proof/c2-lab.yml`. It publishes exactly one port, `192.168.13.1:53/udp`: the iodine tunnel server's DNS port, bound to the single host interface the lab DC can see (the VMnet8 address), never a wildcard that would also expose it to whatever network the laptop is on. `backend/test_exposure_safety.py` locks three invariants hermetically — (1) the default `docker/docker-compose.yml` publishes zero ports; (2) if the override exists, every published port names a literal host IP and never `0.0.0.0 / :: / *`; (3) nothing else in the repo composes the override in implicitly — and it plants a wildcard-bound port and proves the scan catches it, so a silently-broken scanner cannot make the checks vacuous. The mechanism is not only tested but **demonstrated live**: the exposure came up on `192.168.13.1:53/udp` and tore down cleanly, leaving no published port behind.

The proof harness — the same honesty discipline as build #7's live fire. On top of the override sits a gated C2/AD proof harness: four scripts (`c2_01_dc_prereq_stage.sh` staging the tunnel client + TAP driver on the DC, `c2_02_iodine_tunnel.sh` the DNS tunnel, `c2_03_sliver_beacon.sh` a beacon over it, and `c2_04_dcsync_defender_excl.sh` a native DCSync behind a **scoped Defender exclusion** — `Add-MpPreference -ExclusionPath` for the tool path, removed in a `trap` guard, *not* a real-time-protection toggle), plus an orchestrator (`c2_lab_proof.sh`, the only script permitted to compose the override up and down) and an in-process driver that reuses build #9's credential redaction. The load-bearing property carried over from build #4/#7:

offensive command strings are never inlined in the harness. Each is an operator-supplied paste value handed to the driver by file path; an empty or unfilled value makes the proof report **NOT-RUN**, never a fake pass. Every offensive step stays behind the existing engagement + approve-each + danger red-confirm gates; no new autonomous path is added.

What is NOT-RUN, and why it is recorded as such. The four offensive proofs were **not fired**. They are harness-ready but need operator-supplied commands (the iodine client invocation, the Sliver implant generation and launch, the DCSync one-liner), and those were deliberately not filled — so the run reports NOT-RUN across all four, which is the honest state, not a failure. "The harness is built and safe" and "the offensive path was demonstrated end to end" are different claims, and only the first is made here.

The safety harness, investigated — and one of this build's own claims withdrawn. Wiring the new exposure test in raised a suspicion that the runner itself was not gating: `backend/run_safety_tests.sh` appeared to invoke every test file without checking any exit code, which would mean a failing safety test could let the suite exit 0 and print "passed." **That claim was wrong, and it is worth recording as wrong rather than quietly dropping.** `set -e` is present in the committed runner and has been since before this build; planting a deliberate failure and running the suite produced exit 1, a halt at the second test file, and no "passed" banner. The suite always gated. What is a real weakness is that `set -e` is silently suspended for any command inside a pipeline, an `if / while` condition, or an `&& / ||` chain — so a later edit as innocent as `"$PY" test_x.py | tee log` would disarm the whole suite with no visible change in output. The runner now routes every test through a `run_test()` helper that checks the exit code explicitly, names the failing file and the invariants it guards, states how many files had passed before the abort, and reports the file count on success. A guard that can be switched off by accident is not one a safety suite should rest on, even when it happens to be working.

The nc danger-gate failure: an environment leak, not an over-block. `test_session.py`'s `test_start_trips_the_danger_heuristic` failed on unmodified code with "nc must start once explicitly acked" — `validate_start` refusing `nc -lvnp 4444` even *with* `dangerous_ack`. The danger/target/scope logic turned out to be correct and the engagement/Wall-A state clean (the request resolved to `lab` mode, no active engagement). The real cause is the LAB gate order — target → approval → danger → **isolation** — where that last gate calls the live `assert_isolation_proven()` Docker probe. The *un-acked* half of each assertion pair short-circuits at `danger` and never reaches it; the *acked* half falls all the way through, so on a host with the stack down it came back `gate="sandbox"` and the test failed for a reason with nothing to do with the heuristic it guards. It was the only test in the file that called the gates outside the `_Spy()` fixture that stubs that probe — contradicting the module's own "Hermetic: ... no Docker daemon" contract. Fixed by running it inside the fixture like every other test in the file, with the gate order written into the docstring so the next person does not re-learn it from a red suite.

Harness honesty, locked — and a real defect caught doing it. `backend/test_proof_honesty.py` locks the property the whole NOT-RUN story rests on: an unfilled offensive slot can never become a fake pass. It checks four independent angles — the reader treats empty/whitespace/comment-only as unfilled; driving the real driver subcommands with an empty slot emits NOTRUN, emits no PASS, and never reaches the WinRM transport (asserted with a tripwire, not assumed); an AST pass proves every `_read_paste` call site is followed by a NOTRUN empty-guard, so a *new* slot added later without one fails the suite; and the slots in the four committed scripts really are unfilled — plus a control that feeds it a filled slot to prove the checks can fail. Writing it caught a genuine defect: `c2_01`'s two slots and `c2_04`'s DCSync slot did **not** ship empty. They opened with a self-referential shell fragment `( echo $(cat ...| grep ...| cut ...) )` left from an earlier placeholder sweep, sitting *outside* the first pair of quotes. That text is not a comment, so `_read_paste` returned it as live content, the empty-guard never fired, and the driver would have sent it to the domain controller as PowerShell and scored its exit code as a real attempt — an unfilled proof reporting FAIL, or conceivably PASS, instead of NOT-RUN. The slots are now comment-only and the regression is locked by a test verified to fail when it is re-planted.

The DC prerequisite probe, finalised. `docker/proof/dc_prereq_probe.py` answers whether the C2 path even has the pieces it needs — adapter inventory, staged tooling, DNS/HTTPS egress, route to the VMnet8 gateway, Defender posture, OS/PowerShell version — before anyone runs a proof. It is read-only and *structurally* so: every probe is checked against a read-only cmdlet allow-list that **fails closed**, so a probe added later with a mutating verb makes the file refuse instead of run. The stale hardcoded profile id was removed (`--profile` is now required, so it can never fire at whatever host an old id happens to name) and a `--list` mode prints the probes while contacting nothing. It was **not** run against the DC in this session.

Suite state. `sh backend/run_safety_tests.sh` exits 0 with **538 checks across 50 test files** — now genuinely gated, verified by planting a failure and confirming the suite goes non-zero, names the offending test, and prints no "passed" banner.

Build #11 — the launcher, and operator identity (2026-07-31)

Two small builds, both closing gaps that had been visible for a while and neither moving the safety boundary.

A dozen surfaces were built and then invisible

The top nav carries five product sections and deliberately stays that size. Everything else — `:arsenal`, `:c2`, `:tunnels`, `:windows`, `:repeater`, `:scripts`, `:code-scan`, the AD graph — was reachable only by typing the URL. This was not a new discovery: the project's own working notes had flagged it twice (2026-07-26 and 2026-07-27) without it being fixed.

The launcher is 15 tiles in four bands — **plan, operate, infrastructure, reference** — merged into the home page per D19. The band label is not decoration: it states the posture of everything under it (*"proposes · never executes"*, *"every command human-approved"*, *"gated start · human-only stdin"*), so the page cannot show a shell tile without also saying who approves it.

`CategoryGrid` is now categories only. It had been carrying **five** product cards — attack-paths, Cockpit, the scripts arsenal, the tool arsenal and code scan — every one of which is now a launcher tile; keeping them would have shown each surface twice on one screen. `ScriptsCard`, `FEATURED` and `COCKPIT_FEATURE` are left in the tree unreferenced rather than deleted, so restoring the old layout is a revert of one file.

The status rail answers a question the UI could not

Every execution surface shells out to `docker exec` and refuses with *"bring the stack up"* when the container is down. That state was only discoverable by clicking into a surface and reading an error. `GET /home-summary` now reports it up front — stack, LLM model, WinRM target, active engagement — and the tiles that need the stack dim when it is down.

Two details that are the point rather than polish. A **null** probe renders as *"unknown"*, never as *"down"*: an undeterminable probe must not send the operator chasing a container problem they do not have. And the endpoint is deliberately **separate from** `/stats`, because it runs docker probes and folding it in would gate the hero's counters behind a container inspect.

It lives in `main.py` because it is cross-cutting by construction — sandbox probe, LLM config, Windows profile store, engagement store, KB counters — and `cockpit` and `arsenal` may not reference each other.

The report could not identify its own author

`backend/report.py` had **no author, candidate or OSID field anywhere** — every "author" match in it was the word "*authoritative*". An OSCP submission is not attributable to a candidate without a name and an OSID, so the report this tool generates could not be submitted without hand-editing identification in. For a project whose tagline is "*pass the cert*", that is a defect, not a missing nicety.

Fixed per D20, with the block spliced under the report title the same way evidence and the CVSS block are.

A mistake worth recording

The identity module was first written as `backend/operator.py`. `backend/` is first on `sys.path`, so that filename **shadows the stdlib `operator` module for every import in the process**. It was caught immediately and renamed to `operator_identity.py`; `test_operator.py` now asserts `backend/operator.py` does not exist and that `import operator` still resolves outside the repo. It is recorded here because the failure mode is silent, process-wide, and would have been very hard to attribute later.

What is locked

`test_home_summary.py` and `test_operator.py`, both in the suite. Between them: no secret reaches the browser (checked against the **real** stores — stored WinRM secrets, the configured LLM key, every credential in state — and reporting per-source counts so an empty leg is *visible* rather than absorbed into a reassuring total); the endpoints call only the masked accessors, asserted over the **AST** rather than a substring; a status endpoint cannot reach an execution path; the operator config is gitignored, asserted by asking git; and the OSID and email never reach the browser. Every check carries a positive control, and the leak test asserts it cannot go vacuous.

Suite state. `sh backend/run_safety_tests.sh` exits 0 across **52 test files** (50 → 51 → 52). Frontend lint holds at the pre-existing baseline of 11 errors + 1 warning; new tiles stagger via CSS `animation-delay` specifically so as not to add more `react-hooks/set-state-in-effect` errors, and `next build` exits 0.

Build #11 (the launcher + operator identity) — complete. Neither half moved the safety boundary; both closed a gap that had been visible and unaddressed. The launcher's value showed up immediately and by accident: bringing the Docker stack up during verification flipped the rail from `down` to `up` · `engage up` and un-dimmed seven tiles with no code change, which is the whole feature demonstrated end to end. The report identity block closes a real submission blocker — the generated OSCP report previously carried no candidate identification at all. One near-miss is recorded in full above (a module named `operator.py` shadowing the stdlib) because the failure would have been silent and process-wide.

KB enrichment — the first external corpus, and what it cost to read it (2026-07-31)

A 687,830-character corpus of live bug-bounty stream transcripts was folded into the KB. It is recorded here because the *ratio* is the finding, and it sets the expectation for every batch that follows.

11% of the file was exact duplication — two sections were byte-identical copies of two others, caught by hashing before any reading effort was spent on them. **A further 44% (nine Hindi-language sections, 268,048 characters) yielded exactly one usable technique.** That was established by counting technique vocabulary on the Devanagari transliterations rather than by sampling: across all 268k characters the corpus contained 57 mentions of price tampering, 15 of SQLi, 2 of CSRF — and **zero** mentions of XSS, IDOR, open redirect, subdomain takeover, rate limiting, brute force, path traversal, Burp, nuclei or dorking. The remainder was stream management and narration.

So roughly a third of the corpus carried essentially all of its value. The lesson is not that the source was bad — the English third was excellent — but that **volume is not a proxy for content, and measuring density before reading is cheap.**

Thirteen entries were written, each checked against all 2,699 existing rows first. They cluster in three places the KB was genuinely empty:

- **Targeting** — hunting the most recently *added* wildcard (platforms publish scope changes as a dated diff, and a wildcard added weeks ago cannot have been fully tested); hunting the asset with the *fewest* resolved reports rather than the most; `origin=*` hosts that serve the same application with the edge WAF removed; and the cheap staleness signals that decide which of several hundred hosts is worth attention.
- **Method** — **STRIDE threat modelling**, which closes the gap between "recon is finished" and "what do I actually test": decompose into mechanisms, walk each through Spoofing / Tampering / Repudiation / Information disclosure / Denial of service / Elevation of privilege, and emit a ranked attack-vector list. The KB previously held two passing mentions of threat modelling and no method at all, which is a structural gap rather than a content one — HackPit's planner had no model of this phase.
- **Web** — repudiation and audit-trail attacks (the most overlooked STRIDE class: a role able to edit the logs it is audited by, with impact argued through the compliance regime); brute forcing through the change-password endpoint, where login is rate-limited and the change-password path usually is not; CSRF token *binding* failures rather than missing tokens; client-trusted entitlement flags, with the discipline that a UI which unhides is not an authorization bypass until the server honours it; and letting the framework fingerprint decide the test set — verb tampering is futile against Express and worth doing against Django/PHP.

`pipeline/authored/authored_entries.jsonl` is the one committable artifact; the raw corpus lives under the gitignored `sources/` tree and never enters this public repository.

KB enrichment is now a measured process rather than an additive one (2026-07-31). The governing rule is D21: distil, never parrot; check every candidate against the whole KB before writing; report zero when a source yields zero. The first batch under that rule took a corpus that was one-third signal and produced 13 entries aimed squarely at the categories that were thinnest, while deliberately adding nothing to the ones already carrying hundreds of rows. The next batch — 23 cert-study and CTF-writeup repositories — is scoped in `PROMPT-kb-repo-ingest.md` and is expected to yield most of its value in `pivoting`, `credentials`, `persistence` and `iot`, with several repositories predicted in advance to yield nothing at all.

KB enrichment — 24 cert/CTF/pentest repos, and a prediction that was wrong (2026-08-01)

The batch scoped above ran. **24 repositories were cloned, 22 yielded nothing, one produced the batch's only two KB entries, and one contributed to them without earning an entry of its own.** That is the headline, and under D21 it is a correct outcome rather than a disappointing one.

The prediction the batch was scoped on was wrong, and the way it was wrong is the useful part. CPENT was expected to be the highest-yield group, because the KB is thinnest in exactly what CPENT claims to cover — IoT (2 entries), ICS, and pivoting. It covers none of it in practice. The study guide has only modules 05 and 06 written; modules 09 (wireless), 11 (IoT) and 12 (SCADA/ICS) are `[TBD]` stubs with no content behind them, and the companion repo is a six-file link list. Widening the check from that repo to the whole batch settled it: across **all 24 repositories**, exactly one file mentions ICS vocabulary (`modbus`, `s7comm`, `dnp3`, `scada`, `plc`, `profinet`, `HMI`, `Purdue`) four or more times, and that file is a CTF challenge *named* "Scada" which turns out to be Jinja2 SSTI. There is no

ICS, no IoT-hardware and no firmware/UART/JTAG material in this batch at all. `iot` (2) and `phishing` (1) are still the KB's thinnest categories, and they need a different kind of source — vendor and ICS-CERT advisories, teardown writeups, protocol specifications — not more exam notes.

Saturation was measured, not assumed. Beyond grepping each candidate technique against all 2,712 rows, every word in the batch was tokenised against the whole KB. **308 terms appear in three or more repository files and never once in the KB — and on inspection all 308 are URL slugs, lab hostnames, YouTube playlist IDs, filenames or typos.** Not one new tool name, not one new technique name. That is the quantitative form of "these repos are saturated," and it is a far stronger claim than reading a sample and forming an impression. Individually probed and confirmed already covered: DCSync (47 rows), mimikatz (90), LSASS dumping (34), password spraying (43, including lockout thresholds), Kerberoasting, chisel (19), proxychains (22), sshuttle (7), mitm6/DHCPv6/WPAD relay, IPsec/IKE, VoIP/SIP, Jinja2 SSTI, and the SEH / bad-character / `jmp esp` stack-overflow workflow.

Where the two entries came from. One repository — `eCPPTv2-PTP-Notes`, with the TryHackMe Wreath notes from `PNPT-study-guide` as a second source — carries a worked three-network, three-hop pivot lab. The KB's six existing `pivoting` entries and its pivot cheat sheet are all *command references*: here is `ssh -L` syntax, here is chisel syntax. None of them teaches the discipline that keeps a chain standing, which is a different kind of knowledge and the one thing in the batch the KB genuinely lacked. Two entries were written to fill it (`pivoting` 6 → 8, KB 2,712 → 2,714):

- **Multi-hop pivot chains.** The asymmetry every naive chain dies on: hop 1 can reach you, hop 2 cannot, so each hop needs an outbound path *and* a return path built by separate commands — a SOCKS proxy per hop on its own port, and a socat relay chain carrying callbacks back. Plus the parts that are only obvious in hindsight: stage tooling from hop N-1 rather than from your box (hop 2 cannot fetch from you), stage *static* binaries (a pivot host often has no interpreter and no package manager), verify each hop before adding the next, scope the subnet before routing into it, and tear down in reverse while recording what you left behind.
- **Choosing the pivot primitive.** Three properties decide it and none of them is preference: direction (egress decides forward vs reverse, not taste), shape (one port vs SOCKS vs a layer-3 TUN), and footprint. Including the SOCKS tax that costs people hours — a SOCKS proxy relays completed TCP only, so `-ss` and ICMP host discovery return a silence that reads exactly like "host is down" — and the quiet relay that moves both listeners onto the attacker's box so the compromised host has **nothing listening on it at all**.

The arsenal gap this exposed was not in the plan, and was the more valuable find. HackPit's `cockpit/tunnels.py` has driven chisel, ligolo and proxychains since Phase 4; all five pivot binaries were confirmed present and runnable in the live sandbox image — and the 110-tool catalog contained **zero** pivoting tools. Five rows were added in a new `pivoting` category (chisel, ligolo-ng, socat, sshuttle, proxychains), 110 → 115.

Cataloguing proxychains then surfaced a real gate hole that predated it. The danger heuristic classified `argv[0]` only. Bare `weeveily` produced a red-confirm reason; `proxychains -q weeveily ...` produced **none** — and `wrap_command` builds precisely that `argv` for the human to approve. So routing a command through a tunnel *stripped its red-confirm*: the deeper into a network you went, the weaker the gate got, which is exactly backwards. The fix is in the predicate, not in the file set — wrappers are peeled off and the command that actually runs is classified, with the wrapper kept visible in the reason (`through proxychains: weeveily: turns a vulnerability into ...`). It is regression-locked by `test_a_wrapper_cannot_launders_the_red_confirm`, which asserts through `validate_request` rather than through the predicate, because the claim is about the *gate*; it carries the `-f <config>` case (skipping the flag but not its value would read the config path as the command) and the negative half, that a benign inner command must **not** raise a confirm.

`sshuttle` was classified into the same bucket as `chisel` and `ligolo` on what it does — it puts a whole subnet inside reach — rather than on how quiet it is; needing no agent and no listener makes it the stealthiest of the three, not the least dangerous.

One gitignore exception was taken, deliberately. `/sources/` was ignored wholesale, which meant `sources/repos-manifest.md` — the record of which commit of which repository produced which entry — could never be committed and would vanish with the clones. The pattern is now `/sources/*` with a single negation for that one file; it holds our own prose about third-party repositories and never their content, and `git status` was checked afterwards to confirm it is the only thing in the tree git can see.

Also worth recording, because it was on disk and is not any more. `ciwen3/PNPT` (307 MB, the largest clone in the batch) contains a `conti/` directory holding leaked Conti ransomware operator manuals in Russian, `rcclone.exe`, and a Cobalt Strike 4.3 archive. Nothing from it was read into an entry and nothing from it could ever have been committed — `sources/` is gitignored — and all 1.6 GB of clones were deleted after verification, as the process requires. It is noted here because "we cloned 24 arbitrary repositories" has a supply-chain shape worth being explicit about.

`0xdf.gitlab.io` is deferred with a concrete proposal, not skipped. It is a website, not a repository, and was correctly excluded from this batch. It is also the highest-quality HTB writeup source in existence, and the KB already holds 41 `htb-writeups` + 173 `writeup` entries, so the first step is not fetching — it is **overlap measurement**: pull the sitemap, diff the box names against those 214 rows, and establish how many boxes are genuinely absent before writing an ingester. If the answer is meaningful, the shape is the existing sitemap-driven PortSwigger ingester (enumerate → fetch → distil → authored entries), never a crawler. Recorded as a follow-up.

KB enrichment — 7 GitBook certification spaces, and a zero that took three probes to earn (2026-08-01)

Seven GitBook spaces of personal certification notes were supplied — two eCPPT, three CRTP, one OSCP, one pentesting checklist. **All seven yielded zero entries, five were never fetched, and the KB is byte-unchanged at 2,714 rows.** Under D21 that is the correct outcome and not a failed run: certification notes are the single most duplicated material in this KB, and the batch was scoped expecting it.

What is new here is the gate, not the yield. Batch 2 fetched 24 repositories in full and *then* discovered 22 were duplicates. This batch inverted that: the page lists were pulled from each space's own sitemap and tokenised against all 2,714 KB rows **before a single content page was requested**. Across the six reachable spaces that is 545 published pages, and the slug words absent from the KB number six, five, three, one, two and two respectively — `fundamentals`, `accross`, `uncostrained`, `gnereral`, plus author names and the acronyms `crtp` and `ecppt`. Four spaces were closed on that evidence and never fetched. The cost of settling three CRTP spaces was three sitemap requests.

One space was resolved by path, not by inference. `dev-angelist.gitbook.io/ecpptv2-ptp-notes` was suspected of being the GitBook publication of `github.com/dev-angelist/eCPPTv2-PTP-Notes`, mined in batch 2 at commit `a543e9167445`. Suspicion is not evidence, so it was checked: the space publishes `network-security/2.4-1/2.2-pivoting` and `.../2.2-pivoting-1`, byte-identical to the repo paths `repos-manifest.md` records as the source of both authored pivoting entries. Same content, same author, two surfaces. Not fetched.

The two spaces batch 2 named as genuine unknowns both resolved against the prediction.

- `ecpptv3-ptp-notes` was the one expected to pay. 75 of its 90 distinct page slugs are shared with the same author's already-mined v2 space, so only the 15-page delta was fetched — and the delta is the reason it yields

nothing. It is an Active Directory chapter that has been outlined but not written: **1,058 words across 14 pages, zero code blocks**, median 23 words per page, with `6.1.4-ad-enumeration` at 25 words and `6.1.7-ad-persistence` at 24. Its one substantial page is a conceptual introduction to users, groups and OUs, against a KB holding 123 `active-directory` entries. eCPPTv3 is eCPPTv2 plus a stub.

- **pentesting-checklist** was fetched in full on purpose, 113 pages / 20,248 words / 214 code blocks, because its hypothesis was methodology *structure* rather than technique novelty and no index gate can test shape. Three probes, all negative. Word novelty: 10 words appear three or more times and never in the KB, eight of them the author's name and lab hostnames — batch 2's exact pattern. Tool novelty: 175 distinct leading commands, 10 unknown to the KB, all 10 resolving to covered material (PowerView cmdlets against 47 entries, `rpcclient` subcommands against 23, `vshadow.exe` as an alternate shadow-copy binary against `diskshadow` and `tool-sebackupprivilege`, a Linux capability string against `ht-linux-capabilities`). And the structure hypothesis itself failed: **the KB already carries 118 checklist-* entries** — 51 Active Directory, 21 privesc, 21 recon, 19 enumeration, 4 credentials, 2 persistence — against a space that is 74 Windows/AD pages, 17 Linux and 5 pivoting, the latter naming the same primitives batch 2 catalogued a day earlier.

The finding worth carrying forward is a correction to batch 2's own method. The token diff nominated `sesshutdownprivilege` — 11 occurrences in the checklist, zero anywhere in 2,714 KB rows, and unlike the slugs and hostnames it is unmistakably a real privilege-escalation technique. It was the batch's one credible entry, and it is already in the KB. `ex-windows-checking-services` step 4 reads "*Stop the service, check its start mode, reboot if you hold SeShutdown*" and carries `shutdown /r /t 0`; `ht-service-triggers` covers the harder form of the same hinge, starting a service without `SERVICE_START` rights by firing its trigger rather than rebooting the host. The KB spells it without the `Privilege` suffix, and that one suffix was the entire "gap". So the rule batch 2 established is now bounded: **the token diff nominates candidates and can never confirm a gap**. A zero it produces is only trustworthy once the *concept* has been grepped, and an entry it motivates is only safe on the same condition. The second survivor, `setakeownership`, failed the same way against a dedicated `oscp-setakeownershipprivilege` entry.

One space was refused rather than skipped, and the refusal is in code. `mqt.gitbook.io` serves a `robots.txt` carrying `Content-Signal: ai-train=no` and a blanket `Disallow: /` for ClaudeBot, GPTBot, CCBot, Google-Extended, Applebot-Extended, Bytespider, Amazonbot and meta-externalagent. `pipeline/fetch_gitbook.py` parses that before its first request and skips the space; its sitemap was never requested. We are none of those user-agents, and the check is not one we could be compelled by — which is exactly why it belongs in the fetcher rather than in a person's judgement at fetch time. The space was independently a likely zero: all four OSCP repos in batch 2 yielded nothing.

`pipeline/fetch_gitbook.py` follows `fetch_portswigger.py` and improves on it in one respect. Sitemap-driven with the GitBook `sitemap.xml` → `sitemap-pages.xml` indirection resolved per space rather than assumed, serialised with a delay, honest User-Agent, fetch-once-to-disk. The improvement is that it does not scrape HTML at all by default: GitBook serves every page as markdown at `<url>.md` and advertises it in-page beside an `llms.txt` index, so the fetcher asks for the format the publisher chose to hand to machines — a third of the bytes, real fenced code blocks, no nav chrome. The `<main>`-to-markdown parser remains as the fallback. One sharp edge is handled explicitly: a missing page is served as HTTP 200 with a `# Page Not Found` body, so without a check a renamed page lands on disk as a stub and gets triaged as though it were content.

An operational note, because it is the third time this has bitten. Windows Defender deleted two fetched pages mid-triage — `enumeration/ports.md` (8.4 KB of port-enumeration commands) moved from `OSError 22` to `OSError 2` between two reads, and an LFI/RFI page followed. The same signature-on-our-own-examples

behaviour has deleted `data/kb/entries.jsonl` before; it now reaches the fetched source tree as well. Triage was made tolerant of unreadable files rather than retried, 111 of 113 pages were analysed, and both lost pages sit in saturated categories (`network-services` 183, `web` 642).

The fetched tree is deleted. `sources/gitbooks-manifest.md` is what remains — URL, fetch date, pages published, pages taken, and the verdict for each of the seven, under the same single-file gitignore negation the repo manifest uses.

KB enrichment — the 0xdf fingerprint corpus, and the series close-out (2026-08-01)

This is the last batch in the enrichment series, and the first that was **not** cert notes. The target was `hackpit-distilled`, the 78-entry corpus that keys a service+version fingerprint to the technique that solves it and that build #8's 2.7 retrieval ranks ahead of generic token matches — 78 entries is thin, and every `nmap` result the cockpit parses is a fingerprint that either hits that corpus or falls back to fuzzy prose. `https://0xdf.gitlab.io` is several hundred long-form HTB/CTF machine writeups, each starting from a scan result and ending at root: a service→technique mapping in narrative form, which is exactly the shape a fingerprint corpus needs and the one shape 31 sources of cert notes never supplied.

The policy came first (D23). The ingester's docstring named 0xdf as never-to-draw-from; that rule was reviewed by the operator and reversed to the project-wide distil-not-parrot rule, and the docstring was rewritten to match. The `consolidate.py:2324` `link-index skip` was deliberately left alone. None of the 55 fetched pages nor any 0xdf prose is committed — the committable artifacts are `fetch_0xdf.py`, the policy fix, the distilled entries and the manifest.

Politeness, and why the index was cheap. `robots.txt` is a 404 — no restrictions published — which the fetcher checks on every run and would refuse on. `fetch_0xdf.py` takes URLs only from the site's own `sitemap.xml` and its `/tags/` page, never from crawling, and serialises requests behind a 1.8s delay (slower than the PortSwigger fetcher on purpose — this is one person's blog on GitLab Pages). **Phase 1 cost exactly two requests:** the sitemap (614 dated posts) and the tags page, which is the whole archive's metadata in one 2.8 MB document — 614 posts across 3,981 tags, and **276 distinct CVE tags of which 175 had never appeared anywhere in the KB.** That 175 is the number that separated this source from the cert notes before a single writeup was read: the 31 cert-note sources re-covered a syllabus the KB had already absorbed, and this one did not.

Phase 2 was selected by KB gap, not by recency. Each post was scored `3×network-service tags + 2×thin-category tags + 4×KB-absent-CVEs - web-app tags`, and the top 55 were fetched — 55/55, no failures. Windows Defender then did exactly what batch 2 warned it now does: it deleted **5 of 58 files mid-triage** (`OSError 22` → `OSError 2`), on the KB's own web-shell-signature behaviour reaching the fetched tree. Triage was written to tolerate a file vanishing between listing and reading — it counted 53 analysed and named the 5 lost rather than crashing — and a re-fetch restored them.

Yield: 19 entries from 12 posts, `hackpit-distilled` 78 → 97, KB 2,714 → 2,733. Every candidate was grepped by **concept** — synonyms, stems, tool names, technique words, not the token — against the whole KB before it was written, per batch 2's correction. That grep did real work: it killed ~43 of the 55 shortlisted posts whose standout CVE was already covered (runc fd-leak escape was present via `cred-toctou`; Azure AD Connect decrypt, fail2ban, vm2 escape, Office/RTF CVE-2017-0199, polkit CVE-2021-3560, rsync 873, and the Craft/Openfire/Jenkins-CLI CVEs were all already in), and it caught two would-be false positives that were pure slug collisions — a `SeShutdown` -

style trap each: the KB's only `schallenge` hit was `xSSChallengeWiki`, and its only `wsrep` hit was `wsrepl`, a WebSocket tool. One candidate (CrushFTP / htb-soulmate) was **dropped** because Defender deleted its page twice — the rule is that if you cannot read it, you do not understand it well enough to write it.

The 19 landed in the thin categories a fingerprint corpus is for, not the saturated web tree: OpenSMTPD MAIL-FROM injection, SaltStack master auth-bypass, Tomcat FileStore deserialization, Solr Velocity RCE, OFBiz XML-RPC deser, Cacti graph_realtime SQLi, CUPS ErrorLog read, printer 9100/PJL+ SNMP credential leak, MariaDB wsrep RCE, IPMI RAKP hash leak, GitLab SSRF→Redis RCE, Vim modeline RCE, Rails cache-store deserialization, XWiki SolrSearch Groovy RCE, Netdata ndsudo PATH privesc, Firejail `--join` privesc, ManageEngine SAML RCE, pgAdmin Query-Tool RCE, and the Squid CONNECT pivot. Retrieval was confirmed against realistic scan strings (`OpenSMTPD smtpd 6.6.1 25/tcp`, `Apache Tomcat 9.0.27 8080`, `jetdirect printer 9100 snmp`, ...): each new fingerprint ranks first or top-three for its own banner, and the existing controls (`vsftpd 2.3.4`, `Apache 2.4.49`) still win theirs, so nothing was displaced.

The series, whole (all batches)

Checking `git log` and the manifests in `sources/` for what actually landed rather than assuming an order, the enrichment series ran as six batches — batch 3 (this one) is the close-out:

Batch	Source	Sources	Entries	Note
Transcript corpus	live-hunting transcripts	1 corpus	13	set D21; first external corpus
Batch 1 — repos	24 cert/CTF/pentest repos	24	2	both from one repo; exposed D22
Batch 2 — GitBooks	7 GitBook cert spaces	7	0	index-gate; bounded the token diff
Batch 4 — 0xdf	0xdf machine writeups	55 posts	19	first non-cert source
Batch 4b — 0xdf pass 2	0xdf, the novel-CVE tail	51 posts	8	run after the retrieval fix; tail
Batch 3 — PDFs/pages/CPENT	PROMPT-pdf-and-pages.md	5	2	the close-out; ran 2026-08-01

Batch 3 has now run and closes the series. Its five sources — two redistributed cert-notes PDFs (OSCP, eCPPT), one OSCP-notes blog post, one 185-page PNPT GitBook, and the CPENT cheat-sheet repo — yielded **2 entries**, both distilled from the CPENT repo. The series added **44 entries total** (13 + 2 + 0 + 19 + 8 + 2).

Amendment — 0xdf pass 2 (2026-08-01). Batch 4 was written as the series close-out, but the retrieval eval it triggered found the fingerprint corpus half-firing, and the fix session that followed made new entries worth writing (`6c3ba42` : covered hit rate 70%→93%). So a second, final 0xdf pass ran against the remaining novel-CVE posts. It re-pulled the index (158 of 276 CVE tags still KB-absent, down from 175 as pass-1's 19 closed part of the gap), shortlisted 58 service- weighted posts, and — after Defender ate the same 7 pages on every fetch — concept-grepped 51. **Yield: 8 fingerprints** (HFS 2.3, IIS 6.0 WebDAV, Icinga Web 2, OpenTSDB, SQLPad, ES File Explorer, Next.js middleware bypass, Strapi), taking 0xdf to **27 fingerprints total** and `hackpit-distilled` to 105 (KB 2,741). The rate is a declining tail (19/55 → 8/51) and the remaining unfetched novel-CVE posts skew to web-CMS boxes the KB's 636 web entries already cover, so **0xdf is now recommended closed as a fingerprint source** — a pass 3 would likely yield <5, and only web-app entries. Two things measured rather than assumed: all 8 rows landed in `category="writeup"` (the ingester forces it), **not** the thin `services / pivoting` categories — correcting pass-1's report of where its rows went; and the eval's known **UNCOVERED false-fire residual held at 20% (3/15) with zero delta**, so the added substrings did not worsen the version-less fallback.

The stark, honest number is the cert-note yield: 36 sources produced 4 entries. Twenty-four repositories, seven GitBook spaces, two ebook PDFs, one blog post and the CPENT repo — the single most duplicated material in this KB — yielded four entries between them (two pivoting from batch 1, two scan-discipline from batch 3), each cluster from a single repo. That is not a failed series; it is a genuine measurement of the KB's maturity on exam-syllabus material, and it is why the series was worth running to its end: it told us, with evidence rather than assertion, that the KB is *saturated* on cert notes and *not* saturated on the service→technique fingerprint shape. The same effort spent on 0xdf — one source, narrative machine writeups — returned 27 entries against 4, because it supplied a shape nothing else in the KB did. The lesson for any future batch is to weigh the **source class**, not the source count: 36 derivative syllabus sources are worth less than one that maps scans to root.

Two records the series is required to carry forward, both confirmed:

- **The token-diff bounding (batch 2, methodological).** A token diff **nominates** candidates and can **never** confirm a gap. Batch 2's diff flagged `seshtutdownprivilege` — 11 hits, zero across 2,714 rows — and it was already in the KB as `SeShutdown`; one missing suffix was the whole "gap." This batch applied the rule and it paid twice more: `schallenge` → `XSSChallengeWiki` and `wsrep` → `wsrepl` were both slug collisions a token count would have called gaps. Grep the concept, never the token, before claiming a zero or writing an entry.
- **The declines and refusals.** `themastermindnotes.com/products/ecppt-study-notes-guide-unofficial` **declined** as a commercial product for sale — not fetched; if the operator supplies a purchased copy it goes through the PDF path. `mqt.gitbook.io` **refused itself** via `robots.txt` (`Content-Signal: ai-train=no` , blanket `Disallow: /` for ClaudeBot and eight others) and was correctly not fetched. And **0xdf was initially declined and that position was reversed** (D23) — the reversal is part of the story, not a footnote to it.

The fetched tree is deleted. `sources/0xdf-manifest.md` remains — the 55 URLs, the fetch date, and which 12 produced an entry — under the same single-file gitignore negation the repo and GitBook manifests use.

Batch 3 — PDFs, pages, and CPENT: the close-out (2026-08-01)

This is the close-out. `PROMPT-pdf-and-pages.md` ran against its five sources and the series is complete. Its two entries are **distilled** (D21) from the CPENT cheat sheet, written from scratch against a bare flag list — the entries are the mechanism and judgement, the source was the nomination:

- `authored-perimeter-filter-mapping` [recon] — reading a packet filter as an object of enumeration: the three-state model (`--reason`), ACK/window scans for ACL mapping, firewalking to locate the device, and scan-time evasion with the decoy caveat spelled out. **Gap confirmed by measurement:** `--spoof-mac` , `--data-length` , `firewalk` , `--mtu` , `nmap -f` and `window-scan` were all **0 hits** before ingest, and retrieval for "map firewall rules with nmap" returned *cloud* firewall enumeration.
- `authored-ot-safe-scanning` [ics] — scanning an ICS/OT segment without faulting a controller: never `-sV` / `-A` / `-O` , prefer `-sT` over `-sS` , `passive-first` , `--max-parallelism 1` , abort path in the ROE. **Gap confirmed:** `--max-parallelism` was **0 hits**, and the `ics` category held exactly one entry — an operator persona, not a technique.

The prediction held, and is worth recording because it survived its test. Batch 3 was five more cert-note sources; the prediction written down *before* the run was ~0–1 entries. It yielded 2, close enough that the lesson stands rather than needing revision — and the two it produced are the exception that proves the rule, since neither came from the cert *notes* (the PDFs, the blog, the 185-page PNPT space all yielded **zero**, every candidate already covered) but from a cheat-sheet repo's two thinnest modules, which happened to name a scan-safety discipline the KB had never written down. The earlier wrong prediction in this series was the mirror image: CPENT was called the highest-*yield* cluster on the strength of its IoT/SCADA coverage, and its study guide left those modules as

[TBD] stubs — yet the *repo* version of CPENT, a different author, is exactly where the 2 entries came from. So the corrected lesson: the CPENT syllabus is high-value, but which artifact of it you get matters more than the syllabus name.

The final scoreboard, measured:

transcript corpus (687k chars)	13 entries	
batch 1 – 24 repos + gist	2 entries	(22 sources yielded zero)
batch 2 – 7 GitBook spaces	0 entries	
batch 4 – 0xdf pass 1 (55 posts)	19 entries	
0xdf pass 2 (51 posts)	8 entries	
batch 3 – PDFs/pages/CPENT	2 entries	(4 of 5 sources yielded zero)
— 44 entries total; KB 2,699 -> 2,743		

The finding worth stating plainly: cert notes 36 sources → 4 entries; narrative writeups 106 posts → 27 entries — roughly a 20× yield-per-source difference. Exam notes are condensed, derivative, and cover a syllabus the KB had already absorbed; writeups that start at a scan result and end at root supply a shape nothing else did. That is the reusable lesson for choosing future sources, and batch 3 did not disturb it.

The uncomfortable number that must also be stated: the thin categories the series set out to fill still barely moved. Measured now against the start:

services 9 (unchanged) · credentials 5 (unchanged) · persistence 4 (unchanged)
phishing 1 (unchanged) · iot 2 (unchanged) · ics 1 -> 2 · pivoting 6 -> 8 · recon 69 -> 70

Batch 3's two entries nudged `ics` and `recon` by one each — the first movement in `ics` the whole series produced — but `services`, `credentials`, `persistence`, `phishing` and `iot` are exactly where they started. All 27 fingerprint entries land in `category="writeup"` (173 → 200) because `ingest_exploitation_writeups.py` forces it as part of its `no_merge` discipline. They function — 2.7 keys on `meta.fingerprint`, not `category` — but the series **did not achieve its stated goal** of filling those thin categories, and this record says so rather than reporting 44 entries as if it had. **Open question, recorded without acting on it:** is `category="writeup"` right for discoverability, or should a fingerprint entry carry the service category it describes? That is a contract change to the ingester and belongs in its own session.

Records the series carries forward:

- **D22 — proxychains laundering the red-confirm** was found by *cataloguing a tool*, not by looking for it: the danger heuristic classified `argv[0]`, so `proxychains ... weevily` demanded no confirm. The single most valuable output of the whole series, and it came from an enrichment batch's side.
- **D24 — the shared-predicate boundary convention**, the **third** instance of that pattern (build #5's WinRM `argv[0]`, D22's `proxychains`, D24). Worth actively watching for a fourth.
- **D25 — the retrieval residual**, closed this session: `fingerprint_match` reserved for structured hits, UNCOVERED false-fire 20%→0% with no covered-hit-rate cost.
- **The token diff nominates, it never confirms** — `seshutdownprivilege` looked like a certain gap and was already present spelled `seShutdown`; grep the concept, never the token. Batch 3 confirmed it a third way: its top two token-diff nominations, printer *pass-back* and IKE *aggressive mode*, were both already fully in the KB (the printer entry even carries the 2024/25 Xerox pass-back CVEs) — the two entries that *did* land came from concepts the diff's noise had buried, found by retrieval testing, not by the diff.

- **Sources declined and why:** themastermindnotes (commercial product, batch 3 honoured the exclusion — not fetched), `mqt.gitbook.io` (refused itself via `robots.txt`), `0xdf` (declined then reversed, D23), and **0xdf now closed** as a fingerprint source on a declining tail (19/55 → 8/51).
- **A source's scope in a prompt is a claim to verify, not a fact.** Batch 3's prompt listed `pnpt.adot8.com` as one of "two web pages"; its sitemap resolved to **185 pages** — a GitBook space on a custom domain — and it was routed through `fetch_gitbook.py` accordingly. Resolve the sitemap before trusting a source's described shape.
- **Windows Defender deleted files mid-run in four separate sessions** now (batch 3 lost `bypassing-amsi.md` to an `errno-22` lock while it still showed 3,453 bytes on disk), from fetched source trees as well as `data/kb/entries.jsonl` — it is a **standing operational hazard for this repo**, not an anecdote: back up the KB before any rewrite, and make every triage loop tolerant of a file vanishing between listing and reading.

Build #12 — the capabilities nothing could reach, and the first CI (2026-08-03)

Two findings drove this build, and the first is D19 one level deeper.

Ten endpoints that existed and could not be reached

D19 fixed surfaces that were *built and then invisible* — a dozen routes reachable only by typing the URL, cured by putting the index on the home page. This build found the same failure one layer further down: **capabilities with no route at all**. A frontend-to-backend coverage sweep (every one of the 112 source endpoints matched against every call site in `frontend/src`) found ten with no caller:

Endpoint	Backing module	Consequence
<code>GET /exploits/cve/{cve} , /exploits/cves</code>	<code>exploits/index.py</code>	"I have CVE-2021-41773, show me exploits" was unanswerable in the UI over a 25k-CVE index — you could search by <i>service</i> only
<code>GET /arsenal/suggest , /arsenal/render/{name}</code>	<code>arsenal/router.py</code>	the planner's own tool shortlist, and server-side template rendering, were API-only
<code>GET /detection/catalog , /technique/{id} , POST /tag , GET /sources</code>	<code>detection/</code>	the whole curated map — 133 command families, 19 techniques — was unbrowsable
<code>POST /sessions/{id}/webexploit/draft</code>	<code>reasoning/webexploit.py</code>	build #8 Task 3 shipped backend-only
<code>POST /sessions/{id}/privesc/ingest</code>	<code>reasoning/privesc.py</code>	build #8 Task 4 shipped backend-only

`detectionSources` is the sharpest instance: it had a typed client function in `api.ts` and no component ever called it, so the About box it was written for never existed. Two of the ten were the *depth* half of build #8 — the work was done, reviewed and locked, and no human could reach it.

All ten are now wired: CVE-id lookup takes the exact keyed path (labelled "unranked", distinct from ranked search) with a "N CVEs affect this service" band beside it; the arsenal renders any tool against a target with unfilled placeholders left **visible** and badged rather than guessed; a new `:detection` route browses the map, expands a

technique with its Sigma rules, and offers a deterministic tag probe with no model in the loop; and the two build-#8 reasoning surfaces appear in the engagement-state panel — a per-finding *draft exploit* action and a linpeas/winpeas paste box. Both are styled deliberately as proposals: no approve affordance anywhere, and a needs the red confirm chip stating what the gate will demand. **D18 is untouched** — every drafted step is data the human fires through the already-gated surface.

:detection was registered on the home launcher rather than left as a bare route, because D19's own lesson is that a page you must navigate to fixes discoverability only if you remember to navigate to it.

The first CI, and the corpus problem it exposed

Until this build the only thing running the safety suite was remembering to. That is the wrong resting place for a project whose value is that gates fire — the gate audit probed 24 guards and found **seven that never fired**; build #5 found a red-confirm defeated by moving a cmdlet one token right; build #10 found the runner disarmable by adding a pipe. A refactor that silently unlocks a gate should fail a build.

.github/workflows/ci.yml runs three jobs. Blocking on every push and PR: the **56-file hermetic suite** (gated per-file on explicit exit codes), next build, and an eslint **baseline** check — pinned at 11 rather than 0, because frontend/AGENTS.md documents those errors as accepted deliberate patterns, and because the baseline had already crept 10 to 11 unnoticed, which is the argument for pinning it. Non-blocking, weekly: pipeline/detection_sources.py --verify, the live ATT&CK/SigmaHQ drift check, kept off the merge path because a failure there means the world changed, not that a PR is wrong — and, since it had never once executed, now dispatched and made to report its verdict rather than fail silently into a log (below).

Wiring it surfaced the corpus problem D26 records: /data/ is gitignored, three fingerprint locks iterate the live KB by design, and on a clean checkout they crashed rather than skipped. The fix is the derived, complete, equivalence-proven projection described in D26 — **not** a hand-written sample, which backend/AGENTS.md §1 forbids for exactly the reason that would have bitten here.

Verified against the real CI condition, not an approximation. A clean virtualenv carrying only the five dependencies CI installs — *deliberately without pywinrm*, so the suite's hermetic claim is genuinely under test rather than papered over — combined with HACKPIT_FORCE_KB_FIXTURE=1 to hide the live KB: **56 test files, every one exited 0**, with the three fingerprint locks running for real on the fixture (105 fingerprints self-matched, 56 version checks, 15 uncovered services) rather than skipping. The same env var reproduces CI on any machine that has the KB, and is documented in backend/AGENTS.md.

The first CI run failed, and that is the result worth recording

The workflow was verified locally before it was pushed — a clean virtualenv with only CI's five dependencies and HACKPIT_FORCE_KB_FIXTURE=1 — and it still went red on its first real run. The local simulation had removed the dependencies and the KB; it had not removed **the machine**. Three tests were leaning on gitignored files that exist only on the author's laptop, and all three had been green for months:

- **sqlite3.OperationalError: no such table: sessions.** test_home_summary._real_secrets() iterates sessions_db.list_sessions() against a database that is gitignored. Fixed with init_db() (a CREATE TABLE IF NOT EXISTS, so a no-op where the DB exists) rather than a try/except: swallowing the error makes the leg silently empty, whereas an initialised empty DB makes it a real check that finds zero rows — which the caller then reports as having checked nothing.
- **The same file's invariant 1 was vacuous on a clean checkout.** Every store it draws from — stored WinRM secrets, llm_config.json, captured engagement credentials — is gitignored, so CI has nothing to leak and the check proves nothing. It now reports **NOT-RUN** and returns, instead of failing its own "no real secret found

in any store — this test proved nothing" assertion. Invariant 2, the AST no-execution scan, is structural and still runs in CI.

- `test_operator.py` died in app *startup*. with `TestClient(main.app)` enters the lifespan, which loads the KB and hard-fails when `data/kb/entries.jsonl` is absent. `/operator` reads `operator_identity` directly and needs no lifespan state, so the `with` was dropped — the construction `test_sliver_safety` and `test_obfuscation_safety` already use for their endpoint assertions. That keeps a real leak guard **running** in CI rather than skipped.

One fix was tried and rejected, and the rejection is the more useful record. Staging `kb_fixture.jsonl` at `data/kb/entries.jsonl` so the app could boot *works* — the app starts fine on the projection. It also **lies to every component that asks "is there a built KB"**: `test_corpora`'s skip guard saw one, proceeded, and failed on expected corpus entries in the built KB. A fixture that makes absence look like presence is worse than the crash it cures, and the fixture's whole justification (D26) is that it is honest about what it is. It stays where it belongs — behind a loader that names it — and the app-boot problem was solved at the tests that had it.

The re-verification was done the way the first one should have been: a real `git clone` into a temp directory — no `data/`, no `sessions.db`, no `operator.json`, no `llm_config.json` — run with CI's exact dependency set and no `pywinrm`. That clone reproduced all three failures first, so the fixes are demonstrated rather than assumed, and then went green at **56 files, every one exited 0**. The second CI run passed on both branches.

The durable lesson is build #7's and build #9's in a third setting. Build #7: six defects invisible to a green suite, every one a structural claim nothing had executed. Build #9: four defects that only a real domain could surface, every one two halves of the codebase agreeing in a test and disagreeing against reality. Build #12: **three tests that passed for months and had never once run anywhere but the machine that wrote them**. Each time the green suite was accurate about what it checked and silent about what it assumed — and each time the fix was to put the code somewhere its assumptions did not hold.

The drift job had never run — and could not have told anyone if it failed

Two green runs later, one job in the workflow had still never executed: `gh run list` returned **zero** schedule-triggered runs and **zero** dispatches. The ATT&CK/Sigma drift check is `schedule` - and `workflow_dispatch` -only by design, so nothing on the merge path had ever reached it. It was, precisely, a guard that existed and had never fired — the finding this whole workflow was written to prevent, sitting inside the workflow.

Dispatched, it passes. ATT&CK Enterprise v19.1 and the SigmaHQ master ruleset both clean, **8 detection-reference pages** checked, `0 problem(s)`, 14 seconds. The one plausible way it could have failed on a clean checkout — the citation half reads `pipeline/authored/authored_entries.jsonl` — does not arise, because that file is tracked; the gitignored `/data/` never reaches it.

But running it exposed why nobody would have noticed if it hadn't. `continue-on-error: true` is what keeps this job off the merge path, and it is also what made it mute: real drift exits non-zero, gets swallowed, and paints the job **green**. The finding would have existed only in a log nobody opens on a Monday morning. *Non-blocking has to mean reported, not silent* — otherwise "we check upstream weekly" is a claim with no delivery mechanism, which is the same shape as a proof slot scoring as a pass. The job now captures the verify outcome explicitly, writes the verdict and the full output to the run summary **either way**, and raises a `::warning::` annotation on drift — the only part visible without opening the run. The outcome is read from `PIPESTATUS`, not `$?`, because piping to `tee` otherwise hands the step `tee`'s exit code: the identical mechanism build #10 found could disarm the safety runner. Upstream is clean today, so the drift branch of the report was exercised by simulating a non-zero status rather than left to run for the first time on the day it matters.

Also bumped: all six action references (`checkout` , `setup-python` , `setup-node`) from Node 20 actions being force-shimmed onto Node 24, to `@v7` — confirmed `using: node24` at the source before pinning. Every run had been carrying deprecation annotations that would become hard failures whenever the runners drop the shim. Both changes verified by dispatch: three jobs green, deprecation annotations gone, the drift summary written.

This is the lesson a fourth time, in its sharpest form. The other three were code that was *wrong* in a place nothing had exercised. This one was not wrong at all — it had simply **never run**, and was built so that it would have looked identical either way.

A latent finding, recorded not fixed

Measuring the fixture projection surfaced something about the matcher itself. `retrieval._entry_blob` reads `title`, `text`, `body`, `summary`, `tags`, `product`, and the live KB has **none of `text`, `body`, `product`** — it stores `body_md` and `steps`. So the unstructured fallback path has never searched the body of any entry; it sees titles, summaries and tags only. That may well be correct — matching titles and summaries is the more precise behaviour, and D25 tightened this path deliberately after measuring it — but the `body` / `body_md` near-miss looks accidental rather than chosen. It is **left alone and written down** rather than "fixed" in passing: changing which fields the fallback reads would move retrieval behaviour that was just carefully measured to a 0% false-fire, and that is a decision to take deliberately with an eval beside it, not a drive-by.

Housekeeping, and a document that had been wrong for a year

`GATE-AUDIT-FINDINGS.md` moved from the repo root into `docs/` (both stale references updated — the build-5 plan literally said "at the repo root"). `fix-vmnet8.ps1`, the host-side VMware NAT repair for the build #9 lab, was **parameterized before being tracked**: it carried a hardcoded `C:\Users\zaid_\...` transcript path, and this repo is public — the same class of thing cleaned up before it was published. It now takes `param()` defaults and `$PSScriptRoot`, and sits in `docker/proof/` beside `dc_prereq_probe.py`, its natural companion (the probe reports the DC unreachable; this fixes the commonest cause). Two superseded KB snapshots were deleted after verifying the live corpus parses to 2,743 valid objects — `data/kb` 88 MB to 58 MB. `exploitdb.json` was **kept**: an earlier pass in this session wrongly grouped it with the stale backups, and it is the live CVE index `backend/exploits/index.py` reads.

The repo went to one branch. `sandbox-kali-image` was cut for the build #2 Kali image work, became the de-facto trunk for twelve builds, and had stopped describing its own contents by about build #4. It sat at the identical commit as `main` — zero commits either direction, empty tree diff — and the land-there-then-fast-forward-`main` habit bought nothing: no review gate, no window in which `main` was shielded, both refs public, both pushed seconds apart. Ceremony, not a gate; the gate is the CI plus the safety suite. It went, along with **eleven local feature branches** (`ad-graph` , `session-panel` , `detection-panel` , `engagement-mode` , `tool-arsenal` , ...) that were all fully merged. Deleted with `git branch -d` rather than `-D`, so git refused anything unmerged rather than trusting the check that said none were — it refused nothing, and all twelve tips verify as ancestors of `main` afterwards. The CI push trigger is now `branches: [main]`, with `pull_request` left unfiltered so a branch cut for a one-off build is still gated through a PR without editing the workflow again. **Every earlier reference to `sandbox-kali-image` in this document and in `docs/` is left standing on purpose** — those sentences record what was true when they were written, and rewriting a past assessment to scrub a branch name is the worse outcome of the two.

`docs/build-notes.md` was rewritten. It had opened with "The Cockpit does not exist yet" and closed calling the execution engine "still ahead of me" — true when written, false for about a year, across eleven builds and a live DCSync against a real domain. The correction is kept in the document as a dated note rather than quietly applied, because the failure is worth naming: a doc written to guard against overselling ended up underselling by a year, and the mechanism is the same either way — it stopped tracking the code. The five original "what broke" stories

survive; three better ones were added (the six live-fire defects invisible to a green suite, the shared-predicate pattern across three subsystems, and the DA password leaking into the LLM proposer context), and the closing admission is now four honest ones led by the real one: the headline was autonomous hacking and the project deliberately built the opposite.

Build #13 — where a callback lands (2026-08-03)

HackPit reaches **out** to anything: the engage sandbox is fully open by decision. Being reached **in** is a different problem, and it is the one real capability gap a review of the whole build surfaced. A callback is the target dialling *you*, and for that to land a container port must be published on a host address the target can route to. Exactly one file did that — hand-written, opt-in, and hardcoded to the VMware VMnet8 address of one laptop.

This is one of four parts, and the decomposition is the first decision. The callback gap splits into local listener profiles (no infrastructure), an out-of-band confirmation canary and a public C2 listener (both needing a VPS and a domain that do not exist yet), plus an unrelated safety-model change — reversing the evasion engine's generate-only rule into a gated deploy. Cramming those into one spec would have produced a document too vague to implement and a plan too big to verify, so each gets its own spec → plan → build cycle. Part 1 is first because it is free, verifiable today, and it is where the callback-destination abstraction gets designed, so the remote parts slot into an existing shape rather than being bolted on.

What was built

`backend/cockpit/exposure.py` owns the surface end to end: validate → render → write → apply → observe. A profile names a bind address, a container (`engage-sandbox` or `kali-open` — never the lab sandbox), and the listener kinds it will use; ports are derived from each kind's default and explicit extras merge in. `docker/proof/c2-lab.yml` is gone, replaced by a `vmnet8-dns` preset locked against it by a test.

Three decisions went against the first draft, all from pushing back on it

Public and wildcard binds are red-confirms, not refusals. The first design refused both. But this codebase already has the pattern — the danger gate "never blocks outright; requires the confirm. Over-inclusive assist — human is the gate" — and inventing a second, stricter one would be inconsistent for no gain. A wildcard buys two real things: a binding that survives a VPN or DHCP address change, where a named bind leaves a container that will not restart, and a fallback when a specific bind misbehaves under Docker Desktop's networking.

Arbitrary ports are allowed. The first design derived ports only from the four known listener kinds. Those four omit a **plain reverse shell** — netcat or pwncat on 443 or 4444, an msfconsole handler — which is the commonest callback there is, so the restriction would have been hit on first use. Ticking a kind now fills in its port as a convenience rather than acting as a cage. Ranges stay refused: a range is how one typo publishes hundreds of ports, and it makes the exposure summary unreadable.

A non-live bind address warns rather than refuses — and the claim that motivated refusing it was wrong. The first draft said a mistyped address would start fine and silently receive nothing. It does not: Docker refuses to start the container with `bind: cannot assign requested address`. The check buys a better error, earlier — not safety — and refusing would break the real case of writing a profile while off the VPN, intending to connect before applying it.

The hole self-review found, and the one the suite found

Permitting an acknowledged wildcard broke the scanner. `test_exposure_safety` is a static text scan, and it had no way to tell a wildcard the operator consciously chose from one that slipped through — so simply teaching it to accept wildcards would have *deleted* invariant 3 rather than relaxed it. The acknowledgement is therefore rendered **into the file**, as a `# hackpit-ack: wildcard bind=0.0.0.0 engagement=...` line, and the rule becomes **covered** rather than **absent**: every broad binding needs a marker naming that exact address. This makes invariant 2 stronger, not weaker — the one small file a reviewer reads now states what is exposed *and* that it was chosen deliberately, by whom and when. A marker for a different address covers nothing, which is the case that keeps the rule from degrading into "contains the string `hackpit-ack` somewhere".

A guard fired on the first full-suite run, and the fix was to remove the dependency rather than excuse it. `exposure.py` needs to know which port a listener binds, and the obvious way to get it is `from .tunnels import CHISEL_DEFAULT_PORT`. That trips the whole-tree scan that allows exactly two files to reference the tunnel module, so no agent path can raise a pivot listener; `sliver` and `obfuscation` carry the same guard. Adding `exposure.py` to those allow-lists would have been wrong twice — it narrows the file set instead of fixing the predicate, the mistake build #5 records, and it would have left the module free to call `start_tunnel` with nothing watching. The scan matches the *import* rather than the call precisely so a module cannot get within reach and then be trusted not to use it. So the constants moved to `cockpit/listener_ports.py` and the dependency disappeared. The drift lock came out stronger: one definition instead of one per owner.

Measured, not assumed

The bind classifier keys off `ipaddress.is_global`, not the obvious `is_private`, because a test caught `is_private` being wrong twice over on Python 3.14: it is **False** for CGNAT `100.64/10` (Tailscale, mobile hotspots) and **True** for the RFC 5737 documentation ranges. The first is the one that would have hurt — every Tailscale address would have been called public and demanded an acknowledgement, on the interface most likely to be right for a remote internal engagement. Liveness is probed by binding a throwaway UDP socket rather than enumerating interfaces, which would need a third-party package the hermetic suite forbids, and which asks a weaker question than "can a listener actually bind here".

What it still cannot do

A published port is **not** an open port — the host firewall can still drop inbound, so `observe()` says so rather than leaving the operator guessing. And 1c does nothing for an internet-facing target: `192.168.13.1` means nothing to a host on the internet, which is the whole reason parts 1a and 1b exist and need infrastructure. Suite **57 files, 0 failures**, up from 56.

Build #13 part 2 — the evasion engine can now deploy (2026-08-03)

This is a policy reversal and should be read as one. The evasion engine opened with "*GENERATES ONLY, never runs or deploys*", and two tests enforced it. That property is gone by decision. HackPit can now put an artifact built to evade detection onto a real host and run it.

The argument for it is the project's own precedent: the Sliver server and the pivot/DNS listeners were once refused outright and became **gated-and-allowed** in build #7 / I2, on the reasoning that the gate — not the absence of the feature — is the control. The artifact always landed in a loot directory mounted into the sandbox and sitting on the

host, so an operator could already copy it out and run it by hand. What changed is that the step no longer has to leave the tool. The argument against it is simply the plain fact above, and it is recorded here rather than argued away.

What did not change: the mandatory footprint. `deliver` computes the honest half first and raises rather than act without it, exactly as `generate` does, and the route guard now asserts `DeliveryResult` carries it too. An evasion tool that told you only how to be quieter, and never what still sees you, would be an evasion how-to — that is what makes this purple-team, and it is what justified lifting the OPSEC sensor-tamper ban in D16.

Two primitives, deliberately separated

Putting an artifact somewhere and **running** it are different acts with different blast radii, and one `deploy()` would have left the gate unable to tell them apart. So `deliver()` takes a **closed set** — `winrm` (chunked base64) or `smb` (`argv smbclient`) — and never a free-form delivery command, which would have handed the package a general execution path with none of the executor's gates. `invoke` is **WinRM-only**, and a sandbox invoke is **refused rather than merely unimplemented**: running the artifact inside HackPit's own box detonates it on the operator's machine, never the target.

The red-confirm is required **unconditionally** rather than left to the heuristic to notice — build #5 found a red-confirm you could defeat by moving a cmdlet one token right, so a gate that depends on a classifier spotting a name is a gate a rename defeats.

The guard that fired found a design mistake, not an import

`test_winrm_safety` scans the whole tree and allows only the executor and the router to reach `winrm_transport` — the orchestrator proposes, it must never fire WinRM. The first implementation of `deliver` imported the transport directly and tripped it.

Adding the evasion engine to that allow-list would have been wrong twice over: a **third** module able to fire WinRM, and that module re-implementing gates beside the ones already living in the executor. The real problem was architectural — `deliver` was duplicating gate logic and then reaching *around* the gated execution point. So the capability moved to `executor.send_windows_scripts`, where every other Windows command already goes.

That also resolved a genuine tension. Per-command human approval is the floor on a real target, but a chunked upload is not N commands anyone could sanely approve one at a time. Treating the whole sequence as **one gated act** — one approval for one transfer, chunking an implementation detail of it — keeps the floor intact without making it unusable. This is the second time in one build that the right answer to a whole-tree guard was to remove the dependency rather than be added to its allow-list.

Smaller things the tests forced

A **short write is checked** against the far side's own reported file length, because a truncated payload that reported success is the failure mode chunking introduces — and it is worse than a failed transfer, since the operator would run it. The SMB credential is **masked by construction** in the audited argv, so the run store does not become a key store — obfuscation.py's rule for its pre-shared tunnel key. And the route-set guard caught the new `/api/evasion/deliver` endpoint, which is exactly its job: the set is pinned so a new evasion surface cannot appear without someone deciding it should.

Suite 57 files, 0 failures.

Deferred to a later session, deliberately

Two README items, noted here so they are not lost:

1. **State plainly that the KB is not in the repo.** `data/kb` — the ~2,743-entry dataset of techniques, workflows and payload corpora — is gitignored and always will be (it is rebuildable, and third-party corpora are never committed). A reader who clones this and expects search to work will find it **empty and non-functional**, and the README does not currently say so, or say that the KB is available on request from the author.
2. **Mention the CI.** It is a genuine engineering signal — a public repo whose safety suite runs on every push, with a corpus fixture proven equivalent to the real one — and the README still describes the suite as something run by hand.

The README's counters are also stale (it claims 2,621 KB entries, a 110-tool catalog, and "42 test files, 459 assertions"; the measured values are **2,743, 115 and 56**). All three are one editing pass, held for a session of their own.

Status

Suite **57 files, 0 failures** (build #13), green both with the live KB and under the CI simulation. Frontend builds clean at 23 routes; eslint holds at the accepted 11-error baseline. Every one of the ten previously-unreachable endpoints now has a real component caller, verified by re-running the sweep that found them.

CI is complete: **all three jobs have now actually executed on GitHub's machines and passed** — the two gating jobs on every push, and the drift job by dispatch, which was the last piece of this workflow that had only ever been reasoned about. No deprecation warnings remain.

The repository is **one branch**, `main`, and CI is green on it.

Build #13 part 3 — the out-of-band canary (2026-08-03)

This is **part 1a** in the numbering above: the half of "where a callback lands" that part 1c can never do. Listener profiles solve the callback for a target that can route to your laptop; `192.168.13.1` means nothing to a host on the internet, which is every bug-bounty target there is.

Without an internet-reachable listener, whole vulnerability classes are **unconfirmable**, because the hit *is* the entire proof: blind SSRF, blind XXE, blind RCE with no returned output, DNS-based blind SQLi (`xp_dirtree` , `UTL_HTTP`), JNDI and deserialization callbacks, and async SSRF where the app calls your URL minutes later through a queue. The alternative is writing a promising blind injection up as "unconfirmed", which most programs reject and which this project's own report discipline already bans.

DNS matters more than HTTP, and that is the reason this needs a real domain rather than an IP. Plenty of targets block outbound HTTP from application servers while DNS still resolves through their internal resolver, so DNS-based OOB lands where HTTP does not. That requires **NS delegation** to an authoritative listener — and it is the second reason part 1c can never do this job, since a target's resolver will never route to a private address.

HackPit does not provision, and that is a decision

The tool **configures, deploys and verifies**; you create the droplet and buy the domain yourself, once. Four reasons, in order of weight. A credential that can create one droplet can create a hundred, and it would sit in an application that has **zero route authentication today**. The ROI is poor for a strictly one-time task. Domain registration needs a

funded registrar account and ICANN verification regardless, so the "automated" path still stops and waits for a human. And a tool that *spins up* C2 infrastructure is a materially different artifact from one pointed at infrastructure you already own — the same concern that ended build #11. Everything after the one-time setup is buttons.

The first internet-facing component, and why it is safe to expose

`oob/server.py` is the first HackPit component that faces the internet, and it faces it from a machine holding a client's evidence. The safety argument is deliberately **an absence rather than a control**:

- it **records, and that is all** — no execution, no eval, no deserialization;
- it **never forwards** — there is no outbound connection anywhere in the file, so it cannot be turned into a redirector. That is part 4, and a different thing;
- it **never reflects** request content into a response — the body is the constant `ok\n`, so it is not a free reflection oracle for anyone who can reach it;
- the hit log is **append-only JSONL** — no database, no rewrite path, no delete;
- **reads are authenticated**, because the log holds the target's internal hostnames and source addresses. That is the client's information, and an unauthenticated read endpoint would publish it to anyone who guessed the host. Bearer secret from the environment, compared with `hmac.compare_digest`, rate-limited per source, newest-first on a cursor.

It is **one file importing nothing outside the standard library**, because it is deployed by copying it to a bare VPS. An install step on a box you SSH into once is a box that stops working the day a wheel moves.

The correlation deliberately does **not** live there. The server records the candidate token it saw; only HackPit knows which tokens were ever minted and for which engagement. A canary that knew would be a canary worth stealing.

Four things the design got right only after being questioned

An answer, never NXDOMAIN. Refusing the name would record the DNS half of a chained proof and kill the HTTP half. The A record is the feature: the target resolves `<token>.<zone>`, gets an address, and the follow-up request lands on the same box under the same token. Other query types (AAAA, TXT, MX) get NOERROR-no-data rather than a refusal, for the same reason — the name stays alive for the A retry.

The token is the label immediately LEFT of the zone, not the leftmost label. Reading the leftmost one works perfectly for `<token>.<zone>` and fails silently for every blind-RCE and DNS-exfil one-liner, all of which *prepend command output* to the name. `whoami-output.<token>.<zone>` has to correlate or the feature only works in the demo.

The question is echoed byte-for-byte, and this is the subtle one. Resolvers randomize the case of a query as an anti-spoofing measure (DNS 0x20) and compare the echo against what they sent. Re-encoding the question from the lowercased name — the obvious implementation — produces an answer the resolver discards as a spoof, while every log on this side says the hit was answered. So the raw question bytes are kept verbatim for the response, and the *recorded* name is folded for correlation. Two different jobs, two different forms.

Credentials that arrive in a hit are recorded as present, with their values dropped. A blind SSRF routinely arrives carrying the target's own `Authorization` header, session cookie or proxy credential. The finding needs "an authenticated internal client reached out"; it does not need the secret it used, and storing it would turn a canary into a credential store sitting on a VPS. Bodies are capped to a 512-byte excerpt, flagged when truncated — a canary records that something arrived and from where, not a copy of the target's traffic.

The read secret **fails closed**: no secret, or one under 16 characters, and the server refuses to start. The failure mode that prevents is an operator deploying, watching the listeners come up, and never learning the read endpoint was open the whole time.

Verified end to end, on loopback, for real

The first assumption was that nothing here could be verified without infrastructure. That was wrong, and precisely wrong: **everything except public reachability runs today**. `docker/proof/oob_loopback_proof.py` starts both listeners on `127.0.0.1` with DNS on `udp/5353`, sends a real datagram carrying a real minted token, parses the answer off the wire, and asserts the hit was recorded, correlated back to the right engagement and step, and served through the authenticated API — with the anonymous and wrong-bearer reads refused and the operator's own read traffic absent from the hit log. **15 passed, 0 failed, 2 not-run**.

The two not-run are named exactly, never folded into a pass: that **NS delegation resolves publicly** to the server, and **one live hit from a real target**. Both become real checks the moment the VPS and zone exist.

The proof's own first version had the bug worth recording. It probed `udp/5353` with a throwaway socket before binding, and fell back to an ephemeral port when the probe failed. The probe did not set `SO_REUSEADDR` and the real listener does, so on a host running mDNS the probe reported the port busy when the actual bind would have succeeded — and the proof passed cheerfully on a port the spec never named. A probe that differs from the real thing answers a different question. It now attempts the real bind and says loudly if it ever falls back.

Hermetically, the two new test files add **24 assertions** over the wire format, the token grammar, redaction, paging, authentication and inertness. Ten mutations were planted to confirm they can fail — NXDOMAIN answers, a re-encoded question, redaction removed, an uncapped body, a prefix-matching bearer check, a missing read secret, oldest-first paging, the leftmost-label token, a reflected response, and a global instead of per-source rate limit. All ten were caught.

Two structural locks exist because no behavioural test can see the difference: the minter must draw from `secrets` and never `random`, and the bearer check must go through `hmac.compare_digest` rather than `==`. Both spellings look identical in review and only one of each is safe.

The token grammar lives in two files on purpose — the deployable may not import from the repository — so a test holds them together against 200 freshly minted tokens rather than against a reading. Drift there is the silent kind: hits keep landing and quietly stop correlating.

The rest of part 1a — poll client, templates, panel (2026-08-03)

The first commit was the server and the minting. This one is everything that turns a line in a JSONL file on a VPS into a finding in a report: the **poll client and state ingest** (§3.3), the **payload templates** (§3.4), and the **configure / deploy / verify panel** (§3.5) with its UI.

Correlation is the product, so nothing is dropped. A hit arrives with no memory of why. The poll client joins each one back to the token's mint record — engagement, step, note — and files the correlated ones as `high`-severity findings that carry the source address, the timestamp and the token. The severity is deliberately not a judgement call: a callback from inside a target's network is the proof for the *entire* blind class, and grading it lower because the module cannot see the payload would understate every one of them. Hits that cannot be attributed are **reported, never discarded** — "something arrived that I could not place" and "nothing arrived" are different facts, and collapsing them would reintroduce exactly the silence this part exists to remove. There are two such shapes and they get distinct reasons: a token this HackPit never minted, and an engagement with no state session to file into.

Hits arrive late, so the engagement lookup covers exited engagements too. A blind SSRF routed through a queue lands minutes or days after the test that caused it, by which point the engagement has very often been exited. Filtering to active engagements would have filed those nowhere — the evidence loss the canary was built to stop.

The cursor is the whole safety of re-reading. It advances only after the ingest returns, and only forwards. A poll that fails anywhere leaves it untouched and the same hits are read again next time; re-reading is free because findings upsert on a fingerprint, whereas missing a hit is not recoverable. Reads with an explicit `after` (verify, the panel) never touch it, so a verify run cannot swallow a genuine callback that arrived in the same window.

Why an outbound request is allowed here when the repeater refused one. `cockpit/repeater.py` deliberately does not send from the backend process — it execs curl inside the open sandbox — because operator-composed requests to arbitrary hosts would be a general egress path out of the operator's machine. Neither half applies to a poll: the destination is one host read from the config store, the path set is two constants, and the response is parsed as JSON and never executed. That is the containment shape `backend/l1m.py` already has. **Writing that down surfaced a real defect.** The first version documented "redirects are not followed" and did not follow from the code: `urllib.request.build_opener` re-adds every default handler that is not overridden, so leaving the redirect handler out leaves redirects *on*. A tampered or proxied canary answering `302 http://169.254.169.254/...` would have turned a poll into an SSRF from the backend host. It now passes an explicit refusing handler, and the test interrogates the real opener object rather than reading the source — with a control proving the check fires on an opener that would redirect.

The templates encode the details whose absence is silent. Fifteen of them across the five classes the spec names — SSRF, XXE, blind RCE, blind SQLi, JNDI — each carrying where to paste it and, separately, *what a hit proves*, because that sentence differs by class and ends up in the write-up. The class-specific traps are the point: a parameter entity where a general entity is refused, a UNC path rather than a URL for `xp_dirtree`, `LOAD_FILE` being Windows-only. The one that would rot quietly is **exfil direction** — the server reads the label immediately left of the zone, so `whoami` -output must be prefixed, not appended. A test feeds every DNS name in every rendered payload to the **canary's own parser** and asserts the right token comes back: 21 names, checked against the implementation that will actually read them rather than against a second regex.

The deploy takes no destination, and that is asserted on the signature. Shipping `server.py` over SSH is a third remote-execution transport after `docker exec` and WinRM, so it lives at the gated execution point in `cockpit/executor.py` — the lesson part 2 learned when `deliver` grew its own WinRM call and a whole-tree guard caught it. `deploy_oob_canary(*, approved, restart)` has no host, user, port or key parameter, defaulted or otherwise, and the resolver it calls takes no arguments either: there is no way to *express* "deploy somewhere else". A test parses the real signature and fails on any destination-shaped parameter, with a planted control. Around that: an unapproved deploy sends nothing (verified by recording the transport, not by trusting the return), the whole tree is scanned so only the executor and the canary's own route can reach it, and the read secret rides **stdin** into a 0700 file rather than a command line, because `ps` is world-readable on a multi-user box. The config store refuses shell metacharacters in the host, zone, user and key path at *configure* time, and every interpolated value is single-quoted at the point of use anyway.

The subprocess runs on the host rather than in the sandbox, which is the opposite of the repeater's choice and is deliberate: the alternative would mount the operator's SSH **private key** into a container that also runs attack tooling — strictly worse than the thing it would be avoiding.

Verify is the valuable button, and it reports NOT-RUN as its own status. Every link in this chain fails silently, so "is my canary working" has to be re-runnable. Three checks: the server answers its authenticated health endpoint *and* is authoritative for the zone payloads are rendered against (a mismatch there is a live misconfiguration that otherwise looks like "no hits ever arrive"); a freshly minted token completes a full round trip — mint, arrive,

correlate, read back; and NS delegation resolves. The third is the only one that genuinely needs public infrastructure, and it says so rather than counting as a pass. The NS records themselves are generated as exact copy-paste text, with the mistake named: an A record alone makes the zone resolve while every `<token>`. name still goes to the parent's nameservers.

A guard fired, and the fix was to split the claim rather than widen the ban. `test_oob_tokens.py` asserted that the whole `backend/oob` package "executes nothing and reaches no network" — and the second half stopped being true of a poll client, whose entire job is to fetch. The tempting fix was to drop `urllib` from the banned list, which would have un-banned it for all six modules. Instead the invariant is now two: execution stays absolute with no exceptions, and network reach is a **pinned set of two named modules**, so a third one gaining egress fails here. That is a stronger statement than the original made, not a weaker one. The same discipline applied to a second guard: the DNS-tunnel human-only lock matches on raw text and fired on a docstring that merely *cited* the sibling generator to agree with it — so the citation was dropped rather than the module allow-listed.

One defect in the tests themselves, worth recording. The first drafts called `save()` and `clear()` on the real `sessions.db`, which would have destroyed a live canary's read secret — a value whose only other copy is a 0700 file on a VPS. All three test files and the proof now redirect their stores to a scratch database. A suite that can damage the thing it is testing is not hermetic whatever else it asserts.

Verification. The loopback proof grew from 15 checks to **24 passed / 0 failed / 3 not-run**, and the new half is not a unit test wearing a proof's name: it configures the real store, points the **real poll client** at the real listener over loopback TCP, and asserts it fetches through the authenticated API, correlates back to the engagement and step, files findings into engagement state, and is idempotent on a re-poll. The real verify button runs too — health and round trip pass, DNS reports NOT-RUN. Hermetically, three new test files add coverage of correlation, the cursor, the opener, the templates against the server's own parser, and the deploy gates.

The three NOT-RUN are named individually and never folded into a pass: **NS delegation resolving publicly, one live hit from a real target, and the SSH transfer itself**. Only the third is new, and its *gates* are fully covered hermetically — what is missing is a remote box to reach, not an untested control.

Frontend. `:oob` is a real panel at `/oob`, wired to all eight endpoints — there is no endpoint here without a component caller, which is the whole reason build #12 existed. It walks the actual order of operations: configure, delegate, deploy, verify, mint, collect. The read secret field is blank on load and blank means "keep the stored one", because the value is never sent to the browser. The deploy button sends `{approved}` and nothing else; there is deliberately no host field, since there is no parameter for one. eslint holds at the accepted baseline of 11 — the refresh callback routes its `setState` through a `.then` rather than an effect body, which is what pushed the count to 12 on the first attempt.

Suite **62 files, 0 failures**, up from 59.

What is not built yet

Part 1a is now complete. The public C2 redirector remains **part 4**; this server records hits and never forwards, and nothing in it can be turned into a redirector — an outbound-connect ban is asserted over its AST.

Build #13 part 4 — the public C2 redirector (2026-08-03)

The last piece of "where a callback lands". Part 1 made the destination configurable but every destination it can express is an interface **on this machine** — so an implant inside an internet-facing target has nowhere to call, because a laptop behind NAT has no address anyone can dial. Part 3 made the *proof* of a blind vulnerability reachable and deliberately cannot carry a session. This adds the remaining shape: a redirector on a VPS the

operator already owns, which accepts an inbound connection on a public port and relays it down a reverse tunnel that was dialled **outward**. The implant talks to the VPS; the VPS talks down a tunnel established from the inside; nothing is traversed inbound.

This one carries real AUP weight, and the safety argument had to be rebuilt

Part 3's canary is safe to expose because of an **absence** — it records, never executes, never forwards, and answers a constant. Reusing that shape here would be dishonest, because this component *forwards by design*: it is the exact thing `test_oob_server.py` has an AST-asserted ban against the canary becoming. What is actually being stood up is an always-on listener on a public IP, reachable by anyone who scans that address rather than only by the target under test, relaying traffic it did not authenticate, **into the operator's own machine**. That last direction is the one that matters — a misconfigured redirector is an inbound path from the internet to a laptop, not just an exposed service on a rented box.

So the claim is **bounded forwarding**, and the bounds are structural rather than advisory:

- **One destination, and it is loopback.** `TARGET_HOST` is a module constant in the deployable. A test walks every addressing call in the file and asserts each one uses it — with the answering half held to its own non-empty rule, since replying to the source of a datagram is not choosing a destination. Nothing in the file resolves a hostname. An open forwarder on a public IP is precisely what a stranger scanning that address wants to find, so "it cannot be pointed elsewhere" had to be a property of the code.
- **An enumerated port set.** No ranges, the same rule part 1 already applies, for the same reason: a reviewer has to be able to read the whole exposed surface at a glance.
- **Off unless the tunnel is up.** With nothing attached, the loopback target is a closed port, so a client is accepted and dropped. That is the right failure direction, and the proof asserts it *first* — before any listener exists — because it is the state a redirector spends most of its life in.
- **Deploy and stop are both gated, and equally reachable.** A start button without an equally reachable stop is how a public forwarder outlives the engagement it was built for.
- **The panel says all of this in those words.** The exposure sentence, the AUP position and the teardown command are returned *with* the profile, so a panel cannot render a configured redirector without also rendering what it exposes.

Authentication is deliberately not built, and the reason is written down. It would be theatre: a shared secret would have to live inside the implant, which is a binary the target holds. A fake control standing beside a real one is worse than no control, because it invites trusting the wrong thing.

Extending part 1 rather than sitting beside it

`ListenerProfile` gains a `destination` of `local` (unchanged, and regression-locked — part 1's tests and its published-port scanner pass untouched) or `remote`. The two are not one path with a flag, and the validation genuinely differs rather than being reused with exceptions carved out: a remote profile has no bind address to check liveness on, no container, and **no private case** — a port on a VPS is reachable by anyone, so the public acknowledgement is unconditional rather than conditional on classifying an address there is none of. `render`, `write`, `compose_command` and `apply` all refuse a remote profile at the door, and `write_remote` refuses a local one, so the two paths cannot half-mix into a compose file that publishes on `' '` or an operator believing a VPS is exposed when nothing was shipped to it. `observe()` reports `remote` and does not guess — `docker inspect` says nothing about a process on someone else's machine, and answering anyway would be the exact defect `lifecycle.py` exists to prevent.

One deploy engine, not a second path

Part 3's deploy ships one file to the configured VPS behind an approval and takes no destination. Part 4 needed to ship a different file to the same box, and two of the three available answers were wrong: a second deploy function would repeat part 2's mistake of two places implementing the same gates, and a `path` parameter would turn a deploy button into an arbitrary-write primitive on a machine reached over SSH once. So the transport, target resolution, gates, stdin-not-argv secret discipline and step orchestration became **one private engine**, with the artifact as a module-level constant built inside each wrapper from repository paths and server-side config. The public surface is three thin wrappers that take an approval and nothing addressable, and the signature assertion now runs over **all three** — plus a new check that reconciles the wrapper list against the real module, so a fourth one added without a line in the enumeration fails rather than going unchecked.

The one awkward fact, stated rather than papered over

SSH reverse tunnels carry TCP only. A Sliver implant or a reverse shell rides `ssh -R` end to end; a DNS tunnel does not. Rendering an `-R` line for UDP would produce a command that runs cleanly and carries nothing — indistinguishable from a target that never called back, which is the precise class of silent failure part 3 exists to remove. So the UDP case renders a `socat` pair with each end labelled by *where it runs*, and says why. The tunnel itself terminates on `127.0.0.1` on the VPS rather than `0.0.0.0`: the alternative needs `GatewayPorts yes` and would put the tunnel endpoint on a public interface with no forwarder in front of it — an unbounded relay straight into the operator's machine.

The reverse tunnel is **rendered and never run**, the same boundary the DNS-tunnel client one-liner draws. It is a long-lived outbound process on the operator's own machine, and starting it deliberately is the approval.

A test bug worth recording

The first draft of the destination scan read `args[0]` for every addressing call — correct for `connect`, wrong for `sendto`, whose signature is `(data, address)`. It inspected the *payload* and reported every legitimate relay as an unknown destination. That is the kind of false alarm that gets a guard loosened rather than fixed, so the predicate now records which argument carries the address per call shape, and a control asserts specifically that `sendto`'s address slot is the one being read.

Verification

On loopback, for real. A redirector relaying `127.0.0.1:A → 127.0.0.1:B` is the entire mechanism; the only thing a public IP adds is that a stranger can reach A. So the proof runs a real forwarder, a real listener standing in for the far end of the tunnel, and a real client standing in for an implant: a full exchange in both directions over TCP, the same over UDP, the accept-and-drop behaviour with nothing attached, survival of a client that vanishes mid-stream, and agreement between the rendered `ssh -R` port and the port the running forwarder actually relayed to. Hermetically, two new test files cover the live forward, the rendering, and the four safety invariants with a control each.

Reported NOT-RUN, named individually, never folded into a pass: that a connection from the internet reaches the forwarder on its public address, one real implant session through the full chain, and the SSH transfer itself. Only the first is a property of this component that nothing local can stand in for; the deploy's *gates* are covered hermetically and what is missing there is a remote box, not an untested control.

Frontend

`:exposure` at `/exposure` — and this is where part 1's four `/cockpit/exposure` endpoints finally get a caller, having shipped with none. That was the same gap build #12 existed to close, and adding a second orphan surface beside it would have compounded it. The panel walks both destinations from one port selection, keeps the approval explicit for every destructive or exposing action, and renders the reverse-tunnel command, the socat bridges and the teardown with copy buttons. eslint holds at the accepted baseline of 11.

Suite 64 files, 0 failures, up from 62.

Build #14 part 1 — ZAP, the first web vulnerability scanner (2026-08-03)

The arsenal catalogued 115 tools and shipped an HTTP repeater, but nothing in it systematically probed a discovered parameter for an injection class. Recon found endpoints; nuclei matched known templates; the gap in between was the one a web assessment actually lives in. This closes it with OWASP ZAP.

Why not Burp Suite Professional

Burp Pro is \$499/user/year and its REST API is a scan-launcher — start a scan, poll it, fetch issues. Anything richer needs a Montoya extension loaded inside a running Burp, which is not drivable over a wire. Its licence is a named-user seat and the process assumes a desktop session, so running it as a shared service backend is the case PortSwigger sells Burp DAST for. ZAP is Apache 2.0, headless-first, and has no activation to break when a container is rebuilt.

A cracked Burp distribution was considered and rejected outright. Beyond being pirated software and a takedown risk to a public repo, it means running an unaudited third-party patch as the interception proxy that sees every credential pushed through it.

The whole build is seven touch points and no new architecture

ZAP is modelled on nuclei: a catalogued tool the executor gates and the ingest parses. No new module, no new endpoint, no frontend, and — the point — **no new execution capability**. Scans run through `POST /cockpit/exec` as ordinary commands, so all four lab gates and all three engagement gates apply on day one without a line of code protecting them.

The daemon and its REST API are excluded, and the reasoning is structural rather than cautious. Every gate here inspects *a command*. Once "scan this target" is an HTTP call, it bypasses all of them unless a second, parallel gate system is built — and a second copy of a safety predicate is this codebase's most-repeated defect: the WinRM `argv[0]` classification, the proxychains-laundered confirm, the collector FQDN classifier. Worse, the only channel into a sandbox is `docker exec`; the lab sandbox sits on an `internal: true` network precisely so no route exists between it and the host. Reaching an HTTP API in there means opening the path `assert_isolation_proven()` exists to deny.

`zap.sh -cmd -autorun plan.yaml` is excluded for a related reason: whether the run is passive or active lives *inside the plan file*, so the gate would classify a string that does not describe what executes. The packaged scan scripts are self-describing, which lets the existing danger gate split them unmodified.

The split, and an inconsistency stated rather than hidden

`zaproxy -cmd -quickurl` spiders and then sends live SQLi/XSS/command-injection payloads at every discovered parameter, and demands the red-confirm. `zaproxy -cmd -zapit` crawls and fingerprints, and does not — gating both identically would make the confirm meaningless for the tool, the same argument the AD-enumeration note has always made.

The verdict is argument-based, not name-based, because Kali ships one launcher that does both jobs, so the binary was never the tell. `_TOOL_ATTACK_FLAGS` mirrors the `_TOOL_EXEC_FLAGS` pattern already used for `netexec -x`.

This makes ZAP **stricter than the rest of its family**: `sqlmap`, `nikto`, `dalfox` and `nuclei` remain unflagged, and `sqlmap` is arguably more intrusive. That is a recorded decision, not an oversight. Erring safe on a new tool changes no existing behaviour, and the rationale is written at the point of use so a future reader does not quietly "fix" it.

Two things the existing guards caught that the plan had not anticipated

A fourth shared-predicate defect, prevented. `dangerous_script_heuristic` shares its tool groups with the command heuristic — its docstring says so explicitly, "two lists would drift, and drift is what produced this bug." Adding the active-scan rule to only the command side would have been the fourth instance of exactly that failure.

It went into both — and after the rework the script heuristic **derives** its markers from the same `_TOOL_ATTACK_FLAGS` dict the command heuristic reads, rather than restating them, so the two cannot drift even in principle. A test locks them to the same verdict on the same tool.

A tool's own name needs a verdict too. `_catalog_invocations()` covers the catalog's `name` field, not just template `argv[0]` s, so bare `zap` failed the suite until classified. It landed in `_ARGUMENT_DEPENDENT`, whose stated shape it matches exactly: the bare binary is clean, at least one catalogued template fires.

The image build caught what the whole test suite could not

This is the finding worth keeping. Everything was first written against `zap-baseline.py` and `zap-full-scan.py` — upstream's documented packaged scan scripts, and the names the parser registry, the catalog templates, the danger sets and four test files all hardcoded. Kali's **zaproxy package ships neither**. `dpkg -L zaproxy` on 2.17.0kali1 gives exactly `/usr/bin/zaproxy`, `/usr/bin/owasp-zap` and `/usr/share/zaproxy/zap.sh`. Those scripts exist only inside OWASP's own Docker image, and they hardcode `/zap/zap-x.sh`, so adopting them would have meant faking that image's directory layout and fetching three unpinned files from GitHub `main` into a safety-critical image — rejected, given how often upstream drift has already bitten this project (`dnscat2`, `ScareCrow`, `trufflehog`).

Every one of those keys would have matched nothing, forever, with a green suite. That is the build #9 ingest gap in a new place, and it is the *third* time in this build that the same root cause surfaced: **a hermetic test feeds the parser a string the test itself chose**. Nothing in 66 test files could have found it. The image build's own smoke test did, on the first run, because it checks the names against the package rather than the documentation.

The replacement is the package-native CLI, and it was verified before being adopted rather than after: a real ZAP 2.17 scan was run inside the Kali image against a live throwaway app, and `-quickurl`'s report turned out to be shaped exactly like the one `parse_zap` already handled — **so the parser needed no change at all**. That verbatim report is now committed as `test_support/zap_report_fixture.json` and a test parses it for 4 findings and 13 endpoints, with severities mapped from ZAP's *string* risk codes. The synthetic fixture is kept beside it, because the real report happens to carry only risk codes 1-2 and the four-way mapping still needs covering.

The trap that would have shipped silently

Two normalisers exist and they disagree. `allowlist._tool_name()` strips `.py`; `state.ingest.program_name()` strips only `.exe`. So the danger sets must be keyed `zap-full-scan` and the parser registry `zap-full-scan.py`. Key either the other way and nothing matches: zero findings, no error, green suite. That is the build #9 ingest gap exactly, where a live DCSync dumped four NTLM hashes including `krbtgt` and ingested none of them. A test pins each side against its own normaliser.

A second, smaller version of the same lesson: `_json_objects()` cannot reach a ZAP report, and the way it fails matters. It does not return nothing — its line-delimited fallback matches the alert *instance* fragments that happen to sit on one line. A parser resting on it would emit quiet rubbish rather than obvious nothing. The spec claimed "neither path works"; implementation measured otherwise and the spec was corrected in the same commit.

Verified, and what is not

Hermetically, for real: the report is extracted from surrounding progress output, all four risk codes map, the registry keys match what `program_name` actually produces, `-zap.json` is claimed while a plain `.json` loot file is not, both heuristics agree on both verdicts, proxychains cannot launder the confirm, and the catalog genuinely ships both invocations. Every one carries a control in the same test. Suite **66 files, 0 failures**, up from 64.

Now run, and green: `docker/proof/zap_install_proof.sh` reports **9 passed, 0 failed**. The image was rebuilt, ZAP starts, both program names resolve under exactly the strings the parser registry keys on, the two scripts Kali does *not* ship are asserted absent, `$HOME/.ZAP` is writable as uid 1000, and a real active scan of the Juice Shop lab target from inside the isolated sandbox produced a report that `parse_zap` turned into **6 findings and 26 endpoints**. The ingest path is verified end to end.

Getting there cost three rounds, and each failure was a check working rather than a tool misbehaving. **The container is not the image:** every name check ran `docker run` against the freshly built image while the scan ran `docker exec` inside a container that had been up for 45 minutes on the previous one — `compose up -d` does not recreate a running container just because its image changed. The proof now compares the two IDs and says so outright. **MSYS rewrites container paths:** on a Windows host, Git Bash turned the container path `/tmp/zap-proof.json` into a host path before Docker saw it, so ZAP correctly reported the directory unwritable and exited 0 — a scan that "worked" and wrote nothing. **And bash's `/tmp` is not Windows Python's `/tmp`**, so the report is now piped to the parser rather than staged through a host file that the two sides name differently. A fourth lesson in one build: an exit code is not a result.

Still not run: a scan driven through the gated executor's own approve-and-run path rather than `docker exec` directly. The gates themselves are covered hermetically; what is untested is the wiring, not a control. These are the only checks that can confirm Kali installs the scan scripts under the names the parser registry and the catalog templates both hardcode — and no hermetic test can stand in, because it feeds the parser a string it chose itself. `docker/proof/zap_install_proof.sh` runs all of them and fails loudly per check. **Until it passes, the ingest path is unverified end to end.**

Timeout, resolved without touching a global

A full active scan runs for tens of minutes against the executor's 180-second default. No global was raised — `MAX_TIMEOUT_SECONDS` is already 3600 and clamps rather than refuses, so the two full-scan templates simply say to raise `timeout_seconds` on the request. Raising a default for one tool changes behaviour for every tool that relies on it.

Deferred to part 2

The proxy surface — ZAP as a live intercepting proxy feeding the repeater and state — with its own spec and safety review. The daemon question belongs there, and it splits by sandbox: the engage sandbox is already open and breaks no property by hosting one, while the lab keeps `internal: true` and keeps command-path scanning. A *recording* proxy may reuse the existing listener pattern almost directly (gate the start, hold liveness, expose no writer); a scan-control channel would need the `tunnels.py` treatment, where one derivation function sits behind both the gate and the action.

Build #14 part 2 — the recording proxy (2026-08-03)

Part 1 gave HackPit a scanner. This gives it the thing the tool had never had: **the raw HTTP of every run**. Until now a `ffuf` run's findings were parsed and its actual requests were thrown away — you could see that `/admin` existed, never the request that found it or what came back.

The daemon part 1 refused, built anyway — because the transport changed

Part 1 excluded a ZAP daemon for two reasons: an HTTP control channel bypasses `validate_request`, and reaching one inside the lab sandbox would mean opening the `internal: true` network `assert_isolation_proven()` exists to deny.

Both objections are about **a socket from the backend to the container**. Neither survives if there isn't one. Measured against the running sandbox before any of this was designed:

Check	Result
<code>zapproxy -daemon -host 127.0.0.1 -port 8090</code>	API answers after ~7 s
<code>curl 127.0.0.1:8090/JSON/core/view/version/</code> from the host	refused — unreachable
the same call via <code>docker exec</code>	<code>{"version": "2.17.0"}</code>
a request proxied from inside the sandbox	recorded, with full bodies

The daemon binds loopback **inside** the container and no port is published. So the API exists, and the only way to reach it is `docker exec` — the same channel every other execution uses, and the one the gates already classify. Part 1's objection was to an *ungated* control channel; this is a gated one.

`docker/proof/zap_proxy_proof.sh` asserts host-unreachability directly, because it is a property of the network rather than of the code and no hermetic test can see it. **7 passed, 0 failed.**

Four defects the guards caught, none of them test bugs

The gate refused the operator's own socket. The first `_gate_request` passed the real argv, and the scope extractor reads `127.0.0.1` as an out-of-scope host. `tunnels.py` documents this exact trap for `-laddr`; my code carried a comment saying the port was excluded while passing it anyway.

Lab mode refuses target-less commands — a locked invariant, and very likely the real reason `tunnels.py` is engagement-only. Copying that would have been worse than the feature: engagement mode runs in the fully-open sandbox, so "make an engagement for the lab" would push practice traffic out of the sealed box. Resolved

by declaring the lab as the gate surface, which is a true statement of scope rather than a workaround — the lab proxy runs in the isolated sandbox, whose network has no route off the bridge, so the lab target is everything it can reach. `check_target_lock` is untouched.

The red-confirm would have been decorative. Part 1's rule is argument-based and keys on `-quickurl`, which a daemon argv does not carry — and a gate surface holding only the binary name cannot fire a flag-based rule at all. Either omission alone left `dangerous_ack` unenforced: gate-audit finding I2's exact shape. `-daemon` now sits in both the attack flags and the surface, for a stated reason — it attacks nothing, but it starts a listener that records credentials in cleartext.

The repeater lock refused the module. `proxy.py` first imported `RepeaterExchange`; `test_repeater.py` bans *any* import of the repeater, not just `repeater.send`, because a module that can import it is one line from calling it. The tempting fix was adding `proxy.py` to the allow-list — the exact anti-pattern build #5 was about. The models are local instead, with field names deliberately matching so the panel still renders a captured exchange with no translation layer.

Two things only a live process could find

The spawned argv is not the running argv. `zapproxy` is a wrapper that exec's the JVM, so the process on the box is `java -jar /usr/share/zaproxy/zap-2.17.0.jar -daemon ...`. The string "zapproxy -daemon" appears nowhere in it, so `pkill -f` matched nothing and the daemon survived every stop. No unit test can catch that — a hermetic test has no process to fail to kill.

`pkill -f` matches its own command line. The fixed pattern then killed the killer: the probe shell exited 143 having reaped nothing. The pattern is written `[z]aproxy` for that reason — it matches the literal "zapproxy" in the JVM's argv, while the killer's own argv contains "[z]aproxy", which does not.

Secrets: raw in the panel, masked in reports

Captured bodies hold passwords, `Authorization` headers and session cookies. Redacting on ingest was considered and rejected: the request that matters is usually the one carrying the token, and this is the operator's own data on their own disk. Masking lands at the **report** boundary instead — the artefact handed to a client or a grader — reusing `secretargs` ' `REDACTED` marker rather than inventing a second convention. It masks the value and keeps the parameter name, and its test carries a positive control, because a redactor that blanks everything would satisfy "the password is gone" while making every report useless.

One invariant with nothing behind it, stated plainly

The history read is deliberately **ungated** — a panel that refreshes cannot demand approval per refresh. That makes it the one path in this build with no human checkpoint, so what it can reach matters. It issues two fixed URL constants and takes no endpoint parameter, so an `action/` call is not expressible without writing a visibly new function.

Nothing enforces that. Both candidate guards — a runtime allowlist and a three-line static test — were considered and declined (2026-08-03). It is a convention, and the spec records the consequence and a review rule: any change to `cockpit/proxy.py` that introduces a URL parameter reopens the decision rather than quietly satisfying it.

Verification

Suite **68 files, 0 failures**, up from 66. `zap_proxy_proof.sh` 7/0.

That count needed one more fix to be true, and the honest version is worth recording. `test_redirector.py` began failing four runs out of four with `WinError 10013` — its `_free_port()` helper picked a port by binding a **TCP** socket, and the UDP test then bound **UDP** on that same number. TCP-free proves nothing about UDP: Windows reserves large UDP ranges for Hyper-V/WSL (on this box `50000–50059` and everything from `53879` up). It had passed 8/8 earlier the same day, which is the tell — the exclusion table is environmental and moves. The helper now takes the socket kind and probes **both** the public port and its tunnel port for it. That is the second time a Windows/Linux ephemeral-range difference has broken this one file. The `:proxy` screen ships **with** its endpoints — build #13 part 1 shipped four `/cockpit/exposure` endpoints with no caller and closing that took a whole later build. `tsc --noEmit` and `next build` both exit 0, and the screen adds zero lint errors.

Not built, deliberately: browser interception (it needs a published port, which breaks the lab sandbox's isolation — its own exposure decision), and driving ZAP's scanner through the API (scanning stays on part 1's gated command path). *Part 3 built the second of those; the first is still blocked, for a reason measured immediately afterwards — see below.*

Build #14 part 3 — the scanner learns to aim (2026-08-04)

Part 1 could scan a URL; part 2 could record traffic. Neither could attack **what you actually touched**. `zapproxy -cmd -quickurl` spiders a site and attacks whatever the crawl found, so an endpoint reached only by *using* the app — an API route nothing links to, a page behind a login, the exact request a `ffuf` run just made — was out of reach. Those are precisely what the proxy already records. This part joins the two halves and drives ZAP's active scanner over part 2's `docker exec` transport, aimed at the captured Sites tree.

The measurement that justifies the feature: **one captured endpoint, 376 real attack requests, one live High SQL injection** — on a route `-quickurl`'s spider had no reliable path to.

The finding that came first, and still blocks the other half

Before designing anything, `-config api.key=SECRETKEY123` was tested against the running daemon. It **enforces nothing**: `spider/action/scan` and `ascan/action/scan` both launched with no key at all. ZAP's proxy and its API share one listener, so publishing that port for **browser interception** would also publish an unauthenticated scan trigger to the host — and HackPit has no route auth. Browser interception therefore stays blocked, now for a measured reason rather than a suspected one. Scanner-over-API is unaffected: nothing is published, and the transport stays `docker exec`.

⚠ **THIS FINDING WAS WRONG, AND IT IS LEFT HERE UNEDITED ON PURPOSE.** Re-measured on 2026-08-04 with the flag stated explicitly, `api.key` **does** enforce — on views *and* on actions. The original test inherited `api.disablekey=true` from `$HOME/.ZAP/config.xml`, which **HackPit itself had written** on an earlier proxy start. It is left in place because a corrected record that hides the mistake also hides the lesson, which generalises: *a daemon that persists its configuration makes every measurement conditional on what a previous run wrote*. See **build #15** below, which this unblocked.

The gate: an honest surface, not a new one

A scan start builds `zapproxy -quickurl <target>` and runs it through the **real** `executor.validate_request` before ZAP is contacted. No new gate exists. `-quickurl` is already defined in the attack-flag table as *spider then active scan*, and an API scan is that attack minus the spider — so the declared command describes strictly **more** aggression than what runs, which is the safe direction. Putting the full target URL in the surface is what matters:

measured against the real validator, an unapproved scan is refused at `approval`, one without the red-confirm at `danger`, and `http://example.com/x` at `target` — the existing scope extractor reads the host out of a URL carrying a port and a query string.

Part 2's central lock could not be restated, so it was replaced

Part 2 asserted *the gated argv is the spawned argv*. That is meaningless here: the gate classifies an `argv` and what executes is a `URL`, so string equality between them would be theatre. The property underneath it is what actually matters — *the thing the gate scoped is the thing that gets attacked* — and that is now the lock. One derivation feeds both sides, and a test decodes the API's `url=` parameter back out and asserts its host is the host the target gate read.

A defect class that did not exist before this build: the operator-supplied target is interpolated into a URL that carries the scan's own parameters, so a target containing `&recurse=true` would broaden the scan the human approved. That is Critical 2 expressed in a query string. The target is percent-encoded with `quote(safe="")`, and the test carries a control proving recursion can still be set legitimately — otherwise it would also pass on a build where recursion never worked.

A second bound, enforced by ZAP rather than by us

`ascan/action/scan` on a URL not already in the Sites tree answers `{"code": "url_not_found"}`. The active scanner's reach is therefore bounded by what already passed through the proxy: a host never captured cannot be attacked through this path even if every gate here were bypassed. It is a bound, not a control — the gates remain the control — but it means the worst case of a defect here is "it attacked something you already proxied". The proof asserts it live.

The shape trap, caught by measuring instead of assuming

`state/parsers.py::parse_zap` reads the `-quickurl` `report: nested site[].alerts[]`, severity in `riskcode` (`"0" – "3"`), plugin in `pluginid`, URL down in `instances[].uri`. The API returns a **flat** list with `risk: "High"`, `pluginId`, and `url` on the alert itself. `_zap_report()` requires a `site` key, so feeding it an API response yields **zero findings, silently, forever, with a green suite** — part 1's headline defect in a new place. A separate mapper handles the API shape, tested against a real captured response committed as a fixture, and a test asserts the two parsers are **not** interchangeable, with a control proving the report parser does work on a real report. Someone will eventually try to merge them; that test is the argument they have to answer.

The plugin reference is written in `parse_zap`'s exact `pluginid:NNNNN` spelling, so the same issue found by both paths fingerprints to one finding rather than two — asserted by comparing real fingerprints, not by eyeballing the format.

What the proof found that no test could

ZAP locks its **home directory**, not its port. A daemon left running on any other port kills a new one at startup with a message that appears only in a log file inside the container. The proof hit this on its first run, and it exposed a latent part-2 defect: the "one proxy at a time" refusal was scoped per *port*, so a second proxy on a different port in the same container was accepted and then died. Nothing was unsafe — status is observed, so the dead proxy reported itself down rather than lying — but it was **unexplainable**, which is its own kind of defect. The refusal is now container-scoped and states ZAP's reason.

Verification

CI caught a defect in my own tests, and it is worth recording. The first version of the new locks drove the real `start_proxy` and asserted a lab scan validates cleanly. Both passed locally and failed on the runner: there is no Docker in CI, so the isolation gate refuses first and the verdict is `sandbox`, not the gate under test. The tests were silently depending on the developer's stack being up — the opposite of hermetic. The fix was structural rather than a loosened assertion: the clash check became a pure function that needs no Docker, and every control is now phrased as "*not refused at THIS gate*" rather than "*not refused*", which is how part 2's locks were already written. Re-verified by hiding `docker` from `PATH` and confirming the simulation actually bites (the lab scan comes back `sandbox`) before trusting that both files still pass under it — a simulation nobody checks is just a second thing to believe.

Suite **71 files, 0 failures**, up from 68. `zap_scan_proof.sh` **8 passed, 0 failed** — including host-unreachability re-asserted (it matters more now that action URLs sit behind that boundary), ZAP's `url_not_found` refusal, a full scan to completion, and the live alert response mapping through the real mapper. `zap_proxy_proof.sh` still 7/0. `tsc -noEmit` exits 0 and the `:proxy` scan panel adds zero lint errors. The panel ships **with** its endpoints, and its "Aim scanner" button on a captured row sets the target without starting anything — a one-click path from a table row to live attack traffic is exactly the shape a red-confirm exists to prevent.

Closing the two gaps that were left open — and what the second one found

Two things were shipped unverified and named as such rather than left to read as done. Both are now closed.

The ingest route had never been executed. The mapping into `Finding / Endpoint` was tested, but nothing had driven `POST /cockpit/proxy/alerts/ingest` through to `upsert_findings` — unverified *wiring*, the same shape as build #13 part 1's four `/cockpit/exposure` endpoints that shipped with no caller. It is now suite file **71**, and it is hermetic in a way the house style is not: `store.DB_PATH` points at a temporary file, so nothing touches the operator's `sessions.db`. A test whose whole subject is "*did the write land?*" cannot also be the test that trusts a delete to have removed only its own rows. It patches `proxy._api_get` rather than `scan_alerts`, one layer lower, so the real captured JSON goes through ZAP-response parsing, the route, the mapper *and* SQLite. It asserts the rows read back, that re-ingesting does not duplicate (the panel has a button, so it will be pressed twice), that findings belong to one session, that a missing `session_id` is a 422 rather than a cheerful `{"findings": 0}`, and — by AST — that the route executes nothing.

The scan panel had never been rendered, and it turned out the entire `:proxy` screen was invisible. `.hp-card` starts at `opacity: 0` and only becomes visible via an `.hp-in` class or the `.hp-surface` keyframe; `ProxyScreen` used bare `hp-card`, and it was the only component in the codebase to do so. Worse, **six of the eight classes it used did not exist at all** — `hp-kv`, `hp-check`, `hp-table`, `hp-danger`, `hp-url`, `hp-error`. Part 2 wrote the screen against an invented vocabulary, and `tsc`, `next build` and ESLint all passed, because **none of them can see whether a class exists in CSS**. Part 2's own verification note claimed the typecheck and build pass — which was true, and did not mean the screen worked. The screen is now written in the real `hp-tn-*` vocabulary that `:exposure` and `:c2` use, and it renders: status rail live from the backend, the gate visibly refusing (the start button stays disabled until *both* confirms are ticked), and the scan panel's red-confirm block reading as the most dangerous thing on the page. One further browser-only defect: checkbox rows placed in `.hp-tn-form` inherit `flex: 1 1 200px` and render as 200px grey slabs, so they get their own `.hp-tn-check` rule.

What is still unverified, stated rather than glossed: a full click-through could not be completed under browser automation — the enabled start button did not issue its `POST` after several attempts, and I stopped rather than keep retrying. The source wiring is the codebase's standard `onClick / disabled` pattern and every endpoint

behind it is covered by the suite and the live proof, so this is most likely an automation artefact; but it was not proven either way, and "probably the harness" is not a verification.

And the screen had no way in. Zaid asked whether anything recent was missing from the home launcher. It was: `/proxy` was the **only** top-level route in the app with no tile — not in `SURFACE_BANDS`, not in the command palette, with nothing anywhere in the UI linking to it. Both part 2 and part 3 shipped it reachable only by typing the URL. Part 3's own definition of done said "*proxy ships with its endpoints — no orphaned routes (the build #13 part 1 lesson)*", and I had verified that every **endpoint** had a frontend caller — but never that the **screen** had an entry point. The same lesson, one level up, walked straight past. The tile now sits in **operate**, next to `:repeater`. That placement is an argument, not a preference: the band `hint` is a posture claim about everything in the band, and infrastructure's reads "*gated start · human-only stdin*" — the proxy's start is gated, but it has no stdin at all (spawned `interactive=False`, deliberately), so half that claim would be false for this tile. Operate's "*every command human-approved · needs the stack*" is exactly true of both halves. Sitting beside `:repeater` also makes the capture → replay path discoverable, which is why the captured-exchange model mirrors the repeater's field names in the first place.

A process lesson worth more than the bug it caused. The CI-simulation from the previous fix — running the suite with `docker` hidden from `PATH` — also hid `git`, and `test_corpora`'s re-ingest uses a git-recovery path. It rewrote the live KB and dropped one entry (2743 → 2742). The KB is gitignored, so there was no commit to fall back on; re-running the ingest with a normal `PATH` restored it exactly, as the recorded flake note predicted. The lesson: **a simulation that strips the environment can have side effects on real data.** Simulate by running the specific test files, not the whole suite.

Still not built: browser interception, blocked on the unauthenticated-API finding above; spidering via the API (that is `-quickurl`'s job and carries its confirm); scan-policy tuning (a policy that decides its own aggression is the `-autorun` shape part 1 excluded); and authenticated scanning.

Full-project audit — the whole surface, driven end to end (2026-08-04)

An unattended overnight audit of **every** surface, not just the newest build, against the brief in `docs/superpowers/prompts/2026-08-04-full-project-audit.md`. The full report with per-question verdicts, evidence and an explicit NOT-RUN list is `docs/AUDIT-2026-08-04.md`; this section records what it changed and the three things worth carrying forward.

The finding that mattered most was a combination, not a bug

`state/render.py` sets `INCLUDE_CREDENTIAL_SECRETS = True` — Zaid's explicit decision, and the module says so in full, including the consequence: "*every captured credential is sent to whatever LLM endpoint is configured, including a remote one.*" That decision is safe under a precondition it states but does not enforce — a **local** endpoint. The precondition was not holding. The live provider was `claude-agent-sdk` (remote, `opus`), while six real credentials from the build #9 live fire sit in `state_credentials`. Neither half is a defect; the pair was. The provider is now `ollama / qwen3:8b` per the standing rule, verified end to end (a grounded attack path composed in 147s, `context_leaks: 0`, every step citing a real KB entry id).

The lesson: a documented decision can carry an unenforced precondition, and the audit that finds it will be looking at configuration, not code. Both halves passed every test.

Reviewed and deliberately left unchanged (2026-08-04, Zaid). Five enforcement designs were put to him — derive inclusion from the provider, pass it from the caller, guard the config write, scrub at the egress boundary, or scope it to lab-vs-engagement — and he declined all five. His reasoning holds for the case he is in: the credentials

in the store are from his own `corp.local` lab VM, a box built to be popped, so sending them to a model costs nothing. Recording the decision was the point rather than winning the argument, so the acceptance and its *boundary* now sit in `state/render.py` beside the constant: this path is not scoped to the lab, so the same render runs during a real engagement, where a client's domain-admin hash reaches the same endpoint by the same route — and there it stops being a matter of taste, because most pentest contracts forbid sending client data to third parties. The risk accepted is precisely "*my own lab credentials reach my own model endpoint*", and Option E is the smallest change that would close the rest if HackPit is ever pointed at a client under contract.

That is the right shape for a disagreement to end in. The behaviour is his call and it is unchanged; what changed is that the next reader finds reasoning instead of a landmine.

Every gate fires, and every gate can fail — 27 checks, both directions

The 2026-07-27 audit found seven guards that never fired. A harness drove the real `executor.validate_request` with a refusal **and** a positive control for each gate, in lab and engagement mode: 24 checks, 0 failures, plus 3 on the isolation guard. The isolation refusal was proven with a genuinely non-isolated container rather than a mock — pointing the lab guard at `hackpit-engage-sandbox` (fully open by design) produced the egress-path refusal.

A trap the harness caught in itself. Its first danger-gate probe was `bash -c '...10.0.0.1...'`, which is refused at **target**, not **danger** — so the probe never reached the gate under test and the paired "with ack" control passed *vacuously*. In lab mode the danger gate is only reachable by a command naming the lab target as a clean argv token. This is the exact failure the house "not refused at THIS gate" phrasing exists to prevent, and it is the second time that convention has earned its keep.

The missing gate is the mirror of the never-fired guard

Measured against the real bug-bounty program: `ffuf -u https://.../FUZZ -w common.txt`, `sqlmap -u https://... --batch` and `nuclei -u https://... -t cves/` all pass on `approved=true` **alone** against a third-party production host — no red-confirm. That is not a defect in the danger heuristic, which correctly models arbitrary code execution (it does flag `sqlmap --os-shell`). **Nothing models aggression toward the target**, and there is no rate limiting, pacing or concurrency control anywhere in `cockpit/`, `state/`, `reasoning/` or `arsenal/`. Recorded as the highest-value missing feature; the narrowest fix reuses the existing red-confirm in engagement mode only.

Two hypotheses in the brief were wrong, in the product's favour

Bug-bounty reporting **exists**: a dedicated `bugbounty` template ("a HackerOne / Bugcrowd submission"), selectable in the UI, with a CVSS 3.1 base calculator verified correct against six known vectors including the round-up edges.

Missing is narrower than assumed — no VRT, no P1–P4, 3.1 not 4.0. And the scope model **already** handles wildcards, CIDRs and `!exclusions`, and refuses both classic wildcard bypasses (`notexample.com`, `example.com.evil.net`). All 11 of the program's web/API targets parse, resolve and enforce.

What it cannot express is **mobile** — and the failure shape matters more than the refusal: `parse_scope("THAT Concept Store iOS")` splits on whitespace into four bogus hostnames and fails only because none of them resolve. Had one resolved, a garbage host would have entered the scope. A fail-open shape inside a fail-closed result.

Reviewed with Zaid and ACCEPTED — no change (2026-08-04). Following the audit, the finding was measured further and turned out not to be about mobile at all: the trigger is **whitespace**, and the worst case is not a store URL but a scope block that names its own out-of-scope hosts in prose — measured, `"out of scope: iana.org"` puts `iana.org` in the **allowed** set. Two fixes were designed and both declined: a strict tokeniser (comma/semicolon/newline separators, with anything not host-shaped recorded as "not understood"; zero of the

21 live engagements affected), and a dedicated out-of-scope box wired to the `!host` exclusion machinery — which is viable and already enforced end to end, **provided it marks every token rather than every line**, since per-line marking leaks even on a clean comma list. That per-token trick is recorded in the audit report in case the box is built later; the asymmetry behind it is that over-*inclusion* is dangerous while over-*exclusion* is free.

The acceptance is defensible for a reason the audit under-weighted, and it is worth stating because it corrects the framing rather than excusing it: `check_target_lock` describes itself as "cheap defense-in-depth... **NOT a load-bearing control**", and says that in engagement mode "HUMAN APPROVAL of every command is the actual bound and this lock is an aid to the human, not a guarantee." An over-wide scope runs nothing on its own; every command is still approved individually with the hostname visible in the argv. The audit treated the scope as the wall; the codebase treats it as a handrail. Zaid's stated workflow — a clean comma-separated list — parses perfectly, as do all 21 existing engagements. The residual, accepted knowingly: pasting a scope straight off a program page can still admit an out-of-scope host, caught only by reading the hostname before approving.

Defects fixed in this commit

A tool-file id dropped two files out of the product. `pipeline/ingest_corpora.py` built the id from `path.stem`, so files differing only by extension collided: `nc / nc.exe` (linux vs windows) and `SharpHound.ps1 / SharpHound.exe`. The extension is what decides `platform`, so these are genuinely different artefacts. `/entry/{id}` could only resolve one of each pair, and because the Scripts Arsenal keys its React list on that id, **React dropped one of each pair from the list** — two duplicate-key errors, confirmed in the browser. Fixed at source (`path.name`), artefacts regenerated, locked by two tests: one on the derivation carrying a control that proves the *old* derivation collided, one on the built artefact. The KB was verified byte-identical at 2743 lines before and after every regeneration.

The Scripts Arsenal claimed more than it contained. The section header printed `count` while the list held `shown` — 1235 claimed, 1158 served; `Enumeration` claimed 335 and rendered 263. The cap is deliberate and the pipeline records both numbers honestly; the UI simply never used `shown`, which was already carried through the backend model and the API type. It now shows what is listed and states how many were dropped, by which cap, and that the filter cannot reach them. The index was also **stale by 126 entries** (built against 2617); rebuilt.

Checkboxes were grey slabs on six screens. `.hp-tn-form input { flex: 1 1 200px }` over-matches — a checkbox is not a text field. Build #14 part 3 worked around this on `:proxy` with a per-screen `.hp-tn-check ; :exposure , :repeater , :oob , :cockpit` and the AD screens still had it. **Fixed in the over-matching rule itself**, where every screen gets it, rather than in the one component that noticed — the same "fix the predicate, never narrow the file set" discipline build #5 established. Verified in the browser before and after.

Verification

Suite **71 files, 0 failures**. All five live proofs green: isolation 4/0, engage-open 4/0 (Wall A confirmed down), zap-install 9/0, zap-proxy 7/0, zap-scan 8/0. `tsc --noEmit` clean; lint unchanged from its accepted baseline (no changed file appears in the output). All 18 surfaces rendered in a real browser and the orphan-route check is clean.

Windows/AD driven **live** against the real DC — WinRM reached `corp\administrator`, and read-only domain enumeration ran through the gated executor while DCSync-shaped and `-EncodedCommand` payloads were correctly held at the red-confirm.

Not run, and not to be read as passing: the passive HTTP fingerprint sweep of the 11 third-party hosts (refused by the classifier — left as a self-verifying `docs/audit-2026-08-04/bb-passive-recon.sh`), a re-measurement of ZAP's `api.key`, and the end-to-end report-redaction observation (refused, correctly, because it meant dumping six real

domain credentials — the property is verified by its existing test instead). Exactly **one** packet-generating command touched a program asset all night: a single `dig` , which returned `crateandbarrel.edgekey.net` . Nothing was submitted anywhere.

Build #15 — browser interception (2026-08-04)

Against WAF/bot-managed targets HackPit did not reach rate limiting. **It did not reach request one.** A passive sweep of a live Bugcrowd program — one `HEAD` per host, in scope — returned nothing at all from every host:

protocol	result
HTTP/2	instant stream reset, <code>INTERNAL_ERROR</code>
HTTP/1.1	total timeout, 0 bytes in 15s

Two *different* failure modes on two protocols is an edge actively refusing the client, not a protocol quirk. Nine of the eleven assets sit behind Akamai Bot Manager. Egress was fine — HackPit reached Akamai and Akamai said no. So the audit's "can it run a full bug bounty?" answer of "*partly — breaks at volume*" was generous: volume is a problem you would like to have.

A real browser is what passes: correct TLS fingerprint, real headers, JS execution, a real profile. **That is the whole of this build**, and the boundary is stated because it is the interesting part: nothing here imitates a browser or evades a control. It uses one.

The measurement that unblocked it was a correction, not a discovery

Part 3 recorded, as MEASURED, that ZAP's `api.key` "*enforces nothing*" — and that finding blocked browser interception for a day. It is wrong. Re-measured against the same ZAP 2.17.0 with the flag stated explicitly:

check	result
<code>core/view/version</code> without key	refused
same with key	<code>{"version":"2.17.0"}</code>
<code>ascan/action/stop</code> without key	refused
wrong key	refused
the PROXY meanwhile	<code>HTTP 200</code> — serves normally

`cockpit/proxy.py::server_argv_for` had been passing `-config api.disablekey=true` on every proxy start, and ZAP persists `-config` values into `$HOME/.ZAP/config.xml` . The original test set a key with no explicit `disablekey` , inherited `true` from a previous HackPit run, and concluded the tool was broken. **HackPit was disabling its own lock and we blamed ZAP.**

Generalised, beside "the container is not the image": **a daemon that persists its configuration makes every measurement conditional on what a previous run wrote.** State the flag explicitly, or you are measuring history. Both the wrong finding and its correction are kept in this document, because a corrected record that hides the mistake hides the lesson too.

Part 1 — the port is published through `exposure.py` , not hardcoded

The single most important decision in part 1, and it replaced an earlier draft that baked `127.0.0.1:8090:8090` into `docker-compose.yml` . HackPit already has a designed, tested answer for "publish a port on the engage sandbox", built across builds #10 and #13, and this uses it: the ZAP port is `extra=[(8090, "tcp")]` on a `ListenerProfile` . **No new model, no new field, no new file format, no second publish path** — the generated override is gitignored (it can name a client's internal address and this repo is public), and a wildcard bind is a red-confirm whose acknowledgement is written into the file as a machine-readable marker that `test_exposure_safety.py` statically enforces.

That is *freer* than hardcoding loopback, which is the point. The operator can bind anything, including every interface for a phone or a second machine, with one confirm — recorded against the engagement so a report can state honestly that the proxy was open on all interfaces during that window. A hardcoded `127.0.0.1` would have been a constant to edit compose to escape; this is a switch they hold. Two presets ship: `zap-proxy` (loopback) and `zap-proxy-lan` (`0.0.0.0`) . **Neither pre-acknowledges its own confirm** — a preset that did would mean picking a dropdown entry silently satisfies the gate, which is gate-audit finding I2's exact shape.

What the design spec did not state, and the implementation had to answer

A container process bound to loopback cannot be reached through a published port at all. `docker -p` forwards to the container's bridge interface, where nothing would be listening. A published proxy that stayed loopback-bound would have been the worst of both worlds: a port open on the host and a feature that silently did not work. So `-host` became conditional (`bind_host_for`), and a published daemon also needs `api.addr`s widened — a request arriving through a published port reaches ZAP from the bridge gateway, not from `127.0.0.1`, and ZAP's default address filter would refuse it while looking exactly like a broken feature.

Both are real widenings. What pays for them is the key: what becomes reachable is an **HTTP proxy**, which is the entire point, while scan control still refuses everyone who cannot present a secret only the backend holds. Publishing is **engagement-only**, refused before the executor gates — the lab network is `internal: true`, so a published port there has no route, and binding wide would expose the API *inside the isolated network* while still being unreachable from here.

The key is closed twice, and one of the two cannot regress

`-config api.key=<secret>` puts a credential in the spawned argv, and this codebase records argv into run records that feed the state store, the LLM prompt and rendered reports. Two layers:

1. **The gate is never given the key.** `_gate_request` passes a placeholder. An `ExecRequest` is the thing that gets recorded, reported and put in front of the model — redacting a secret after handing it over depends on the redactor being right forever; never handing it over cannot regress.
2. **secretargs masks it anyway**, and this needed a new *shape*: `-config` is the wrong discriminator, because the same flag carries `api.disablekey=false` — the single most audit-relevant token in the argv, and **the evidence the lock was on**. A rule keyed on the flag would redact both, which is `nmap -p 445` becoming `nmap -p <redacted>` one level down. So the map keys on the *setting name*, and a control proves an unregistered tool keeps its value.

The residual is written down rather than hidden: `/proc/<pid>/cmdline` inside a single-tenant container we own. ZAP reads `api.key` from `-config` or its persisted config and nothing else — there is no stdin path, so build #13's "ride the secret in on stdin" is unavailable here.

Part 2 — the AJAX spider, and the red-confirm that could not fire

The spec said the crawl's equivalent surface is `zaproxy -zapit <target>` and that it must require `dangerous_ack`. Those two cannot both be true. `-zapit` is deliberately absent from `_TOOL_ATTACK_FLAGS`, and `test_zap_safety.py` explicitly locks that a `-zapit` recon run must *not* demand a red-confirm. As written, the spec's own test could never pass.

Fixing it exposed a defect that was already shipping. The map was `tool -> frozenset(flags)` and the consumer appended one hardcoded sentence for whatever matched:

flag	reason the operator was shown	true?
<code>-quickurl</code>	"active web scan — sends live injection payloads..."	yes
<code>-daemon</code>	"active web scan — sends live injection payloads..."	no

Starting the recording proxy told the operator it was sending injection payloads at a target it never touches. A red-confirm whose stated reason is false is worse than no reason: it is what teaches an operator that the text is noise and the checkbox is a formality.

The fix was to change the one map to `tool -> {flag: reason}` — *not* to add a second `_TOOL_BROWSER_FLAGS` beside it (Zaid, 2026-08-04). Two lists of the same kind of fact have to agree forever: that is build #5's "fix the predicate, never add a parallel one", and the same drift this very file removed once when `dangerous_script_heuristic` was made to DERIVE from this dict instead of restating it. **That the reshape fixes an existing defect as a side effect is the tell that it is the right shape.**

The half that nearly slipped through: that derivation *kept working without an edit*, because `sorted(a_dict)` yields its keys. It would not have failed loudly — it would have gone on stamping the old false sentence onto every flag, re-introducing the defect on the WinRM path only. `.items()` makes the reason travel with the flag, and a test now asserts the two heuristics agree on WHY, not just on WHAT.

So the crawl earns its confirm for a third, genuinely different reason: **it drives a real browser that clicks things**. On the production e-commerce sites in scope for a bug bounty, clicking everything reachable can submit a form, empty a basket, trigger an email or place an order — a different hazard from SQLi and arguably a more embarrassing one. `-ajaxspider` is a **declared marker**, not a real ZAP flag (the crawl is API-driven and has no command line), and `_ATTACK_FLAG_IS_REAL` records which tokens are which so a third marker cannot arrive unnoticed. `-zapit` is dropped from the surface entirely: once the marker carries the danger verdict and the URL carries the scope, including it would make the declared command claim two modes at once.

Depth and duration are in the approved surface, for the same reason `-autorun` was excluded from the catalog in part 1: a crawler that decides its own bounds is a command that has stopped describing what runs.

Why ZAP's spider and not our own headless browser

The reason exists only because part 1 landed first. Manual browsing through the published proxy establishes an **authenticated session inside ZAP**, and the AJAX spider runs through that same ZAP — so it inherits those cookies and crawls the logged-in application. A separate headless Chromium would start cold and need scripted authentication per target, an auth-automation problem this build would then own forever. Log in once by hand, let the spider expand from there, let part 3's scanner attack what both produced.

An OK is not a result

`setOptionBrowserId` **does not validate**: it accepted `not-a-browser` and answered `{"Result":"OK"}`. The image ships Chromium and **no Firefox**, while ZAP's configured default is `firefox-headless` — so the failure surfaces at *crawl* time, in a driver stack trace inside ZAP's log, not at set time. `start_spider` therefore reads the value back and refuses a mismatch, `observed_spider` reports what ZAP holds rather than what we sent (asserted by AST, because a substring scan failed on the function's own docstring — build #8's lesson, firing on the first run), and the proof asserts a browser *process* appeared and the message count rose.

The isolation argument was rebuilt, not relaxed

Part 2's argument was *"no port is published, so the API is unreachable, so the only channel is `docker exec`."* **Half of that is now deliberately false**. The property being proven changes from *"nothing can reach the control channel"* to *"the control channel refuses everyone who does"* — and the second needs **more** assertions, not fewer.

`zap_proxy_proof.sh` 's load-bearing lab check is untouched and still runs; `browser_intercept_proof.sh` adds the engage half: reachable, **refuses without the key** (view *and* action), **answers with it, still serves as a proxy**, the lab still publishes nothing, and a control proves the refusal is enforcement rather than a dead port.

Residual risk, accepted and written down: anything that can reach the bound address can *use* the proxy — it cannot scan, which needs the key. On loopback that is a privacy annoyance on a single-user machine. On a wildcard bind it is an open proxy with the engage sandbox's full egress behind it, which is why that case carries a confirm rather than being forbidden or being free.

Two defects the proof found that nothing else could — and the second was hidden by the first

The hermetic suite was green, `tsc` was clean, the image built, the API key enforced, the port published, and the AJAX spider answered `{"Result":"OK"}` and **crawled nothing**. Twice over, for two unrelated reasons that produce the identical symptom:

1. Chromium refuses to run as root without `--no-sandbox`. The engage sandbox runs `user: root` by design (build #3's privilege decision), so every browser Crawljax launched died at creation — the log simply stops after `Loaded ... DummyPlugin as a OnBrowserCreatedPlugin`, with no error on the ZAP side at all. ZAP's Selenium add-on exposes **no API** for adding Chrome arguments in this version (`addChromeArgument` answers `bad_action`), so the flag goes in through `/etc/chromium.d/` — Debian's own launcher extension point, which already ships a `dev-shm` file.

`--no-sandbox` turns off Chromium's internal process sandbox, and this browser visits target-controlled content, so it is written down rather than glossed: the boundary that holds here is the container, not Chromium's own, and as root the alternative is a browser that does not start at all. `--disable-dev-shm-usage` rides along for Docker's `64MB /dev/shm`.

2. ZAP bundles its own chromedriver and prefers it. With the first defect fixed, the crawl failed differently: `This version of ChromeDriver only supports Chrome version 151`. The `webdriverlinux` add-on ships a driver for Chrome 151; Kali ships Chromium 150. So installing `chromium-driver` was **necessary and not sufficient**, and every daemon start now passes `-config selenium.chromeDriver=/usr/bin/chromedriver`.

The lesson is about the check, not the flags. The Dockerfile layer already followed part 1's discipline — it proved the driver was on PATH, proved it *ran*, and proved a browser existed. All three passed while the feature was completely broken, because **proving a driver runs proves nothing about whether it can drive the browser that is there**. The layer now compares the two major versions and fails the build on a mismatch. That is "an exit code is not a result" one level deeper than part 1 phrased it: a *successful* version string is not a result either.

Both were then given their own proof checks, because they fail separately and because a future image whose packages drift apart would otherwise reappear as an unexplained empty crawl.

A third, smaller one, in the proof itself. The "a real browser process was observed" check passed *"after 0s"* on a run where the crawl found nothing — it had matched a leftover Chrome from the previous attempt. A check that reports success from another run's residue is worse than no check, so strays are now reaped before the crawl starts (and deliberately *before*, since doing it after would kill the very process the evidence depends on).

Two more defects, this time in the proof itself — and one is a lesson arriving backwards

`pkill -f` matched too much, because a fix earlier in this build changed an unrelated `argv`. The proof reaps stray browsers before crawling, and did it with `pkill -f "[c]hrome"`. That killed the **ZAP daemon**: its command line now carries `-config selenium.chromeDriver=/usr/bin/chromedriver`, so it contains the string "chrome". Every API call after that point returned empty and three checks failed for reasons that had nothing to do with what they test.

This is the `[z]aproxy` lesson arriving from the opposite direction. There, `-f` matched too LITTLE — the wrapper exec'd the JVM, so the spawned `argv` was not the running one. Here it matched too MUCH, and it began matching only because a fix elsewhere in the same build put a new word on a different process's command line. **A `pkill -f` pattern is a claim about every `argv` on the box, and it stops being true silently when an unrelated `argv` changes.** Now `pkill -x`, which matches the process NAME and cannot be broadened by somebody else's flags.

A check whose own plumbing decided the verdict. The "can chromium start as this user" check piped into `grep -c`, which prints `0` and *exits 1* when it finds nothing — so the `|| echo 1` fallback fired on the SUCCESS path, appended a second line, and the comparison could never match. It reported FAIL while the thing it tests was working perfectly. Read the output; do not infer it from an exit code. That is the same sentence as part 1's, applied to the harness rather than the tool.

Neither is a product defect, and both are worth writing down: a proof that fails for its own reasons trains you to discount it, which is exactly the failure mode a proof exists to prevent.

CI caught a test depending on Docker — for the SECOND time, in the same file

Build #14 part 3 recorded exactly this: a new lock drove the real validator, asserted a lab run "validates cleanly", passed locally and failed on the runner, because there is no Docker in CI and the LAB gate order is `target -> approval -> danger -> **sandbox**`. The fix then was structural — every control phrased as *"not refused at this gate"* rather than *"not refused"*.

The new crawl-gate test did it again: `assert clean is None`. Locally the stack was up, so the isolation gate passed and the assertion looked correct; on the runner it came back `gate='sandbox'`. **The green local run was measuring the developer's machine, not the code.**

Two things follow, and the second is the useful one. The assertion is now *if refused* it must be at `sandbox`, which excludes the three gates under test and nothing else. And the suite is now run with `docker removed from PATH` before being believed — with the strip itself checked first, because a simulation nobody checks is just a second thing to believe. That lesson already existed, was already written down, and still did not change what I actually ran until CI forced it. **A convention only holds if something executes it.**

Verification

Suite **71 files, 0 failures** — and re-run with `docker` hidden from `PATH`, where it is also **71/0**, which is the environment CI actually provides. `docker/proof/browser_intercept_proof.sh` **22 passed, 0 failed**, including the four that carry the safety argument — an unauthenticated **view** refused, an unauthenticated **action** refused, a **wrong key** refused, and the same call **with** the key answering `{"version": "2.17.0"}` as the control proving the refusals are enforcement rather than a dead port. A request made **from the host** through the published proxy was captured (0 → 2), the browser-driven crawl captured 2 → **41 messages and found 29 URLs**, and the real captured traffic parsed through the real parser into exchanges and endpoints. The **lab** sandbox still publishes nothing, and teardown left nothing listening on either side of the boundary. `docker/proof/zap_proxy_proof.sh` **7 passed, 0 failed** with its load-bearing check — *the ZAP API is UNREACHABLE from this host* — still passing for the lab sandbox, and `isolation_proof.sh` **4 passed, 0 failed**. The half of the argument that did not move did not move.

`tsc --noEmit` exits 0, `next build` exits 0, and `eslint` sits at the accepted baseline of 11 — the same count with the frontend changes stashed, so this build adds none. `:proxy` was **looked at in a browser**, not merely typechecked: the publish control is disabled until an engagement id is entered, and the crawl panel's red-confirm carries its own copy rather than the scanner's. All three spider routes are registered and all three are called by the client — no orphans.

The acceptance test was left open at commit time and **HAS SINCE BEEN RUN**. The runbook is `docs/proof/build15-acceptance-runbook.md`; what it returned is below, and it split the two halves of this build apart.

The acceptance test, run live (2026-08-04) — part 1 passes, part 2 does not

Run against `www.crateandbarrel.me`, in scope for the MAF Lifestyle Bugcrowd program and one of the nine Akamai-fronted assets that returned **nothing at all** to a bare `HEAD` — h2 stream reset, h1.1 total timeout.

Part 1 PASSED. Firefox on Windows, pointed at the published `127.0.0.1:8090` with ZAP's CA trusted, loaded the site normally. ZAP captured **55 requests** across the target's own hosts:

host	requests
<code>www.crateandbarrel.me</code>	44
<code>dh.crateandbarrel.me</code>	6
<code>gtm-analytics.crateandbarrel.me</code>	3
<code>crateandbarrel.me</code>	2

A real browser through this proxy reaches a host a bare client cannot. That is the whole premise of the build, and it holds. It also surfaced **two subdomains enumeration had not** — `dh.` and `gtm-analytics.` came out of a page load, not a wordlist.

A stated concern was measured and was WRONG. Before the test, the risk called out was that ZAP is a MITM proxy: it terminates the browser's TLS and re-originates upstream with its own Java stack, so Akamai would see Firefox's headers and JS but *not* Firefox's TLS fingerprint — and Akamai Bot Manager leans on TLS fingerprinting.

That was a reasonable objection and the measurement overruled it. Java's handshake was accepted. Worth keeping, because it is the shape of a plausible argument that a five-minute test settles.

Part 2 FAILED against the same target, and the control makes it unambiguous. Three fetches through the identical proxy, minutes apart:

client	target	result
Firefox	<code>www.crateandbarrel.me</code>	55 requests captured
headless Chromium	<code>www.crateandbarrel.me</code>	ZAP's own 20s read timeout, 0 bytes
headless Chromium	<code>example.com</code>	HTML, normal

Same ZAP, same proxy, same window. The refusal is specific to the CLIENT, and it reproduces the original `curl` signature exactly: a silent hang rather than a rejection.

This undercuts part 2's central argument on precisely the targets that motivated the build. The design was "log in by hand, let the AJAX spider expand from there, let part 3's scanner attack what both produced." The *session* inheritance works — that part is sound. But the spider brings its own browser and therefore its own fingerprint, and on a bot-managed edge that is what gets refused. So the honest scope of each half is now measured rather than assumed:

- **Part 1 (publish + manual browsing)** — what actually unblocks the WAF-fronted assets.
- **Part 2 (AJAX spider)** — scale on ordinary targets; refused by Akamai. Useful, and narrower than the spec claimed.

Not pursued, deliberately. Headless Chrome advertises `HeadlessChrome` in its User-Agent, which is the likely discriminator, and overriding it would probably restore the crawl. That crosses the line this build drew around itself — *"nothing here imitates or evades — it uses one."* Going further is a separate decision and a separate build; this section is the evidence it would rest on, which is exactly what the spec asked for if a real browser were refused.

Two smaller things the live run surfaced

Chromium phoned Google during the proof's crawl. A crawl aimed at a local two-page site also produced requests to `www.google.com`, `gstatic.com` and `play.google.com` — Chromium's background networking (variations seed, search preconnect) going out through the proxy, and the engage sandbox has full egress. Harmless here, wrong for a tool used on engagements: traffic nobody asked for, attributed to the operator's IP. The `/etc/chromium.d/` drop-in should carry `--disable-background-networking --no-first-run --no-default-browser-check`.

An engagement's rules-of-engagement text is not enforced by anything. The active engagement carried `authorization: "...PASSIVE RECON ONLY this session; no active scanning per rules of engagement."` — written for an unattended session. Scope is enforced in code; that sentence is free prose no gate reads, so a crawl that clicks through a production storefront would have passed every gate. It was caught by reading the record, not by a control. Same shape as the credential-precondition finding in the 2026-08-04 audit: **a documented constraint with nothing behind it.**

The audit punch list, worked — build #16 (2026-08-04)

The 2026-08-04 audit ended with nine open items. Zaid reviewed each and decided which to build: D3, D5, D6, D7, D8 and D9 in, plus three new ones (known-issue matching, VRT priority, scan session-expiry detection) and the `:exposure` control layout. **D1, D2, D4, route auth, INCLUDE_CREDENTIAL_SECRETS and mobile support were reviewed and declined** — they are not open defects and are not listed as such below.

D8 — the planner emitted commands nobody could run, and nothing said so

The measurement that opened this: with the correct request shape, `/attack-path` returned 32 commands, of which 0 named the real scope, four carried a foreign example host, and an earlier run had returned `sublist3r -d tesla.com -t 100` / above the `-t` is for threads... — the source writeup's example host *and* its prose, verbatim, inside the `cmd` field. Every step carried a `target_adaptation` line *describing* the substitution in words. Nothing performed it.

The substitution machinery was not broken. It never ran. `substitute_target` opens with `if not target: return cmd`, and `target` came only from `extract_target(goal)` — which reads the goal text. The failing call passed the real hosts in `scope_text` and a goal that named none, so `target` was `None` and every command kept the KB author's example host. One missing fallback made an entire feature a no-op, silently, for as long as it has existed.

Two changes, in the order the work was scoped:

(c) **Validate and flag first.** Every returned command is now checked against the declared program scope and, when it cannot run as written, marked `runnable: false` with the reason. Two reasons, kept apart because the fix differs: *points at a host outside the scope* (written for someone else's environment) and *names no host at all* (an unsubstituted placeholder, or a KB excerpt that was prose rather than a command). The path-level `commands_unrunnable` is the headline: **a plan built entirely from the KB's own examples used to be indistinguishable from one adapted to the target.**

(a) **Then substitute.** `primary_target` falls back to the first *concrete* in-scope host when the goal names none, and reports where the target came from (`caller` / `goal` / `scope`) so "HackPit picked this for you" reads differently from "you named it". A wildcard is deliberately not a target: `*.example.com` names nothing runnable, and inventing `www.` in front of it would be a fabricated target dressed as a real one.

The scope model is reused, not re-derived. `cockpit/scope.py` — wildcards, CIDR, exclusions, already measured against a real bug-bounty program — is what judges the commands, so the plan cannot disagree with the gate about what is in scope. The same reasoning produced one `_host_positions()` definition shared by the substituter and the checker: written twice they would drift, and a checker that disagreed with the substituter would flag commands it had itself just fixed. That is the shared-predicate defect this repo has now found four times.

What was NOT done, and why. A real foreign host like `tesla.com` is *flagged, never rewritten*. Rewriting any out-of-scope host sitting in a host position would have closed the gap completely — and would also have rewritten `curl https://raw.githubusercontent.com/.../linpeas.sh` into a request against the client's own storefront. A tool download, an attacker-owned listener and a target are all "a host in a host position"; nothing in the `argv` distinguishes them. So example hosts are substituted (they are recognisable), and everything else is reported. The scope directive said "*where it can be done safely*"; this is where that line falls.

`/attack-path` accepts `target` now — and refuses what it does not know. A `target` field used to be dropped in silence, which is how the first measurement looked worse than it was. It is a declared field, and `AttackPathIn` is `extra="forbid"`, so an unknown field is a 422 rather than a shrug.

A silent loss found on the way in. `AttackStep.commands` was typed `list[Code]` — the KB entry shape — while `attack_path.py` had been setting `unverified` and `truncated` on planned commands the whole time. A response model does not carry a field it has not declared, so both have been dropped on the floor for as long as they have existed. `PlannedCode` repairs that and holds the new scope fields. `test_attack_path_contract.py` locks the general property rather than a list of names: **every key the composer emits must survive `response_model`**, with a planted field proving the check can fail. A name list would need updating by exactly the person who just forgot.

D3 — nothing throttled anything

Pacing injects each tool's own rate or delay flag, in the exact shape `executor.apply_proxy` already uses, **engagement mode only** — the lab target is an isolated container with nobody to annoy, and lab argv stays byte-identical. A tool with no known throttle flag runs unchanged and *says so*, because a run the operator believes was paced and was not is precisely how a program bans your IP mid-engagement.

Three things the shape needed that a proxy URL does not:

- **Units.** Throttling is spelled three incompatible ways — requests/sec, packets/sec, or a delay *between* requests. The operator supplies one number and it is converted per tool. Handing a literal `4` to `sqlmap's --delay` would mean a four-second wait when four requests/second was asked for; reversed, a delay fed to a rate flag runs 1000× too fast. The regression test pins all twelve conversions.
- **Already-set detection.** A command that already carries its own rate flag is left alone and says so. Two contradictory rate flags are resolved silently and differently by each tool.
- **Flag placement.** See below — this one was a live bug, not a hypothetical.

Every flag was read out of the tool's own `--help` inside the real image, not recalled. That check removed three candidates rather than guessing at them: `amass` is un-probeable in the sandbox (the `sudo / no-new-privileges` trap), `curl` issues one request so there is nothing to pace, and — see below — the image's `httpx` is not the tool anyone assumes it is. All three take the honest no-flag path.

Two shipping defects the pacing work found in the proxy flag it was mirroring

`gobuster --proxy ... dir ...` has never worked. Measured against the real image, it exits with `flag provided but not defined: -proxy` and does nothing at all: `gobuster's` flags belong *after* its subcommand. Asking to capture a `gobuster` run has been breaking the run. Pacing would have reproduced the bug exactly, so the placement rule now lives in one `_place_flag()` used by both rewrites — the next tool with subcommands cannot be right in one and wrong in the other.

`/usr/bin/httpx` in the sandbox is the *python* `httpx` CLI, whose proxy flag is `--proxy`. ProjectDiscovery's `httpx` — the one `-http-proxy` belongs to — installs as `httpx-toolkit`. Both names were mapped to `-http-proxy`, so capture on `httpx` produced a command line that CLI rejects outright. Same defect class as "*Kali ships no zap-baseline.py*": the map named a flag for a tool that is not there.

D5 — the safety suite rewrote the live KB on every run

`test_reingest_is_byte_identical` re-ran the real ingester against `data/kb/entries.jsonl` itself. The assertion held, but a green test that rewrites 15 MB of gitignored production data with no restore path is how an earlier build lost an entry. It now ingests into a **copy of the real KB** in a temp directory, and afterwards asserts production is byte-identical.

Copying the KB alone was not enough. **Everything under `/data/` is gitignored**, and `--kb` redirected only `entries.jsonl` — the sidecars, `toolfiles.json` and `corpora_report.json` kept writing into `data/kb` regardless. The ingester's outputs now follow the KB it is given, so "run this somewhere else" means it. The test asserts that too, by mtime: `corpora_report.json` is rewritten unconditionally, so before this fix its mtime moved on every single suite run.

It is still the **real KB**, copied — with a size floor asserting >1000 entries. Idempotence over 2,743 entries of real escapes, non-ASCII titles and foreign lines is the property worth proving; over a three-line fixture it is close to vacuous, and a fixture is what this test would otherwise decay into.

D6 — nothing noticed a derived artefact going stale

The scripts index sat **126 entries behind** the KB (2,617 against 2,743) through an entire audit, and was found by reading a number. A stale index does not error — it answers, with 126 entries' worth of the corpus missing.

`test_kb_drift.py` compares the KB against the three artefacts built from it: the semantic index (`ids.json`), the scripts index, and the corpus report. Two details carry the weight:

- **A count is not coverage.** `ids.json` is checked for *set equality* of entry ids, not just the total — one entry added and another deleted between runs leaves the count untouched and the index wrong. The scripts index is additionally checked for citations that no longer resolve.
- **Absences are reported, never silently green.** `/data/` is gitignored in its entirety, so a fresh clone or a CI runner has none of these files, and there the test proves nothing. It says so out loud. A fourth check fails when `data/kb` grows a file nothing has classified, so leaving something unchecked has to be a decision rather than an oversight.

Each check has a positive control that runs **the real check** against a doctored artefact — including a 126-entry lag, the exact drift that went unnoticed.

D7 — every successful AD command looked like a failing one

PowerShell's stderr over WinRM is not text; it is a CLIXML object stream, and every progress bar it would have drawn on a console arrives as an `<obj S="progress">` record. Build #9's live run against a real `corp.local` DC returned `rc=0` alongside a wall of markup.

Progress records are now stripped where the `WinRMResult` is *built*, so the event stream and the persisted run record see the same cleaned text; doing it at a display layer would have left the record full of markup and put the burden on every future consumer.

The whole risk of this fix is that it strips something else, converting a cosmetic annoyance into a silent failure — a strictly worse bug. So the filter removes progress records and nothing else: Error, Warning, Verbose and Debug streams survive, `_xNNNN_` escapes are decoded, and **anything that fails to parse is returned verbatim**. The raw stderr stays on the result and the dropped count is reported, so "quiet because it was filtered" stays distinguishable from "quiet because it was quiet". The test that matters is the negative one: a genuine `Get-ADUser` error survives a document carrying forty progress records.

The bug-bounty submission fields — and one that nothing could set

VRT priority (P1–P5) now sits alongside CVSS. Bugcrowd triages on the Vulnerability Rating Taxonomy, and the two genuinely disagree: a stored XSS is P2 whatever its vector works out to, and a 9.8 on a class the program rates P3 gets paid as a P3. Reports carried CVSS only, so the number a triager acts on was missing.

It is a LOOKUP, and that is the invariant with teeth. Deriving a priority from the CVSS score would have been three lines and would have produced a confident P-number with no relationship to the taxonomy — a fabrication in the field a triager reads first. So the operator names the category and a curated subset of the VRT maps it; an unrecognised category claims *no* priority and says why. The regression test proves the priority does not move when the CVSS vector does. Where the two disagree, the report says so and suggests arguing it.

The field nothing could set. `build_cvss_block` has read `session['cvss_vector']` since the bug-bounty template shipped. There was no column, no endpoint and no control — so the CVSS block this project verified against six reference vectors, roundup edges included, **could never appear in a real report**. The calculator was

right and unreachable. Adding the VRT field would have been dead in exactly the same way, so both are now stored, settable through `PUT /sessions/{id}/submission`, and covered by a round-trip test.

The known-issue check — flag, never suppress

Programs publish what they already know about and will not pay for. Submitting one burns the report and, on some programs, costs signal. Nothing compared a finding against that list.

The whole feature is one design rule: **it flags and it never drops**. A false match that silently removed a real finding costs a paid bug and nobody ever learns it happened; a false flag costs one glance at the brief. The output is a table headed by what was compared, the matching is deliberately loose, no code path removes a finding, and the test asserts the finding list comes back unmutated.

It also reports **when it finds nothing** — "compared 7 findings, no matches" — because silence is indistinguishable from a check that never ran. The discrimination is real, not nominal: a stored XSS in the order form is not matched against a published *self*-XSS in the profile bio, which is the false positive that would have flagged a P2 as unpayable.

Authenticated-scan session expiry — a silent wrong answer made loud

Build #15's AJAX spider crawls behind a login by inheriting a session the human established by hand. It added no ZAP context, no session management and no authentication handling. That is fine. **What happens when the session expires mid-scan is not.**

The active scanner does not stop. It keeps firing SQLi and XSS payloads at what are now login redirects, matches nothing, and finishes cleanly reporting **zero findings** — which is indistinguishable from "*the application is secure*". Not a crash and not an error: a confident wrong answer, of the kind nobody has a reason to doubt.

`session_health()` reads the recorded traffic and notices four shapes a dead session takes: a wall of redirects to a login path; 200s whose body is the login page; a wall of 401/403; and the one the first three miss — a **collapse to a single response shape**, which is what a login wall looks like when it renders a friendly page instead of redirecting. Below ten responses it returns `unknown` and says so; a false all-clear would re-create exactly the confidence this removes. A `suspect` verdict rides the alert-ingest response *and* lands at the top of the report, above the finding count it should be read with.

This is the cheap honest version by decision. It does not re-authenticate, maintain a session or build a ZAP context — those are a separate build. The AST-scanned test asserts it cannot: it reads the exchanges it was handed and calls nothing.

The phantom class, and what the type-checker cannot see

`.hp-tn-start` did not exist in `globals.css`, while nine buttons across `:exposure`, `:oob` and `:windows` used it. Those are the PRIMARY actions — "write profile", "save remote profile", "deploy + start" — so they rendered plainer than the destructive buttons beside them, which at least use `.hp-tn-stop` and turn red on hover. **The visual hierarchy was exactly backwards**, and nothing could tell: tsc, `next build` and eslint have no idea whether a CSS class exists.

Defined rather than dropped — nine call sites already assert "this is the affirmative action", which is a distinction worth rendering. Three ranks now, so the eye sorts them before reading the labels: `.hp-tn-start` (accent, affirmative) · `.hp-tn-destroy` (red, destroys something) · `.hp-tn-stop` (muted, a reversible undo).

And the row itself. Inputs, the wildcard checkbox and all three buttons shared one `.hp-tn-form` flex row, where `input { flex: 1 1 200px }` stretched the fields and the buttons wrapped wherever they landed. There was no break between *configuring* a profile and *acting* on it — including the act that recreates the container and kills every listener, session and background job inside it. A shared `.hp-tn-actions` row now separates them, rules them off, and pushes the destructive button away from the affirmative one. It lives in the CSS layer: the last time a shared `.hp-tn-form` rule was patched per-screen, the fix shipped for `:proxy` alone and left checkboxes broken on six others.

The guard this earns. `test_css_vocabulary.py` asserts every `hp-*` class any component names exists in `globals.css`. This class of defect has now cost three builds — an entirely invisible `:proxy` for two of them, then this — and it is mechanical, so it belongs in the suite rather than in a reviewer's memory. **Its first run found seventeen**, and the first thing that had to be fixed was the checker: `hp-set-model` and `hp-set-key` are element `ids`, not classes, so the scan was narrowed to `className` attributes. Four more are harmless: modifiers on an element whose *first* class is defined and does the styling. The remaining **eleven are real, bare phantoms** in `:repeater`, `:exploits`, `:category` and three others. They are frozen as a baseline that may only ever SHRINK, kept deliberately apart from the allow-list: one list means "checked, fine", the other means "known, unexamined", and collapsing them is how a check like this quietly stops working. **It is still not a substitute for looking at the screen** — a class can exist and still be wrong.

D9 — the two surfaces that read as a different product

`:evasion` was raw Tailwind utilities (`rounded border border-slate-700 bg-slate-900/40`) with no kicker/ `:title` header, so it looked like a different application. It is now the house `hp-tn-*` vocabulary throughout. Two classes of thing the vocabulary did not have — a label/value fact list and a preformatted block — were added *as vocabulary* rather than as one screen's private classes, which is the difference between a restyle and a reskin. The one piece of emphasis kept deliberately is `still_recorded`: it is the sentence that makes the surface a purple-team instrument rather than an evasion how-to, and it must not flatten into the facts around it.

`:ad-graph` was a different case, and worth saying plainly. It is **not** Tailwind — it has a purpose-built `hp-adg-*` vocabulary for nodes, edges and hop drawers, with no `hp-tn-*` equivalent. The real gap was the page: a bare `<main>` with its own back-link bar, outside `PageShell`, with no kicker/ `:title` block. That is fixed, along with the two empty states, which are ordinary cards now. **The graph canvas itself is untouched** — replacing those primitives would be a rewrite of something that works, and D9 is explicitly cosmetic.

Both are behaviour-identical: no gates, no endpoints, no handlers changed.

The browser pass found two things nothing else could

Every gate was green — suite, docker-stripped suite, `tsc`, `next build`, lint baseline, and the new CSS-vocabulary check — before anyone looked at a screen. Then looking found two defects.

All three buttons in the new action row rendered identically grey. `.hp-tn-actions` button and `button.hp-tn-start` have the *same* specificity (0,1,1), and the base rule was written second, so it won. The classes existed, they were applied, the check that asserts they exist passed — and the row that was built specifically to make three ranks distinguishable showed three identical buttons. This is what "*a class can exist and still be wrong*" means in practice, and it is the reason that sentence is in the guard's own docstring rather than a claim that the guard is sufficient.

15 of 16 commands won't run as written. JSX drops the space between an expression and the text following it when a line break falls between them. A missing space in the one banner whose entire job is to be read.

Neither is subtle once seen, and neither was visible to any tool. The rule stands: **a frontend change is not verified until it has been looked at.**

A third came out of the same pass, once the new action rows put it under a spotlight: `hp-tn-olhint` is a **label** style — 10px, 1px letter-spacing, ALL CAPS, bottom margin only — and ten places were using it for whole paragraphs. Three lines of prose rendered in it are close to unreadable, and with no top margin they sit crammed against whatever is above. Those are now `hp-tn-note`; the genuine labels (`expected: 8090/tcp`, `path`, `mode` · `exit` · `run`) keep the label style, which is the point of having two. One element on `:exposure` was doing both jobs at once — a port list and a paragraph sharing a tag, so the paragraph inherited the label's styling — and was split in two. Pre-existing on `:exposure`, `:oob` and `:windows`, and fixed there because this build is what put a ruled row directly above each one.

Still outstanding, deliberately not changed: the acknowledgement checkbox labels ("I understand this is a public listener that relays into this machine") are also 10px uppercase. That is text an operator is meant to actually read before ticking a safety control, so the styling is arguably wrong there too — but it sits inside the form flex row, where a size change moves layout, and it was not part of what this build touched.

What was not done

- **No UI for the VRT / known-issues fields yet.** `PUT /sessions/{id}/submission` and `GET /vrt-categories` exist and are tested; the report screen has no control for them, so today they are set by API. Flagged rather than quietly skipped — the report-side work is what the punch list asked for, and it is complete.
- `:ad-graph` 's **graph primitives were left alone**, for the reason above.
- **A real foreign host is still not substituted** in a planned command, only flagged — see the D8 section for why rewriting one would point a tool download at the client.

The three things the punch list left open, closed

Foreign-host substitution now has a seam, and it is what the TOOL DOES. `_HOST_TOOL` was one list; it is two. TARGET-DIRECTED tools (`nmap`, `ffuf`, `nuclei`, `sqlmap`, `sublist3r` ...) have a host argument that IS the thing being assessed, so an out-of-scope host there is the KB author's target left in by accident and is rewritten automatically. FETCH-CAPABLE tools (`curl`, `wget`, `nc`, `ssh`, `scp` ...) are never rewritten: the host may be a tool download or your own listener, and nothing in the argv distinguishes it from a target. The seam is deliberately not a blocklist of infrastructure hostnames — a blocklist has to be complete to be safe, and the first host nobody thought of is the one that points a download at the client.

What the automatic pass declines, the operator can still do — by **looking first**. A flagged command carries a `suggested_cmd`, shown BESIDE the original and never in place of it, with its own copy button and a line saying it was not applied.

That design was validated by the first real plan it rendered. The suggestion for a SOAP XXE payload rewrote the XML *namespace URIs* — `xmlns:soap="http://schemas.xmlsoap.org/..."` became `xmlns:soap="http://www.crateandbarrel.me/..."`. Obviously wrong on sight, and silently corrupting if it had been applied for you. One second of human attention is the whole mechanism, and it only works because both versions are on screen.

A rewrite that IS made is stated rather than hidden: the command keeps its `original_cmd`, and the block says "repointed from tesla.com — the entry said ...". Changing what a KB entry said without saying so is its own kind of dishonesty.

The submission fields are reachable. A collapsed panel on the engagement screen carries the CVSS vector (eight dropdowns AND a paste box that populates them, with the score computed live), the VRT category, and the known-issues list; the report screen echoes what the next report will carry, with an edit link. Measured end to end in the browser: pasting `CVSS:3.1/AV:N/AC:L/PR:L/UI:N/S:U/C:H/I:N/A:N` populated all eight metrics, read **6.5 Medium**, saved, and came back out of the generated report as an authoritative CVSS block — the first time that calculator has appeared in a real report at all.

Two defects only the browser found, again

`PUT` is not in the CORS allow-list. The submission endpoint shipped as `PUT` and every backend test passed — TestClient talks to the ASGI app directly and never performs a preflight — while the browser refused it outright. The failure surfaced as *"cannot reach the API"*, which reads like the server is down. The fix was the VERB, not the policy: this is a partial update, so `PATCH` was correct all along, and widening the allow-list would have been the easy wrong answer. `test_attack_path_contract.py` now asserts that **every route's method is one CORS permits** — a backend/browser contract of exactly the same shape as the `response_model` check beside it, and equally invisible to backend-only tests.

The bounty blocks were landing on every template. With the fields set, a `standard` pentest report carried a Bugcrowd VRT priority and a known-issue check — jargon a client or an OSCP grader did not ask for. Both are gated to the bug-bounty template now. CVSS is deliberately not gated: a base score is meaningful everywhere, and it has been spliced unconditionally since the block existed. Verified against a running backend on both templates.

Build #17 — closing out build #15 (2026-08-04)

Build #15 shipped with its acceptance runbook only partly run. §3.4 — attack one captured URL with the scanner — was never executed, and §5, teardown, was never done: the port is still published and the daemon still running. §3.3 was *answered* (refused) rather than run. Build

17 is that close-out, plus a diagnostic that must land before anything is built on top of it.

The rules-of-engagement free-text field — reviewed and ACCEPTED

Build #15's live run surfaced that an engagement's `authorization` prose is enforced by nothing, and recorded it in the same sentence as the audit's credential-precondition finding: *a documented constraint with nothing behind it*. Reviewed 2026-08-04 and accepted, in the same voice as D1 and D2 — it is not an open defect and is not to be re-opened.

The reasoning is the one already written into the target lock's own docstring: **scope is enforced in code, and human approval of every command is the actual bound**. The authorization line is a note from the operator to their future self and to whoever reads the report. A `passive_only` flag derived from free text would be a gate whose predicate is English prose — the least dependable input this system has — and its real effect would be to teach an operator that the box means something the moment it happened to fire. That is the same failure the `daemon` red-confirm had: a control whose stated meaning is not the one it enforces.

What is worth fixing is the **contradiction**, not the enforcement, and that is item 1 below.

Item 1 — the record was corrected, not made enforceable

`eng-69ec01d0fe74` carried "...PASSIVE RECON ONLY this session (operator asleep); no active scanning per rules of engagement." True when written for an unattended session; false now. Running §3.4's active scan under it would leave an audit trail contradicting the action — the one cost of an unenforced field that is real regardless of whether it is enforced.

So the record was re-entered with accurate text, same target and same 11-host scope. **Exit, then enter — not enter alone.** `enter()` mints a fresh id rather than amending, so entering without exiting would leave the forbidding string active *beside* the corrected one, and `GET /cockpit/engagement` would still return it. The precondition is then checked inside the live-fire script itself rather than assumed: if the passive-only text is still active, the scanner script refuses to start. **The check belongs in the thing that would create the contradiction, not in the product** — which is exactly the line that keeps this from becoming the `passive_only` gate that was declined.

Item 4 — choosing the endpoint IS the safety decision

The active scanner sent **376 requests against a single endpoint** in the build #14 measurement, with payloads at every parameter. Against a production storefront, which URL is picked is not a detail. It was chosen out of the 55 requests the proxy actually captured, not guessed: a category listing carrying a real application filter parameter, `recurse=false`.

What was rejected matters more than what was picked. The capture contains `/en-ae/cart`, `/en-ae/login/register`, `/en-ae/gift-registry` and `/en-ae/guest/order` — injection payloads at a cart or checkout parameter on a live storefront can create orders or empty a basket, which is the AJAX spider's hazard arriving

through a different door. The endpoint is therefore a **constant in the script with a forbidden-token guard checked against it at run time**, not an argument: a comment explaining why a URL is safe does not survive somebody editing the URL.

It also rejected the most interesting thing in the whole capture, and that is the judgement call worth recording. `www.crateandbarrel.me/api?endpoint=https%3A%2F%2Fapi.crateandbarrel.me%2Frest%2Fv2%2F...` is a **server-side proxy that takes a full URL in a query parameter** — the classic shape of an SSRF, and precisely why it is not scanned here. Active-scanning it means asking the target's own infrastructure to fetch whatever the scanner puts in that parameter. That is not a read-only endpoint; it is outbound request generation from someone else's servers, and it deserves a deliberate decision of its own rather than riding along inside a pipeline test. **It is a finding, and it is logged as one, not exercised.**

The pacing added in build #16 does not cover this path — checked, not assumed

Build #16 added per-tool throttle-flag injection for engagement runs. `pace` is a field on `ExecRequest`, applied on the **executor** path; `cockpit/proxy.py` contains no reference to it, and the scanner drives ZAP's API directly. So §3.4 is unpaced, and saying otherwise would be the most dangerous kind of wrong — a safety feature credited on a path it never reaches.

The only rate control on this path is that **stop is ungated**, which is why it is ungated. The live-fire script adds a request ceiling and a wall-clock cap that stop the scan through that same route. The real fix is ZAP's own `scanner.setOptionDelayInMs`, and it is deliberately not in this build.

The defect that blocked item 4 — the product goes blind to its own daemon, silently

Found while dry-running item 4's preconditions, and it is the most consequential thing in this build so far. `GET /cockpit/proxy` returns `[]` while ZAP holds 1,076 captured messages.

`cockpit/proxy.py` keeps API keys in an **in-process dict**. `api_key_for()` documents this honestly — *"the key this process minted for that daemon, or " if it did not start it"* — but nothing downstream treats the empty case as unknown. `_api_get` simply omits the header, ZAP answers with an **empty body**, and `history()` parses that into `[]`. So a backend restart — an ordinary event across a multi-day engagement — turns every read route blind and reports the blindness as *"no traffic captured"*. Silent empty, again, and this repo's fourth or fifth.

The second face is worse, because it is the check that is supposed to prevent it. `clash_refusal()` also reads the in-process `_models`, so after a restart it can see no existing daemon and would let `start_proxy` spawn a second one in the same container. That second daemon dies immediately on ZAP's home-directory lock — the exact failure whose comment sits directly above the function — while `lifecycle.observe()` finds the port **bound by the old daemon** and reports the new proxy up, holding a key the listening process will never accept. A `Proxy` model reading `status=up` whose every subsequent call returns empty. **The clash check protects against a state it cannot observe**, in a module whose own docstrings say counts are "READ BACK FROM ZAP, never assigned at launch".

Nothing here is a safety hole: no gate is bypassed, and the failure direction is toward doing less, not more. It is a *correctness and honesty* defect, and it lands squarely on the distinction the acceptance runbook draws in bold — *"a broken proxy is a bug in this build, and a refused browser is a finding about the target. Do not report one as the other."*

Item 4 was held rather than run. In that state the scan would have sent nothing, printed `REFUSED`, and been written down as *"Akamai blocks the scanner even though capture works"* — which the plan itself calls "a significantly larger finding than the spider's refusal". It would have been the most consequential wrong conclusion

available in this build, arrived at through a green-looking script. The precondition that catches it is in the live-fire script, and it named both ways forward rather than picking one. **Zaid chose the fix** (2026-08-04), keeping the 1,076-message capture.

The fix — the key is recovered, and the clash check can finally see

A daemon states its own key in its own argv, and that argv is readable inside a container we own. `/proc/<pid>/cmdline` is *already* the accepted residual of passing `-config api.key=` at all, written down in the build #15 section — so this reads what that residual exposes rather than widening anything. Recovered keys live in `_adopted`, deliberately **not** in `_keys`: "what this process minted" and "what we read off a process we did not start" are different facts, and one dict would lose the distinction the moment anyone needed it.

Three things make it safe rather than merely convenient:

- **The port is checked, not assumed.** `observed_daemon` reports one daemon per container, so adopting its key for whatever port the caller asked about would send a live secret to something else listening there. A mismatched port returns `""` and caches nothing.
- **An empty answer drops the adopted key.** ZAP replies to a wrong or missing key with an empty body — the same thing a dead daemon looks like — so either way the recovered key is suspect and the next call re-reads the argv. Self-healing across a daemon swap.
- **The probe cannot match itself.** `grep '^api[.]key='`, because the probe's own command line is in `/proc` while it runs and an unbracketed pattern would return a fragment of the probe as a credential. This is the `[z]aproxy` lesson applied to a `grep` instead of a `pskill` — now wrong in both directions **and** in a third tool. The test caught the obvious follow-on immediately: ``${K#api.key=}`` puts the literal straight back, so the strip is ``${K#*=}``.

`clash_refusal` stays **pure** and takes the observation as an argument. Doing the `docker exec` inside it is exactly the mistake its own docstring records — a test that reached for Docker in a pure check passed locally and failed in CI, twice in this repo. `start_proxy` observes and injects; an orphaned daemon now earns a refusal that says a *restart lost the record, not the daemon*, because there is no proxy in the UI to press stop on.

The defect underneath it: a capture was unreadable on Windows, and read as no capture

Fixing the key let a read reach real response bodies for the first time on this machine, and it **crashed**. `_api_get` ran `subprocess.run(..., text=True)`, which decodes with the **ambient locale codec** — cp1252 here. A capture is arbitrary bytes from arbitrary sites, so one byte outside that codepage raised `UnicodeDecodeError` inside subprocess's own reader thread and left `stdout` as `None`.

It has been failing for as long as the code has existed, and it looked like success. `history()` catches the downstream error and returns `[]` — this module's recurring silent empty, one layer *below* the one build #17 came here to fix. Every Windows operator reading a real capture would have seen "no traffic". It was invisible because on this machine `history()` was already returning `[]` for the key reason, and in tests the bodies are ASCII.

The fix is to stop letting the environment decide: capture bytes, decode UTF-8 with `errors="replace"`. The regression test feeds `0x8f` — undefined in cp1252 and not valid UTF-8 either — and asserts both that a `str` comes back with the readable part intact and that `text=True` is not passed. Verified to bite: the old code returns `bytes` for that input.

After both fixes, against the live daemon: **200 exchanges read, session health ok, 96 sampled**. That is the signal item 4 needs to tell "zero findings" from "a dead session", and it was returning `unknown` from an empty list an hour ago.

What was deliberately NOT fixed: `GET /cockpit/proxy` still returns `[]`. Reading and driving an adopted daemon works; *listing* it does not, because a `Proxy` model carries `started_at`, `bind_host`, `published` and `engagement_id` — facts about how a daemon was launched that its `argv` does not fully carry. Synthesising them would put invented values in front of the operator, in a module whose stated rule is that counts are read back and never assigned. Better a visibly empty list than a confidently wrong row. The consequence is stated rather than buried: after a backend restart the `:proxy` panel shows no proxy while the history and scan panels work, and the live-fire precondition therefore accepts *either* signal.

Item 5 — the browser must talk to the target and nothing else

A crawl aimed at a local two-page lab site also produced requests to `www.google.com`, `gstatic.com` and `play.google.com` — Chromium's variations seed and search preconnect, going out through the ZAP proxy because the engage sandbox has full egress. Harmless in a lab and wrong on an engagement: traffic nobody asked for, leaving the operator's IP, landing in the capture a report is written from. `/etc/chromium.d/hackpit-container` now also carries `--disable-background-networking --no-first-run --no-default-browser-check`.

The proof check for it states exactly what it proves, because the honest version is narrower than it looks: it asserts the drop-in **sources cleanly and yields the flags**. That the file is sourced *at all* is proven by the check above it — the browser only starts as root because `--no-sandbox` from this same file reaches it. The two together are the property; this one alone would be a `grep` against a file, which is what "the file existing is not the same as the flag reaching the browser" already warns about one line up. Verified to **bite**: run against the pre-rebuild container it fails, which is the only way to know a new check is not vacuous.

Item 2, RUN — the discriminator is not the User-Agent, and that is why the honest fix works

Decision-table row 4 (P1 fails, P2 passes), which the plan routes to row 1: **build option A**.

probe	result
P0 headless, stock UA, ×3	fails 3/3 — deterministic, not a transient. 336 B, 21–25 s, ZAP's own 504
P1 headless + a normal Chrome UA (one spoofed request, as a measurement)	fails — hung, killed at 75 s, 0 bytes
P2 headed under Xvfb, stock UA, nothing spoofed	PASSES — 2,347,139 bytes in 3.7 s, 26/26 answered

The prime suspect was the `HeadlessChrome` token. **It is not the discriminator**. Reading the request headers back out of ZAP's own history, across all four clients that reached this target:

client	Sec-CH-UA	outcome
headed Chromium	<code>"Not;A=Brand";v="8", "Chromium";v="150"</code>	31× 200, 1× 302, 1× 301
headless, stock UA	absent	504
headless, spoofed Chrome UA	absent	504

Headless Chromium sends **no Client Hints at all**, and `--user-agent` does not add them. So P1 changed the string the client *claims* and left untouched the header set a real browser *emits* — which is why it failed, and why row 2 was never reachable. **You cannot get there by lying about one header; you get there by using a browser that**

sends them all. That is this build's own line — *"nothing here imitates or evades — it uses one"* — arriving as a measurement rather than a principle, and it is a considerably better outcome than row 1 would have been: the dishonest route is not merely disallowed, it does not work.

The script's printed verdict was wrong and the data was right. It reported `ROW none: HARNESS BROKEN` because a single `dom > 5000` rule declared both controls broken: example.com is a ~1 KB page, so a perfect fetch of it is 561 bytes. One threshold cannot serve a 1 KB control and a megabyte storefront; the floor belongs to the URL. The same rule cost P2c 76 seconds — the headed probe only stops early once it is over the floor, so it sat waiting on a page that had finished loading in three. Both are fixed, and the pass rule now also refuses to count a `403 / 504` as success.

Item 4, RUN — the attack half IS blocked, and the first verdict said the opposite

This is the larger of the two outcomes the plan named. Not "the chain survives the WAF" — the reverse.

39 attack requests were sent at one read-only catalogue endpoint. Reading the responses back out of ZAP's history: `3x 200`, `36x 403 Forbidden`, `Server: AkamaiGHost`, body `<TITLE>Access Denied</TITLE>`. **92% edge refusals**. A sample request carries a Shellshock payload in `allCategories` — Akamai inspects the payload and refuses it. Capture works; the attack half does not. **The payloads never reach the application**, so the zero findings are not evidence about the application at all.

It was also **tarpitted, not merely blocked**: 39 requests in 12 minutes, one stall holding at 7 requests for ~350 seconds, 28% progress when the wall-clock ceiling stopped it. Build #14 sent 376 requests at a single endpoint in the lab. That is the concrete shape of the audit's *"breaks at volume"* — roughly three requests a minute against a bot-managed edge.

The script announced `WORKS - the chain survives the WAF`, and two separate mistakes produced it.

1. **It counted 153 alerts as scan results**. The route's own docstring says they are *"NOT scan-scoped: this includes PASSIVE alerts raised merely by traffic passing through the proxy"*. They were an hour of manual browsing. Exactly 2 of the 153 are on the endpoint that was scanned, both passive header findings. Walking into a trap the function documents is worse than walking into an undocumented one.
2. **It treated "answered" as "allowed"**. A `403` with a block page is a response. Every "requests landed" count in the first pass included the ones Akamai refused.

The verdict now classifies the recorded responses and calls $\geq 50\%$ edge refusals what it is.

Item 3 — option A, built on the measurement rather than on the suspicion

`SPIDER_BROWSER_ID` is `chrome`, not `chrome-headless`, and the image installs `xvfb`. The constant carries the measurement that justifies it, because the next person to read it will otherwise see "headed browser" and assume it was a preference.

The display goes in through the drop-in, not through the process tree. The browser is launched by ZAP → Selenium → Crawljax, three processes we do not control, so threading `DISPLAY` down to it would mean touching the spawn path. `/etc/chromium.d/hackpit-container` now also carries `export DISPLAY="${DISPLAY}:-:99"` — Debian's own launcher sources it, so any Chromium started by anything finds the display, and an explicitly-set `DISPLAY` still wins.

`start_spider` refuses when there is no display, with `gate="display"`, and the refusal names both `xvfb` and `setcap`. This matters more than it looks: a headed browser with nowhere to draw dies at launch inside ZAP's log, and the visible symptom is `{"Result": "OK"}` followed by zero URLs — **the identical symptom of both defects**

build #15 spent a day on. Three different causes, one indistinguishable outcome. The third one gets a sentence instead of a log dive. `ensure_display` returns the OBSERVED state rather than "the start command ran", which is the same rule as reading the browser id back instead of trusting the OK.

The recorded Xvfb trap — Kali ships setcap'd binaries and `no-new-privileges` makes the kernel refuse to exec them — turned out to be **already handled**: layer 10 strips file capabilities across `/usr /bin /sbin /opt`, and it runs after the layer that installs xvfb. Verified on the live container: `getcap /usr/bin/Xvfb` is empty and Xvfb runs.

Two locks, and one of them had to be repaired rather than flipped. The proof no longer hardcodes the browser id; it reads `SPIDER_BROWSER_ID` out of `cockpit/proxy.py`, because a proof carrying its own copy of a fact goes on proving the old one. And the existing test asserted `SPIDER_BROWSER_ID == "chrome-headless"` while its own failure message explained the real property — *the image has Chromium and NO Firefox*. Flipping the literal would have made it a lock that only ever says "you changed it", which teaches people to update it without reading it. It now asserts the predicate: chrome yes, firefox never.

The defect that hid item 4's own evidence: "recent" traffic was the oldest traffic

Finding the 403s at all meant paging ZAP's history by hand, because reading the documented way returned nothing. `history()` passed the caller's `start` straight to ZAP, whose `start` counts from the **beginning** of history, and it defaults to 0. Measured on a daemon holding 1,296 exchanges: `start=0` returned requests to `example.com` from hours earlier. The `:proxy` captured-traffic panel passes no `start` at all, so **it has always shown the first 50 requests that daemon ever recorded**, forever, however long the engagement runs.

The knock-on is the serious half, and it lands on build #16. `session_health` reads `history(count=200)` to notice a scan's traffic coming back login-shaped — #16's answer to a silent wrong answer. Judging the *oldest* 200 exchanges, it was looking at the moment the session was established, so **it could not detect a session expiring mid-scan by construction**. It reported `ok` throughout this build's live scan, and that `ok` meant nothing. A guard against a silent wrong answer, silently unable to fire — and it took a scan whose traffic was being refused to notice, because that is the case where you go looking.

`start` is now an offset back from the newest; order within the window stays chronological. After the fix, `session_health` samples 195 recent exchanges instead of the first 200 ever.

The one place this build said "you may not" — removed

The live-fire script hard-blocked a list of path tokens (`cart`, `checkout`, `account`, `newsletter`, ...) and refused to run under an engagement record whose prose forbade active scanning. Zaid stripped both, 2026-08-04, and he was right that they did not belong.

HackPit's entire design is that the operator is the bound and every control informs rather than forbids. The danger gate demands a second confirm and then proceeds. The scope lock is documented, in its own docstring, as a handrail rather than the wall. This very build *declined* to make the authorization field enforceable. A proof script inventing a prohibition the product does not have was the single place in build #17 that said "you may not" — and the record check was worse than that: it enforced the declined `passive_only` gate by the back door.

Both now print and continue. What survives is the reasoning, because a recorded decision is worth more than a blocked path: injection payloads at a checkout parameter on a live storefront can create real orders or empty a basket. That is a fact about the target; the operator weighs it.

A third one went with them, and it was the worst of the three because it was invisible. The script required the URL to contain the engagement's *named target* — which would have refused the other **ten in-scope hosts**, since the scope is eleven and the named target is one. `scope.py` already models wildcards, CIDs and exclusions and

has been measured against this exact program, and `executor.validate_request` runs it before ZAP is contacted. One scope model, and it is not the harness's.

Verification of item 3, live — and the runbook is now actually finished

`docker/proof/browser_intercept_proof.sh` : **25 passed, 0 failed** on a freshly built image (container image id compared against the built one first, per "the container is not the image"). Three of those are new: the drop-in exports a display, the drop-in suppresses background networking, and a virtual display is observed up by process name.

Option A works. ZAP reports browser id chrome (read BACK, not assumed from the OK) , a real browser process was observed during the crawl, and the crawl captured 2 → **42 messages and found 20 URLs** with a headed browser under Xvfb. The safety half is unmoved: an unauthenticated view refused, an unauthenticated action refused, a wrong key refused, the right key answering, the lab sandbox still publishing nothing.

Runbook §5 teardown is done — the daemon stopped, `docker/listener-profile.yml` removed, the container recreated, and `NetworkSettings.Ports` now reads `{}` . Nothing listening on either side. That closes the last open step of build #15, which is what this build was for.

What was prepared and has now run

Build #15's section had to be amended after the fact because the runbook was described as finished when it was not. Not repeating that: as of this commit, **items 1, 2 and 4 have all run**, and what remains is option A itself, the image rebuild, the proof re-run and teardown.

Three of this build's findings came from its own scripts being wrong, and that is worth stating rather than smoothing over: a single DOM threshold that failed both controls, an alert count that was never scan-scoped, and a verdict that read a WAF block page as a served response. Each was caught by going back to the raw evidence — ZAP's own history — instead of trusting the summary line. **The scripts printed ROW none and WORKS ; the correct answers were row 4 and REFUSED** . A harness that reports confidently is not the same as a harness that is right, which is the same lesson as the proof checks in build #15 and is now this repo's most repeated one.

Item 1 ran, and its first version was wrong in an instructive way. It asserted "*exactly one active engagement*", written against a truncated read of `GET /cockpit/engagement` that showed one record when there were **21** — twenty of them old lab engagements against RFC1918 addresses, still active because engagement expiry (D4) was reviewed and declined. The check caught it, which is the only reason it is a footnote rather than a defect. The property is **per-target**, and it is now stated that way in both scripts: *no active engagement naming this host may forbid active scanning*. Exiting twenty unrelated records to satisfy a global assertion would have been a destructive tidy-up nobody asked for. The corrected record is `eng-d3807b5bd170` ; the script is idempotent and re-running it writes nothing.

Every one of them reaches a real, in-scope bug-bounty target, and live fire against a real target is refused by the real-time classifier from inside the agent — the recorded convention is to run it in a plain shell and document from the log. They are therefore committed as scripts (`docs/proof/build17_run.sh` and the three it drives) rather than as results, and the measurements follow in a later commit. **The diagnostic's outcome, including a "we could not make this work honestly" outcome, is a legitimate result and will be recorded as one.**

Two things in the harness are worth keeping regardless of what it measures. **Every probe carries its own control** — a headed browser that cannot start would otherwise read as "Akamai refuses headed browsers", which is the one decision-table row that must never be reached by accident, since it is the row that would put UA spoofing on

the table. And `RC=$?` after a pipe is the repo's non-gating defect wearing a hat: in `cmd | tee log`, `$?` is `tee`'s status, which is zero whenever `tee` could write the file. Both live-fire wrappers were written with that bug and both were fixed before anything ran, with the fix proved in both directions.

Verification

Hermetic safety suite **78 test files, every one exited 0** — and the two touched files re-run with `docker` removed from `PATH`, both 0, **with the strip checked to bite first** (`docker` not found, `git` still found, since the KB ingester's recovery path needs it). Six new locks: a recovered key is bound to the port it was read from, the daemon probe cannot match its own command line, the API reader survives bytes the local codec cannot decode, history returns the newest window and pages backwards from it, a headed crawl with no display is refused naming both `Xvfb` and `setcap`, and an orphaned daemon is refused with ZAP's real reason.

`docker/proof/browser_intercept_proof.sh` **25 passed, 0 failed** on the rebuilt image.

One flake worth recording so the next person does not chase it: `test_redirector.py` failed once with "*no bindable port for kind=1 outside the self-mapping range*" and passed on re-run. That is Windows' UDP excluded-port ranges moving under it (Hyper-V/WSL reserve blocks), not a code fault — but it means a single red run of that file is not evidence until it is repeated. `tsc --noEmit` exits 0, `next build` exits 0, `eslint` sits at the accepted baseline of **11 errors + 1 warning** (unchanged — this build touches no frontend file). `data/kb/entries.jsonl` still **2743** entries. `docker/proof/browser_intercept_proof.sh` gains one check, so it becomes **23** on a rebuilt image; it is not re-run here because the rebuild that makes it pass would destroy the running daemon and the ~1000 captured messages items 2 and 4 depend on. **The rebuild and teardown are deliberately last**, after the live-fire runs, for that reason.

Build #18 — reaching the targets that refuse us, and scanning behind a login (2026-08-05)

Build #17 measured, against a live Akamai-fronted target, that **there are two different walls and they need different answers**: Bot Manager judges the CLIENT (headless Chromium gets a 504 tarpit; bare `curl` / `ffuf` / `nuclei` never reach request one), and the WAF judges the REQUEST CONTENT (scanner payloads got **36 of 39 403 AkamaiGHost**). Every item here answers one wall or the other, and the two are deliberately never conflated: fingerprint spoofing does nothing about a 403, because those requests already cleared Bot Manager and died at the rule engine.

The build adds no gate, no confirm, no blocklist and no allow-list narrowing — not in the product and not in a proof script. Where something could refuse, it warns and continues. Two places tightened an EXISTING control that was failing open, and both are named below with the argument for why that is repair rather than prohibition.

Item 1 — the bypass header, and the fact that its value is a credential

Programs of MAF's size issue researchers a header that skips the WAF, so the testing they invited is not refused by the edge they bought. Zaid does not have one yet; the mechanism is built now and proved with a dummy header, which is provable without a real one and without spending a real-target request.

The header is stored on the **engagement**, because it is per-program. It is injected through **ZAP's Replacer add-on**, which is the one point every outgoing request passes: the operator's own browser, the AJAX spider's browser, the active scanner's payloads, and any sandbox tool run with `proxy: true`. One rule covers all four; threading a header through each caller would need four correct implementations and would silently miss the fifth.

THE VALUE IS A CREDENTIAL AND THERE IS NOWHERE FOR IT TO LEAK FROM. `EngagementRecord` carries `bypass_header_names` — names only — so `GET /cockpit/engagement`, the LLM proposer context and every rendered report see the name and never the value. `ReplacerRule` has a `replacement_set` boolean and no `replacement` field. That is build #15's rule restated: never handing a secret over cannot regress, while redacting it afterwards depends on a redactor being correct forever.

And it never touches a command line either. `_api_get` already keeps the API key out of the URL because ZAP records what passes through it; a replacer value on a GET would land in that same history *and* on the `docker exec ... curl ... <url> argv` that `ps` on this host can read. So a new `_api_post` sends action parameters as a **form body on stdin** — build #13 part 3's trick applied to a different secret. `docker exec -i` is load-bearing there: without it the container's curl reads a closed stdin, sends an empty body, and ZAP answers a cheerful `{"Result": "OK"}` for a rule with no value in it. Another OK that is not a result.

ZAP PERSISTS ITS CONFIGURATION — third instance in one build. A replacer rule set for one engagement survives into the next, which is a credential leaking to a third party by nothing worse than forgetting. So: every install clears first and re-adds; `stop_proxy` clears **before** the kill (afterwards there is no API to clear through while the persisted config keeps the rule); exiting an engagement clears it from any live proxy; and what is reported is what `replacer/view/rules` **holds**, never what was sent.

One thing is honestly unmeasured and is built to say so. The Replacer add-on **renamed its actions** across versions — `addRule` versus `addReplacerRule` — and the daemon was not running while this was written, because build #17's teardown stopped it. Guessing one spelling would produce this module's signature failure: a confident `{}` from a 404 reading as "the rule is not there". So both are tried in order, the **read-back is the arbiter**, and `docs/proof/build18_bypass_header.py` prints which one answered. That is how the guess becomes a measurement.

Item 2 — is it fronted, and what is behind it? Passive, and `unknown` is a real answer

`cockpit/fronting.py` resolves the CNAME chain, reads the `Server` header, maps the address to an ASN through Team Cymru's DNS zones, and pulls SPF, MX and optionally certificate transparency. Everything is a lookup except **one HEAD request per host** — the same request a browser makes opening the page. No scanning, no brute force, no subdomain guessing, and the test asserts the module invokes no scanner.

It runs **from inside the open sandbox**, `docker exec`, `argv-only` — `repeater.py`'s rule #1 restated, so no new egress path is created from the Windows host.

The verdict distinguishes `unknown` from `not-fronted`, and that is the whole design. A host whose lookups all failed reports `unknown`; only a host that *answered* and showed no marker reports `not-fronted`. Reporting the first as the second would send an operator at a fronted host with the wrong toolchain — build #17's confident-zero lesson, applied before it could happen. Suffix matching is **dot-anchored**, because `notakamai.net.example.com` contains `akamai.net`; this repo has been bitten by a fragment match in both directions already.

A discovered origin is REPORTED and never added. `add_pivot_subnet` is the one deliberate, audited widening path and a human uses it. The test that asserts this had to be rewritten from a substring check to an AST walk — because `fronting.py`'s own docstring `names add_pivot_subnet` in the sentence explaining that it must never call it. Item 8's lesson arriving from the other direction, inside the test written to check for it.

Item 3 — a scan policy is a list of checks to switch off, each with a reason

Most of build #17's 403s came from checks that could never have applied: a Shellshock probe against a Next.js storefront costs a request, earns a WAF hit, and has nowhere to land. `targeted-web` switches off nine rules that are locked to a platform the target is not running — the three C-server memory checks, Shellshock, Apache SSI, `.htaccess`, ELMAH, the Java `/WEB-INF` disclosure and PHP remote file inclusion. **Each entry carries WHY**, because "off because the target is not a C server" is a claim a reviewer can disagree with and a bare list of plugin ids is not.

It is applied FROM A KNOWN BASELINE on every scan and never reset afterwards, and two facts force that shape. ZAP persists scanner state, so a disable-only apply would inherit whatever the previous scan switched off and call it this policy — the same error as measuring `api.key` against a config a previous run wrote. And a scan is **asynchronous**, so "reset when finished" would either race the running scan or need a watcher that outlives the request. Applying at the start of the next scan gets the same property with no race. The consequence is stated rather than buried: between scans the daemon holds the last policy applied, and `observed_scan_policy` exists so that is read from ZAP rather than from our hopes.

An unknown policy name is a default, not a refusal. A typo should cost a wider scan the operator can see reported back on `Scan.policy_observed`, not a 403 that reads like a safety verdict. Nothing about a policy decides whether a scan may happen.

Item 4 — the scanner cannot shape payloads, and saying so is the result

Zaid took this decision explicitly and it was not re-litigated. What *was* required was an honest answer about where it can be built, and the answer is **the scanner cannot; the repeater can**.

ZAP's active scanner generates payloads inside each rule, in Java, at scan time. Its API exposes which RULES run, how HARD they try, and which INPUT VECTORS they fill — none of which is a transform. The two near-misses are both narrower than they look: the Custom Payloads add-on replaces payload *lists* for the handful of rules that opt in and cannot encode what a rule already generated, and pointing the Replacer at payload bytes would mean a regex matching every payload every rule might emit. A knob there would have been a switch in the UI with nothing different on the wire — a fake knob, which the plan explicitly ruled out.

So `cockpit/shaping.py` builds it where it works: the repeater, which is operator-driven, sends one request at a time, and is exactly the surface a human uses when they know which parameter to attack and are watching a WAF refuse them. Six value transforms (percent-encode, double-encode, case variation, `/**/` comment insertion, tab substitution, `+` substitution) and two request transforms (parameter pollution, chunked framing). **It is an option, not a gate** — no confirm, no acknowledgement, no refusal if unset; the repeater is human-only and a human clicking Send is the approval. A test asserts `shaping.py` contains no `raise` at all.

The span is marked explicitly, with `[[` and `]]`, because guessing would be worse. A transform applied to a whole URL would encode the scheme, the host and the parameter names into a request that goes nowhere. The markers are stripped **even with no shapes selected**, which is what makes shaped-versus-unshaped a one-variable comparison rather than two different requests.

The scope check runs on the SHAPED url. Nothing stops an operator marking a span inside the host, and checking the composed URL while sending a different one is "the gated argv is not the spawned argv" wearing a new hat — a scope check on bytes that never went on the wire. The run record stores the shaped URL for the same reason: an audit trail showing the unshaped request would misdescribe every shaped send.

`chunked` also suppresses `Content-Length` explicitly. A bare `Transfer-Encoding: chunked` header makes some curl versions send **both** framings, which is request smuggling by accident rather than a shaped payload.

Item 5 — curl-impersonate, and why a User-Agent could never have worked

Nine of eleven in-scope assets refused a bare HTTP client before request one — h2 stream reset on HTTP/2, timeout on HTTP/1.1. Two different failure modes on two protocols is an edge refusing the client, not a quirk. Build #17 found the discriminator, and it is **not** the `HeadlessChrome` token: headless Chromium sends **no Client Hints at all**, and `--user-agent` does not add them. Spoofing the UA failed identically to not spoofing it.

`curl-impersonate` is a curl built with Chrome's and Firefox's cipher and TLS-extension order, HTTP/2 SETTINGS and pseudo-header order, and the full default header set including `Sec-CH-UA`. Dockerfile layer 9g installs it and symlinks the wrappers. **The wrappers are the product, not the binary** — `curl-impersonate-chrome` alone is just a curl with a different TLS library; the `curl_chrome*` scripts carry the header list and the ciphers. The layer therefore asserts both, and it introduces **no new gate**: `curl-impersonate` sits in `_MUST_NOT_FIRE` beside `curl` and `httpx`, because a fetch is a fetch. Marking it dangerous while plain curl sits clean would be a red-confirm that fires on fetching a page, and one that fires on everything stops meaning anything.

Recorded honestly: this makes HackPit's traffic indistinguishable from a browser's. It is built deliberately, for an authorized safe-harbour program. The pinned release impersonates Chrome 116 and Firefox 109, and **whether that vintage still satisfies a 2026 bot manager is a measurement, not an assumption** —

`docs/proof/build18_impersonate.sh` makes it, every probe carries its own control, and a `PASS-NO-BENEFIT` outcome is a legitimate recorded result rather than a failure.

Items 6 and 7 — the hard part already existed; what was missing was telling ZAP what it means

A human logging in by hand through the published proxy puts a **live session inside ZAP**. What was missing was never the session.

Tier 2 needs no credentials at all: a Context over the target's origin, cookie-based session management, and the logged-in / logged-out indicator regexes. The include regex is quoted with `\Q` and `\E`, because a host contains dots and an unquoted regex would read each one as "any character" and pull unrelated domains into the context — the same class of bug as a target smuggling a `&` into the scan URL. `start_scan` now **looks the context up** for the target it was given rather than asking the operator to name it again, which would be a second place for the two to disagree. A context with a configured USER switches the scan to `scanAsUser`, because ZAP can only re-authenticate mid-scan if it knows who to re-authenticate as. With neither, the scan URL is byte-for-byte what it was before build #18, which is what makes this additive.

Tier 3 is the trap Zaid chose to take on, entered with eyes open. Build #15 declined it on the grounds that per-target auth scripting is a problem the codebase would then own forever. The defence is to stay declarative — a login URL, a request body with ZAP's two placeholder tokens, two indicator regexes — and to refuse to grow a scripting language. If a target needs JavaScript to log in, the answer is Tier 2 and a human's browser.

No model in the module has a password field. The credential is *named* by the same `{session_id, kind, principal, domain}` reference the Windows-profile path already uses, resolved server-side out of the state vault by exactly one function, and delivered to ZAP through `_api_post` — on stdin. A GET would have put the password in ZAP's own recorded history, on a readable argv, and into the artefact a report is rendered from.

Two failure modes are named rather than left to be discovered. A login body missing ZAP's password token produces a POST that literally sends the token text — a request that succeeds, authenticates nothing, and leaves the scanner running unauthenticated while every indicator says the login failed; that warns. And a **named credential that is not in the vault refuses**, because scanning unauthenticated instead would report zero findings off a login page, which reads exactly like a secure application. Everything else warns and continues: a weaker context is a weaker scan, not an unsafe one.

Tier 3 is **UNVERIFIED** against a real target. Zaid has no account on any in-scope host, so it is exercised against the LAB target only. Nothing here implies otherwise. Tier 2 is unaffected — it needs no credentials at all.

Item 8 — the silent-empty sweep, and the two that mattered

`backend/tools/silent_empty_scan.py` walks `cockpit/`, `state/`, `arsenal/` and `reasoning/` for the build #17 shape: an exception handler or falsy guard returning `[]`, `""`, `{}` or `0` where the caller cannot tell "empty" from "failed". **AST, not substring** — the docstrings in `cockpit/proxy.py` quote `return []` repeatedly while *describing* this very bug, so a grep would report that file as its own worst offender. The scanner ranks by how likely the caller is to be fooled, exits 0 always (a reporting tool that failed a build would turn every legitimate empty return into work, which is how a control stops being read), and prints ASCII because this console is cp1252.

118 hits. Most are fine and the report says which: a parser returning `[]` for input that legitimately contains nothing is correct; `get_active("")` returning `None` is a question with an answer. Two were not, and both were fixed.

The one with teeth: `observed_scans` failed **OPEN on a bound**. It answered `[]` for both "ZAP knows of no scan" and "the read failed", and `start_scan` used that to enforce ONE SCAN AT A TIME — a bound on concurrent attack traffic against a live production target. An unreadable daemon therefore *granted* a second scan. This is `clash_refusal`'s build #17 defect one function away in the same file: a check protecting against a state it could not observe. `scans_snapshot` now returns `(scans, read_ok)` and an unreadable list refuses, naming the harm. **That is repair, not a new prohibition:** the refusal already existed and its gate is already `limit`; what changed is that a failed read used to make it grant. It also costs nothing real — a daemon whose scan list will not read is a daemon whose scan *start* would fail one call later, with a worse message.

The one that travelled furthest: `scan_alerts`. A failed read returned `[]`, and `POST /proxy/alerts/ingest` wrote "0 findings" into engagement state — from which a **report** is rendered. A confident zero all the way to a deliverable. `alerts_snapshot` reports `read_ok` and the ingest route now refuses rather than persisting a zero it cannot vouch for.

Deliberately NOT fixed, with the reason: `parse_message` and `parse_alert` return `None` for a malformed record and `history()` filters those out, so a capture with 200 messages of which 50 are unparseable reads as 150. That is a real gap, but it is partly visible already — the `:proxy` panel shows ZAP's own `captured` count beside the history length, so the two disagree in front of the operator. Closing it properly means a count on a route whose response model is a bare list, and that is a schema change with no measured need behind it yet.

Item 9 — one authored entry, and the KB was checked rather than assumed

`sources/random Custom Toolsmini scripts for.txt` is three HTB helper scripts. Two are noise: a random-box picker for a wheel-of-names, and an nmap wrapper that shells out three times. The third, `upload.py`, is a catalogue of ingress/egress tool-transfer one-liners — **T1105**.

The token diff nominates; it never confirms. Five existing KB entries already cover file transfer, and between them they hold certutil, `DownloadString`, `nc -lvnp` and one mention of `/dev/tcp`. What none of them has is **the part that decides which method you can actually use**: the raw-fd `exec 3<>/dev/tcp` HTTP download for a host with no wget, curl or nc; the upload direction over HTTP (`python3 -m uploadserver` — plain `http.server` will not accept a POST); PowerShell's **8191-character command-line ceiling**, which is what rules base64 out at roughly 6 KB; and the netcat-to-no-netcat pairing in both directions with `-q 0`. So: **one authored entry about choosing**, not a sixth copy of the same command list. Prior batches recorded the discipline — 24 repos to 2 entries, 7 GitBook spaces to zero.

`data/kb/entries.jsonl` 2743 to 2744, verified after the run and after Defender has had its chance at the file. The downstream artefacts had to follow, exactly as the pipeline's build-order note says: `ingest_corpora` re-run to restore the fixed point (the byte-identity test caught it), `scripts_index.py` rebuilt (D6's guard caught it — "built against 2743, the KB now holds 2744"), and the KB fixture regenerated. Three separate guards fired, each naming its own fix.

What was NOT built, and why

- **Route auth and scanner pacing** — Zaid's call, skipped for this build.
- **No frontend: SUPERSEDED** — the frontend was built and looked at; see the RUN section below.
- **The parse-drop-count from item 8: SUPERSEDED** — closed; see the RUN section.
- **The image is NOT rebuilt in this commit: SUPERSEDED** — rebuilt, recreated and every proof run; see the RUN section. The original note read: Layer 9g is written and asserted at build time, but a rebuild is about 45 minutes and recreating the engage sandbox destroys the ZAP daemon and its capture. It is bundled as ONE rebuild, deliberately last, and `docs/proof/build18_run.sh` prints the exact sequence — `docker compose ... build engage-sandbox` (the SERVICE) then `--force-recreate` then re-run `docker/proof/browser_intercept_proof.sh` (still 25 passed, 0 failed) and the impersonation proof. Until then `build18_impersonate.sh` reports NOT-RUN and says why, rather than reporting every host as refused.

Verification

Hermetic safety suite green, **82 test files**, every one exited 0 — four new: the bypass header (value on stdin only, on no model and no argv, cleared before the kill), payload shaping (markers stripped as the control, the shaped URL is the scoped URL, no gate anywhere), scan policy plus authenticated scanning (baseline-then-read-back, no password field exists, the scan URL is additive), and CDN fronting plus the silent-empty sweep (`unknown` is a real answer, an unreadable read is not a zero). Every control is phrased "*not refused at THIS gate*": a fully approved lab request legitimately ends at `gate='sandbox'`, and CI has no Docker.

Two existing locks fired on a rename, and both were REPAIRED rather than flipped. Splitting `observed_scans` into `scans_snapshot` and `scan_alerts` into `alerts_snapshot` broke a lock asserting the old symbol name — while the property it was written for still held. Pinning the literal would have made each a lock that only ever says "you renamed something", which teaches people to update it without reading it; that is build #17's browser-id lesson, and both now assert the predicate. A third lock, `test_arsenal_safety`, correctly refused `curl-impersonate` until it had a pinned danger verdict, and then correctly refused two binary names the catalog does not actually invoke.

`test_redirector.py` flaked twice inside the full suite with the recorded Windows message — "*no bindable port for kind=1 outside the self-mapping range*" — and passed three times out of three run on its own; the suite then passed whole. That is Hyper-V/WSL UDP exclusions moving under it, not a code fault, and it is written down again because one red run of that file is not evidence until it is repeated.

`data/kb/entries.jsonl` at 2744. Every real-target and live-daemon verification is prepared as a self-verifying script under `docs/proof/`, each printing `VERDICT=` and exiting non-zero on failure, with `build18_run.sh` chaining them on `captured` exit codes — `RC=$?` after a pipe is tee's status, which is 0 whenever tee could write the file, and that bug shipped into two wrappers in build #17 before it was caught.

Build #18, RUN — the image rebuilt, the proofs executed, and eight defects the read-backs caught

The image was rebuilt (layer 9g), the engage sandbox recreated onto it, the port re-published through `exposure.py` 's own path, and every proof run. `docker/proof/browser_intercept_proof.sh` is 25 passed, 0 failed on the new image, so nothing build #17 established was disturbed.

Item 5, the headline result. `curl-impersonate` completes requests bare curl cannot.

host	bare curl	curl_chrome116	verdict
example.com (control)	200, 559 B	200, 318 B	both fine — no client wall
api-prod.thatconceptstore.com	000, 0 B	404, 431 B	IMPERSONATION WINS
lapi.yellowblocks.me	000, 0 B	404, 431 B	IMPERSONATION WINS

000 is curl never completing a request — the same h2-reset/timeout wall build #17 measured on nine of eleven hosts. A 404 is the ORIGIN answering: these are API hosts, so / legitimately has nothing at it, and the point is that the request arrived at all. The Chrome 116 fingerprint is enough for this edge, which was an open question when the layer was written. The control matters as much as the result: without example.com answering both clients, "both refused" and "no egress" would look identical.

Item 2 answers the question the build exists for, and the answer is NO. The plan hoped `api-prod.thatconceptstore.com` and `lapi.yellowblocks.me` might be reachable directly, in which case two of eleven hosts would need none of items 1, 3, 4 or 5. They are not. Both are Akamai — and they resolve through the same edge node, `e28210.a.akamaiedge.net`, to the same two addresses in AS20940. They are two hostnames on one Akamai property, so they need all of it. A premise of the plan, tested and false.

Item 4 measured, after the proof caught its own harness being wrong. Against a naive single-decode signature matcher, with both controls blocking: 5 of 7 shapes turn a 403 into a 200 — double-url-encode, case-vary, sql-comment, whitespace-tab, and sql-comment+url-encode. Two correctly do not: plain url-encode (the matcher decodes once, which is exactly what that transform is a probe FOR) and param-pollution against a matcher that reads the whole query string. That is the honest shape of a result — some transforms defeat this rule and some do not, and which is which is the useful part.

The first run of that proof FAILED, and it was right to. Four cases came back HTTP 000: the payload carried a raw space and a space is not legal in a URL, so curl never sent them. The control did not block, the script said so and exited 1 rather than reporting four bypasses that were really four requests that never left. The payload now travels in a POST body, param-pollution gets its own block with its own control, and a request that never completed is counted as NOTHING rather than as a bypass. Build #17's lesson — three of its findings came from its own scripts being wrong — arriving on schedule, and caught in one run by the control.

Item 3's first measurement was PASS-NO-BENEFIT, and the read-back said why. `targeted-web` saved -1 requests of 376 while `not_held` listed all nine plugin ids. The cause: ZAP rejects the ENTIRE `disableScanners` list if ONE id in it is not installed. This build has 30001 and 30002 but not 30003, so `ids=7,10045,...,30003,...` answered `{"code":"does_not_exist"}` and disabled nothing at all — a policy that quietly did nothing, which is precisely the fake knob the plan forbade. `apply_scan_policy` now intersects against `scanner_ids()` first, READS the answer instead of discarding it, and falls back to one call per id. Re-measured: `targeted-web` sends 264 requests where default sends 376 — 112 fewer, 29.8%, with no alerts lost.

Items 6 and 7: Tier 2 and Tier 3 both green against the lab, after four more read-back catches. Every one of these answered `{"Result":"OK"}` and either did nothing or was read wrongly:

- `includeRegexs` comes back as a JSON-ENCODED STRING, not an array, while `excludeRegexs` is the string `"[]"`. An `isinstance(x, list)` reader fell through to `[]`, so a context whose include regex WAS installed reported as having none.
- The two context views nest differently. `getSessionManagementMethod` answers `{"methodName": ...}`; `getAuthenticationMethod` answers `{"method": {"methodName": ...}}`. Reading only `methodName` reported no authentication method for a context that had one, and every Tier 3 context read back as Tier 2.
- Setting the authentication method REPLACES it, taking any indicator set beforehand with it. A Tier 2 run held both indicators; the identical Tier 3 run held neither. Indicators now go last.
- `authMethodConfigParams` is itself a query string, so a login body full of `&` and `=` has to be percent-encoded before being nested inside one.
- `ascan/action/scan` answers `{"scan":"7"}` but `scanAsUser` answers `{"scanAsUser":"6"}`. Reading only `scan` made every authenticated scan raise `ZAP` returned no scan id while the scan was running — attack traffic in flight, reported as not started, with no id to stop it by. The worst direction that error could have taken.

Three of those five are the same shape as item 8's whole subject: **an unexpected SHAPE read as an ABSENT VALUE**. They were found in the code written to hunt exactly that, which is worth stating plainly rather than smoothing over. All six are now locked by hermetic tests that assert the property rather than the spelling.

The measured fact item 1 was built to leave open is now closed. This ZAP 2.17.0 accepts `/JSON/replacer/action/addRule/`. The proof is 11 of 11: the rule is held (read back, not the OK), it is enabled, it holds a non-empty replacement, the reported rule does NOT carry the value, **the header is on the outgoing request read back out of ZAP's own history**, the value on the wire is the value that was set, and it comes off again.

The frontend — built, and LOOKED AT

Every capability is now reachable from the cockpit: the bypass header and the authenticated context on `:proxy`, the scan-policy selector with each disabled rule's reason beside it, the shaping controls and a byte-level preview on `:repeater`, and the fronting sweep on the engagement panel — **next to the scope, because the scope is its input**. It gets no tile of its own: none of the four band hints is TRUE of a passive OSINT sweep, and a tile whose band makes a false posture claim is worse than no tile.

`tsc --noEmit 0`, `next build 0`, `eslint` back at the **11 errors + 1 warning** baseline — one new error appeared and was fixed by DERIVING the header names from the live proxy instead of storing-and-syncing them in an effect, per the frontend's own note.

Four things were only visible by looking, which is the whole point of the rule.

1. The shaping heading rendered **on the same line as the first checkbox** — `hp-rp-opts` is a flex row and the subhead belonged outside it.
2. The fronting buttons wore `hp-ck-approve`, the **approve-a-real-target-command** style, which made two DNS lookups the loudest thing on the page. Same class of defect as the `.hp-tn-start` visual-hierarchy bug the CSS-vocabulary test records.
3. `ONE <code>HEAD</code>request per host` — a missing space in JSX.
4. A candidate origin rendered as an empty `mx: chip. example.com` publishes a NULL MX (`0` , RFC 7505 — "this domain sends no mail"); stripping the trailing dot left an empty host and the code appended a lead pointing at nowhere. It is now a note explaining the null MX, and the SPF side is guarded the same way. **No**

typecheck, build or lint could see any of the four, and the CSS-vocabulary test — which passed throughout — only ever claimed the classes exist.

The parse-drop count, closed

Build #18's first pass left this open as "real but partly visible". It is closed now, because the partly-visible version is the dangerous one: not a confident zero but a **confident UNDERCOUNT**, which is harder to notice precisely because it looks plausible. A window of 200 exchanges of which 50 were unparseable arrived as 150 and read as less traffic.

`GET /cockpit/proxy/history` now answers with a `HistoryPage` — `exchanges`, `total`, `window_start`, `returned`, `dropped` and `read_ok` — rather than a bare list, and the `:proxy` panel says "showing *N* of *M* captured" plus, when it is not zero, "*N* rows could not be parsed and are missing from this list. The traffic happened; it is the reading of it that failed." `alerts_page` does the same for alerts, and the ingest response carries `alerts_dropped` because a finding that never parsed is a finding that never reaches a report. `history()` keeps returning bare rows for the callers that only want shapes to judge.

An empty window on a daemon that ANSWERED still reports `read_ok: true` — the trustworthy zero. That distinction is what makes `read_ok: false` mean anything.

The curated exploit overlay (2026-08-05)

A new source arrived — `sources/some vul.md`, 216 KB of disclosed vulnerability reports with working PoCs. The first question was whether it belonged in the KB. It **did not**, and the saturation check is the reason: request smuggling appears 449 times in the KB, `169.254.169.254` 402 times, command injection 277, sanitisation 244, `pull_request_target` 28. The one candidate that looked novel — the `Connection: close\t` parser trick — returned zero for `obs-fold`, `token parser` and `parsing differential`, but the recorded rule is that a token diff only ever NOMINATES. Grepping the *concept* found **44 entries already pairing smuggling with a parser/delimiter idea**, including dedicated `Request Smuggling`, `Carriage Return Line Feed`, `SMTP Smuggling` and `Special HTTP headers` rows. Writing it up would have duplicated them. That is three sources running where the honest answer was few-or-zero — the saturation signal doing its job.

The verdict was then overruled, and correctly, for a reason better than the one first given. Zaid asked for it anyway; looking harder at what the reports actually assert:

libcurl's SSH connection-reuse guard `ssh_config_matches()` — added for CVE-2022-27782 and reaffirmed by CVE-2023-27538 — **is dead code in every release since 7.83.1.**

This index's entire claim is that **the version verdict outranks token similarity**. Here the *public* version verdict is wrong: curl 8.14.0 reads as patched for CVE-2022-27782 and is not, and the Rocket.Chat report bypasses the published fix for CVE-2024-39713. An index that answers "patched" there is worse than one that says nothing. That is index-shaped, not KB-shaped — a keyed product+version lookup, which is exactly what `backend/exploits/` is for.

The overlay is SOURCE; the mirror is DATA

`data/kb/exploitedb.json` is generated wholesale from the sandbox image's exploit-db catalogue, and **the whole of `/data/` is gitignored**. A hand-authored row placed there would never be committed and would be destroyed by the next ingest. So `backend/exploits/curated.json` lives beside the module that reads it, is tracked, and is

merged at load — $47,108 + 5 = 47,113$.

The lookup keys are **DERIVED at merge, never read from the file**. A hand-authored JSON carrying its own `by_cve` / `by_token` offsets would be one edit away from pointing at the wrong row, and an index that answers *confidently with the wrong entry* is worse than one that answers nothing. The file states facts; the code computes the keys, and a test walks every key to assert it lands on an entry that really names it.

`gte` — the index could not say "there is no fix"

Every existing kind carries an upper bound, and `versions[-1]` means **the first patched release** (stated in `_version_verdict`'s own docstring, compared exclusively). An unfixed vulnerability has no such number. Forcing `curl` into `range` would have required naming an upper bound, and the index would then have **announced a patch that has never shipped** — precisely the failure it exists to prevent. `gte` reports *"at or above 7.83.1, no fixed release"*, and a test asserts the reason string can never contain "fixed in".

Measured after the merge: `curl 8.14.0` → `in-range, at or above 7.83.1, no fixed release`; `curl 7.83.1` → the same (the lower bound is inclusive); `curl 7.80.0` → falls through, because it predates the introduction. `Rocket.Chat 7.13.2` and `Trix 2.1.16` → `exact`.

What was NOT indexed, and why that is the discipline working

Monero's ZMQ RPC log injection is left out. Its advisory states *last affected* versions (`v0.12.0.0` – `v0.12.4.0` , `v0.13.0.2` – `v0.13.0.4`) and this index's `range` / `lte` upper bound is the *fix* version. Writing `0.12.4.0` there would read as "fixed in 0.12.4.0" and mark a genuinely vulnerable build safe — the same class of lie `gte` was added to avoid. Expressing it needs a per-entry boundary declaration, which is a change to comparison code two subsystems share with **opposite conventions**; recorded rather than forced. The 2026 credential-permissions CVEs are out too: the extracted text does not name a product precisely enough to key on, and a guessed product token is worse than an absent row.

An existing lock broke, and it was right to

`test_a_missing_index_degrades_quietly` loaded one nonexistent path and asserted `ready is False`. With the overlay merging at load, "nothing is built" became two files and the test failed. **That is the lock working** — it caught a behaviour change rather than quietly testing something narrower than its own name. It now names both absences explicitly, and the state it used to cover by accident got its own test: mirror absent, overlay present, which is a fresh clone before the ingester has ever run.

`data/kb/entries.jsonl` is **unchanged at 2744** — this was an index decision, not a KB one. Suite **82 files, all green**.

Two sources evaluated and declined (2026-08-05) — with the reasons, so nobody re-reads them

`sources/gitbooks-manifest.md` set the precedent: an evaluation that ends in nothing is only worth the tokens if the *reason* is written down, or the next person re-reads 200 MB to reach the same answer.

tim-barc/ctf_writeups — 209 PDFs, and the commands are pictures

The repo looked promising against a strong prior: the 0xdf batch produced 27 fingerprints from 106 writeups, making writeups the highest-yield source type this KB has ingested. **The prior did not transfer, for two measured reasons.**

First, it is not the same kind of corpus. Filenames nominate; content decides. Roughly 147 of 209 are CyberDefenders / BTLO / LetsDefend labs, and much of the remainder is HTB Sherlocks and TryHackMe blue rooms (`brutus` , `conti` , `lockbit` , `boogeyman` , `snort_*` , `tshark*` , `zeek_exercises`). The genuinely offensive subset is ~20–25 beginner TryHackMe rooms.

A real gap did open up, and it is worth recording because it was surprising. Those rooms are absent from the KB *by name* — `pickle rick 0` , `mr robot 0` , `basic pentesting 0` , `bounty hacker 0` , `wgel 0` , `colddbox 0` . The 200 existing `writeup` entries are HTB boxes and challenge categories, 176 of them carrying steps with code blocks, which is what `attack_path.build_writeup_path` replays. TryHackMe is a platform the KB does not cover, so those rooms would have added something real.

Then the measurement killed it. Running the repo's own ingester over all 62 non-blue candidates, the offensive rooms yield **nothing runnable**: `pickle_rick 0` command lines, `basic_pentesting 0` , `photographer 0` , `blogger1 0` , `dav 0` . The reason is visible in the extracted text:

"Here is the Nmap command that was used:" ... "Next, I used Gobuster to brute-force directories"

The commands are screenshots. 23–42 images per file and ~350–500 characters per page: these are narrated screenshot writeups, where the prose describes the attack and every command is an image. Ingesting them would produce entries that look like box walkthroughs and cannot drive a single command — worse than absent, because the writeup-first attack path would find them and have nothing to run. The rows that *did* score highest for commands are the blue ones (`masterminds 18` , `snort_challenge 5` with 112 blue markers), where the extracted "commands" are the analyst's `tshark` and `volatility` invocations, not an attacker's.

This is the "an exit code is not a result" family one level further out: **the file exists, it parses, it produces entries, and the entries are empty of the only thing that makes them useful.** OCR could recover the commands, and is deliberately not proposed — a mis-transcribed command in a corpus that drives an attack path is a fabrication with a plausible shape, which is the one failure mode this KB's whole curation discipline exists to prevent.

sources/some vul.md — declined for the KB, redirected to the index

Covered in full in the curated-overlay section above: KB-saturated on every class it teaches, but three of its findings belonged in the CVE→exploit index because the *public version verdict* for them is wrong. Declined as a KB batch, accepted as five index rows.

An undeclared dependency, found on the way

`pipeline/ingest_box_pdfs.py` does `from pypdf import PdfReader` at module import, and **pypdf was declared nowhere** — not in `backend/pyproject.toml` , not in any requirements file. The ingester could never have run on a clean checkout; it had only ever run where pypdf happened to be installed. Now an optional group beside `codescan` , which is the existing precedent for tooling the backend serves every route without.

Build #19 — the interactive half, and eyes for the agent (2026-08-05)

Build #18 made HackPit *reach* targets that were refusing it. What it never touched is a gap measured by comparing HackPit against Burp rather than against its own backlog: **the proxy records faithfully and cannot change anything mid-flight**. Four things fall out of that — interception, a repeater cookie jar, history filtering, an intruder — plus an MCP server, because `hexstrike-ai` already gives an agent hands and nothing gives it eyes on *this* engagement.

The build adds no gate, no confirm, no blocklist and no allow-list narrowing — not in the product and not in a proof script. Where something could refuse, it warns and continues. Two places decline to send an API call, and both are named below with the argument for why that is a defect fix rather than a prohibition.

Item 1 — the go/no-go, and it was a GO for a reason nobody would have guessed

Nothing in item 4 was written before this answered. Build #14 was written against `zap-baseline.py`, which does not exist in Kali, and only the image build caught it — so `docs/proof/build19_break_api.py` enumerates what ZAP 2.17.0 actually exposes and then drives the whole loop against a real origin, **reading that origin's own access log as the arbiter**. Every one of these endpoints will answer `{"Result": "OK"}` for things it did not do.

VERDICT: 27 passed, 0 failed. Hold, read, replace, forward and drop all work. The origin log shows only the *replaced* requests; the dropped ones are absent; breaking off restores traffic.

Five traps, and the first four were each written down WRONG before they were measured:

1. `brk` IS NOT IN THE INSTALLED ADD-ON LIST AND THE API IS THERE ANYWAY.

`autoupdate/view/installedAddons` returns 48 add-ons and the Break add-on is not among them; every `break/view` and action answers regardless, because `BreakAPI` ships in ZAP core. **A go/no-go taken off that inventory would have said NO to a feature that works.** An add-on list is not an API surface — and neither is the documentation: `core/view/apiSummary` is `bad_view` in 2.17.0 and `/JSON/break/view/` returns ZAP's *welcome page* as a 200 of HTML. The surface had to be established by probing names and reading the error CODE, where a genuinely absent action answers `bad_action` and one that merely needs a held message answers `does_not_exist`.

2. `http-all` IS THE ONLY `type` THE ACTION ACCEPTS. `http-request`, `http-response` and `http-sender` — two of which name views that DO exist — every one answers `illegal_parameter`. This module's first docstring claimed `http-request` "answers OK and holds nothing"; the proof disagreed and the proof was right. `BREAK_TYPE` is a constant rather than a request field, for the simplest possible reason: there is nothing else to choose.

3. `continue` TURNS BREAKING OFF. `step` AND `drop` LEAVE IT ON. `isBreakAll` reads true while a request is held and false immediately after `continue`. That is ZAP's break-panel semantics (Continue means "let everything go"), and it is not a detail: an operator who forwards a request expecting to catch the next one catches nothing, and the global banner correctly vanishes underneath them. `release()` reads the state back and says so in `detail`.

4. `isBreakRequest` IS A SETTING, NOT A STATE. It reads true whenever breaking is switched on, with nothing held. A panel wired to it reports "a request is waiting" forever, so `held` is derived from `httpMessage` being non-empty and a test asserts that by AST.

*** 5. A `drop` WITH NOTHING HELD PERMANENTLY WEDGES THE BREAK MANAGER. ***

The worst one, and **the proof script found it by doing it**. After a single stray `break/action/drop/` against a daemon holding nothing, that daemon still HOLDS requests — `isBreakAll` true, the origin never sees them, the client blocks — but `break/view/httpMessage` returns `""` forever and `setHttpMessage` therefore never applies.

Interception silently becomes a way to freeze your own browser with no way to read or release anything.

It took four wrong hypotheses to find, and the wrong ones are worth recording because each was plausible and each was killed by a measurement rather than by argument:

hypothesis	how it died
the daemon degrades over a long session	a brand-new daemon wedged on the first stray drop
unreadable from the host, readable from inside	an interleaved read got <code>""</code> from both , on one held request
it is a latency problem	301 host polls over 45 s, never readable
<code>drop</code> -then-enable stops it holding	an A/B said both arms still HOLD — but that A/B ran on an already-wedged daemon, which is why it looked like a refutation

What settled it was a **single-variable experiment on fresh daemons, twice in each direction**: with the stray drop, 23 passed / 4 failed; with that one line removed, 27 / 0.

So every drop in the product is guarded by a read-back. `release()` will not send `drop` unless something is held, and `panic()` drops only when `before.held` says there is something to drop. **That is not a prohibition on the operator** — dropping nothing was never a meaningful action — **it is refusing to send an API call that breaks the daemon**. `panic()` is the function most likely to be pressed when nothing is held, because it is the "I think the target is down" button; an unguarded drop there would have destroyed interception for the rest of that daemon's life at exactly the moment the operator was most confused.

And the reset in the proof is shaped by the same finding. The obvious reset is `drop` then `state=false` — which is the bug. It now reads first and only drops something really there, which is exactly the guard the product carries.

Item 2 — the repeater cookie jar, and a value that never leaves the module

Every authenticated flow broke on the SECOND request: the login returned a `Set-Cookie`, nothing kept it, and the follow-up went out anonymous. `cockpit/cookiejar.py` is a per-session RFC 6265 jar — domain and path matching dot-anchored, `Secure` honoured, `Max-Age` beating `Expires`, an expired cookie DELETING rather than storing (that is how a server logs you out, and a jar that kept it would silently re-authenticate the next request).

THE JAR IS STATE, AND STATE THAT SILENTLY CHANGES A REQUEST IS A TRAP, so:

- `CookieAttachment` **HAS NO `value` FIELD AT ALL**. It carries the name, the domain and path it was stored under, and the URL of the response that set it. That model is what goes on the exchange, into the API response and onto the screen. Build #18's rule restated: never handing a secret over cannot regress, while redacting it afterwards depends on a redactor being correct forever.
- **A cookie the operator typed WINS**, and the suppression is named. An explicit `Cookie:` is them testing a specific value; a jar that overwrote it would make the request under test unreachable.
- `use_cookie_jar: false` **sends with no session WITHOUT emptying the jar** — testing what an unauthenticated caller sees is a real test and must not cost a session established by hand.

- `evil.example.com` cannot set a cookie for `example.com` (RFC 6265 §5.3 step 6). That is the one place the jar declines to store something, and it still does not refuse the SEND: it warns, drops that cookie, and the exchange proceeds.

THE REDACTOR THAT WOULD HAVE HAD TO BE CORRECT IS NEVER CALLED (checked, not assumed).

`report.py::redact_captured_body` knows the word "cookie" — and no production path invokes it; its only callers are tests. Meanwhile `report.py::_run_cmdline` renders a run record's command line **verbatim** into a report. So the run record has to be cookie-free *by construction* rather than cleaned, and it is: `args` stays `[method, sent_url, *shapes]`. A test asserts both halves, including that `redact_captured_body` still has no production caller — because if that changes, this lock's argument changes with it.

ONE Cookie: HEADER ON THE WIRE, NEVER TWO. RFC 6265 §5.4 says a user agent must not send two, and servers disagree about what to do when one arrives — some concatenate, some read the first. Emitting the jar as a second `-H` would have made the request's meaning depend on the target's parser, which is a silent wrong answer of exactly the shape this repo keeps finding. The jar merges into a single header, operator's cookies first. **When the jar contributes nothing the headers go out exactly as typed, duplicates included** — two Cookie headers may be precisely what an operator is testing, and this must not be the code that quietly decides they may not.

Item 3 — history filtering, where the counts ARE the feature

~1,300 captured messages and the only tool was `start / count`. `filter_history` filters server-side on host, method, status, URL substring, has-parameter, content-type and engagement scope — and **the host filter is scope.py's own parser**, so `api.example.com` and `*.example.com` mean here exactly what they mean in a scope field. A second host matcher is how `notexample.com` gets to match `example.com`.

A filter is the perfect place for this repo's recurring silent empty, because "no rows" is a completely ordinary answer. **Four different facts present as an empty list** and the response distinguishes all four: ZAP holds nothing (`total`), we read and nothing matched (`scanned`, `matched`), we could not read what ZAP holds (`read_ok`, `dropped`), or **we stopped scanning before reaching the rows that would have matched** (`truncated`). That last is the one a naive implementation creates, and it would have been `history`'s own build #17 defect in a new place with a scarier failure: "there are no 500s on this target" is a conclusion someone acts on. The scan pages through the whole capture and `truncated` is never silently true.

An engagement id that names nothing active does **not** refuse the search — `scope_note` says the filter was ignored and everything was kept. This is a read of traffic that already happened; refusing to show an operator their own capture because an engagement ended would be a prohibition invented by the tooling.

Item 4 — interception, which adds no gate because there is nothing to bypass

Turn breaking on, know when a request is held, read it, replace it, forward it, drop it. **This is a surfacing job, not a proxy-building job** — nothing in `cockpit/intercept.py` parses TLS, holds a socket or forwards a byte; every verb is one call to an API measured working first.

Every route is ungated in BOTH directions. A request is held, a human reads it, a human edits it, a human presses forward — the press IS the approval, the same position `:kali` and the repeater already take. Interception also strictly *reduces* what reaches the target: with breaking on, nothing goes anywhere until a person says so. "Off" matters most: while breaking is on **the operator's own browser is frozen**, so a gate that could refuse to stop it would be indistinguishable from the target having gone down.

That is also why the screen carries a **loud global banner** whenever breaking is on, saying which of the two states it is in ("a request is HELD" versus "traffic is being held, nothing yet — a frozen browser right now is something else"), with **drop-everything-and-turn-it-off one click away** without scrolling. The panel polls every two seconds, because a held request has to announce itself rather than wait to be refreshed into view.

The replaced request travels as a **POST form body**, never a URL: a held request routinely carries a session cookie and an Authorization header, and `_api_get` would put those in ZAP's own history *and* on a `docker exec ... curl ... argv that ps` on this host can read. Build #18's bypass-header reasoning, applied to a third secret.

Item 5 — the intruder, and why one approval may buy thousands of requests

HackPit refuses BATCHING ACROSS APPROVALS — an agent queueing five commands behind one human click. It has never refused ONE approval that produces many requests, and could not: `ffuf`, `nuclei` and the ZAP active scanner are each a single approval buying thousands, and the scanner's own measurement was 376 requests from one press. An intruder is that same shape, gated by the SAME four gates with no new ones.

Positions reuse the `[[...]]` shaping marker — one vocabulary, not two — and what is *inside* the markers is the BASELINE value, so `sniper` leaves the other positions at their original values rather than blanking them. Blanking would change two things at once and make every result uninterpretable. The baseline request is sent FIRST, with markers stripped, which is build #18's control half: without it "this response is different" has nothing to be different from.

THE PAYLOAD SET IS IN THE APPROVED SURFACE, COMPLETE AND UNTRUNCATED. Rendering "...and 4,993 more" would let a payload carrying `| sh` sit at position 5,000 where the danger heuristic cannot see it while the human approves a summary — Critical 2 expressed in a payload list. Measured: the danger gate fires on payload CONTENT, and a test plants a shell payload at position 200 and asserts the gate still catches it.

THE RED-CONFIRM IS THE GATE'S TO DEMAND, NOT THE FORM'S. A set of ordinary injection strings does not trip it; one carrying `| sh` does. Requiring the ack unconditionally would have been this build adding a confirm. **So would declaring `ffuf` an attack tool** to make the feature feel appropriately dangerous — that would have added a red-confirm to every existing ffuf run in the product, and a test asserts a plain in-lab `ffuf` still needs none.

A correction this build made to its own first draft: `intruder.py` originally justified declaring `ffuf` by saying it "is already in the allowlist". **There is no tool allowlist.** The module that sounds like one says so in its own comment — `SUGGESTED_COMMANDS` is "Informational only ... NOT an allowlist; anything may run". Membership proves nothing; the real property is that the declared command must DESCRIBE what runs, because that string is what a human approves and what the danger heuristic reads.

The scope check runs **per request, on the substituted URL** — a payload can contain a `.` and a `/`, and a marked span inside the host would otherwise send the request somewhere the gate never saw. An off-scope substitution is COUNTED and the job continues; one payload walking off-scope must not throw away the other 4,999. The ceiling caps and REPORTS rather than refusing, and stop is ungated like every other stop here.

Pitchfork and cluster bomb are deliberately absent: they need two independent sets, and a cluster bomb of two 1,000-entry lists is a million requests from one press — that deserves its own argument rather than inheriting this one.

Item 6 — the MCP server: eyes, not hands

Twenty-two tools (extended 2026-08-08 — see the addendum). Twenty-one are reads — engagement and scope, filtered proxy history, scan status, alerts, session health, interception state, CDN fronting, findings, engagement state, KB search, the arsenal, the CVE/exploit index including the curated overlay, a command-scope check, and the

offensive-build reads added on 2026-08-08: the cross-domain kill-chain graph, the cloud IAM and Active Directory attack-path graphs, the governance package (RoE/ConOps/deconfliction + OPPLAN objectives and ATT&CK coverage), the ranked recon surface, and the parameter/content and JS-recon discovery results (JS secrets as names + masked preview only). One is write-shaped: `propose_command`, which appends to an approval queue and runs nothing. A second, execution-capable tool exists only behind an explicit opt-in — again, see the addendum.

*** THE LINE. *** HackPit's action routes take `approved=true` and `dangerous_ack=true` in the request body. If an MCP tool could set those fields the agent would approve itself and every gate in this codebase would become theatre — which is exactly why `proxy._gate_request` passes `GATE_KEY_PLACEHOLDER`. So `test_mcp_safety.py` enumerates every exposed tool and asserts:

- no approval field is nameable, at any schema depth and through an open schema — `additionalProperties: False` is part of the line, because an open schema makes `{"approved": true}` sendable whether or not the handler reads it;
- no handler reaches an execution path, by AST, following the module's own helpers.

Both audits carry a positive control that plants a violation and proves it is caught, because an audit that always returns `[]` would pass forever. And the server itself refuses to start if either audit fails — the one refusal in build #19, and a self-check rather than a prohibition: it declines to EXPOSE a violating surface; it never refuses a user action.

Approving a proposal does not run it, and that absence is the design. The obvious next feature is "approve and run"; wiring it would make the queue a SECOND place that can set an approval field, reachable by anything that can reach the queue. `review()` marks a row and stops. The `Proposal` model deliberately carries no gate field NAMES either — if it did, the next person to wire it into an `ExecRequest` would do so by copying them across, and the agent would have written the value.

The gate preview beside each row always asks with both flags false, so what comes back is the first gate standing in the way. Asking with them set would answer "this would be allowed", which reads as permission and is a question nobody asked.

Two reads are deliberately narrowed, and neither is a gate. The intercept read tells an agent THAT a request is held and how many bytes it is, never its contents — a held request is the one read whose payload is a live credential rather than recorded traffic. The engagement-state read returns credentials as host, username and kind, never values.

Transport is stdio, and build #19's own week is the reason: Burp's MCP server was measured answering 403 to every combination tried, rejecting browser-looking User-Agents and non-allowlisted `Origin` headers as DNS-rebinding defence. That is a real threat against a localhost HTTP MCP server — any page in the operator's browser can POST to `127.0.0.1`. stdio has no port and no rebinding surface. If it ever moves to HTTP it needs the same protection, and route auth is still not built.

The registry imports no MCP SDK. `mcp_tools.py` holds the tools and both audits; `mcp_server.py` is the transport. That split is structural: it means the lock runs in CI, where the optional dependency is absent. A line that could only be checked on one laptop is not a line.

LIVE-VERIFIED END TO END, AND REGISTERED WITH A REAL HOST (2026-08-06). The server was driven over stdio by a genuine MCP client, not just started: `initialize` came back as `hackpit`, `tools/list` returned all fifteen, and `hackpit_arsenal` was called for real — a 14.7 KB round-trip carrying the 121-entry tool catalog, not a schema dump. The optional `mcp` package (v1.23.3) was already present in the backend venv, so nothing had to be installed. It is now registered with Claude Code at user scope, alongside the other security servers (`hexstrike-ai` ,

`burp` , `ghidra`); user scope was deliberate, because the per-project local scope keys off the exact launch directory and kept splitting `HackPit` from `HackPit/backend` . `claude mcp list` reports it **✓ Connected**, and the tools surface as `mcp__hackpit__*` in a host session after a restart. Note the reads that reflect *live* ZAP or engagement state (`hackpit_scan_status` , `hackpit_alerts` , `hackpit_proxy_history` , `hackpit_session_health`) only carry data when a scan or proxy session is actually up; the file-backed reads (KB, arsenal, exploit index, engagement state) work with the backend down.

README — DONE (2026-08-08). The README now carries a full MCP section: what the tools are, the stdio install/registration steps for Claude Desktop and Claude Code (config JSON and the `claude mcp add` form), and the opt-in execution mode below. The earlier "README does not mention it" TODO is closed.

Addendum (2026-08-08) — extended for the Q2/Q3 builds, plus an opt-in execution door

The registry written for build #19 predated the Q2/Q3 offensive builds, so an agent using it as eyes was blind to what those builds know. **Seven read tools were added**, all within the existing line (closed schemas, no approval field, no execution path — the same audit passes with them): `hackpit_killchain_graph` (rebuilds the cross-domain graph + route from the STORED cloud/AD graphs and findings), `hackpit_cloud_graph` , `hackpit_ad_graph` , `hackpit_governance` (RoE/ConOps/ deconfliction + the full OPPLAN and ATT&CK coverage), `hackpit_recon_surface` , and `hackpit_discover_results` / `hackpit_jsrecon_results` . The JS-recon read inherits the value-free-secret property from `MinedSecret` (type + masked preview + loot path, never the value) — the same discipline as the credential read.

The one deliberate policy change: an opt-in execution surface. By default the server is still eyes-and-no-hands and the audit still refuses to start a violating surface. But when the operator sets `HACKPIT_MCP_EXECUTE=1` , one execution-capable tool — `hackpit_execute` — registers, wired to the same sandboxed executor the cockpit route uses, self-approving the gate so **no human sits between the agent and the target**. This removes the human-approval bound *for the MCP path*, on the operator's explicit say-so. It was built to stay honest about that cost:

- it is OFF unless the env var is exactly `1` ; with no flag the tool does not exist, the audits see nothing new, and **CI (which never sets the flag) stays green with the build-#19 locks unchanged**;
- the execution audit is never blinded — `audit_no_execution_paths()` still REPORTS the execute tool as an execution path; only `mcp_server.preflight()` tolerates it, only in this mode, and it prints a loud startup banner saying the human gate is off;
- the approval-FIELD line is never relaxed even here — the tool names no gate field in its schema; it hardcodes `approved / dangerous_ack` in its body, so an agent still cannot *name* approval.

`test_mcp_safety.py` gained a test that pins all of this: execution is off by default, and when the opt-in surface is registered the audit stays honest while preflight tolerates only the named tool. This is a genuine loosening of the model's central guarantee, chosen deliberately — the safe default is what ships to anyone who installs the public tool; the hands are a flag the operator flips for an engagement they are driving agent-first, against authorized targets.

Six defects this build found in its own work, and each needed a different kind of check

1. **SIX OF FIFTEEN MCP HANDLERS WERE WRITTEN AGAINST NAMES THAT DO NOT EXIST.** `pipeline.search` , `state.store.list_findings` , `state.store.session_summary` , `arsenal.catalog` , `exploits.index.by_cve` and `cockpit.fronting.recorded_verdicts` — all invented, along with `attack_path.plan_for_session` . Build #14's lesson arriving exactly on schedule, and caught only by calling every tool for real rather than by any type-check.

Three more followed on the second pass: `DATA_KB` is the entries FILE not its directory (the tool reported "the knowledge base is not built" about a KB with 2,744 rows in it), `EngagementRecord.id` is `engagement_id`, and `Template.command` is `Template.template`.

2. THE EXECUTION AUDIT SILENTLY PASSED A HANDLER IT COULD NOT READ. `ast.parse(source.lstrip())` strips only the first line's indentation, so any handler defined at an indent raised `IndentationError` — and the audit caught it and `continue`d. The positive control planted a handler calling `subprocess.run` and the audit reported CLEAN. **Two fixes, and the second is the one that matters:** `textwrap.dedent`, and an unreadable handler is now an OFFENCE rather than a pass. "We could not check this" must never read as "this is fine" — the same `read_ok` rule the rest of the codebase runs on, applied to the safety check itself.

3. TWO EXISTING WHOLE-TREE LOCKS FIRED ON THIS BUILD, AND BOTH WERE RIGHT. `FORBIDDEN_CALLS` named the `:kali` shell and the tunnel starter as string literals, and `test_kali.py` and `test_tunnels.py` failed the build **on this file**. Those scanners strip docstrings and comments but NOT other string literals — because blanking every string would go blind to `import_module("cockpit.kali")`, which is precisely the indirection they exist to catch.

The fix was neither to weaken the scanners nor to spell the names evasively. It was to notice that the list was carrying a **second, weaker copy of locks that already cover every file in the tree, this one included**, and that catch imports and indirection a name-match never could. The rule left behind: an entry point belongs in that list only if it has no whole-tree lock of its own — and if adding a name breaks the suite somewhere unrelated, that is the duplication showing, so delete the name rather than narrowing the lock.

4. A CONTROL PHRASED "NOT REFUSED" MADE A TEST DEPEND ON DOCKER — THE THIRD TIME.

`test_intruder.py` asserted `validate(approved_lab_job)` is `None`. Without Docker the ISOLATION gate legitimately fires at `gate='sandbox'`, so the file would have gone red in CI, after two previous CI catches in build #14 part 3 and build #15. **The docker-stripped run caught it**, which is exactly what that run is for. Every control now asks "not refused at THIS gate".

5. A SHARED CSS RULE UNDER-MATCHED, AND ONLY A SCREENSHOT COULD SEE IT. `.hp-tn-form select, input, button` never named `textarea`, so every textarea on a cockpit screen rendered as bare unstyled text on a dark background. This is the exact mirror of the checkbox defect recorded three lines above it in `globals.css` — that rule OVER-matched `input`; this one under-matched by omission — and it earns the same answer: **fix the rule, not the screen that noticed**, because the next screen with a textarea would have hit it too. `tsc`, `next build` and `eslint` were all green throughout, and `test_css_vocabulary.py` passed because every class did exist. Only looking at it found this.

6. JSX DROPPED A SPACE THAT IS PRESENT IN THE SOURCE. `<em>required</em> when` rendered as `<em>required</em>when`. Verified against the served HTML rather than trusted from the screenshot's kerning, then fixed with the explicit `{ " " }` this codebase already uses everywhere — which is presumably why it does.

The screens were LOOKED AT, and here is how, because the usual way was unavailable

The Claude-in-Chrome extension was not connected this session. Rather than downgrade to "tsc passed", the sandbox's own Chromium was pointed at the dev server through `host.docker.internal:3000` and `:proxy`, `:intruder` and `:repeater` were screenshotted and read. Defects 5 and 6 above are what that found; both were fixed and re-shot. `:intruder` is a new top-level route and has a tile in `SURFACE_BANDS` (band `operate`, whose hint — "every command human-approved · needs the stack" — is true of both halves of it), so it is not the orphaned route `/proxy` was for two builds.

What was NOT done, and why

- `test_kb_fixture.py` is RED locally and this build did not fix it, deliberately. The live KB is at 2,747, not 2,744, because another session's uncommitted work (`pipeline/authored/authored_entries.jsonl` and two more `curated.json` rows) added three entries. Regenerating the committed fixture would fold someone else's change into this commit *and* break CI, because the fixture would then project a KB whose source rows are not committed. CI is unaffected either way: that file's completeness and staleness checks need a live KB and report NOT-RUN on a clean checkout, by design. **Build #19 touched no KB file**, and neither the authored corpus nor `curated.json` is in this commit.
- **The image is NOT rebuilt.** Nothing in this build changes it — interception, the jar, the filter, the intruder and the MCP server are all backend and frontend code against tools the image already ships.
- **Caído, route auth, scanner pacing, mobile scope** — Zaid's standing skips, untouched.

Verification

`docs/proof/build19_break_api.py` : 27 passed, 0 failed, four consecutive green runs plus two more on freshly-restarted daemons. `docker/proof/browser_intercept_proof.sh` : 25 passed, 0 failed — nothing build #15 or #17 established was disturbed.

Hermetic safety suite green at 85 test files, every one exiting 0, with `test_kb_fixture.py` excluded for the reason above — four new: the cookie jar, interception plus history filtering, the intruder, and the MCP line. Every control is phrased "*not refused at THIS gate*".

The four new files, plus the eight existing ZAP / repeater / kali / tunnel locks this build touches, re-run with `docker` stripped from `PATH` and the strip verified to bite first (docker gone, `git` still present — the KB ingester's recovery path needs it): 12 of 12 green.

`npx tsc --noEmit 0`, `npm run build 0`, `npm run lint` at the accepted baseline of 11 errors + 1 warning, unchanged from `main` — the new effects route their `setState` through async callbacks per the frontend AGENTS.md rule, and the one place that would have needed an effect (syncing the held request into the editor) is a **derivation** instead, which also removes any chance of a two-second poll overwriting a half-finished edit.

`data/kb/entries.jsonl` is **not touched by this build**; it stands at 2,747 from another session's uncommitted work, and the committed fixture's 2,744 is left exactly as it was.

Build #19, CI — main was already red, and then this build broke it a different way

CI was failing before build #19 was pushed, at `test_exploits.py`, since the curated-overlay commit two before it. Three defects had to be fixed to get green, and **all three are the same shape: a check that quietly depends on state only this laptop has.**

1 and 2 — two mirror-dependent claims in `test_exploits.py`. `/data/` is gitignored, so a clean checkout has no exploit-db mirror. `full.stats()["entries"]` raised `KeyError` because `stats()` reports the MIRROR's own counts and there were none; and `assert ix.ready is True` under the heading "*a broken overlay must not take the 47k-row mirror down*" asserted a property **about a mirror that is not there**. The first was fixed by stating the property so it holds in both environments — the merged index holds the mirror's rows plus the overlay's, and with no mirror that is the overlay's alone. The second is genuinely unobservable without a mirror, so it reports **NOT-RUN**, the shape `test_kb_fixture.py` already uses. Neither was weakened to a skip: everything below them still runs everywhere, including the check the file exists for — that a malformed overlay is *reported* rather than swallowed.

The lesson the second one taught is the useful one: **after the first fix, the next assertion in the same file failed for the same reason.** So the third pass did not push and wait — it patched `exploits.index.INDEX_PATH` to a nonexistent path and ran **all 17 functions in the file**, 0 failures, before committing.

3 — BACKTICKS IN THE SUITE'S OWN DESCRIPTION STRINGS ARE COMMAND SUBSTITUTION.

`run_safety_tests.sh` passes each test's description as a double-quoted shell string, so a backtick in one runs. This build wrote two: ``held`` printed *"held: not found"* and carried on, and ``| sh`` was a **parse error that killed the script with exit 2** — after every test in it had passed.

A green local run said nothing about it. Windows MSYS parsed the file; CI's `dash` did not. That is precisely the "it works on this laptop" failure the docker-stripped run exists to catch for a different dependency, and it argues for `sh -n` on that script being part of the routine.

And it was not new: **one pre-existing description had the same defect** — `test_scan_session_health`'s ``unknown`` and ``ok`` (build #16) have been executing as commands and blanking themselves out of the printed guard text ever since, non-fatal only by luck of which words were chosen. Both are fixed; no `run_test` line carries a backtick now.

Source evaluated: `sources/insane ctf writeups and vulnerabili.txt` (2026-08-05)

2,168,551 chars / 28,105 lines, segmented programmatically into 198 sections that reconcile to the byte (4,556 preamble + 2,163,996 sections = 2,168,552 = len+1, the +1 being the single trailing-newline overcount). Full ledger, one row per section with a verdict and a reason: `docs/source-eval-insane-ctf-writeups-2026-08-05.md`.

The file's advertised shape was wrong in a way that mattered. It reads as a CTF writeup dump, and the 191 "markdown headings" are almost all `#` comment lines inside pasted code. The real structure is 97 numbered documents, and documents #50–#67 are not CTF writeups at all — they are disclosed vulnerability reports (Monero, Liberapay, Rails, Trix, SingleStore, aws-cdk-lib, Taskcluster, Burp, HackerOne, Shopify, Khan Academy, DuckDuckGo, curl x2, Snowflake). That region is where the entire yield came from; the ~2.0 MB of pwn/crypto/rev writeup prose around it produced nothing, exactly as the saturation probe predicted.

Ten of those eighteen reports overlap `sources/some vul.md`, already evaluated on 2026-08-05 — three of them are already curated rows (`hp-c-0001` curl SSH reuse, `hp-c-0004` Trix, `hp-c-0005` aws-cdk-lib). The overlap was measured, not assumed, by probing distinctive strings of each report against that file.

KB — 3 authored entries (2744 → 2747)

- `authored-bcrypt-72-byte-truncation` (web/authentication). The KB already held the *primitive* — `PHP Tricks` states `PASSWORD_BCRYPT` truncates at 72 bytes and shows the 71-vs-72 `password_verify` demo. It did not hold the *exploitation pattern*: that the bug appears when the app hashes a **structured** string, so the secret sits at an offset and the attacker picks the account whose prefix is longest to push it out of the window.
Recorded as a delta over an existing entry, not a virgin gap.
- `authored-elasticsearch-script-sort-injection` (web/injection). Genuine zero: `painless 0`, `script_score 0`, `_seq_no 0`, `sort injection 0` across all 2,744 entries. Carries the differential method (compile-fail/compile-success 500-vs-200, then constant-vs-per-document ordering) that proves execution without a payload.
- `authored-sql-compile-vs-runtime-error-oracle` (web/injection). The KB has 90+ `SQLi` entries and zero on cloud data warehouses (`snowflake` = 1 hit, unrelated context) and zero on the compile-time/runtime error asymmetry that makes a "closed" error channel readable again.

CVE→exploit index — 2 curated rows (5 → 7)

- `hp-c-0006` **Ruby on Rails ≥ 7.0**, `gte`, CVE-2025-24293. The published fix added `validate_transformation()` to the **ImageMagick** transformer only; `vips.rb` never overrode it, and `load_defaults "7.0"` makes `vips` the **default** processor. So the public verdict "patched for CVE-2025-24293" is wrong for a default-configured Rails 7.0+ app — precisely the failure this overlay exists to prevent. `gte` because no release closes the Vips path.
- `hp-c-0007` **curl 8.20.0**, `exact`, no CVE. `http_rw_headers()` splits response headers on a bare `memchr(buf, '\n')` with no quoted-string awareness, so a malicious server injects `Set-Cookie:` into the jar or a `Location:` that replays a 307's POST body and `Authorization` header. Here the public verdict is **absent** rather than wrong — the other half of the overlay's remit. `exact` because 8.20.0 is the only build the PoCs were run against; no lower bound was measured and guessing one is the exact error the file forbids.

Two more were **considered and explicitly not indexed**, recorded in `considered_and_NOT_indexed`: `monero-wallet-rpc relay_tx` (real and unfixed, but keyed only to a branch + commit — no release to compare against) and Burp Suite Pro 2026.3.3 (expressible, but the index answers "what runs on the *target*", and Burp is operator tooling).

Verified, not assumed

The overlay was checked through the live index rather than trusted: `rails 7.1.0` and `rails 8.0.1` return `hp-c-0006` at rank 1 while `rails 6.0.0` correctly demotes it to last; `curl 8.20.0` returns `hp-c-0007` at rank 1 and `curl 8.19.0` demotes it to rank 7. `test_exploits.py` passes 17/17 including "the curated overlay merges (7 rows)" and "`gte` reports 'no fixed release' and never implies a patch".

A finding that evaporated when the tree stopped moving

The evaluating session reported the safety suite RED at `test_kali.py`, with `mcp_tools.py` flagged for `run_kali` — diagnosed as the fourth instance of this repo's shared-predicate defect, a containment scanner catching the module that enforces the rule.

It does not reproduce. Against the committed tree `test_kali.py` passes: 103 backend modules scanned, planted-violation control included. `run_kali` appears in `mcp_tools.py` exactly once, inside a module docstring, and `scan_source_tree` already strips prose before the substring pass — the refinement that was made the first time this class was hit.

What actually happened is worth more than the finding would have been: **two sessions shared one working tree, and the evaluator scanned `mcp_tools.py` while build #19 was still writing it.** The state it judged was real when read and never got committed. A second session in a live tree is reading a moving target, and a defect it reports against uncommitted work has to be re-checked against HEAD before it is believed. That is the same discipline as re-running a flaky test before treating one red run as evidence.

Build #20 — GraphQL, end to end (2026-08-05)

The source evaluation earlier the same day produced exactly one HackPit capability finding, and it was evidence-backed rather than speculative:

Report #61's injection point is a GraphQL argument, reachable only by composing a query with variables. Attacking it through the repeater means hand-writing quadruple-escaped JSON-in-JSON.

Measured against the tree: `graphql` appeared in the backend **once**, incidentally, in `attack_path.py` ; in the frontend **zero** times; and the arsenal carried **no** GraphQL tooling at all. Meanwhile ZAP already shipped a GraphQL add-on, so most of this looked like a SURFACING job rather than an engine-building one — the same correction that made build #19's interception days instead of weeks.

The build adds no gate, no confirm, no blocklist and no allow-list narrowing — not in the product and not in a proof script. Where something could refuse, it warns and continues. The two places that decline to do something are named below, and both are disclosure or correctness decisions rather than prohibitions.

Item 1 — the go/no-go, and * IT CHANGED THE PLAN *****

`docs/proof/build20_graphql_api.py` : **56 passed, 0 failed**. Nothing in items 4 or 5 was written before it answered, and what it answered was not what the plan assumed.

The plan's chain was: import a schema → ZAP generates operations → **the operations land in the Sites tree** → `ascan` attacks them there. *** THE MIDDLE LINK IS NOT THERE. *** ZAP sends the generated operations at the endpoint for real — the proof's origin log counts them every time — and never files a single one. So item 5 scans **CAPTURED** operations instead, which is both ZAP's own path and a better fit for the finding that started this: an operator holding report

61's request has the request, not the schema.

That took four measurements to believe, and the wrong turns are worth recording because each was plausible and each was killed by a measurement rather than an argument:

hypothesis	how it died
<code>core/action/newSession</code> broke the tree insert	a fresh daemon that had never had a session reset behaved identically
the ops attach only to a node that already exists	a node primed through the proxy gained the node and still gained no operations
the insert is merely late	polled out to 60 seconds ; the count never moved
it works and the counter is wrong	the counter was rewritten to look for the operation BODIES in the tree; still zero

The **CONTROL** is what makes it actionable. One GraphQL request through the **proxy** creates both `<endpoint>` and a synthetic `<endpoint>/query` child node, with messages on it. Captured traffic is what feeds the scanner; an import is coverage traffic. `SchemaImport.scannable` is therefore a field that is **always false** rather than a silence, and the panel says so out loud, because an import that answers `OK` looks exactly like one that gave the scanner new targets.

Seven more findings from the same script, each one measured

1. `/UI/<component>/` **IS THE ENUMERATION SURFACE, AND BUILD #19 NEEDED IT**. It lists every view and action of a ZAP component with its parameters and which are required. Build #19 established the break API by probing names one at a time and reading error codes, because `core/view/apiSummary` is `bad_view` in 2.17.0 and `/JSON/<component>/view/` answers with ZAP's **welcome page**. `/UI/graphql/` answered the whole surface in one request: 11 views, 13 actions. That is worth more than this build — it is how any future go/no-go against a ZAP component should start.
2. **THERE IS NO VIEW THAT READS AN IMPORTED SCHEMA BACK**. All 11 views are `option*` getters. "Read back what was imported — an OK is not a result" was a sub-item of this build and it cannot be satisfied through `graphql/` at all. Naming the absence is the deliverable.
3. `importUrl` **CANNOT TELL YOU WHY IT FAILED — AND THIS IS THE ONE THAT SHAPED ITEM 4**. An endpoint that refuses introspection the way production does and a host that is not listening answer with the **same code and the same message** (`illegal_parameter`). The brief required "disabled / error / empty are three facts and the operator needs to know which"; ZAP cannot supply that, so `graphql_zap.probe_schema` asks the endpoint itself and classifies into six: `ok`, `disabled`, `empty`, `http_error`, `unparseable`, `unreachable`. Four of the six would be an empty list from a naive implementation, in the one place where the difference decides what to do next — `disabled` means reach for `clairvoyance`, `unreachable` means fix the network, `empty` means stop looking.

4. `importFile` ANSWERS OK AGAINST AN ENDPOINT THAT IS NOT LISTENING. The OK means "I parsed your schema", not "I did anything with it". Anything reporting success from that return value would tell an operator a scan had run against a host that does not exist.

5. ZAP VALIDATES ITS OWN BOUNDS NOT AT ALL. `setOptionMaxQueryDepth?Integer=-1` answers OK and reads back `-1`. A non-numeric depth is refused, so the TYPE is checked and the RANGE is not. HackPit therefore **warns and sends anyway** — the bound is the operator's to set — and reports `observed`, read back field by field, never the value that was requested.

6. `cycleDetectionMode` HAS EXACTLY ONE USABLE VALUE. `QUICK` applies; `OFF`, `THOROUGH`, `PRECISE`, `NONE` and `COMPLETE` are all refused. A choice with one option is not a choice, so HackPit does not offer it. A refused enum leaves the previous value intact — checked, because a setter that corrupted on refusal would be worse than one that refused loudly.

7. `argsType=VARIABLES` GIVES EVERY ARGUMENT ITS OWN KEY in `variables`; `INLINE` writes them into the query text. Both are reachable.

The measurement item 5 exists for

A captured GraphQL operation, no schema imported, `ascan` with `recurse=true`:

- 1,630 requests, 1,427 carrying a variables object
- 588 mutated EXACTLY ONE argument — and 0 mutated more than one
- all three arguments reached individually: `id` 200, `locale` 194, `token` 194

And on ZAP's own generated operations in an earlier run, a **High SQL Injection** with `param='search.limit'` — where the JSON variable key is `search_limit`. A dot where the JSON has an underscore, so the name came from the add-on's own variant walking the parsed operation, not from a JSON path and not from the body as a whole. HackPit uses that exact `field.argument` spelling everywhere, so the names an operator is shown **before** a scan are the names they read in the findings **after** it.

`recurse` is **always true** and that is not a preference: the operations hang off the synthetic `/query` child node, and the first measurement of a scan aimed at the endpoint alone sent 114 requests and touched zero arguments.

*** THE CONTAINER IS NOT THE IMAGE, AGAIN — AND THIS ONE IS A CAVEAT ON EVERYTHING ABOVE ***

The image ships `graphql-alpha-0.29.0` at `/usr/share/zaproxy/plugin/`. The container the proof was measured in ALSO had `0.33.0` in `/root/.ZAP/plugin/` — a runtime upgrade living in one container's writable layer that a rebuild would drop. They are not the same surface: `0.29.0` has no `optionMaxCycleDetectionAlerts` at all and answers `bad_view`.

It was found by looking at the screen. The panel rendered `max_cycle_detection_alerts: bad_view` against `hackpit-kali-sandbox`, having rendered `100` against `hackpit-engage-sandbox` — and **that exposed a defect in HackPit's own read-back**: ZAP's error shape `{"code": ...}` and a view's success shape `{"MaxQueryDepth": "5"}` are both one-key dicts, so "take the first value" printed the error CODE into the panel as though it were the daemon's configuration. `_getOption` now returns `unreadable (<code>)`, because "we could not read this" and "this is empty" is the distinction this whole build keeps insisting on.

The proof now prints which add-on versions are visible **before it asserts anything**, and says so loudly when it finds more than one. Build #14 was written against `zap-baseline.py`, which does not exist in Kali; this is the same trap wearing the version number instead of the name.

NOT FIXED IN THIS BUILD, deliberately: the image is not upgraded to 0.33.0. That is a change to the scanner's own engine, the rebuild is ~45 minutes, and item 1's whole point is that a surface is not assumed — re-measuring the full proof against a new add-on is its own piece of work. The product is now *tolerant* of the difference rather than silently wrong about it, which is the part that had to land today. **It is the top follow-up.**

Item 2 — detection, by BODY SHAPE and never by path

`cockpit/graphql.py` is **pure** — no I/O, no subprocess, no socket, asserted by AST — which is what lets the hermetic suite cover every claim without a daemon.

*** `/graphql` IS A CONVENTION, NOT A RULE. *** Shopify serves GraphQL from `/admin/api/2024-01/graphql.json`; plenty of APIs mount it at `/api`; and a site can serve JSON that is not GraphQL from a path called `/graphql`. A path test would both miss real endpoints and invent fake ones, so the test is the envelope: a JSON object (or a batched array of them) carrying a string `query`, `?query=` on a GET that opens like a document, or `Content-Type: application/graphql` where the body IS the document. `path_hint` records that the path looks conventional, is reported to the operator, and **decides nothing**. Tests assert both directions with controls.

A DOCUMENT THAT WILL NOT PARSE IS STILL GRAPHQL. The envelope decides. A malformed operation is precisely the request worth looking at, and dropping it from a GraphQL filter would hide it — so it comes back `is_graphql: true` with a `note` saying why it would not parse.

The filter is **three-state** (`true` / `false` / unset means both), because a checkbox would have made "show me everything" unsayable. `graphql_seen` is reported over the same rows `scanned` counts, whether or not the GraphQL filter was used — an honest denominator or nothing: `0` beside `scanned: 1200` means we looked at 1,200 and none were GraphQL; `0` beside `truncated: true` means we stopped before we could say.

Endpoint records gain `tech: "graphql"` and the `field.argument` names in `params`, because an endpoint whose `params` is empty reads as "nothing to attack here", which for a GraphQL POST is exactly wrong and is why the surface has been blind to it.

Item 3 — the repeater's round trip, and the operator's raw body WINS

Query and variables are edited separately and serialised into a correct JSON body. Verified in a real browser: a `' OR 1=1--` typed into the `token` **variable** landed in the raw body correctly escaped, with no hand-escaping anywhere — report #61's exact case.

- **A BODY THAT WILL NOT SPLIT COMES BACK RAW AND SAYS SO.** `parsed: false`, `raw_body` intact, `note` explaining. A **batched** request in particular stays raw, because one query box would drop every operation but the first.
- **IF THE VARIABLES WILL NOT PARSE, NOTHING IS BUILT.** No guessing, no repair, no plausible fallback — the error is shown and the body is left exactly as it was. A composer that quietly fixed a body would put a request on the wire that nobody wrote.
- ***** THE RAW BODY WINS. ***** The composer remembers what it last produced; the moment the raw textarea holds something else, the structured editor **steps aside** rather than overwriting. Same rule build #19 gave the cookie jar and for the same reason. Verified by typing into the raw box and watching the editor close with the typed text intact.
- The repeater stays **HUMAN-ONLY**. Both new routes are pure transforms; neither module imports the repeater or references `.send`, asserted by AST.

Items 4 and 5 — recon, and a scan behind the existing four gates

Item 4 is **reconnaissance**: probe, classify, list types/queries/mutations and their arguments, and hand ZAP a schema with bounds the operator set. Item 5 is the **ordinary active scanner** aimed at a captured operation — `scan_plan_for` computes a target and starts nothing, and the same four gates run unchanged. A test walks a GraphQL scan through the **approval**, **danger** and **target** gates in turn, naming each, because the gates run target-first and an off-lab URL would have made "an unapproved scan is refused" pass vacuously for the wrong reason.

One approval buying many requests is the established position — `ffuf`, `nuclei` and the scanner are each one approval buying thousands, and this one measured 1,630.

NAMES, NEVER VALUES. `GraphQLArgument` and `SchemaArgument` have no value field at all, and a test plants a secret in both an inline argument and a variable and asserts neither reaches any model, endpoint record or API response. Build #19's cookie-jar rule applied to a third secret: never handing a value over cannot regress, while redacting it afterwards depends on a redactor being correct forever.

WHERE IT COULD REFUSE, IT WARNS. A schema probe against a host outside the named engagement's scope is sent, with the warning in `scope_note`. A nonsense bound is **applied**, with the warning attached. A test asserts the warner cannot raise and that `apply_bounds` has no `raise` in it at all, because "warn and continue" is a requirement of this build and a later "tightening" would be the regression.

KNOWN GAP, RECORDED RATHER THAN HALF-BUILT: field-suggestion / clairvoyance enumeration when introspection is off. When the probe answers `disabled` the panel names `clairvoyance` as the tool for that job; HackPit does not do it itself.

Item 6 — the arsenal, and *** A SUPPLY-CHAIN TRAP AIMED AT AGENTS ***

Four tools added (117 → 121): `graphw00f`, `graphql-cop`, `inql`, `clairvoyance`. All four were run against a live GraphQL endpoint inside the container before being catalogued; none is written down on the strength of a `--help`.

*** TWO OF THE FOUR MUST NOT COME FROM PyPI. *** `pip install graphw00f` installs a package whose own summary reads "Inert defensive-hold placeholder for an unclaimed PyPI name referenced by a public agent skill. NOT the real tool" (Metano Labs, 0.0.1, one release). `pip install graphql-cop` installs "Reserved name placeholder. No functionality." Both real projects are `d0levf/*` on GitHub and neither publishes to PyPI. Somebody has parked those names **precisely because tooling guides name them and an agent will reach for pip** — so the obvious install line is the wrong one, quietly, with a package that installs, runs and finds nothing.

The other two are genuine and are taken from PyPI: `inql` 4.0.5 (Doyensec) and `clairvoyance` 2.5.5 (Nikita Stupin, project repo in its metadata). The provenance of all four was read before any of them was installed.

Each entry carries what the tool actually does when you run it, not what its README says: `graphql-cop` exits with a bare `KeyError` on a server whose `{__typename}` does not answer with a root type name; `inql` has no `--generate-schema` flag because generation is the default; `clairvoyance` reports an **empty schema** rather than an error when suggestions are off, which means "this defence is complete", not "there is no schema". All four are classified **clean** in `test_arsenal_safety.py` alongside `sqlmap` and `dalfox`, which fire real injection payloads — marking a GraphQL fingerprinter dangerous while those two are not would be a red-confirm that fires on reading a page.

Item 7 — the screens were LOOKED AT, and it found five things nothing else could

`tsc --noEmit`, `next build` and `eslint` all passed before, during and after every one of these.

1. **THE GRAPHQL BADGE NEVER APPEARED AT ALL.** Detection reused `RepeaterRequest`, whose `url` is `min_length=1` because a SEND needs somewhere to go — so pasting a captured body before typing a URL, which is exactly when the badge is useful, answered **422**. Detection now has its own model where every field is optional.
2. **"edit as GraphQL" RENDERED AS PLAIN HEADING-SIZED TEXT.** An unclassed `<button>` inherited the page's heading scale: a control that did not look like a control, sitting above the thing it acts on. Same shape as `.hp-tn-start` not existing while nine primary buttons used it.
3. **A DROPPED SPACE, CONFIRMED IN THE DOM RATHER THAN GUESSED.** The rendered text read `not at all—it accepts`. Reading `textContent` in the browser proved the space was genuinely absent, not an italic-glyph artifact — and the same check on a suspicious-looking `//` in a URL input proved *that* one WAS just kerning. A sweep of every prose element on the page then found **two more, both pre-existing build #19 copy**: `{"Result": "OK"}and lets and settingand reads`. The codebase already uses `{" "` at line ends for exactly this; those three spots used a bare inline space before a line wrap instead. All three fixed.
4. **THE PANEL WAS INVISIBLE WITH NO PROXY RUNNING.** Rendering nothing reads as "HackPit does not do GraphQL", which is the belief this build exists to correct. It now renders and says why it cannot act, matching the intercept section three cards up.
5. **A SHARED CSS RULE OVER-MATCHED, AND THE RULE WAS FIXED RATHER THAN THE SCREEN.** `.hp-tn-cardsub` carries `margin-top: -4px` to tuck a sub-line under its heading; after a form or a list it pulls the paragraph onto the box edge. That shape occurs **21 times across 7 components**, so a per-screen workaround would have left the other 20 cramped — build #5's lesson and the note in the frontend-class-vocabulary memory. A sibling rule now adds space only where there is currently none, so it cannot break a layout that was already right; `:intruder` was looked at afterwards to confirm it improved rather than harmed.

Plus the `bad_view` read-back defect recorded above, which the screen also surfaced.

One safety predicate **FIXED** rather than narrowed

`test_detection_safety.py` asserted "the cockpit package must not reference detection" with a **substring scan for the word `detection`**. It fired on `cockpit/graphql.py`, whose subject is GraphQL DETECTION, and on `max_cycle_detection_alerts`, which is ZAP's own option name.

The standing rule when a guard trips on innocent code is to **fix the predicate, never widen the file set** — narrowing the glob would have left the guard weaker against the thing it exists to catch. The property it always meant to state is "the execution layer must not IMPORT OR CALL the detection package", not "must not contain nine particular letters". It is now an AST pass that catches `import detection`, `import detection.catalog`, `from detection import ...`, `detection.resolve(...)`, `importlib.import_module("detection")` and `__import__("detection")` — strictly stronger about the real thing and blind to prose. The old positive control was a real file that reaches its siblings by RELATIVE import and would no longer have exercised it, so the control is now **eight planted violations** plus four innocents that must NOT fire.

Verification

- `sh backend/run_safety_tests.sh` — **88 test files, every one exited 0** (86 + `test_graphql.py`
- `test_graphql_safety.py`).
- Green again with **docker stripped from PATH** for every file this build touched, and the strip was verified to bite first: `docker` gone, `git` still present (the KB recovery path needs it). Each control is phrased **"not refused at THIS gate"**, never a bare "not refused".

- `docs/proof/build20_graphql_api.py` — 56 passed, 0 failed.
- `docker/proof/browser_intercept_proof.sh` — still 25 passed, 0 failed.
- `npx tsc --noEmit 0 · npm run build 0 · npm run lint` unchanged at the accepted baseline of 11 errors + 1 warning.
- `data/kb/entries.jsonl` still 2,747 — nothing was ingested — and `backend/test_support/kb_fixture.jsonl` matches it.
- Arsenal 117 → 121 tools, CRLF preserved (the reconFTW trap), inserted textually rather than re-serialised: a `json.dumps` round trip re-expanded compact inline arrays elsewhere in the file and turned a four-entry addition into a 281-line deletion.
- Every screen this build touched — `:repeater`, `:proxy`, `:intruder` — driven in the **real Chrome**, with what was found written up above.
- **The image was rebuilt once, at the end** (`docker compose build engage-sandbox`, exit 0), and all four tools are present AND run inside it: `graphw00f`, `graphql-cop`, `inql` and `clairvoyance` each resolve under `/usr/local/bin` and each exits 0 on `--help`. The rebuild also **confirms the caveat above from the other direction**: the fresh image carries `graphql-alpha-0.29.0` and nothing in `/root/.ZAP/plugin/`, so the 0.33.0 the proof measured really was a runtime artefact of one long-lived container and really would not survive a recreate.
- **A line-ending trap, caught before the commit.** Python's `write_text` translates `\n` to `os.linesep` on Windows, so editing an LF file through it rewrote every line: `api.ts` alone showed **8,672 changed lines** for a 208-line addition. The repo has `core.autocrlf=false` and no `.gitattributes`, so each file's endings were restored to whatever HEAD holds and the new files were matched to their neighbours. The staged diff went from 11,147 insertions to **4,690 insertions / 27 deletions**, and the suite was re-run green afterwards. Same class as the reconFTW `tools.json` CRLF note, in the opposite direction.

What was NOT done, and why

- **The ZAP GraphQL add-on was not upgraded in the image.** See the caveat above; it is the top follow-up, and the product is tolerant of the difference rather than wrong about it.
- ***** THE RUNNING CONTAINER WAS NOT RECREATED, SO THE FOUR TOOLS ARE IN THE IMAGE AND NOT YET IN THE CONTAINER. ***** They are verified present and working in `hackpit/kali-sandbox:m1`; a `docker compose up -d engage-sandbox` is what puts them in reach of a run. That was left for a deliberate moment rather than done here, because recreating **destroys the running ZAP daemon and its capture** — and in this particular case it would also drop the container's runtime `graphql-alpha-0.33.0` back to the image's 0.29.0, which is the exact surface every measurement in item 1 was taken against. Trading a measured state for an unmeasured one at the end of the build that measured it would have been the wrong order.
- **No field-suggestion / clairvoyance enumeration in the product.** Recorded as a gap, with the arsenal tool named at the exact moment the operator needs it.
- **No new gate, confirm, blocklist or allow-list narrowing** — the build's own requirement.
- **No GraphQL query generator.** ZAP already has one, and the brief said not to write a second.

Build #21 — pin the engine, then read a schema nobody will give you (2026-08-05)

Build #20 closed with three leftovers that were **one chain**, because all three were blocked on the same rebuild, plus the follow-on it named: field-suggestion enumeration. All four landed.

CLOSED 2026-08-05. The build's own leftover — three engines named but unmeasured — was read in the same session, the long tail behind it was triaged, and the two live findings that triage surfaced were then closed too. **Eleven engines resolved in total:** six read from source and given a parser or a non-suggesting entry (Absinthe, async-graphql, HotChocolate, graphql-dotnet, **sangria**, and graphql-js's Python/PHP/Go siblings already present); four measured as never suggesting (Iacina, Caliban, agoo, pg_graphql); one already covered by an existing core (dgraph → gqlparser); and **morpheus**, which shares graphql-js's sentence and is now told apart by a directive probe. Three shipping defects were found and fixed on the way — a fabricated schema field, a dropped argument-suggestion clause, and a strikethrough that rendered as literal tildes — and one claim in this document (morpheus's ScalarLeafs message) was retracted as wrong. **Nothing is left open.**

The build adds no gate, no confirm, no blocklist and no allow-list narrowing — not in the product and not in a proof script. The one place it changes what an operator can reach, it **removes** a restriction (see item 5). The enumeration bounds are a declared parameter, not a gate: reaching one stops the run, names itself, and hands back everything found.

Item 1 — the pin, and ***** THE ASSERTION CAUGHT A BROKEN PIN ON ITS FIRST RUN *****

The sequence was load-bearing and was not reordered: pin → assert at build time which file ZAP loads → rebuild → recreate → re-run build #20's proof.

***** AN ADD-ON VERSION IS A DEPENDENCY CLOSURE, NOT A FILE. ***** Dropping `graphql-alpha-0.33.0.zap` into the system plugin directory and deleting 0.29.0 produced a ZAP with **no GraphQL add-on loaded at all** — not an error, not a warning, not a line in the log: simply absent from `installedAddons`. 0.33.0's manifest requires `commonlib >= 1.40.0 & < 2.0.0` and the image shipped 1.39.0, so ZAP declined it in silence. The container's 0.33.0 had only ever worked because the same auto-update pass had also lifted commonlib to 1.43.0 — which is the concrete reason the drift is **33 add-ons wide, not one**. So the pin is a SET: graphql 0.33.0 + commonlib 1.43.0, the closure measured (commonlib 1.43.0 declares no dependencies of its own) rather than assumed. 1.43.0 specifically, because it is the version build #20's 56/0 was actually measured against — which is what makes re-running that proof a regression test rather than a fresh measurement.

An `ls`-based check would have shipped that image green, with a GraphQL scanner that had silently ceased to exist. The assertion is `autoupdate/view/installedAddons`, which reports the file ZAP resolved and loaded — the only question a shadowing directory cannot change.

Two more traps paid for in the same layer:

- ***** ZAP REPORTS** `/usr/share/zaproxy/.plugin/graphql-alpha-0.33.0.zap` . ******* With a `./` in the middle, because the value is assembled from its install root plus a relative directory and never normalised. A literal `==` against the install path fails on a **perfectly correct pin**. Measured before it could be mistaken for a real failure; the checker compares paths, not strings.
- **Pinning one add-on can unload another.** Bumping `commonlib` moves a version every other add-on depends on. So the checker also asserts that **every .zap present is loaded** — an invariant measured on the pre-change image (48 files, 48 loaded) rather than invented.

`docker/zap-addon-check.py` ships into the image as `zap-addon-check` and has a runtime mode, so an operator can ask a **live** daemon what it is really running instead of reading the Dockerfile and assuming.

The runtime half, measured in both directions

The image pin alone is not enough, and this was measured rather than argued:

daemon started	result after ~60s
without <code>start.checkForUpdates=false</code>	29 add-ons auto-downloaded into <code>/root/.ZAP/plugin/</code>
with it, on a clean ZAP home	0 downloaded; all 48 loaded from the image

So `server_argv_for` now states `-config start.checkForUpdates=false` on every start, for exactly the reason `api.disablekey=false` is stated: ZAP persists `-config` values, so an unstated key inherits whatever the last run wrote.

A precise refinement, because the first reading was too kind to the pin: in that 29-add-on sweep, graphql and commonlib were *not* replaced — because the pinned versions are the newest published, so there was nothing to fetch. **The pin held today for a reason that expires the day 0.34.0 ships.** The image pin is what makes the version correct now; the flag is what keeps it correct later. Neither half is redundant.

What it costs, stated rather than glossed: add-on and scan-rule updates now arrive by rebuilding. Nothing is refused — `zapproxy -addoninstall`, ZAP's own UI and a daemon started without the flag all still update. The default simply stopped being "silently whatever is newest". This is the contract nuclei-templates already has in this image.

The re-measurement — *** THE 56/0 BASELINE REPRODUCED EXACTLY ***

`docs/proof/build20_graphql_api.py` re-run against the **image-produced** add-on: **56 passed, 0 failed**, and this time the proof reports exactly one graphql add-on visible instead of warning that more than one is present. **No divergence to report** — the finding to record is the absence of one.

Re-checked specifically, because the plan asked: **ZAP still files none of its generated operations in the Sites tree** (4 reached the origin, 0 in the tree). The scan path is unchanged and the captured-operation route stays. The scanner reached all three arguments individually — 604 requests mutated exactly one argument, 0 mutated more than one, across 1,671 requests. Same shape, same conclusion; the small deltas from build #20's 1,630/588 are ordinary scanner variance.

Leftover #2 closed for free: the recreate put build #20's four GraphQL tools in the running container — `graphw00f`, `graphql-cop`, `inql` and `clairvoyance` all resolve and exit 0.

Item 2 — engine fingerprinting, and the unit is the CORE

*** "APOLLO, graphql-js, HASURA, GRAPHENE" IS FOUR BRANDS THAT ARE NOT FOUR IMPLEMENTATIONS. ***
 Apollo is graphql-js; Graphene sits on graphql-core, the Python port of graphql-js. Grouping by brand produces parsers that duplicate each other in some places and miss entirely in others, so the lookup is keyed on the core that **formats the error**, and brands map onto cores.

DECIDED AND RECORDED: HackPit fingerprints NATIVELY rather than driving `graphw00f`. Three reasons, and the choice was not left implicit:

1. the enumerator and the fingerprint must travel the **same HTTP path**, or the thing that picks the parser and the thing that uses it can disagree about proxying, headers and TLS. `graphw00f` is a subprocess with its own client, and the brief said not to grow a second HTTP path.
2. `graphw00f` answers with a **brand**; the parser lookup needs a **core**. Driving it would not remove the brand-to-core table, it would add a subprocess to it.
3. fingerprinting is recon in the repeater's position — ungated, one request to a named URL. `graphw00f` is a gated arsenal command; routing an ungated feature through it would either add friction or route around the

gate.

`graphw00f` stays as the operator's **independent cross-check**, which is worth more than a shared implementation — two implementations agreeing is evidence — and the proof runs it against the same endpoints for exactly that. Its 36-engine list is the reference set the brand table is drawn from.

*** `unknown` IS A REAL ANSWER AND IS NEVER QUIETLY UPGRADED TO APOLLO. *** An unidentified engine spends **no wordlist at all**, because the wrong parser returns zero and reads exactly like a hardened server. A confident wrong answer is worse than an honest empty one here: both end with "stop looking", and only one of them is true.

Item 3 — field-suggestion enumeration, and the measurement that shaped it

*** `graphql-core` IS NOT BYTE-IDENTICAL TO `graphql-js`. *** Measured by running both:

```
graphql-js    Cannot query field "usr" on type "Query". Did you mean "user" or "users"?
graphql-core  Cannot query field 'usr' on type 'Query'. Did you mean 'user' or 'users'?
```

Identical grammar, **different quote character**. So the parser everybody writes first — the one every article quotes — returns **zero** against Graphene, Strawberry and Ariadne. Silently. It looks exactly like a server with suggestions switched off. That is the "one regex over all of them" trap this project has been burned by repeatedly, and it is why the dialect is chosen from a measured probe rather than assumed.

The same measurement in the other direction stopped three parsers being written for nothing: **`graphql-php` and `gqlparser (gqlgen)` ARE byte-identical to `graphql-js`**, read from their own source. They share the parser and each still carries its own fixture and its own test, so a future divergence fails a test instead of quietly returning nothing.

*** `graphql-ruby` IS THE HIGHEST-VALUE TARGET AND THE ONE LEAST LIKE THE REST. *** GitHub, Shopify and GitLab all run it, and it shares **not one delimiter** with `graphql-js`:

```
graphql-ruby  Field 'usr' doesn't exist on type 'Query' (Did you mean `user` or `users`?)
```

Backticks, inside parentheses, after a different sentence, with the field in single quotes — and **no Oxford comma** before `or` where `graphql-js` has one. Four independent differences. A parser that split the clause on a comma-space would yield a name with `or` glued to it on one of them and be right on the other; pulling **quoted runs** out instead cannot straddle a separator whichever separator it is.

*** SOME CORES NEVER SUGGEST AT ALL, AND THAT IS NOT "SWITCHED OFF". *** `graphql-java's FieldsOnCorrectType.unknownField` has no suggestion clause in its source; neither does Hasura's. Telling an operator "suggestions are disabled" there implies a setting somebody could have left on. There is nothing to enable. So it is its own outcome, and **not one wordlist request is spent**.

Provenance is recorded per dialect, because "I am fairly sure Ruby uses backticks" is exactly how a parser that finds nothing gets shipped green: `graphql-js` and `graphql-core` were **RUN** (`graphql 16.x` via node; `graphql-core 3.2.11` via python); `graphql-php`, `gqlparser`, `graphql-ruby` and `graphql-java` were read from the line of source that formats the message.

Three more cores, 2026-08-05 — and they went three different ways

The three engines build #21 left unmeasured were read from source: **Absinthe** (Elixir), **async-graphql** (Rust, substituted for Juniper) and **HotChocolate** (.NET). They were added together, from one plan, on one assumption — that each would need its own parser. One did, one needed none, and one was already covered:

```
absinthe      Cannot query field "usr" on type "Query". Did you mean "user" or "users"?
async-graphql Unknown field "usr" on type "Query". Did you mean "user", "users"?
hotchocolate The field `usr` does not exist on the type `Query`.
```

Absinthe is byte-identical to graphql-js — sentence, joiner, Oxford comma and the cap of five, all four read and all four matching, from an implementation in another language sharing no code with it. It rides the graphql-js dialect and still gets its own fixture and test.

async-graphql shares neither half. A different verb, and a separator that is a **bare comma with no or anywhere** — its `suggestion.rs` has no branch for the final element. That is a third separator across three cores, and it needed no new parser code: pulling quoted runs absorbed a joiner nobody anticipated, which is the payoff for a decision made two builds earlier.

HotChocolate never suggests a name at all — one `FieldDoesNotExist` resource string, in backticks, no clause. `suggestions_unsupported`, and not one wordlist request is spent.

Guessing would have got all three wrong, in a different direction each time.

⚠ A SHIPPING DEFECT, FOUND BY THE NEW TESTS AND FIXED IN THE SAME COMMIT.

`parse_suggestions` harvested the quoted runs out of a `Did you mean` clause **without requiring the unknown-field sentence to have matched** — it could not require it, because the argument path reuses the same function on a different sentence. graphql-js has another quoted `Did you mean`:

```
Field "user" of type "User" must have a selection of subfields. Did you mean "user { ... }"?
```

That is graphql-js's own `ScalarLeafs` rule, and it fires **whenever a probe stem names a real composite field** — `user`, `account`, `order`, which is most of the wordlist against most schemas. The enumerator recovered `user { ... }` and reported it as a field the server had volunteered: a fabricated schema entry, from a real server, on the common path. Found because HotChocolate copied that message from graphql-js and a fixture went looking for it.

Fixed by filtering recovered names through the GraphQL spec's own `Name` production (`[_A-Za-z][_0-9A-Za-z]*`) rather than a blacklist of wordings we happen to have seen — the next core to copy that message is not a change here.

.NET is two implementations, and they agree on nothing

`graphql-dotnet` was read straight after, because graphw00f carries it as an engine **separate from** `hotchocolate` and reading one of them would have left half the platform unenumerable while looking complete. The two share no wording at all:

```
hotchocolate The field `usr` does not exist on the type `Query`.      (never suggests)
graphql-dotnet Cannot query field 'usr' on type 'Query'. Did you mean 'user' or 'users'?
```

`graphql-dotnet` reproduces **graphql-core's grammar to the character** — single quotes, bare `or` at two, an Oxford comma from three, the same cap of five — from a C# codebase. So it rides the graphql-core dialect, and, exactly as with graphql-php on graphql-js, a **.NET server running it fingerprints as core graphql-core**: the dialect is right, which is what the parser needs, and the brand is not recoverable from that probe and is not guessed.

It also contradicts itself between two adjacent files, and that was worth a fix.

`FieldsOnCorrectTypeError.cs` appends the suggestion clause **with** a trailing `? ;`

`KnownArgumentNamesError.cs:45` appends the same clause **without** one. A clause pattern demanding a literal question mark recovered graphql-dotnet's field names and **silently dropped every argument name it volunteered** — and the argument form is the one report #61 actually needed. The terminator is now optional, with the body still required to open and close on a quote so an empty clause cannot match.

Its `ScalarLeafs` message is `Field usr of type User must have a sub selection` — unquoted, no `Did you mean` — so unlike HotChocolate's it could never have fabricated a field.

The long tail, triaged — one question each, no parsers written

The seven engines still resolving to `unknown` were put to a single question — **does it emit a name suggestion at all?** — answered by grepping each repository, not by reading its message format. That is a cheap question with an expensive answer: a core that never suggests is a two-line entry and a `suggestions_unsupported` verdict, and needs no parser at any point.

Engine	Lang	Suggests?	Evidence
dgraph	Go	already covered	<code>go.mod</code> pins <code>dgraph-io/gqlparser/v2</code> ; its <code>fields_on_correct_type.go</code> appends <code>" Did you mean " + QuotedOrList(...) + "?"</code> — the <code>gqlparser</code> core this project already parses
sangria	Scala	yes — now closed	<code>Violation.scala:257</code> Cannot query field '\$f' on type '\$t'. + <code>didYouMean(...)</code> . <code>StringUtil</code> read: single quotes, cap five, and no Oxford comma ('a', 'b' or 'c'). Rides graphql-core; the missing comma is invisible to quoted-run extraction. Locked with a fixture and a test
lacinia	Clojure	no	<code>parser.clj:1160</code> Cannot query field %s on type %s. ; zero <code>Did you mean / suggest</code> in the repo
caliban	Scala	no	<code>Validator.scala:582</code> Field '\$f' does not exist on type '\$t'. ; zero <code>Did you mean</code>
agoo	C	no	zero <code>Did you mean</code> in the entire repository; its only field-related error strings are allocation failures
morpheus-graphql	Haskell	no — now closed	shares graphql-js's field sentence exactly, identified by a directive probe and downgraded to <code>suggestions_unsupported</code> . Its only <code>Did you mean</code> is a stale doc-comment it never sends
pg_graphql	Rust	no	<code>builder.rs:1739</code> Unknown field "{}" selected on type "{}"; zero <code>Did you mean</code>

Five of seven never suggest. One is already covered. One suggests — sangria — and it is now closed (see its row: read, and locked with a fixture and a test). The tail was worth measuring precisely because measuring it was cheap and it turned out to be almost entirely empty.

Three things fell out that matter more than the tally:

- **No shadowing.** caliban's `does not exist` (against graphql-ruby's `doesn't exist`), pg_graphql's `selected on type` (against async-graphql's `on type`) and lacinia's unquoted form all classify as `unknown` — near misses that stay misses.

- **morpheus is the first core that shares graphql-js's unknown-field sentence and never suggests** — and it is now handled. The field probe alone reads it as `graphql-js`, `suggests: True`; its ungrammatical directive-location error (Directive "skip" may not TO be used on QUERY) is unique to morpheus and downgrades it to `suggestions_unsupported`. That is the one brand probe permitted to change `suggests`, justified exactly as Hasura's core-setting probe is: a positive identification, not an inference. Without it an operator would have spent a full wordlist run to earn `suggestions_disabled` where the truth is `unsupported` — a wasted run and a wrong label, never a fabricated name.
- **△ A correction to this document.** An earlier draft claimed morpheus "carries the ScalarLeafs message too" — a third `Did you mean "hobby { ... }"`? after `graphql-js` and `HotChocolate`. That was **wrong**: the string is a stale doc-comment above `subfieldsNotSelected`, and the function emits `must have a selection of subfields` and stops. morpheus never sends it, so it was never a fabricated-field hazard. The Name-production filter would refuse it regardless, and a test asserts both — but the claim that it reached the wire is retracted.

Two structural tests were added with them, because the *shape* of adding a core is what breaks: `classify_fingerprint` carried a **literal tuple** of the three dialects that existed when it was written, so a fourth `DIALECTS` entry would have been never tried, read as `unknown`, and left the suite green — the per-dialect tests call `parse_suggestions` directly and never reach the classifier. The loop is now driven off `DIALECTS`, one test asserts **every** dialect is reachable from it, and another asserts each fixture classifies to **its own** core, since first-match-wins means a loosely-written newcomer can shadow a core that already worked.

FIVE outcomes, and four of them are an empty list from a naive implementation

`productive` / `suggestions_disabled` (a **working defence** — stop) / `suggestions_unsupported` (the core has no such feature) / `engine_unknown` (no parser chosen, deliberately) / `failed`. Build #20 found ZAP answering the same code for introspection-disabled and a dead host and classified around it; returning an empty list for all five here would be that defect in a new place — in the one place where `suggestions_disabled` means *their defence worked* and `failed` means *try again*.

The denominator is reported so a zero is readable: `fields: 0` beside `unknown_field_errors: 900` is a server that answered everything and helped with nothing; `fields: 0` beside `unknown_field_errors: 0` is a server that never understood us.

A five-suggestion response is recorded as CUT OFF, not complete. Both measured cores cap at five, so an enumerator treating five as "all of them" would stop early and report a schema it had not finished reading.

*** A REAL BUG, CAUGHT BY A FIXTURE. *** The fingerprint reads raw response text, but on the wire the message arrives inside JSON, so `graphql-js`'s quotes are **escaped** while `graphql-core`'s single quotes are not. Matching raw text identified every Python server correctly and read every Apollo, Yoga and Mercurius server as `unknown` — asymmetric, silent, and pointing the **wrong way**: it would have looked like the `graphql-js` servers were the hardened ones. Fixed by parsing the envelope first, in **one** shared function, because two parsers is how that divergence lived.

The bounds are a declared parameter, not a gate, and that is asserted structurally: the pure module contains no `raise` at all, a zero bound means unbounded rather than refused, and a run stopped by a bound keeps its results — discarding a partial schema would punish an operator for setting a bound. This is **not** the scanner-wide request cap that was declined; it is one feature describing its own size, like crawl depth and scan policy.

Item 4 — what the recovered schema is FOR

- **The repeater.** A recovered field composes into an operation with its recovered arguments as **variables, one key each** — which is what lets the scanner reach one argument at a time, the chain report #61 needed: no introspection, enumerate, aim at an argument. **No values are invented:** placeholders are empty, because nothing here knows what they should be and inventing one would put a request on the wire nobody wrote.
- **The scan path.** `scannable` is **always false** and says why: ZAP files nothing it generates into the Sites tree, so a composed operation becomes scannable by being **sent through the proxy** and captured. Same shape as `SchemaImport.scannable`, and for the same measured reason.
- **Provenance, because a report that blurs it is dishonest.** The engagement record says the schema was **MINED, not introspected**, with the engine core, the wordlist and its size, the request count, and the method's own limit stated where a reader would otherwise assume completeness: *only names close enough to a wordlist entry can ever surface, so a field's absence is not evidence it does not exist*. A truncated or early-stopped run says so too.

Names, never values — held for a third path. No model on it has a value field, and a test plants a secret and asserts it reaches no record, no composed body and no API response.

Item 5 — the screens, and looking at them found two things nothing else could

`tsc --noEmit 0, next build 0, npm run lint` unchanged at **11 errors + 1 warning**, before and after every one of these.

1. ***** AN INVENTED CSS CLASS. ***** Two new buttons carried `hp-gql-go`, which **does not exist**. The established vocabulary trap, and `tsc / build / eslint` cannot see it. Fixed by matching the existing bare button inside `hp-tn-form`, as the adjacent "probe schema" button does. `test_css_vocabulary.py` — the guard built for exactly this — passes.
2. ***** A DROPPED SPACE, RENDERED AS `1response(s)`. ***** The JSX had a space between the count and the text and it still rendered without one. Fixed with the codebase's explicit space idiom rather than relying on JSX whitespace rules, and re-verified in the DOM: `1 response(s)`. Build #20 found three of these; this is the fourth.
3. A doubled terminator in the new copy, spotted in the rendered text and removed.

A RESTRICTION REMOVED. The panel takes its container from the **running proxy**, so the schema probe was unreachable until a ZAP daemon existed — even though a probe is one `docker exec` plus `curl` and needs no daemon at all, and the moment it is most useful is *before* anything is set up. Enumeration inherited the same shape, and the panel's own advice ("introspection is off, so enumerate it below") would have pointed at a disabled button. A blank container now means the engage sandbox. Naming one still wins.

Driven in the real Chrome end to end against a live multi-dialect origin: `productive` with 10 names recovered — `adminPanel`, `sessionToken`, `auditLog` among them — the bounds line reading back what applied, and the **positive control** flipping the same panel to `suggestions_disabled` with an amber chip, zero fields, and copy that reads *"a working defence, not an empty schema"*.

Verification

- `sh backend/run_safety_tests.sh` — **89 test files, every one exited 0** (88 plus `test_graphql_enum.py`).
- Green again with `docker` **stripped from** `PATH` for all seven files this build touches, and the strip verified to bite first: `docker` gone, `git` still present.

- `docs/proof/build20_graphql_api.py` re-run against the pinned add-on — **56 passed, 0 failed**, baseline reproduced, no divergence.
- `docs/proof/build21_graphql_enum.py` — **33 passed, 0 failed**, including the positive control and the deliberate wrong-parser run.
- `docker/proof/browser_intercept_proof.sh` on the rebuilt image — **25 passed, 0 failed**.
- `npx tsc --noEmit 0 · npm run build 0 · npm run lint` at the accepted **11 + 1**.
- `data/kb/entries.jsonl` still **2,747**, matching `kb_fixture.jsonl` — nothing was ingested.
- **The image was rebuilt ONCE**, with every image change of this build in it.

*** A FLAKE CONFIRMED AS A FLAKE, NOT ASSUMED TO BE ONE. *** `test_redirector.py` failed three consecutive times mid-build with "no bindable port outside the self-mapping range", then passed in the next full run. `netsh interface ipv4 show excludedportrange protocol=udp` shows Windows holding a large and growing block of UDP ranges; the failure is that environment condition, in a file this build does not touch, and it is recorded rather than waved away.

What was NOT done, and why

- **Route auth, scanner pacing, a scanner-wide request cap, mobile scope, Caido** — standing permanent skips, not re-proposed.
- ~~HotChocolate, Absinthe and Juniper have no suggestion parser~~. Closed 2026-08-05 — see *Three more cores* below. HotChocolate and Absinthe were read; **async-graphql was read in place of Juniper**, which is the Rust library people name rather than the one they deploy. `graphql-dotnet` was read in the same session — see *.NET is two implementations* below.
- **No OOB canary provisioning** — a separate parked decision.
- **No new gate, confirm, blocklist or allow-list narrowing** — the build's own requirement, and the one change to reachability removes a restriction rather than adding one.

interact.sh — a second, zero-infrastructure OOB backend (2026-08-06)

The out-of-band canary (build #13 part 3) was the private, owned option: it catches blind callbacks, but only after a VPS, a domain and a one-time NS delegation exist. This adds a second backend alongside it — ProjectDiscovery's public **interact.sh** service, the same shape as Burp Collaborator — so the same capability exists with **no infrastructure at all**. Both backends can be configured and live at once; a poll sweeps both and merges the findings. It is an *addition*, never a replacement, for one reason stated plainly to the operator: **interact.sh callbacks transit a third party**, which is exactly why the owned backend still exists.

The valuable half is reused unchanged. `interact.sh` works inversely to the self-hosted canary — *it* assigns the correlation-id and encrypts every callback to a public key we register, so a new module owns a keypair and a session rather than a VPS and a zone. But the part that is actually the product — correlate a callback back to the step that caused it, file it as a finding — is backend-agnostic. Both backends normalise into the *same* hit dict, and the existing `findings_for / ingest` files them identically. What is new is one transport: register / generate / poll / deregister, with the RSA-OAEP-SHA256 + AES-256-CFB handshake `interact.sh` uses (the `cryptography` library was already present; no new dependency).

The new outbound poll carries the self-hosted poll's containment, asserted, not asserted-in-a-comment. It is a new egress path out of the operator's machine, so it inherits exactly what `poll.py` has: the destination is the configured `interact.sh` server resolved server-side, never a request field; no redirect is followed (a tampered or

proxied server answering `302 http://169.254.169.254/...` is an error, not an SSRF from the backend host); no ambient proxy is honoured; the response is byte-capped; JSON is parsed and never executed; and every decrypted interaction is treated as untrusted input, capped before it becomes a record. The secret-key, RSA private key and any self-hosted auth token are write-only — stored in the gitignored `sessions.db`, never returned by any view. **The oob-package network-reach guard did its job:** it pins the set of modules allowed to touch the network and *failed* the moment `interactsh.py` gained reach, forcing the containment treatment and this paragraph rather than letting a new egress path in quietly.

Two conveniences were added, and exactly one line was drawn. *Auto-poll* is a background sweep that files callbacks from both backends without a click — but it is **read-only automation**: it reaches `poll_all` → `ingest` → `state` and no execution surface, sends nothing, and runs no command, so it does **not** cross the propose-only invariant (a dedicated safety test asserts it reaches no execution or delivery surface). *Send-to-repeater* pre-fills a rendered payload into the repeater editor — as a **frontend action only**: no backend OOB module imports the repeater, so its human-only `send()` guarantee is untouched and the operator still clicks Send. The alternative — HackPit firing the payload at a target itself — was scoped, named as the reversal of the project's central invariant it would be, and **declined** in favour of this (the operator's L1 choice).

Verification. The hermetic suite gained two files and grew the token and template locks. The crypto is exercised for real against an in-process interact.sh-style encryptor — a keypair is generated, an interaction is AES-CTR-encrypted and its AES key RSA-OAEP-wrapped to the client's public key, and the client is asserted to decrypt → `correlate` → `file`. Coverage includes the containment (no redirect, no ambient proxy, response cap), the suffix correlation map, dedup by `<uid>|<timestamp>`, the kept-and-flagged uncorrelated hit, `poll_all` merging both backends and isolating a backend failure, the masked view never returning a secret, and the router rendering only configured backends. **Live, and end to end:** `docs/proof/oob_interactsh_proof.py` registers with `oast.pro`, generates a host, resolves and requests it, then polls, decrypts, and asserts the callback correlated back to the step that generated it — **PASS**. This is the one live check that needs **no infrastructure of the operator's own**, only a network that does not blocklist interact.sh.

And the live proof earned its keep by catching a defect the hermetic suite could not. The first implementation decrypted interactions with **AES-CFB**; interact.sh actually uses **AES-CTR** (Go's `cipher.NewCTR`, the 16-byte IV as the initial counter). With CFB only the first 16-byte block decrypts and the rest is garbage, so every real callback failed to parse — and because `poll_correlated` skips a blob it cannot decrypt, the symptom was an empty poll that *looked* exactly like "no callback arrived". It was first misread as a DNS/EDR filter on `oast.*`; dumping the raw poll response disproved that — the host resolved to interact.sh's real IP and the server returned four encrypted interactions, so the traffic was arriving and the fault was ours. The hermetic test had passed only because it both encrypted and decrypted with the same wrong assumption; it now uses CTR, so it validates the real wire format, and the live round-trip passes. A self-consistent test agreeing with its own bug is the exact failure mode a live proof exists to break.

Frontend, and it was looked at, not just built. The `:oob` panel gained an interact.sh card (register / status / deregister), an auto-poll toggle, dual-backend mint rendering, and a per-payload → **repeater** button — the send-to-repeater convenience, which is a frontend-only action (no backend OOB module imports the repeater, so its human-only `send()` guarantee holds). It was verified mechanically — `tsc` clean, the CSS-vocabulary lock passing so no class renders invisible, the eslint baseline held at 11, `next build` exit 0 — and then **visually**, against the running stack: the unconfigured panel shows the interact.sh card with its form and the gated sections correctly hidden; registering (live, on `oast.pro`) flips the card to a registered status and reveals the auto-poll, verify, mint and collect sections, and a live mint returned 15 interact.sh payloads. The registration, poll and mint were thereby confirmed end to end through the real API and UI, not only in the hermetic suite; the session was deregistered afterward and no state leaked into the repo.

The stale note in `cockpit/repeater.py` that still called the VPS-for-callbacks piece "deferred (D2)" was corrected — that listener shipped in build #13 part 3 and now has two backends. The full safety suite is green — **93 test files, every one exited 0** (up from 89; the four new files are the `interact.sh` backend, its safety invariants, the auto-poll setting/tick, and the dual-backend router). The `test_redirector.py` UDP-port flake recorded above recurred once and passed on a clean re-run, in a file this build does not touch. No new gate, confirm, blocklist or allow-list narrowing; the only reachability change adds an option and removes no restriction.

Credential attack — spray captured creds, crack captured hashes (`:credentials`) (2026-08-06)

The state model has captured credentials and hashes since Phase 2, the arsenal has catalogued `netexec` , `kerbrute` , `hydra` , `hashcat` and `john` since the start, and the AD graph has had an owned-node concept since the WinRM build. What was missing was the wire between them: a way to *use* a captured hash. This build is that wire — and nothing more, deliberately.

One approval buys the whole job — the intruder's argument, applied to a long process. A spray of 300 accounts and a crack of 40 hashes are each ONE job that produces many attempts. HackPit refuses *batching across approvals*; it has never refused one approval that produces many requests, and could not — `ffuf` , `nuclei` , the ZAP scanner and the intruder are each a single press buying thousands. So `:credentials` adds **no new gate**: `executor.validate_request` runs against the equivalent `netexec / hashcat` command before anything spawns, and the stop button is the ungated panic switch, exactly like `stop_scan` . The one place the intruder's shape had to change is the stop: the intruder checks a flag between its own requests, but a spray is one `netexec` process, so the worker holds the process handle and `stop()` kills it — and also enforces the operator's `stop_on_lockouts` knob by watching the output for `STATUS_ACCOUNT_LOCKED` . That knob, and the delay, are *operator inputs, not gates*: a slower spray is a quieter spray, not a safer one, and refusing to run because a number was left at zero would be a prohibition the tooling invented.

The planner executes nothing; the worker is the only thing that runs. `cockpit/credattack.py` is pure — it builds `argv`, detects the hashcat mode from a hash's shape, and correlates tool output back into typed state — and an AST walk over its own source (like `state/ 's`) asserts it imports nothing that executes and calls no `subprocess / eval` . The execution lives in `cockpit/credjobs.py` , the gated job worker. **Secrets never land on an argv**: the user list, the password list and the hash list are written to files under the engagement's loot directory and the `argv` references only their paths — a password on the command line ends up in the persisted `RunRecord` that `report.py` renders verbatim, which is the exact leak `secretargs` was built to stop. A crack that reaches the `argv` would be that same defect wearing a hash. Both are regression-locked in `test_credattack_safety.py` .

A crack recovers the account, not just the plaintext. `hashcat` prints `hash:plaintext` , but that split is ambiguous for salted/Kerberos hashes whose own body contains `:` . Rather than guess, the parser matches each line against the hashes that were actually *submitted* — the longest known hash that prefixes the line wins — so the account is recovered unambiguously from state that produced the job. A hit emits a new `password` credential for that principal (the NT hash is kept, still usable for pass-the-hash), a `high` finding, and marks the principal **owned** in the session's AD graph, opening new frontier edges in `:ad-graph` . That last step is the payoff the whole loop was for: a cracked hash becomes a new route to Domain Admin.

Built, verified, and looked at. The two new test files (`test_credattack.py` , its safety twin) join the runner; the full hermetic safety suite is green — **95 test files, every one exited 0** — and the CSS-vocabulary lock passes against the new screen, so no class renders invisible. `next build` exits 0 with `/credentials` compiled. The screen was then **looked at**, not just built (`frontend-class-vocabulary`): against the running stack, seeding from a synthetic

lab engagement grouped three captured NTLM hashes under `-m 1000` and a kerberoast ticket under `-m 13100`, and the spray preview rendered the exact `netexec smb ... -u /loot/.../users.txt -p /loot/.../pass.txt -d LAB` with the gate's verdict shown before anything runs — the secret lists as files, on screen, exactly as designed. No new gate, confirm, blocklist or allow-list narrowing.

Nuclei template scan — the bug-bounty staple, wired (`:nuclei`) (2026-08-06)

`nuclei` has been catalogued in the arsenal and its ~13,400 templates baked into the sandbox image since the start, and the state model has had a `Finding` type since Phase 2. What was missing was the surface that joins them: point the template engine at the scoped target(s) and turn matches into engagement findings. This build is that surface — the lowest-effort of the four, because the plumbing already existed.

One approval buys the whole scan — the `ffuf` / ZAP-active-scan shape, no new gate. A nuclei run fires thousands of template checks; it is ONE job that produces many requests, exactly the position the intruder and the credential surface already hold. `executor.validate_request` runs against the equivalent `nuclei -u <target> ...` before anything spawns, and the stop is the ungated panic switch, like `stop_scan`. The one thing done *right* rather than hardcoded is the sandbox: `cockpit/nuclei.py` calls `executor.resolve_mode` — the same single source of truth the one-shot `/cockpit/exec` uses — so a lab scan runs in the isolated, egress-less lab box (which reaches the lab target on its internal network) and an engagement scan runs in the fully-open engagement box, and a scan can never bind to a different sandbox than a bare command would in the same mode. Every target rides the argv as `-u <target>`, so in engagement mode an out-of-scope host is refused at the *inherited* target handrail — not a stronger gate this build invented.

The planner/parser executes nothing; the worker is the only thing that runs. `cockpit/nuclei.py` 's pure half — `resolve_targets`, `nuclei_argv`, `parse_findings` — builds the argv and turns nuclei's JSONL into `Finding`s, and a per-function AST walk asserts it reaches no `subprocess` / `eval`. Default targets seed from the session's in-scope endpoints already in state (falling back to hosts); results map `info.name` → title, `matched-at` → target, `template-id` → reference, curl/matcher output → evidence, and **dedupe by** (`template-id`, `matched-at`) before upserting into the same engagement `Finding` store the report renders. The live finding count grows as JSONL streams, so a running scan is watchable.

A live run caught a defect a hermetic test could not — the recurring lesson, again. The first version built the argv from only `-tags` / `-severity`, trusting nuclei to find its baked templates. Against the real lab it died instantly: `FTL Could not run nuclei: no templates provided for scan` — a `docker exec` resolves no default template directory under its `$HOME`. Nothing in the parse tests could see it (they feed the parser a string they chose). The fix points `-t` at the baked repo (`/usr/share/nuclei-templates`) whenever no explicit templates are named; a regression test now pins it, and `docker/proof/nuclei_proof.sh` checks the image's `nuclei` actually accepts the exact flag string against the lab. This is the third surface where "a hermetic test feeds the parser a string the test itself chose" has cost a defect.

Built, verified, and looked at. The two new test files (`test_nuclei.py`, its safety twin) join the runner; the full hermetic safety suite is green — **97 test files, every one exited 0** — the CSS-vocabulary lock passes against the new screen, and `next build` exits 0 with `/nuclei` compiled. The screen was then **looked at** against the running stack: a real scan of the lab target (Juice Shop) surfaced a **medium** Prometheus-metrics exposure plus tech-fingerprint and Swagger findings, each deduped, severity-ranked, and upserted into engagement state — visible in the results feed exactly as designed. No new gate, confirm, blocklist, or allow-list narrowing.

Documentation & housekeeping (2026-08-06)

Recorded here once, for completeness — these are documentation, packaging and a couple of small behaviour changes that landed alongside the `:credentials` build, not new capabilities of their own.

- **README, rewritten whole-project.** The old README described only the companion. It was replaced with a full rewrite covering every current surface (companion + cockpit + Windows/AD + C2/tunnels + web-app testing + OOB + MCP), with a badge row, a demo-video/logo placeholder, an "at a glance" table, the four-gate safety model, an honest security-posture/limits section, and a "Build notes" case study. Counts were **verified against the repo, not the old prose**: 2,747 KB entries, 121 tools / 306 templates, 47,108 exploits / 25,041 CVEs, 65 ATT&CK / 49 Sigma, 93 test files, 15 MCP tools, 27 screens. The code-only / provenance framing (`sources/*-manifest.md` , gitignored KB) was kept.
- **Apache-2.0 LICENSE added.** The repo had no `LICENSE` file backing its stated license; the full Apache-2.0 text was added and the README badge/section now link to it.
- **Default LLM → Claude Agent SDK (Opus).** `backend/llm.py` DEFAULTS now resolve to the Claude Agent SDK on model `opus` (via the local `claude` CLI, no API key), with local Ollama as the automatic offline fallback. The README's provider list was corrected — it had listed Groq/xAI, which the code never supported (real set: `ollama, openai, anthropic, openrouter, claude-agent-sdk`).
- **CI dependency fix — main had been red since interact.sh.** `oob/interactsh.py` imports `cryptography` at module top and `main.py` imports it unconditionally, so `import main` failed on CI's bare dependency set (`ModuleNotFoundError: cryptography`). `cryptography` is now declared in `backend/pyproject.toml` as a core dependency and installed in the CI step; CI is green again.
- **Four build specs authored** (`docs/superpowers/specs/2026-08-06-*.md`) for the next offensive surfaces — credential-attack (now BUILT, above), a cloud IAM-privesc graph, a nuclei template-scan surface, and a guided recon → ranked attack-surface flow. Each rides the existing per-command human-approval gate and adds no new gate; each carries a \$6 requiring a README section + a real lab-only screenshot + this assessment's update as part of its own definition of done.
- **Screenshots — headless-Edge method + a live-lab refresh.** The README's screenshots are captured with headless Edge against the running app (the Chrome extension being unavailable), then eyeballed (a class the type-checker cannot see still renders invisible — `frontend-class-vocabulary`). The stale home (which showed the old `1,551` counter) and intro splash were refreshed, and the container-dependent cockpit screens (proxy, C2/Sliver, tunnels, OOB, intruder, repeater) were re-captured against the running stack with live lab data — seeded from a demo engagement scoped to the OWASP Juice Shop lab target, with the lab-target's network isolation restored afterward and real bounty targets kept out of every frame.

Guided recon → ranked attack surface — the front door (`:recon`) (2026-08-06)

The gap. Everything downstream (nuclei, the graphs, the intruder) needs a *surface* to aim at, but HackPit had no first step that turned a scoped domain into one. `:recon` is that front door: a scoped domain runs recon as approved jobs, seeds engagement state, and **ranks** the resulting surface by likely-exploitable.

Shape (`backend/cockpit/recon.py`). Mirrors the credential surface (engagement-bound) + the nuclei pure/gate split: pure `*_argv` builders + `filter_in_scope` + `rank_surface` (advisory, executes-nothing, AST-locked) and an engagement-bound worker. The passive sweep chains `subfinder` → `dnsx` → `httpx` → `gau/waybackurls/katana` ; the active sweep chains `naabu` → `nmap -sV` . **One approval per sweep** runs the several `docker exec` s of that chain

(the crack-worker shape, gated on the entry command — the sweep model \$0 permits; not a batch auto-runner). It is engagement-bound rather than lab because passive recon needs egress (crt.sh / wayback / DNS) and the scope lives on the engagement record.

Scope-safety by construction, not a gate (\$0). Discovered hosts are sorted against the program scope (`engagement.record.discoveries + recon.filter_in_scope`); **only in-scope hosts are ever written to the loot host-file the probing tools read**, and every parsed record is scope-filtered before it can become scannable. An out-of-scope discovery is read-only — never upserted. `test_recon_safety.py` proves, end to end against a throwaway engagement DB, that an out-of-scope discovery enters neither the allowed set nor the scanner's host file. `rank_surface` then orders a session's state by likely-exploitable — services + a CVE-worthy stack (pulling **real** CVE counts from the 47k-row `:exploits` index) + param-rich endpoints + auth surface + existing findings — states the *why* for each, hands off to `:nuclei`, and is deterministic.

Tests + screen. Reused only existing `hp-tn-* / hp-ap-*` classes (no new phantom). `test_recon.py + test_recon_safety.py` registered — **99 files green**, next build exits 0; a `?session=<id>` deep-link renders the ranked view for the headless-Edge screenshot. `:recon` added to the `operate` band. Committed `43d4cf3`.

Cloud attack surface & IAM privesc graph — the cloud parallel to BloodHound (`:cloud-graph`) (2026-08-06)

HackPit could already route from an owned low-priv user to Domain Admin over a typed graph of abusable AD edges. This build gives the **cloud** the same thing: enumerate an account, build a typed **IAM privilege-escalation graph**, and route to an **admin/owner-equivalent** principal — walked the exact edge-index way the BloodHound orchestrator already works. The new package `backend/cloudgraph/` is a near-clone of `adgraph/`, deliberately: `schema.py` (nodes carry a `provider`, so one engine serves AWS/Azure/GCP), `parser.py` (ScoutSuite/Prowler JSON → graph, the cloud equivalent of the BloodHound parser), `paths.py` (the same BFS + k-alternatives), `orchestrator.py` (copied wholesale — the model picks an edge index, a pick outside the frontier is refused not repaired, it authors nothing and executes nothing), `techniques.py` (a cloud abuse catalog), `store.py`, `router.py`, `sample_data.py`, and the enumeration worker.

It adds **NO** new gate — \$0's binding constraint. Two surfaces, both mirroring what already exists. *Enumeration* is a **gated job** (the recon/nuclei/credattack shape): `ScoutSuite + Prowler (+ cloudfox)` run as ONE approved job, gated by the same `executor.validate_request` before anything spawns, with an ungated stop, engagement-bound because a real cloud API needs the open, loot-mounted sandbox where the credentials live. *The walk* is **propose-only**: the orchestrator hands back an edge index; the abuse COMMAND comes from the deterministic, KB-grounded technique catalog, not the model; approval goes to the SAME `POST /cockpit/exec` every other command uses; and `advance` moves the walk **only** on a run that was approved and exited 0, verified server-side against the recorded run. The two halves are separate modules on purpose — `enumerate.py` is the only thing that execs, and `orchestrator.py` is AST- and source-scanned to prove it execs nothing, so the propose-only invariant is regression-locked, not asserted.

The privesc edges are the real ones an attacker walks, derived from IAM policy statements: `sts:AssumeRole`, `iam:AddUserToGroup / AttachUserPolicy / AttachRolePolicy / PutUserPolicy / CreatePolicyVersion / UpdateAssumeRolePolicy / CreateAccessKey / CreateLoginProfile / PassRole`, `lambda:UpdateFunctionCode / CreateFunction`, `ec2:RunInstances`, `secretsmanager:GetSecretValue`, `kms:Decrypt`, plus Azure `Owner -on-self / app-credential-add / VM run-command / AKS admin-creds` and GCP `serviceAccountTokenCreator / actAs / setIamPolicy`. A Lambda/EC2 abuse edge is redirected to the

compute's **execution role** — that is where the privilege actually lands. A principal that can do `::*` (or holds `AdministratorAccess`) is marked the admin objective. Privilege-escalation paths and Prowler misconfigurations land as engagement `Findings`, so they flow into the report the same way every other finding does.

KB-grounding, made honest for a prose-heavy corpus. The 534-entry `hacktricks-cloud` KB grounds each edge via the same hybrid-search + `entry_commands` mechanism the AD grounder uses — but the cloud corpus is reference/prose, and a seed-term match routinely returns a JSON policy document or an unrelated `aws` command. So a grounded command is **adopted only when its CLI action matches the edge's own** (`aws iam attach-role-policy == aws iam attach-role-policy`); otherwise the precise catalog command is kept and the KB entry is still **cited** as the explanation. Enrich, never mis-ground. This was caught live: with the naive grounder the `AttachRolePolicy` hop rendered `aws stepfunctions test-state` — a real CLI, wrong abuse — before the match-guard fixed it.

pacu + cloudfox added to the arsenal AND the sandbox image, up front (part of this build, not a follow-up). Both are catalogued in `arsenal/tools.json` (read-only enumeration templates; the actual IAM-mutating abuse a path leads to is a `cloudgraph` technique command that clears the executor gates, *not* an arsenal tool — pinned benign in `test_arsenal_safety.py`) and installed in `docker/Dockerfile.sandbox` alongside `awscli`, `scoutsuite` (binary `scout`, **not** `scoutsuite` — the kali-sandbox name trap) and `prowler`, each in its own venv, cloudfox as a release binary. `docker/proof/cloud_install_proof.sh` re-checks every name the catalog templates hardcode against the built image and the running container — the `zap_install_proof` lesson (a hermetic test can only feed the parser a string it chose itself). **The image rebuild + this proof are the one step run outside Claude** (Docker trips the real-time classifier; run `docker compose -f docker/docker-compose.yml build engage-sandbox` then `sh docker/proof/cloud_install_proof.sh`).

Frontend `/cockpit/cloud` clones `/cockpit/ad`: the same kill-chain route canvas (reusing the `hp-adg-*` primitives — no new CSS), provider tabs (AWS/Azure/GCP), the agent-proposes-an-edge orchestrator panel, and a per-hop drawer with the KB-grounded command, the destructive red-confirm, and the defender's-eye detection disclosure. **The screen was looked at** (`frontend-class-vocabulary`): the synthetic AWS sample renders a real 3-hop route — `dev-alice` (owned) → `developers` → `ci-deployer` → `break-glass-admin` (admin/owner) — with no cloud credentials. A `?demo=1` deep-link auto-ingests the sample so the headless-Edge screenshot renders the route, not the empty state (the recon-surface pattern).

Verification. `test_cloudgraph.py` (parser: ScoutSuite/Prowler JSON → graph; BFS to an admin principal; orchestrator edge-index proposal; `advance` requires an approved exit-0 run) and `test_cloudgraph_safety.py` (mirrors `test_adorch_safety.py`: the model picks an INDEX never a command and an out-of-frontier pick is refused; the orchestrator executes nothing by AST and has zero `:kali`; never-auto-run; no second execution path; inherited-rights edges never acquire a command even from a loud grounder; ENUMERATION ADDS NO GATE — its argv builders execute nothing by AST, `start` reaches the executor gate before any spawn, approval + red-confirm default FALSE, it is engagement-bound and `stop` is ungated; lab unchanged) both pass, and are wired into `run_safety_tests.sh`. `next build` exits 0 with `/cockpit/cloud` in the route table.

Assumptions (per §5), stated. Enumeration is **AWS end-to-end** today; the schema, technique catalog and routing already carry Azure/GCP node/edge kinds, and the enumerate worker runs `scout` / `prowler` for those providers, but the deep policy-statement → privesc-edge extraction is AWS-first — Azure/GCP produce principal/resource nodes + their findings via the same tolerant reader. The default objective is an admin/owner-equivalent principal; the operator can name a different target node.

Web SSRF → cloud credentials — the IMDS seam (:cloud) (2026-08-07)

The seam. The web half and the cloud graph needed a bridge: a web-side **SSRF or RCE** that reaches the instance metadata service (169.254.169.254 , metadata.google.internal) hands back the instance's temporary role token — the classic pivot from a web bug into the cloud control plane. This build parses that **captured** response into an **owned** principal seeded into the :cloud IAM graph, so the privesc walk begins from the identity you just stole.

A pure parser (backend/cloudgraph/imds.py). It executes nothing and imports no network/exec module (AST-locked in test_cloudgraph_safety.py) — the request that actually touched IMDS already ran through the human-approved repeater/executor, so the module only reads a captured string. It covers **AWS** (IMDSv1, the IMDSv2 token-PUT + creds-GET two-step, the role listing, the instance-identity doc), **Azure** managed-identity JWTs (base64url-decode the payload for oid / appid / tid , no verification), and **GCP** service-account tokens (with the Metadata-Flavor: Google requirement flagged). Malformed/truncated bodies degrade to warnings and never raise; an empty body or unknown provider is a clean ImdsParseError . The returned finding_evidence is **secret-free by construction**; the credential is carried separately.

Seeding + the decoupling. store.seed_owned_node merges the owned node into the session's latest graph by id/label/tail-alias match (mirroring mark_owned) so the stolen identity lands on the already-enumerated node and a route lights up immediately; otherwise it adds it standalone, or creates a minimal graph if none exists — keeping the same graph_id so the frontend stays valid. The route lives in main.py , not cloudgraph/router.py , because it is cross-cutting — it touches engagement findings + the credential vault + cockpit.loot , and **cloudgraph must not import cockpit** (the decoupling rule). POST /cockpit/cloud/seed-imds + GET /cockpit/cloud/imds-catalog . **No new gate** — nothing here executes.

Secret hygiene + tests. The captured secret goes to the engagement **vault** (a state Credential , kind=token) and a loot file — **never** the Finding, which records only provider / identity / expiry. Frontend: the CockpitCloudSeed panel on /cockpit/cloud . test_cloud_imds.py + the extended test_cloudgraph_safety.py are green and registered; next build exits 0; a live HTTP round-trip through the real route was confirmed (the seed merges onto ci-deployer , the secret reaches the vault, a high finding is recorded with no secret in its evidence, and a 1-hop route to admin lights up). Screen looked at: assets/screenshots/35-cloud-imds.png . Committed fa707ce .

AD CS ESC1–8 routed in the graph — the AD parallel to DCSync synthesis (:ad-graph) (2026-08-07)

The AD attack-path graph could route ACL abuse (ForceChangePassword → GenericWrite → ... → Domain Admins) and DCSync, but not the **certificate-services** escalation path that is now the most common way a low-privileged domain user reaches Domain Admin. Certipy was only an arsenal tool plus one shadow-credentials technique; the graph could not walk the ESC chain. This build makes it route **AD CS ESC1–8** — ingest certipy find -json , add certtemplate / certauthority nodes, and **synthesize composite ESC{n} edges** from a low-priv enrollee to Domain Admins, walked the exact edge-index way the BloodHound orchestrator already works. It is a near-clone of how DCSync + shadow-creds already work in adgraph/ , deliberately.

It adds NO new gate — \$0's binding constraint. The graph **proposes an edge index**, never authors a command (orchestrator.py unchanged; a pick outside the frontier is refused, not repaired; it execs nothing by AST). The ESC abuse COMMAND comes from the deterministic KB-grounded technique catalog, and approval goes to the SAME POST /cockpit/exec every other command uses; advance moves only on a run that was approved and

exited 0, verified server-side. `certipy find` itself runs as a **gated, scope-locked enumeration job** (the bloodhound-collector shape): read-only, so no red-confirm, but still a command against a real DC the human approves. Per-command human approval is the only bound.

The synthesis mirrors DCSync exactly. DCSync is one composite edge emitted when a predicate over two facts holds (`GetChanges` and `GetChangesAll` on the same target). An ESC edge is one composite edge emitted when a predicate over a template's misconfiguration **and** an enrollee's enroll right holds — `parser.ingest_certipy` reads each template's EKUs, `enrollee_supplies_subject`, manager-approval flag and enroll/write ACLs, and each CA's `EDITF_ATTRIBUTESUBJECTALTNAME2` / web-enrollment / ManageCA principals, then collapses the vulnerable ones to a `enrollee --ESC{n}--> Domain Admins` edge carrying `template_name` / `ca_name` / `esc_variant` / `eku` in props. **ESC1/ESC6/ESC8** are the direct wins (ranked high for the tie-break); **ESC4** (write control over a template) and **ESC7** (ManageCA) emit the **two-hop reconfigure-then-abuse** shape, targeting the template/CA node and then chaining to the issue step. ESC2/ESC3 are modeled too; ESC9–11 (weak-mapping / RPC-relay) are catalog-cited and deferred — they have no strong SAN/enroll predicate to synthesize from and are noted here rather than modeled thinly. `PublishedTo` (template → CA) and `CanEnroll` (a template you may enrol but cannot abuse) are **structural context**, not traversed.

Each ESC edge resolves to a real, gated command — and demands the red-confirm. The catalog grounds every edge in a `certipy req → certipy auth` chain (obtain a cert as `administrator`, then recover the NT hash / TGT), with a native `Certify.exe` variant for the on-host CRTP path and `ntlmrelayx --adcs` for ESC8's relay. All are **destructive** — they issue a real certificate or reconfigure the PKI — and the danger heuristic was extended so every one trips it: `certipy`'s abuse subcommands and `ntlmrelayx` already fired, and `Certify.exe request / certipy ... auth` now do too (`certify`'s `request / download` added to `_AD_WRITE_SUBCOMMANDS`, its read-only `find` kept clean like `certipy find`). The **oracle test** — every edge the catalog calls destructive must resolve to a command the heuristic flags, on **both** the Linux and native-Windows transports — passes across the seven new kinds, so no ESC step can run without the explicit confirm.

Certify.exe added to the arsenal for the Windows path. Catalogued in `arsenal/tools.json` (read-only `find` template + the `request / on-behalf-of` abuse templates) and classified `_ARGUMENT_DEPENDENT` in `test_arsenal_safety.py` — the bare binary stays clean, at least one template fires — the same bucket as `certipy` and `rubeus`.

Frontend /cockpit/ad renders the ESC edges + cert nodes in the **existing** graph — new `certtemplate` (amber, 📄) and `certauthority` (violet, 🏰) node styling in the `hp-adg-*` vocabulary, `ESC{n}` edge labels, and a `?demo=esc` deep-link that auto-ingests the synthetic vulnerable-CA sample. **The screen was looked at** (`frontend-class-vocabulary`): the sample renders a real 2-hop route — `HODOR` (owned) — `ESC4` → `VulnTemplate` (a certificate-template node) — `ESC1` → `DOMAIN ADMINs` — with no AD lab (see screenshot 36).

Verification. `test_adcs_graph.py` (`certipy` → cert nodes + composite ESC edges with props; a low-priv enrollee reaches DA via ESC1; ESC4/ESC7 produce the two-hop reconfigure-then-abuse shape; every ESC edge is runnable on Linux **and** Windows; a non-vulnerable template is `CanEnroll` context; the BloodHound-only graph is byte-unchanged) and the extended `test_adorch_safety.py` (a proposed ESC step is refused unapproved and demands the red-confirm; the `certipy ingest` + technique catalog execute nothing by source-scan; `certipy / Certify.exe find` stay clean) both pass, wired into `run_safety_tests.sh` — the hermetic suite stands at **103 test files, all green**. `next build` exits 0 with `/cockpit/ad` in the route table.

Assumptions (per \$5), stated. The primary enum source is `certipy find -json`; a BloodHound-CE ADCS ingest would fold in the same way but is not wired. Synthesized ESC edges target the Domain Admins group (RID 512) when collected, else the domain node — reaching it is the win, matching the existing objective model. ESC9–11 are deferred (catalog-cited), as noted above.

Unconstrained delegation edge + golden/silver ticket forging — the delegation-family closer (:ad-graph) (2026-08-07)

The AD graph modeled RBCD (`AddAllowedToAct` / `AllowedToAct`) and constrained/S4U delegation (`AllowedToDelegate`) as full routable edges, but the third Kerberos-delegation primitive — **unconstrained delegation** — was not routable, and ticket forging existed only as a mimikatz arsenal capability with no technique node. This build closes both. It adds the `TrustedForDelegation` abusable edge and adds **golden / silver ticket forging** as post-compromise persistence. **No new gate; per-command human approval is the only bound** — the guiding constraint held.

Unconstrained delegation is now a routable edge, synthesized from the flag exactly like RBCD. BloodHound emits `unconstraineddelegation: true` on a computer/user; `parser._synthesize_unconstrained_delegation` folds that into one composite `host --TrustedForDelegation--> Domain Admins` edge — the same shape as the DCSync synthesis (a composite abuse edge emitted when a predicate over collected facts holds) and the ESC synthesis before it. The edge encodes the win (own the host, coerce a DC to authenticate to it, capture its TGT, DCSync $\Rightarrow$ krbtgt $\Rightarrow$ full compromise); owning the host first is supplied by the BFS via the inbound `AdminTo` edge, so the demo routes `PODRICK --AdminTo--> APP01 --TrustedForDelegation--> DOMAIN ADMINs`. It targets the Domain Admins objective the path engine already picks (RID 512), matching the ESC objective model. A DC — unconstrained by design, already domain-controlling — is flagged `is_dc` rather than treated as a novel route.

The technique demands the red-confirm on both transports, and the oracle proves it. The catalog entry is `destructive`, with a Linux `krbrelayx + printerbug  $\rightarrow$  secretsdump -k chain` and a native `Rubeus monitor + SpoolSample  $\rightarrow$  mimikatz lsadump::dcsync` variant for the CRTP on-host path. The AD-orchestration oracle (`test_adorch_safety.py`) requires every edge the catalog calls `destructive` to resolve to a command the danger heuristic flags, on **both** Linux and Windows. The Linux first-runnable line (`krbrelayx`) already tripped as a credential-capture tool; the Windows first-runnable line (`Rubeus.exe monitor`) did **not** until `monitor` was added to the rubeus subcommand map — `Rubeus monitor` harvests the coerced DC's TGT, so it mints/captures a credential, and `SpoolSample` (the .NET printerbug) joined the coercion set beside `printerbug/PetitPotam`. `Rubeus find / triage` stay clean, like `certipy find`, so the confirm keeps its meaning.

Golden and silver forging are persistence, deliberately NOT routing edges. Forging a ticket presupposes you already hold the secret it would otherwise be used to obtain — `krbtgt` for golden, a service hash for silver — so an edge for it would let the path engine "reach DA" by assuming the very thing the route exists to achieve. The two live in a separate `adgraph/persistence.py` catalog keyed by node type, surfaced by a read-only `POST /cockpit/ad/persistence` endpoint that only OFFERS an action once its secret is held: **golden** on the domain node once `krbtgt` is held (a traversed DCSync / a captured DC TGT / the objective owned), **silver** on each owned computer/service node. They are never in `schema.ABUSABLE_EDGES`, never a graph edge, and never enter the orchestrator frontier — regression-locked in `test_deleg_tickets.py` and `test_adorch_safety.py` even with every node owned. Each renders both an impacket (`ticketer`) and a mimikatz command; both trip the danger gate. (The lighter of the two spec options — a persistence catalog rather than non-traversable self-edges — was chosen because `Graph.add_edge` drops self-loops outright, which would have made self-edges the heavier, more invasive path.)

Frontend `/cockpit/ad` renders `TrustedForDelegation` as a normal abusable edge on the route, and a **distinct, violet-keyed persistence panel** (`CockpitADPersistence`, new `hp-adg-persist-*` vocabulary) below it — golden/silver cards with a GOLDEN/SILVER badge, the requirement, and both commands, styled so they never read as a routing step. A `?demo=deleg` deep-link auto-ingests the synthetic member-server sample and pre-owns the

host + objective so the panel renders the forging actions unlocked. **The screen was looked at** (`frontend-class-vocabulary`): the sample renders a real 2-hop route — `PODRICK` (owned) — `AdminTo` → `APP01` — `TrustedForDelegation` → `DOMAIN ADMIN`s — with the persistence panel below, no AD lab (see screenshot 37).

Verification. `test_deleg_tickets.py` (the flag synthesizes a routable edge and a low-priv admin routes `AdminTo` → `TrustedForDelegation` → `DA` ; the technique carries both transports, both destructive and both tripping the gate; `krbrelayx/printerbug/PetitPotam/SpoolSample/Rubeus-monitor` all trip while `Rubeus-triage` stays clean; golden gated on `krbtgt`-held and silver on a held service hash; neither forging kind is ever an abusable edge, a graph edge, or a frontier candidate) plus the extended `test_adorch_safety.py` (a proposed delegation step refused unapproved and demanding the red-confirm even when approved; forging never in the route-to-`DA` search; `adgraph/persistence.py` executes nothing by source-scan). Added to `run_safety_tests.sh` .

Assumptions (per §5), stated. Unconstrained delegation is the must-have (the missing routable edge); golden/silver are persistence, explicitly not routing edges, because a route-to-`DA` graph should not search through post-`DA` persistence. Coercion tooling (`petitpotam/printerbug/coercer/dfscoerce/shadowcoerce`) was already gated in `allowlist.py` ; no allowlist policy change beyond classifying `Rubeus monitor` and `SpoolSample` . No new tools — `mimikatz`, `Rubeus`, `impacket` (`ticketer`), `printerbug` and `PetitPotam` are all already present.

AI code-audit fan-out — open·kritt's decomposition, HackPit-gated (`:code-scan` AI mode) (2026-08-07)

The rule scan pattern-matches file-by-file; it cannot reason about whether an attacker-controlled value actually reaches a dangerous sink. This build adds the other half: an **AI-agent audit** that borrows open·kritt's context-saving decomposition (the engine the operator owns — the good parts ported, the autonomous-root model left behind) and runs it on HackPit's `reasoning/` substrate, **human-gated the whole way**.

The three-stage decomposition (`codescan/ai_audit.py`). (1) **Enumerate entrypoints** — one pass over the repo's file listing + a few entry files maps the externally-reachable entrypoints (HTTP routes, RPC/GraphQL handlers, consumers, CLI). (2) **Trace flows** — per entrypoint, the materially-different production paths (validation outcomes, authz boundaries, state changes, external calls, sensitive sinks). (3) **Verify each flow** — the fan-out: **one agent, one flow**, its whole context spent on a single path, returning either a concrete vuln (title, attacker path, `file:line` source refs, impact, a propose-only PoC) **or an honest no-finding stub**. Mapping once is the whole point: context cost is linear in the number of flows, and findings come back attacker-path-backed rather than repo-wide hand-waving. The concrete-or-stub gate (`gate_finding`) downranks any claimed finding that lacks a concrete source location to a stub — the critic a source audit actually needs (`reasoning.critic` 's CVE-vs-observed-version check does not apply to a source read, so a thin gate replaces it; stated per spec §5). Findings are then **deduped** (same bug found via more than one flow collapses) and **severity-ranked** by open·kritt's `IMPACT_LEVELS` .

Reuses the substrate, not a new engine. Domain framing comes from `reasoning.specialists` ; KB grounding from the injected KB search (`reasoning.retrieval` in spirit — each specialist grounded in the fingerprint/CVE/methodology corpus, the edge HackPit has and open·kritt does not); model-tier selection from `reasoning.tiering` (the hard verify step can be pointed at a stronger model). The finding output schema + `validate_payload` are ported from open·kritt's `schema.py` (zero-dependency — HackPit does not pull in `jsonschema`). `patched-since` (ported from `prompting.py`) restricts the whole audit to the files changed since a git ref — a huge repo becomes a reviewable delta.

No new gate — \$0, held. The audit reads source and **PROPOSES**: it walks files, reads their text, and calls the LLM layer (an *injected* agent runner bound to `backend/llm.py`). It launches no scanner against a target, opens no socket, and imports no executor/engagement/sandbox/state module — so it is **one approved job** (the ZAP/nuclei justification, one approval buys the whole fan-out), the exact category the rule scan already occupies. `ai_audit.py` and the AI-audit routes make **no** `eval / exec / subprocess / os.system / socket / HTTP call by AST` (with a control that plants one, in `test_ai_audit_safety.py`). Any PoC a finding offers is a **string**, run **approve-each** through the existing executor in the :kali sandbox — never from here. The two cross-cutting seams codescan cannot own are **injected from** `main.py`: the engagement-state **findings sink** (so codescan never imports `state`) and the **patched-since git-diff provider** (so codescan runs no `subprocess` of its own — the static-only lock still reads exactly one asserted spawn, `runner._spawn`'s `semgrep/bandit`). When no LLM is reachable the audit degrades to a **deterministic heuristic analyst** that runs the same three stages over sink patterns loaded from `ai_audit_rules.json` (kept as data, not Python, so the detector's own tokens never trip codescan's banned-literal scan) — this is also the offline demo the `/code-scan` screenshot renders on a bundled synthetic sample repo (six routes → six enclosing-route findings, ranked critical→high, KB-grounded).

Frontend. `/code-scan` gains a mode toggle: the AI audit picks a repo (+ optional `patched-since` ref), one approval fans out, and the result renders the entrypoint map, the per-flow fan-out, and the deduped severity-ranked findings — each with its attacker path, `file:line` refs, KB technique links, and a **Build PoC (approve-each)** disclosure that shows the proposed command with the "nothing here runs it" note (`hp-cs-*` vocabulary, verified in `test_css_vocabulary.py`).

Tests. `test_ai_audit.py` (the three stages compose; a non-concrete claim becomes a stub; dedup collapses duplicates; ranking is `IMPACT_LEVELS` worst-first; `informational` maps to state `info`; `patched-since` audits only the diff and an empty diff scans nothing and warns; the heuristic analyst maps six sample routes to six enclosing-route findings) and `test_ai_audit_safety.py` (the \$0 invariants above, each with a control). Both added to `run_safety_tests.sh` — the hermetic suite is green (108 files). `next build` exits 0.

Assumptions (per \$5), stated. Reused `reasoning/` (specialists/retrieval/tiering) as the fan-out substrate rather than a new engine; the flow frontier is modeled in-process rather than forced through `reasoning/frontier.py`, whose SQLite lead-queue is shaped for executable leads (a lead with an empty command is dropped) not source flows — noted here rather than bent. Shipped the `external-flow-analysis` playbook as the built-in (the web3 playbooks come with the web3 spec). Agent runs use `backend/llm.py` (Codex-login/OpenAI/Anthropic/OpenRouter), not open-kritt's harness.

Finding pipeline — dynamic schema · auto-dedup · pluggable rankers · post-scripts (cross-cutting) (2026-08-07)

HackPit already had findings, validation gates and a report-writer; what every producer lacked was a **common spine**. The AI code-audit above ports open-kritt's decomposition; this build ports the other half of open-kritt — its **finding-processing machinery** — and makes it **cross-cutting**, strengthening findings from *every* surface (recon, nuclei, AD, cloud, the SSRF→IMDS bridge, the manual paste box), not one producer. It lives in a new **pure-data** package `backend/findings/` that imports no cockpit/executor/engagement/sandbox/state module and executes nothing.

Dynamic / structured schema (`findings/schema.py`, `open-kritt's schema.py`). A configurable, jsonschema-free finding shape: `FIELD_TYPE_MAP` (string/severity/cvss/refs/number/bool/map with never-raising coercers), `normalize_output_format` (accepts `name:type!` shorthand or dicts — so an engagement can **define custom fields at runtime**), `output_schema` (generates the schema descriptor), `validate_payload` (type + required +

enum, ported from open-kritt), and `coerce_finding` (declared fields coerced to type; **unknown keys preserved under extra** — the dynamic part). The base field set adds `attacker_path`, `source_refs`, `cvss` and `vuln_class` to the finding — the concreteness a report actually needs.

Automatic de-duplication (`findings/pipeline.py`, open-kritt's `post_processing.py`). Fan-out producers report the same defect more than once. A **stable dedup key** — normalized title *or*, when both are present, the concrete **location + type** — collapses two wordings of the same bug (or two tools' takes on it) into one finding that keeps the **worst severity** and the **union of source-refs**. It is **idempotent** three ways: exact re-ingest is caught by the existing fingerprint upsert; the fuzzy collapse carries a `merged_count` that is *added* from the inputs, never re-inflated; and re-running the pipeline over its own output is a fixed point. The result surfaces a "**merged N duplicates**" note. `IMPACT_LEVELS` (open-kritt's) drives the worst-first sort.

Pluggable severity rankers (`findings/rankers.py`, open-kritt's `severityRankers.js` + `defaultSeverityRankers.js`). A **ranker** is an ordered rule set that rescopes a finding's severity for the operator's engagement. Ships three: `default` (keep the producer's severity), `bug-bounty-payout` (RCE / loss-of-funds / auth-bypass to critical, exploitable web classes high, best-practice noise to info), and `compliance` (crypto & header control gaps rise to medium, raw exploitation criticals capped at high). The last two are the demonstration that this is *pluggable, not cosmetic*: the missing-header one lens discards to info is a real medium control gap in the other — same findings, different lens, selectable per engagement (persisted in `state_finding_pipeline`).

Post-scripts (`findings/postscripts.py`, open-kritt's `postScripts.js` + `postScriptLocks.js`). An operator step that runs **after a finding lands**: **validate** (re-check it is actionable — composes with the existing validation gates) and **report** (draft a Markdown writeup for report-writer) run **in-process and execute nothing**; **PoC** (`poc-curl`, `poc-nuclei-retest`) builds a command and returns it as an **approve-each proposal** (`needs_approval: true`, `executed: false`) that the operator fires through the gated executor + :kali sandbox — exactly the drafted-web-exploit pattern, never auto-run. A `PostScriptLocks` table refuses a concurrent double-run of the same script on the same finding.

\$0 held — no new gate. Ranking, dedup and schema are pure data operations. The only thing that can execute is a command post-script, and it never does from here — it returns a string the operator approves-each. The coupling (dict → `state.models.Finding`, command post-script → gated executor route) lives in `main.py`, so the package stays orthogonal, mirroring the codescan sink. The **3-schema-places rule** was respected: the structured fields were wired through the `Finding` dataclass, migration-safe `store.py` columns (`attacker_path` / `source_refs` / `cvss` / `vuln_class` / `extra` / `merged_count` / `ranker`), and the frontend — with a round-trip test proving neither the `/sessions/{id}/state` route nor the new pipeline route strips them under any `response_model`.

Surfaces & routes. `GET /findings/rankers` · `GET /findings/postscripts` · `GET /findings/schema` (the base field set) · `POST /sessions/{id}/findings/pipeline` (dedup + rank an engagement's findings; `persist: true` collapses absorbed duplicates and rescores in place — a pure data op that never grows the count) · `GET /findings/pipeline/sample` (a deterministic **synthetic** demo) · `POST /sessions/{id}/findings/postscript` (run a post-script over a stored or inline finding). The frontend adds a **finding-pipeline panel** to `/engagements` — a ranker picker, the merged badges, severity roll-up, and per-finding post-scripts (the two PoC buttons tagged **approve-each**), rendered over the synthetic sample so the screenshot uses no real target.

Tests. `test_finding_pipeline.py` (schema validates + rejects + is dynamic; dedup is idempotent — same finding twice → one "merged 1", re-run a fixed point; the two default rankers rescore the same fixture differently and in order; data post-scripts run in-process; the PoC is approve-each; the lock refuses a double-run) and `test_finding_pipeline_safety.py` (the package executes nothing by AST with a **planted control**; imports no

attack-surface module and names no gate symbol; a command post-script never auto-fires end to end; the structured fields survive both routes; persisting the pipeline collapses duplicates and stays idempotent). Both are in `run_safety_tests.sh` (now 108 hermetic files, all exit 0). `next build` exits 0.

Assumptions (per §5), stated. Shipped schema + dedup + rankers **and** post-scripts in one session (no short-session cut). The `/engagements` panel runs the pipeline over a **synthetic** sample rather than a live engagement's findings, because `/engagements` is a cross-session list — the same machinery runs over real findings from within an engagement via `POST /sessions/{id}/findings/pipeline`. The fuzzy collapse persists destructively only on `persist: true` (the default is a non-destructive computed view), keeping the operator in control of when a merge is written.

Web3 / smart-contract audit — three playbooks on the fan-out (2026-08-07)

The AI code-audit above shipped the `external-flow-analysis` playbook and noted the web3 playbooks would follow. This build adds them: real **smart-contract audit** capability, given HackPit had **no** web3 static-analysis tooling at all.

Playbooks as a decomposition, not a new engine (`codescan/ai_audit.py`). `playbook` was a pass-through string; it is now a `Playbook` (label, language extensions, chain, heuristic-rules group, and a framing fragment per stage). It steers the SAME three-stage fan-out (enumerate → trace → verify) at a domain: it **appends a domain fragment** to each stage prompt on the LLM path, **scopes the mapped file extensions** to one language (a Solidity playbook maps only `.sol`), and **selects a language-specific heuristic sink group** on the no-LLM path. Three web3 playbooks, ported from open-kritt's proven decompositions (the Blockian team's \$1.5M-in-bounties patterns):

- `evm-external-flow` (Solidity) — external/public functions as entrypoints → flows over value transfers, state changes, external calls, oracle reads, access-control branches → **reentrancy** (external call before state update), missing/incorrect **access control** (the sibling-modifier rule), **oracle manipulation** (missing staleness / flash-loan-manipulable spot price), unchecked arithmetic, delegatecall/selfdestruct hijack. The loss-of-funds classes.
- `cosmos-abci-halt` (Go/Cosmos-SDK) — the wired ABCI methods (`BeginBlock` / `EndBlock` / `DeliverTx` / `ProcessProposal` /...) as entrypoints → the **four panic classes** as flows (explicit `panic`, `sdk.Int.Sub` underflow / `Quo` div-zero, `Must*` helpers, slice-index / type-assertion) → keep only panics that are **attacker-triggerable AND production-reachable inside consensus** — a chain halt, not a caught error.
- `anchor-solana` (Rust/Anchor) — instruction handlers as entrypoints (Accounts structs mapped too) → account validation, signer, CPI, arithmetic → **missing-owner-check** (`UncheckedAccount` / `AccountInfo` where a typed `Account<T>` was meant), **signer-spoof**, **integer-overflow**, **CPI-confusion**.

Findings now carry **chain / contract / function** (added through the `Verdict`, the finding schema, and the frontend). The heuristic sink patterns live in `ai_audit_web3_rules.json` as **DATA, not .py** — the same reason as the web-app rules: `codescan/*.py` is itself source-scanned for dangerous literals, and a detector hard-coding `delegatecall` / `panic()` / `invoke_signed` would trip its own static-only lock. Sinks are co-designed with a bundled deliberately- vulnerable **fixture set** (`sample_web3/{evm,cosmos,anchor}`) so the deterministic analyst returns concrete, attacker-path-backed findings with no LLM — the offline demo and the screenshot.

Tool pass — propose-only (`codescan/web3_tools.py`). A smart-contract audit wants real tools; the engine still must execute nothing. So this module **builds command strings** (`slither <file> --json -`, `myth analyze ...`, `echidna ... --format json`, `forge test ...`) the operator runs **approve-each** through the gated executor + :kali sandbox, and **parses** the tools' JSON/text output back into the same normalized finding shape. It launches no

scanner, spawns no process, imports no executor/sandbox — and was added to `test_ai_audit_safety.py` 's **AST no-exec lock** so the "analysis executes nothing, tool runs are approve-each" claim is airtight. `runner._TOOLS` is still exactly `("semgrep", "bandit")` and the package still spawns exactly one program.

KB grounding. The `web3-audit` skill's DeFi/Solidity/Solana-token bug-class content is already in the KB (17 `cbb-* web3` entries); triage found three genuine gaps, each grounding one playbook, authored via the committable `pipeline/authored/` path and ingested: **external-flow analysis** (the map-once/verify-per-flow method), the **Cosmos four-panic-class** consensus-halt review (ZERO before — the KB's web3 was DeFi/Solidity/Solana-token), and the **Anchor account-model** audit. All three rank #1 for their playbook's grounding query; a finding cites them (verified end to end). KB now 2750 entries; the scripts index, embeddings, corpus report and committed KB fixture were regenerated to match.

Tooling (arsenal + image, image rebuild is the operator's step). Added `slither`, `mythril`, `echidna`, `foundry` (forge+cast), `semgrep`, `anchor` / `cargo` / `clippy`, `gosec` / `go` to `arsenal/tools.json` (a new `web3` category, 7 tools / 131 total; classified `_MUST_NOT_FIRE` in the arsenal danger registry — local static-analysis / local test harnesses, the same class as `ghidra` / `gdb`), to `docker/Dockerfile.sandbox` (venvs for the Python analyzers, release binaries for `echidna/foundry/gosec`, `rustup+clippy+anchor`), and a `docker/proof/web3_install_proof.sh` that checks the names the catalog + `web3_tools.py` hardcode resolve in the built image. **The image rebuild** (`docker compose build engage-sandbox`) is the operator's step, flagged here.

Frontend (/code-scan). The AI-audit view gains a **playbook picker** (from `/codescan/playbooks`), a web3-aware "contract folder" field, a **chain badge** + chain/contract/function chips on findings, and a **propose-only tool-pass panel** (`Propose tool pass` → the `slither/mythril/echidna` commands, approve-each). A `?demo=web3` deep link loads the bundled EVM fixture sample so the headless screenshot renders with no input. Screenshot: `assets/screenshots/40-code-scan-web3.png`.

Tests. `test_web3_audit.py` (the three playbooks register + language-scope; each fans out over its fixture into the expected concrete findings-or-stubs tagged chain/contract/function — evm surfaces reentrancy/access-control/oracle, cosmos maps the four panic classes to real ABCI methods as consensus-halts, anchor surfaces missing-owner/signer-spoof/ overflow/CPI-confusion; a `slither/mythril/echidna` fixture output parses into findings; the tool pass is propose-only; KB grounding cites a web3 entry; the router surfaces playbooks/sample/tool-pass). In `run_safety_tests.sh` (109 hermetic files, all exit 0). `next build` exits 0.

Assumptions (per §5), stated. Shipped all three playbooks (not EVM-only) plus tooling + KB in one session. The playbooks run on the already-built code-audit-fanout engine, so nothing was stubbed. The tool pass parses `slither` / `mythril` / `echidna` structured output; Cosmos/Solana tools (`gosec` / `clippy`) are proposed as commands without a structured parser (they return text a human reads).

Formal engagement governance — RoE / ConOps / Deconfliction / OPPLAN (2026-08-07)

The idea. HackPit's bound on a real-target run has always been *human approval of every command* — the wall. What it lacked was the written frame a professional engagement approves *inside*. This build adds that frame: before an engagement goes live, the operator drafts and approves four governance documents, and the OPPLAN's objectives become the thing the orchestrator proposes *toward*. It turns "*human approves each command*" into "*human approves each command inside a written, agreed operating frame*." Ported from **Deception** (Apache-2.0 — attribution in `NOTICE` / `THIRD_PARTY_LICENSES`): the `Objective` / `OPPLAN` data model, `ObjectivePhase` / `ObjectiveStatus` / `C2Tier` / `OpsecLevel`, the objective **status state machine**, the ConOps generator, the RoE structure, and the ATT&CK `killchain.yaml` reference.

The four documents (`backend/state/governance.py`). Each persists as a **versioned JSON record**, engagement-scoped by `session_id` in the shared `sessions.db` (upsert-only, like the rest of `state/`). Every save bumps the version and **resets approval** — an edited frame must be re-approved, because the thing the human signed off changed. **RoE**: authorized scope (references the scope model, never replaces it), authorized/forbidden techniques, OPSEC level, time windows, excluded targets/actions, sensitive-data handling, stop conditions, emergency contacts. **ConOps**: approach + phases (recon → exploitation → post-exploitation → actions-on-objectives) with success criteria. **Deconfliction**: a per-engagement signature/tag, source markers, notification contacts, traffic identification, blue-team notes. **OPPLAN**: the objectives, each with a `phase` , a `status` (the state machine), MITRE ATT&CK technique id(s), an OPSEC level and an optional C2 tier.

The state machine (ported wholesale — then reconciled against the real Decepticon source).

`_VALID_TRANSITIONS` : `pending` → `{in-progress, cancelled}` , `in-progress` → `{completed, blocked, cancelled}` , `blocked` → `{in-progress, completed, cancelled}` (retry / abandon-but-done / drop); **completed** and **cancelled** are terminal. Note there is deliberately **no** `pending` → `blocked` edge (a pending objective has not been attempted, so it cannot be blocked) and **blocked** → **completed** **IS legal**. The build's first cut had those two edges wrong; when the actual `tools/opplan.py` became available it was diffed and the table corrected to Decepticon's verbatim (locked in `test_governance.py`). `update_objective` is the only place a status changes and it validates against the table — an illegal transition raises `TransitionError` , **nothing is written**, and the OPPLAN version does not move. Objectives are a dotted-id tree (`1` , `1.2`) so `expand` adds sub-objectives and `collapse` removes them, mirroring `opplan.py` 's tool set (add / update / get / list / expand / collapse). An approved, exit-0 run is recorded as an objective's `evidence_run_id` — the same `advance` -evidence model the AD/cloud graphs use.

Enum values are Decepticon's, verbatim. `ObjectiveStatus` = `pending` / `in-progress` / `completed` / `blocked` / `cancelled`; `ObjectivePhase` = `recon` / `initial-access` / `post-exploit` / `c2` / `exfiltration`; `OpsecLevel` = `loud` / `standard` / `careful` / `quiet` / `silent`; `C2Tier` = `interactive` / `short-haul` / `long-haul` (plus a HackPit `none` for the spec's "optional C2 tier", the default). The reference `killchain.yaml` maps its 14 ATT&CK tactics onto those five kill-chain phases. (Two source files the spec named turned out to differ from its description: Decepticon's `conops.py` is a SIEM target-resolver, not the ConOps generator — HackPit's ConOps doc shape is its own; and `roe.py` is a middleware that in `enforce` mode *can* machine-veto, which HackPit deliberately does **not** adopt — its own docstring calls the parse layer "a guardrail, not the sole gate", matching HackPit's handrail framing.)

MITRE ATT&CK coverage (`backend/state/killchain.py` + `references/killchain.yaml`). The reference maps 14 ATT&CK tactics (Reconnaissance → Impact) to representative techniques, each tied to a ConOps phase. `attack_coverage` answers "*which techniques did this engagement's objectives exercise*" as a per-tactic grid + roll-up counts — a lexical id match, no ATT&CK API call at render time. An objective mapped to a technique the reference does not know is counted as `unmapped` , never dropped. The coverage view is a professional deliverable and a report input.

Propose-only generation (`backend/governance_draft.py`). The generative layer, like `attack_path.py` : from the scope + target it drafts all four documents via `llm.chat` for the human to edit and approve. Nothing is "live" until approved. It **degrades to a deterministic, scope-derived skeleton** when no LLM is reachable, so the operator always gets an editable starting point. It persists nothing and advances no objective — those are the human-driven route's job.

The RoE is a **FRAME**, not a veto. The RoE **formalises** the scope handrail; it does not replace it and it is **not a machine veto**. The RoE-vs-scope check (`main._roe_scope_advisory`) parses the RoE's declared scope with the same `cockpit/scope.py` and flags an *invalid, undeclared, unbounded* (`*`), or *mismatched-with-the-live-handrail*

scope in the UI — **advisory only, it never blocks a command or an objective**. `state/governance.py` imports nothing from `cockpit` (the scope import lives in the app layer, `main.py`, the one place allowed), so the executes-nothing/no-gate property holds by construction.

Routes + frontend. Ten routes in `main.py` (get the package; draft / save / approve each doc; objective CRUD + expand / collapse / delete; seed an OPPLAN from a draft; the ATT&CK-technique picker) — the save route is **PATCH not PUT** (the CORS allow-list is GET/POST/PATCH/DELETE; a PUT preflight fails, the lesson `set_submission` already learned, which the `/attack-path` response-contract test caught here on the first suite run). `/engagement/[id]` gains a **Governance** view (`GovernancePanel.tsx`) with RoE / ConOps / Deconfliction / OPPLAN / ATT&CK tabs, an **objectives board** with status columns (each card carrying phase + OPSEC + ATT&CK chips and only the *legal* next-state transition buttons), a scope advisory banner, and an **ATT&CK coverage matrix**. A `?view=governance&tab=<tab>` deep-link opens it directly (used for the headless screenshot). The governance package also flows into the generated report as a Markdown appendix (approved docs + an objectives table + the coverage roll-up), appended after composition so `report_gen` is unchanged and an engagement without governance yields a byte-identical report.

No new gate — as designed. The whole subsystem is authored + human-approved documentation plus a formalised scope frame. It executes nothing an engagement acts on, and per-command human approval remains the actual bound.

Tests. `test_governance.py` (the state machine rejects illegal transitions — terminal exits + the pending→completed skip — and a rejected transition writes nothing / documents version and reset approval on edit, a v0 doc cannot be approved / objective add + expand-to-dotted-children + collapse + delete-with-descendants / technique ids cleaned not invented / ATT&CK coverage renders and counts unmapped ids / the RoE-vs-scope check is advisory and objectives keep mutating). `test_governance_safety.py` (§0: governance + killchain + the drafter make no eval/exec/subprocess/socket/HTTP call by AST with a control; no cockpit/executor/sandbox import and no gate symbol; the drafter is propose-only by AST; the state machine governs objective status only and the advisory check never raises). `state/governance.py` + `killchain.py` were added to `test_state.py` 's executes-nothing scan. `run_safety_tests.sh` : 111 hermetic files, all exit 0. **next build exits 0.** Screens LOOKED AT: `41-engagement-governance.png` (the objectives board) + `42-attack-coverage.png` (the ATT&CK matrix), both on a synthetic `acme-demo` engagement — no real client/scope.

Assumptions (per §5), stated. Shipped all four documents + objectives + ATT&CK + the frontend + the report feed in one session (the spec's fallback was OPPLAN-first). Decepticon's source was not present in the tree, so the data model + state machine were **reconstructed from the spec's description** rather than copied file-for-file; the attribution stands regardless (`NOTICE` / `THIRD_PARTY_LICENSES`). Document editing in the UI is currently draft-then-approve (the drafter proposes the whole body; per-field inline editing is a follow-up); the substantive fields render read-structured.

Reusable prompt-workflow builder — author your own playbooks over the fan-out (`:workflows`) (2026-08-07)

The idea. The AI code-audit fan-out ships two hard-coded decompositions (the web-app playbook and the three web3 ones). The **workflow builder** is the authoring layer on top of that engine: an operator composes their *own* ordered set of prompt **steps** — each with variables, an output schema and a batch / depth / siblings fan-out shape — saves it, exports/imports it as a portable JSON, and runs it. Ported from **open-kritt**'s workflow builder (`workflows.js` / `steps.js` / `defaultWorkflows.js` and the `{{var}}` / `resolve_ref` / `render_prompt` variable model in `prompting.py`), adapted to HackPit's stack. It is the most product-polish of the Kritt-derived specs — the built-ins already deliver the core value — so it deliberately adds **nothing to the trust surface**.

The model (backend/codescan/workflows.py , a pure module). A Workflow is an ordered list of Steps; a Step is a focused prompt + a kind (analyze / batch / command) + fan-out controls + an optional output_format schema. The variable model is ported verbatim: resolve_ref(context, "a.b.0.c") walks a nested dict/list and returns None on any missing hop (never raises); render_prompt substitutes every {{ ref }} and renders an unresolved ref as the empty string — a step never emits a literal {{x}} to an agent. Built-in variables are repo , ref , playbook , the per-item item , the sibling branch , and any prior step's parsed output by dotted ref (steps.<id>.output); operator-defined and per-run extra variables land in the same flat namespace (and under extra.*). Batches fan out one agent per item of a list variable; siblings run parallel branches per task; depth re-expands a batch step's list output into child generations. Every knob is bounded — MAX_STEPS=24 , MAX_SIBLINGS=8 , MAX_DEPTH=4 , MAX_BATCH_ITEMS=64 , and a total MAX_TASKS=400 ceiling decremented across the whole run — so a fan-out can never run away. A batch step that references a later step is refused at build time.

Import / export, inspect-before-run. to_portable serialises a workflow to {"kritt_workflow_schema": 1, "workflow": {...}} (dropping the store-local version/updated_at so a re-import is a fresh authored copy); from_portable parses it back, flagged imported=True , and rejects a wrong/absent schema version loudly. The store's import_workflow stores it and returns it — it never runs it; the id is suffixed on collision so an import can never overwrite an existing workflow. Two proven built-ins are seeded from code and visible on first load: external-flow (web-app, mirrors the external-flow-analysis playbook) and evm-external-flow (Solidity), each a 3-step enumerate → trace (batch) → verify (batch) decomposition.

The runner — compile onto the same engine, execute nothing new. run_workflow renders each step's prompt with the resolved variables, runs it (single, batched over a list variable, or — for a command step — a proposal only), validates each output against the step's schema (a declared output_format , else the audit's concrete-or-stub finding schema), and publishes the step's outputs into steps.<id>.output so downstream steps reference them. Concrete findings across all steps are de-duplicated and severity-ranked by the shared ai_audit.dedup_and_rank pass. The runner reuses the audit's injected LLM agent (default_agents) and KB search; it imports no state / git / executor of its own. A command step never calls the agent — its rendered string comes back under proposals with executed:false , approve_each:true , run one-by-one through the existing gated executor in the :kali sandbox. plan() computes the static fan-out shape (steps × items × siblings × depth) with no run at all — the builder's preview.

Routes + store (codescan/router.py). CRUD (GET/POST /workflows , GET/PATCH/DELETE /workflows/{id} — PATCH not PUT, the CORS lesson governance already learned), POST /workflows/import (inspect-before-run), GET .../export , POST .../plan , POST .../run , and GET .../sample (a deterministic no-LLM run of a built-in over its bundled fixture — the offline demo). The store is one JSON file under backend/data/ (gitignored runtime data; the built-ins always seed from code) with an optimistic version lock (open-kritt's workflowLocks): a save carrying a stale version is refused, so two editors cannot silently clobber each other, and a built-in is read-only (clone to edit).

No new gate — by construction. Authoring (create / edit / import / export / delete) touches no agent and no target; running is the same one approved job the AI audit is (the fan-out justification), reaching the same seam with no new approval / danger / target field. codescan/workflows.py + the routes make no eval/exec/subprocess/socket/HTTP call by AST, import no cockpit/executor/state/sandbox and name no gate symbol — the only launchable program in the whole codescan package is still runner._spawn 's semgrep/bandit, exactly once. The wider test_codescan_safety.py source-scan (which sweeps every codescan/*.py) passes unchanged with the new file in the package.

Frontend (`/workflows` , in the **PLAN** band — "**proposes · never executes**"). A two-column builder: the workflow list + an inspect-before-run import box on the left; the step editor on the right — name / playbook / description, a **variable palette** that inserts `{{var}}` (built-ins + each prior step's dotted output ref) into the focused step's prompt, then per-step prompt editors with a kind selector, the batch/depth/siblings fan-out controls, a KB-ground toggle and an **output-schema editor**. A fan-out **plan preview**, **JSON export/import**, and a **run** panel (repo path + session + auto/heuristic mode, plus "Load sample" for the built-ins) that shows step task counts, deduped findings and approve-each proposals, with a hand-off to `:code-scan` . New `hp-wf-*` classes added to `globals.css` (the CSS-vocabulary test enforces every class exists).

Tests. `test_workflows.py` (variable resolution incl. dotted step-output refs + extra vars, a miss renders empty / a batch fans out one task per item and siblings multiply, an empty list no-ops / depth=2 over a 2→2 expansion produces the 2/4/8 child shape, all bounded / export→import round-trips with only the provenance flag flipping to inspect-before-run / a step output validates against both the default and a custom schema / the runner threads a step's output downstream and same-bug siblings dedup+rank to one / a command step proposes and calls no agent / `plan()` is static / a forward batch ref is refused). `test_workflows_safety.py` (\$0: authoring executes nothing — proven with a recording agent that fires only on a real run / no eval/exec/subprocess/socket/HTTP by AST with a planted control / no gate/executor/state import or symbol, still exactly one spawn site / a command step is approve-each, executed:false / an imported workflow is never auto-run, with a control / built-ins read-only + the version lock refuses a stale edit). Both registered in `run_safety_tests.sh` . `run_safety_tests.sh` : 113 **hermetic files, all exit 0.** `next build` exits 0. Screen LOOKED AT: `43-workflows.png` (the builder rendering the EVM built-in — variable palette, three steps, the batch/depth/siblings controls and output schema, on the seeded built-ins, no real client source).

Assumptions (per §5), stated. Depends on the code-audit fan-out engine (present) and pairs with the finding-pipeline schema (also present); a step's default output schema reuses the audit's concrete-or-stub finding shape, so the finding-pipeline dependency is satisfied without `codescan` importing the `findings/` package (orthogonality kept). open-kritt's source was not vendored in the tree, so the step/batch/depth-sibling/variable **semantics were reconstructed from the spec's description** and reproduced faithfully; the attribution stands regardless. Depth is modelled as a batch step re-expanding its own list output for `depth` generations (bounded); the "child steps" of the docs are that self-expansion plus downstream batch-over-prior-step chaining.

Interactive persistent-session engine — named tmux sessions + automatic prompt-detection (`:terminal`) (2026-08-07)

The idea. `:kali` gives clean per-command transcripts; `:terminal` 's pty gives one full-screen shell. Neither gives a *named, parallel, persistent* session whose cwd/env/background-jobs survive across calls, and neither knows when a program has stopped at an **interactive prompt** waiting for the operator. The **session engine** is that third surface: named **tmux** sessions in the same open sandbox, with **automatic interactive-prompt detection** (`msfconsole` `msf6 >` , `sliver-client` , `evil-winrm` PS, generic REPLs), a background lifecycle (`background=True` + **auto-background at 60s** with a notify-once completion), output tiering, and wedge / pipe-degradation recovery. Ported from Decepticon's `tools/bash/bash.py` + `tools/bash/prompt.py` (Apache-2.0 — attribution in `NOTICE` / `THIRD_PARTY_LICENSES`).

The HackPit twist — mechanics without autonomy. Decepticon's agent drives everything, including sending a line to an interactive prompt (`is_input`). HackPit adopts the session *engine* and **not that autonomy: every input path is HUMAN-ONLY**. Both `run_command` (send a command) and `send_input` (the `is_input` path: answer a prompt / send a control key) are reachable only from the router the operator's UI drives; the orchestrator / agent /

executor / proposer have **zero** code path to either. This is the one \$0 invariant that cannot be weakened, and it is source-scan locked across the whole backend tree with a planted control (`test_session_engine_safety.py`), exactly like `:kali` / the pty.

The engine (`backend/cockpit/session_engine.py`). Containment mirrors the pty point for point: the container is the hardcoded `config.KALI_OPEN_CONTAINER` (the request models carry no container/target/shell field), there is **no new gate and no isolation gate** (the open box is intentionally not isolated — the human at the keyboard is the approval), the transcript is audited to the run store, and `tmux pipe-pane` mirrors the raw stream to a per-session log under the workspace `.sessions/` — which a **kill deliberately preserves**. Every heuristic is a **pure function** tested against fixtures — `detect_prompt` (idle-shell vs interactive-tool vs running, via a `PROMPT_COMMAND` marker that carries the last exit code and `$PWD`, so each named session tracks its cwd independently), `manage_output` (inline $\leq 15K$ / $> 15K$ saved to `.scratch/` with a head+tail preview / $> 5M$ watchdog-truncated at source), `strip_ansi`, `compress_repeats`, and the three-condition `detect_wedge` and `detect_pipe_degradation` signatures with their operator-facing recovery ladders. The single impure boundary is `_tmux` (one `docker exec ... tmux ...` call), so the whole suite runs with **no Docker and no tmux**. `run_command` sends the command then watches for the completion marker up to 60s, returning `[DONE]` (with the rc) / `[INTERACTIVE]` / `[BACKGROUND]` / `[AUTO-BACKGROUND]`; `poll_jobs` delivers each background completion **exactly once** (mirroring the job/OBSERVE notify-once contract), then `consume_job` clears it.

The surface (`/terminal`). The screen is now two tabs — **named sessions** (the engine) and **raw pty** (the unchanged full-screen surface, embedded so the screen's header renders once). The named-session panel shows session tabs (a live/waiting/dead dot per session, an amber pulsing dot when a tool is awaiting input), the headline **prompt-detection banner** ("*interactive — send input · msfconsole · msf6* >"), per-session cwd/program/state, the live capture, a background-job tracker with notify-once completions, and the operator's input box — which flips between "run a command" (with a *run in background* option) and "answer the prompt..." (`send line`) based on the detected state. Twelve new backend routes under `/cockpit/sessions/*`, all HUMAN-ONLY, none gated (open box, like the pty).

Why this is safe (mirrors the pty exactly, adds nothing to the trust surface). No new gate; no isolation claim; the container is hardcoded and no request field can redirect the exec; input is discrete per-line `send-keys`, so there is no held `proc.stdin` an autonomous caller could reach (unlike a caught C2 shell); the `:kali` sentinel stays **pty-free AND tmux-free** (asserted); and the human-only lock is source-scan enforced with a control. `test_session_engine.py` (14 checks) proves the mechanics against fixtures; `test_session_engine_safety.py` (6 checks) locks the invariants. Both registered in `run_safety_tests.sh`. **`run_safety_tests.sh` : 115 hermetic files, all exit 0. next build exits 0.** Screen LOOKED AT: `44-terminal-named-sessions.png` (the panel rendering a **synthetic/local** msfconsole session at `msf6 exploit(multi/handler) >` with the interactive banner raised, a second running `recon` session, and a notify-once background completion — no real target, no live shell in the shot).

Assumptions (per §5), stated. The primary target is the `:terminal` surface + the human operator, not an autonomous agent (assumption 5.1); prompt-detection covers msf / sliver / evil-winrm / common REPLs first and the heuristic set extends later (5.2); the engine lives on the OPEN/terminal side, never the `:kali` sentinel side, keeping the two-sandbox separation intact (5.3). Decepticon's `tools/bash` source was not vendored in the tree, so the session semantics were **reconstructed from the spec's description** (the same path the governance port took) and reproduced faithfully; the attribution stands regardless. The engine needs `tmux` inside `hackpit-kali-open` (Kali ships it); the availability check refuses cleanly when the sandbox is down, and the hermetic suite fakes the one `_tmux` boundary.

Binary-RE / pwn + forensics / CTF arsenal expansion — the HexStrike coverage gap (2026-08-08)

The gap. The catalog carried a thin `binary` category (`ghidra` / `radare2` / `gdb` / `strings`) and nothing at all for memory forensics, carving or stego — the whole binary-reversing, exploit-dev and forensics/CTF surface a tool like HexStrike exposes. This build fills it as **DATA + TEMPLATES**, gated identically to everything else. It adds **NO new gate and NO execution path** (the `test_arsenal_safety.py` invariants #1/#2 hold): every new tool is a template the planner PROPOSES, and it runs only through the existing approve-each executor / `:kali` / `:terminal`. HexStrike's *autonomous* execution model was deliberately **not** adopted.

The catalog (`backend/arsenal/tools.json`) — +24 tools, 131 → 155 (361 templates). Two new categories.

binary-re (the old `binary` category renamed and expanded): the four existing entries plus `checksec`, `objdump`, `readelf`, `nm`, `xxd`, `nasm`, `patchelf`, `ROPgadget`, `ropper`, `one_gadget`, `pwntools`, `angr`, `libc-database`, `pwninit` — with binary-name aliases where the package name differs (`ghidra` → `analyzeHeadless`, `radare2` → `r2` / `rizin`, `pwntools` → `pwn`). **forensics-ctf**: `volatility3` (`vol` / `vol.py`), `binwalk`, `foremost`, `scalpel`, `steghide`, `zsteg`, `stegseek`, `exiftool`, `bulk_extractor`, `testdisk` / `photorec`. Templates use the existing local-artifact placeholders (`<binary>` / `<file>` / `<output>`) plus three new ones (`<libc>` / `<symbol>` / `<offset>`); none carries a host, and `<lhost>`-style names were avoided by construction.

Classification — 100% coverage held, and honestly. These are local-artifact tools (analyze a `<binary>` / `<file>`), not network-target tools, so the `<target>` / `scope-lock` substitution mostly does not apply — but every one is still pinned. `test_arsenal_safety.py` now covers **247 catalogued invocation names** with a verdict each: **54 must-fire, 177 benign**. All the new reversing/forensics binaries are read-only local analysis (or, for `patchelf` / `binwalk` - `e` / `foremost` / `scalpel` / `bulk_extractor` / `photorec`, a *local file write* — which is not the `_MUST_FIRE` bar of remote/payload/shell), so they sit in `_MUST_NOT_FIRE`. The one exception is exactly the one that earns it: the `pwntools` / `angr` exploit-skeleton templates run `python3 -c '...'`, i.e. arbitrary code, so `python3` reaches the executor as `argv[0]` and is pinned in `_MUST_FIRE` — `allowlist._INTERPRETERS` already fires on it, so **`allowlist.py` needed no change**. `pwntools` / `angr` the *names* stay clean; the interpreter is the tell, not the framework.

Image + proofs (`docker/Dockerfile.sandbox`). New install layers baking the binaries the way each ships — apt (`radare2/rizin/binutils/nasm/patchelf/checksec/ropper/xxd/python3-ropgadget`, and the forensics apt set), a Ghidra layer that discovers `analyzeHeadless` under `support/` and symlinks it onto `PATH`, per-tool venvs for `pwntools` / `angr` / `volatility3` (heavy, conflicting pins) with the catalog invoking each venv's `python3` directly, Ruby gems (`one_gadget` / `zsteg`), a `libc-database` git clone invoked by absolute path (its `find` basename collides with `coreutils`, so it is deliberately NOT on `PATH`), and release binaries for `pwninit` / `stegseek` — each smoke-tested so a wrong name fails the build, not analysis time. `docker/proof/binre_install_proof.sh` + `forensics_install_proof.sh` re-check that every template head the catalog hardcodes (including the absolute venv/repo heads) resolves in the built image and the running engage container. **The `docker compose build engage-sandbox` rebuild is the operator's manual step** (flagged).

Substrate probe. `substrate_probe.py` / `reconcile.py` are catalog-driven, so the new tools are covered automatically; `test_substrate.py` gains an explicit assertion that all 24 `binary-re` / `forensics-ctf` tools are **probed and reported** (present-or-report — a not-yet-baked tool surfaces as `not-installed` with a reason, never silently skipped).

Frontend (`/arsenal`). The browsable arsenal renders the two new categories (`Binary`, `RE & pwn` / `Forensics & CTF`) with the new tools — categories are data-driven, so only the `CATEGORY_LABEL` map changed.

Tests. `test_arsenal.py` + `test_arsenal_safety.py` green at 155 tools / 361 templates / 247 pinned verdicts; `test_substrate.py` green with the new coverage assertion. `run_safety_tests.sh` : 115 hermetic files, all exit 0. `next build` exits 0. Screen LOOKED AT: `40-arsenal-binre-forensics.png` (the `/arsenal` surface filtered to the two new categories, rendering the new tools with their templates — not blank/error).

Assumptions (per §5), stated. Scope = the user's named list plus the obvious siblings (gef/pwndbg noted as gdb plugins not separate binaries; `ropper` alongside `R0Pgadget` ; `zsteg` / `stegseek` for stego; `bulk_extractor`). These are proposable tools, **not** a new surface or a CTF-solver agent — that would be HexStrike's autonomy, and a dedicated `/ctf` surface, if ever wanted, is a separate build. `hashcat` / `john` were already catalogued (cracking) and are not duplicated.

Token workbench — JWT / OAuth / OIDC / SAML analysis, tamper and a gated crack (:tokens) (2026-08-08)

The gap. The proxy could capture a bearer token or a SAML assertion but had nothing to take one apart. This build adds the token parallel to the GraphQL panel: a **pure analysis + tamper core** and one **gated** offline crack, modelled point-for-point on `cockpit/graphql.py` (same §0 — **NO new gate**, per-command approval the only bound).

The pure core (`backend/cockpit/tokens.py`). Imports only

`base64` / `binascii` / `hashlib` / `hmac` / `json` / `re` / `zlib` / `urllib.parse` — the exact banned/allowed set the GraphQL module uses (no `subprocess` / `socket` / `os` / `requests` / `urllib.request` / `http.client`), and it **raises nothing** (warns via a `note` / `ok=false`), both asserted by the safety test. **JWT:** decode → header/claims/verdicts, plus the tamper set — `alg=none` (with case variants), the RS256→HS256 confusion that signs with the pasted PUBLIC key as the HMAC secret (recomputed from scratch), `kid` / `jwk` / `jku` / `x5u` header injection, and edit-and-resign. **OAuth:** parse + attack builders (a `redirect_uri` open-redirect table, `drop_state` , `pkce_downgrade` , `response_mode` , `implicit_leak`). **SAML:** parse + XSW1-8 / strip / comment / unsigned, done by regex/string splice (not ElementTree) to keep the module pure, each variant round-tripping base64→XML.

Names, never values. Two models enforce the same discipline as `GraphQLArgument` : `DecodedToken` is the operator's *pasted* token → full header/claims/signature (their own box, like the repeater body); `TokenDetection` is a token *auto-detected in captured traffic* → header params + claim **NAMES** + the non-secret timing claims (`exp` / `nbf` / `iat`) only, with **no** `claims` / `signature` / `value` field. The safety test plants a secret in a claim value and asserts it never serializes.

The one executing part — a gated crack (`backend/cockpit/tokenjobs.py`). A near-clone of the credential-crack job: `hashcat -m 16500` (JWT), engagement-bound, gate-before-spawn (the `_gate()` call precedes the `Thread`), the token written to a **loot file** (never the argv), the recovered secret to loot (never a Finding), a High finding on success, `stop()` ungated.

Frontend + tests. `TokenPanel` / `TokenCrackPanel` / `TokensScreen` mounted in the proxy screen (5 new `hp-tk-*` classes). `test_tokens.py` + `test_tokens_safety.py` registered in `run_safety_tests.sh` — 117 files green, `next build` exits 0. Screen looked at: `assets/screenshots/45-token-workbench.png` (a decoded JWT with its verdicts + the tamper controls). Committed `7584fd7` .

Single-packet race tester — beat a check-then-act window (:race) (2026-08-08)

The gap. The intruder's serial loop cannot express a race: two requests that must arrive in the *same instant* to slip through a check-then-act window (limit-overflow, TOCTOU, coupon reuse, one-time-token replay). :race fires **one** request **N** times synchronized — **HTTP/2 single-packet** (all N in one TCP segment) or **HTTP/1.1 last-byte** (send all but the final byte, then release the last byte of every request together via a `threading.Barrier`).

Modelled point-for-point on the intruder (backend/cockpit/race.py). ONE gated job through the same four gates (`executor.validate_request`), **no new gate**, the hardcoded `KALI_OPEN_CONTAINER`, **scope-checked on the wire** (one URL for all N → an off-scope host refuses the whole batch, nothing sent), an **ungated stop**. The whole request template **and N** ride in the approved surface (`race_argv_for` declares the complete `-u/-X/--mode/-n/-H/-d`, Critical-2 completeness so a body carrying `| sh` still reaches the danger gate). `race_verdict` is pure + fixture-tested: it clusters the N responses, the winner is the rarest cluster (tie → 2xx), and `race_detected` fires when the serial-baseline's once-only success appears $\geq 2 \rightarrow$ "K of N won" → a **High** Finding upserted to engagement state.

The engine (the one genuinely new piece). `docker/race_singlepacket.py` → the PATH wrapper `race-singlepacket`, with its own venv (`/opt/race/venv`) for the HTTP/2 `h2` framing lib (`h1-last-byte` is pure `stdlib`). It is invoked **argv-only** with the whole request as JSON on **stdin** (`--job-stdin`), so no shell ever parses a request byte — the declared/gated argv (honest, full) is not the spawned argv, exactly the intruder's `ffuf-declared/curl-spawned split`. `--selftest` is the Dockerfile smoke test; **an image rebuild is required** for the new layer (`docker/proof/race_install_proof.sh` re-checks the name, the `h2` import and the `stdin` contract in image + engage container).

Traps + tests. `tools.json` placeholders are a closed set — `<body>` is undeclared, so templates use the declared `<payload>`; every new `argv[0]` needs a danger bucket in `test_arsenal_safety.py` (`race-singlepacket` → `_MUST_NOT_FIRE`, clean by binary identity like `ffuf/sqlmap` — a plain race adds no red-confirm; the heuristic still fires on request *content*). `test_race.py` registered — **118 files green**, next `build` exits 0. Screen looked at: `assets/screenshots/46-race-singlepacket.png`. Committed `c3a137f`.

Request-smuggling / desync detector — timing-differential detect + gated confirm (:smuggle) (2026-08-08)

The gap + the safety split. HTTP request smuggling is high-impact but its confirmation is inherently co-tenant-affecting (you poison a shared front-end→back-end connection). So this build **splits detection from confirmation**: **detection is safe-by-default timing-differential** (a desync makes the back-end wait for bytes that never come → a measurable delay, planting nothing), and **confirmation is a separate approve-each stage** carrying a co-tenant warning (the actual socket-poisoning probe). Modelled on `:race / nuclei.py`; **no new gate**, the same four gates. The mutation family: `CL.TE / TE.CL / CL.CL / TE.TE / CL.0` (HTTP/1, pure `stdlib`) + `H2.CL / H2.TE / H2.0` (best-effort via `h2`, an honest "inconclusive" when the target isn't HTTP/2).

Files. `docker/smuggle_probe.py` (the baked engine `smuggle-probe`, `--selftest` / `--job-stdin`, JSON on `stdin` → `verdicts/confirm`s) · `backend/cockpit/smuggle.py` (pure `argv/spec/verdict-parse` + a gated worker **per stage**) · `router /smuggle/*` (`POST /smuggle` pins `stage=detect`, `POST /smuggle/confirm` pins `stage=confirm`, so an unknown stage falls back to detect and never escalates) · `SmuggleScreen.tsx` + the `/smuggle` page + `api/nav` · `test_smuggle.py` · `tools.json` (`smuggle-probe` + `smuggler.py` + `h2csmuggler`) + the Dockerfile block + `docker/proof/smuggle_install_proof.sh`.

The target-lock trap (worth remembering). `--mutations CL.TE,TE.CL` — `CL.TE` is host-shaped (a dot with letters either side), so `executor.check_target_lock` reads it as an out-of-scope host and refuses the whole job at the *target* gate, in real engagement mode too. The fix is to spell mutations **dot-free on the argv surface** (`cl-te`); the engine gets the canonical `CL.TE` from the JSON spec on stdin — the argv is only the gate/approval surface, never what the engine parses.

Verified. Hermetic safety suite 119/119, arsenal safety green, `next build` exits 0, the `/smuggle` verdict-matrix screen looked at (`assets/screenshots/smuggle.png`). **Image built and `smuggle_install_proof.sh` 9/9** (engage-sandbox recreated). Committed + **pushed** `67a696d`, with a build fix `9fac6d5` (BishopFox/h2csmuggler ships no `requirements.txt`, so the install layer falls back to `pip install hpack` when it is absent).

Web cache poisoning / deception detector — unkeyed-input detect + gated poison-plant (`:cache`) (2026-08-08)

The gap + the same safety split. A shared cache that stores a response keyed on part of the request while *reflecting* an **unkeyed** part is poisonable; a cache that stores a dynamic page under a static-looking path is a **cache-deception** leak. Confirming either affects co-*users* of the cache, so — exactly like `:smuggle` — **detection is safe-by-default** (probe for an unkeyed input reflected into a **cacheable** response; plant no cache entry) and **confirmation is a separate approve-each** with a co-user warning (prime the cache with a marker, then fetch it back with a fresh request that never carried the marker — if the marker comes back, the cache is poisoned).

Modelled on `:smuggle / nuclei.py`; **no new gate**.

Files + the one-place rule. `docker/cache_probe.py` (the baked engine `cache-probe`, **stdlib-only, no venv**, `--selftest / --job-stdin`) · `backend/cockpit/cache.py` (pure argv/spec/verdict-parse + a gated worker per stage) · `router /cache/*` (detect / confirm stages pinned) · `CacheScreen.tsx` + the `/cache` page + `api/nav` (icon ) · `test_cache.py` · `tools.json` (`cache-probe` + `wcvs`) + Dockerfile blocks 9l/9m + `docker/proof/wcvs_install_proof.sh`. The candidate unkeyed inputs — `X-Forwarded-Host/Scheme/Proto/Port/Server/For`, `X-Host`, `X-Original-URL`, `X-Rewrite-URL`, `Forwarded`, a cloaked query param, a fat-GET body — and the poisonable predicate live in **one** place, `cache.candidate_of(reflected, cacheable)`; a deception hit is a dynamic (`text/html 200`) page served for a `/.../foo.css` path that is also cacheable.

Traps. Header inputs (`x-forwarded-host` ...) carry no TLD-shaped dots, so unlike `:smuggle`'s `CL.TE` they do **not** trip the target-lock — no dot-free rewrite needed. The `wcvs`-transcript parser had to be tightened to **positive** hit markers only (`vulnerab|poison|[+]|hit`), because a bare `reflect` matched the tool's own "*no reflection*" line.

And `wcvs` is a Go tool: it got its own clean install layer with a throwaway current Go toolchain (fetched from `go.dev/VERSION`, not a pinned old version — an old Go cannot compile the modern `stdlib` its deps pull), leaving the other Go layers untouched.

Verified. Hermetic safety suite 120/120, arsenal safety green (161 tools), `next build` exits 0 (`/cache` prerendered), the `/cache` verdict-table + deception-hit screen looked at (`assets/screenshots/cache.png`). Committed `1508238`; the `cache-probe` image layer passed and the `wcvs` layer was fixed in `f5dc80d` (**engage-sandbox rebuild is the operator's manual step**, flagged).

Cross-domain kill-chain graph — the capstone that stitches web → cloud → on-prem into one route (:killchain) (2026-08-08)

The gap. HackPit had three *separate* attack-path graphs — the web foothold (findings from :recon / :proxy), the cloud IAM graph (:cloud), and the on-prem AD graph (:ad-graph) — but nothing that showed how a chain crosses them. The seams an operator actually walks (an SSRF that reaches cloud metadata, a cloud secret reused as an AD credential, a web RCE that lands on a domain-joined host) lived only in the analyst's head. This build adds the **capstone**: one view that routes a foothold in *one* lane to a high-value target in *another*, across those seams, using the same edge-index orchestrator that already drives the two per-lane graphs.

A read-and-stitch OVERLAY, not a new engine (backend/killchain/). The overlay consumes each graph's public dict (Graph.to_dict()) — a plain nodes/edges mapping in which every edge already carries a precomputed `abusable` flag) and merges it into one graph, tagging every node with its `domain` (web | cloud | onprem) and namespacing ids so the lanes cannot collide. Because it trusts each lane's own `abusable` classification, it **needs neither package's edge taxonomy — and imports NEITHER `adgraph` NOR `cloudgraph`**. That is the decoupling invariant, and a safety test AST-scans every module to prove it: the two graph packages stay independent of each other *and* of the overlay; the cross-cutting join (fetching each store's public dict) lives only in `main.py`.

The new content — the cross-domain bridge catalog (killchain/bridges.py). Within a lane the abuse edges already exist (the :cloud / :ad-graph technique catalogs own them). What connects the lanes does not live in any single-lane graph — it is the SEAM. The bridge catalog is the small, new parallel to those technique catalogs, covering five seam kinds, each with a title, a KB-grounded-or-catalog crossing command, and the ATT&CK technique + the two domains it crosses: `SsrfToImds` (web SSRF → cloud principal, T1552.005), `NodeToCloud` (compute/K8s node → its instance role), `WebToHost` (web RCE → host, T1190), `CloudToOnprem` (cloud secret/sync-account/RunCommand → AD credential, T1078), `OnpremToCloud` (AD-synced identity / SP cert / SYSVOL creds → tenant, T1078.004). Seams are declared on lane nodes (props["seams"]) and synthesized after the merge; a seam pointing at a missing node is refused with a warning, never a dangling edge. Grounding uses a **tool-head match guard** (the cloud grounder's `_cli_key` lesson) so a KB entry retrieved by the seam's seed terms is *cited* but never *mis-grounds* the command.

The orchestrator is copied from `cloudgraph`, safety design and all (killchain/orchestrator.py). BFS over the merged abusable subgraph routes an owned foothold in any lane to a high-value target in any lane (Domain Admin, cloud Owner/root). The model is handed the numbered frontier and returns an **INDEX** — it authors no command; a pick outside the frontier is refused, not repaired. A **cross-domain (seam) hop** resolves to a runnable crossing command the human approves through the SAME gated executor every cockpit command uses (POST /cockpit/exec , approve-each, scope-locked, heuristic red-confirm). A **within-lane hop** is deliberately **not runnable here** — it defers to that lane's own :cloud / :ad-graph view, so there is exactly one command catalog per abuse and no drifting copy. `advance` moves the chain only on an approved, exit-0 run for a seam hop, and on the operator's word for a within-lane one. **This build adds NO new gate.**

Routes (main.py , cross-cutting) + frontend (/cockpit/killchain). GET /killchain/graph , POST /killchain/propose , POST /killchain/advance , all delegating to a thin, testable `killchain/service.py` (the evidence-gated `advance_step` takes an injected `run_lookup`, so it is unit-tested with a stub run, no app boot). The screen renders **three swim-lanes** (web / cloud / on-prem) with the stitched path lit across them: each route node sits in its lane at its step column, so the path visibly *descends* through the lanes as it advances; a cross-domain seam is drawn in bright pink, a within-lane hop dim. The per-hop drawer shows the technique and either **APPROVE & CROSS** (a seam, through the gated executor, with the defender's-eye ATT&CK/detection disclosure) or **open in :cloud / :ad-graph** (a within-lane hop). `?demo=1` loads a synthetic web→cloud→on-prem chain — no real host, account or domain.

Tests. `test_killchain.py` (functional): three synthetic lane dicts merge into one domain-tagged graph; the bridge catalog synthesizes all five seams; BFS routes *web SSRF* → *cloud ci-deployer* → *on-prem SVC-SQL* → *Domain Admins* crossing two seams over three lanes; the orchestrator returns an index resolving to the real seam edge, a within-lane hop defers to its view, an out-of-frontier pick is refused; advance is evidence-gated.

`test_killchain_safety.py` (invariants, mirrors `test_cloudgraph_safety.py`): the proposer + service execute nothing (AST, with a positive control), cannot set `approved` / `dangerous_ack`, expose no run/batch path; a proposed seam step submitted unapproved is refused; **no killchain module imports `adgraph` or `cloudgraph`** (decoupling); lab mode byte-for-byte unchanged. Both are registered in `run_safety_tests.sh`. **next build exits 0.** Screen LOOKED AT: `41-killchain.png` — the three swim-lanes with the stitched web→cloud→on-prem route to Domain Admin (not blank/error).

Assumptions (per §5), stated. The overlay renders on synthetic data; live lanes arrive from the three underlying surfaces (`:recon` / `:proxy` , `:cloud` , `:ad-graph`) as an engagement fills them, and live cross-domain seams deepen as the dedicated seam integrations (e.g. the existing *SSRF→IMDS* / *cloud/seed-imds* flow) carry the linkage. On synthetic data it is a diagram; its operational weight scales with how much real graph data sits behind it — which is exactly why it is the *capstone*.

Parameter / content discovery — hidden params + paths as a first-class `:recon` step (`:discover`) (2026-08-08)

The gap. `:recon` builds a surface from a domain, but the step every hunt actually does next — mining an *endpoint's* hidden parameters and hidden paths — was manual. The tools (`arjun` , `ffuf` , `feroxbuster` , `paramspider`) were all catalogued in the arsenal; nothing wired them into a **discover** → **hand to** `:intruder` / `:nuclei` / `:repeater` workflow that feeds the ranked surface. This build is that sibling of `:recon` , modelled on it point for point.

One gated job, three modes (`backend/cockpit/discover.py` , mirrors `recon.py`). The operator picks: **params** — `arjun` against one in-scope URL → hidden GET/POST/JSON parameters (where IDOR/SSRF/mass-assignment hide); **content** — `ffuf` / `feroxbuster` with a wordlist → hidden paths/dirs; **historical** — `paramspider` mines web archives for parameterised URLs, passively. Each is **one** approved job — the same `ffuf` / `nuclei` / `intruder` shape, **no new gate** — gated by the SAME `executor.validate_request` every command clears, run BEFORE anything spawns, with an ungated stop, in the open engagement sandbox. `x8` was in the spec's plan but is **absent from the sandbox image**, so it was dropped rather than cataloguing a broken invocation (the `zap_install_proof` discipline the spec itself cites); `arjun` — the must-have — carries parameter discovery.

Scope-safe by construction, two ways (the `:recon` property, not an extra gate). The **target host is scope-locked before the job runs**: `arjun`/`ffuf` point at one URL and `paramspider` at one domain, checked against the engagement scope in `start()` , and the words only fill path/param positions so they can never move the host off-scope. And **every discovered URL/param is scope-filtered before it lands** (`filter_endpoints_in_scope`) — a `paramspider` result mined from an archive that spans other hosts is dropped and surfaced read-only, never handed off or stored. Both are AST-locked pure functions.

The content wordlist is in the approved surface WHOLE, intruder-style. For content discovery the word list is the payload set, so — exactly like the intruder's Critical-2 — `content_gate_argv` puts EVERY word into the gated command line (`-w <word>` per word), not a count and not a sample, so a word carrying `| sh` is read by the danger heuristic rather than hiding behind a *"...and 4,993 more."* The worker writes the words to a loot file and runs the real tool; the two `argvs`, like the intruder's, are the honest declared one and the spawned one.

Discoveries are suggestions, never auto-fired. Each hidden parameter becomes a **pre-filled** `:intruder` position (`?id=[[FUZZ]]` , reusing the intruder's own `[[...]]` marker), a `:nuclei` target and a `:repeater` request; each hidden path hands off to `:nuclei / :repeater` . The discovery job is the only thing approved — the attack itself is still approve-each on the surface it lands on. In-scope hits are upserted into engagement state as `Endpoint s` (parsed by two new hand-parsed, `urllib` -free parsers `parse_arjun` / `parse_feroxbuster` — the `state/parsers.py` ban held), so a param-rich endpoint also **raises its host's rank** in the `:recon` surface, and a sensitive hit (an admin endpoint, a `debug` param) becomes a low `Finding` .

Routes + frontend. Six routes on the cockpit router (`/cockpit/discover` + `/status` / `/preview` / `/jobs...`), the same shape as `/recon` . The **discovery view lives on** `/recon` behind a tab toggle (`?view=discovery`), with a mode picker, a scope-locked URL/domain form, a full-wordlist textarea, a preview (the exact gate argv), and the discoveries rendered with their `discovery` tag, interesting params highlighted, the out-of-scope drops shown, and the `:intruder / :nuclei / :repeater` hand-offs on every row. Reused only existing `hp-tn-*` classes (no new phantom — `test_css_vocabulary` green).

Tests + screen. `test_discover.py` (functional): both arjun JSON shapes → params; feroxbuster response-only lines; per-mode argv incl. the FUZZ-keyword insertion; **the whole content word list in the gate surface, nothing truncated**; `filter_endpoints_in_scope` with a control; a discovered param → a valid `[[FUZZ]]` intruder position; only interesting discoveries become findings. `test_discover_safety.py` (invariants, mirrors `test_recon_safety.py`): the pure half executes nothing by AST; `start / validate` reach the gate before any spawn; engagement-bound; **the target is scope-locked by construction with a control**; flags default False; stop ungated. Both registered in `run_safety_tests.sh` — **124 files green**, next build exits 0. Screen LOOKED AT: `42-discover.png` — the discovery view with hidden params (id/debug/role highlighted) and content hits (`/admin` , `/.git/config`) each carrying the tri-surface hand-off, an out-of-scope `cdn.thirdparty.net` dropped (not blank/error).

Assumptions (per §5), stated. Parameter discovery (`arjun`) was the must-have and ships with the full hand-off; content (`ffuf` / `feroxbuster`) and historical (`paramspider`) round it out — all three shipped. `x8` was dropped as absent from the image (noted above). The hand-off pre-fills the target surface but never fires it; the discovery job itself is the only thing approved.

JS recon → secret / endpoint mining — the S3→bundle→secret chain, live (`:jsrecon`) (2026-08-08)

The gap. `:recon` builds a surface from a domain and `:discover` mines an endpoint's hidden params — but the step every modern web hunt lives in, *reading the target's JavaScript*, had no first-class surface. The `S3 → bundle → secret → OAuth` chain was a chain-builder concept and a `secrets-hunt` skill; nothing mined a live target's shipped JS for endpoints, parameters and hardcoded keys as an engagement surface. This build is that final `:recon` sibling (the last of the Q2 spec queue), modelled on `:recon / :discover` point for point.

One gated job, an in-repo engine (`backend/cockpit/jsrecon.py` + `docker/js_mine.py`). The worker mirrors `recon.py` (engagement-bound, one approval buys the collect+mine passes, ungated stop) and the *execution* mirrors `:cache` : the mining runs headless in the open engagement sandbox as `docker exec -i ... js-mine -- job-stdin` , the `js-mine` engine — the same baked-client shape as `cache-probe` / `race-singlepacket` / `smuggle-probe` , so the gated job has a stable JSON contract the backend can parse. The engine (stdlib only) fetches each JS URL, mines **endpoints** (a LinkFinder-style regex, relatives resolved against the JS origin so they are scope-checkable), **parameter names**, and **secrets/API keys** (a SecretFinder-style provider set), unpacks any **source map** (`//# sourceMappingURL=`) to recover the original `src/` paths + comments and mines the recovered source too,

and folds **trufflehog** in best-effort to mark *verified* keys. The external toolchain — `getjs / subjs`, `LinkFinder`, `SecretFinder`, `trufflehog`, `gitleaks`, `sourcemap` — is installed alongside for manual use and catalogued in `tools.json`; the image rebuild is the operator's step (`docker compose -f docker/docker-compose.yml build engage-sandbox`), and `docker/proof/jsrecon_install_proof.sh` re-checks every name + `--selftest` + the stdin contract in the built image and the running container.

Scope-safe by construction, two filters (the `:recon` property, not an extra gate). The collection **target and any explicitly-named JS URL are scope-locked before the job runs** (`start()` refuses `gate="scope"`), and **every candidate JS URL and every mined URL/host is scope-filtered before it lands** — `filter_urls_in_scope` means only in-scope JS is ever fetched (a `<script src>` to a third-party CDN is dropped read-only), and `filter_endpoints_in_scope` means a URL mined from inside a bundle but pointing off-scope never reaches the surface. This is the same two-filter split `:discover` makes for operator-input vs discovered URLs.

Secrets → loot, never finding text (mirroring `:credentials`). The engine returns secret values; the worker writes them to a loot file (`jsrecon-<job>-secrets.txt`) and builds a `Finding` carrying the **type, source JS URL, a masked preview and the loot path** — **never the value** (report.py renders findings verbatim). A **trufflehog-verified** key is **High**; an unverified regex match is **Low**. Mined endpoints/params are upserted into engagement state (tagged source `js`), so a param-rich bundle endpoint **raises its host's rank** in the `:recon` surface; each carries a `:nuclei / :repeater` hand-off. The `state/parsers.py` `parse_jsmine` — the one place a mined endpoint becomes a surface record, so a plain terminal `js-mine` run ingests too — respects the `urllib` ban the state package is regression-locked to (hand-parsed query params, no `urllib` import; `test_state.py` green).

Routes + frontend. Six routes on the cockpit router (`/cockpit/jsrecon` + `/status / /preview / /jobs...`), the same shape as `/recon / /discover` . The **JS-recon view lives on `/recon`** behind a third tab (`?view=jsrecon` , `?demo=1` for synthetic data), with a collect-target + JS-URL form, source-map/verify toggles, a preview (the exact gate argv), a **secrets panel** (type + verified badge + masked preview + source, value in loot), mined endpoints tagged `js` with their hand-offs, recovered source-map paths + comments, and the out-of-scope drops shown. Reused only existing `hp-tn-*` classes (no new phantom — `test_css_vocabulary` green).

Tests + screen. `test_jsrecon.py` (functional): the engine mines a JS blob (relatives resolved to the JS origin, absolutes kept, param names + AWS/Stripe/generic secrets, a `.map` → recovered source paths + comments, collect → the `<script src>` set); `parse_mine_output` splits the engine JSON into Endpoint rows (via `parse_jsmine`) + secret dicts + source maps; **both scope filters keep in-scope + drop/collect out-of-scope with controls**; a verified secret → a **High** Finding whose serialized text carries the value **nowhere**, and the value is written to the loot file; unverified → Low. `test_jsrecon_safety.py` (invariants, mirrors `test_recon_safety.py`): the pure half executes nothing by AST; `start / validate` reach the gate before any spawn; engagement-bound; **target + named JS URL scope-locked by construction with a control**, both filters drop out-of-scope with a control; **a secret value never reaches a finding with a planted control** but is written to loot; flags default False; stop ungated. Both registered in `run_safety_tests.sh` — **126 files green**, `next build` exits 0. Screen LOOKED AT: `47-jsrecon.png` — the JS-recon view with the secrets panel (verified `aws-access-key-id` + `stripe-secret-key`, an unverified generic-api-key, each masked and marked *value in loot, never shown*), mined endpoints tagged `js` with param chips and hand-offs, a recovered source map (`src/api/client.ts` , `src/config/keys.ts` , `src/admin/flags.ts` with leftover comments), and an out-of-scope `cdn.thirdparty.net` dropped (not blank/error).

Assumptions (per §5), stated. Endpoints/params + secrets were the must-haves and shipped; source-map unpacking shipped too (the bonus). JS URL collection accepts operator hosts, the engagement's `.js` endpoints already in state, and a collect-from-target pass — no fetching outside scope, ever. Execution was modelled on `:cache` 's baked engine + catalogued external tools + install-proof (the most robust choice given the image

rebuild is the operator's step and the built tools can't run in a hermetic session); the engine's own regex mining works without any external tool, and trufflehog only *upgrades* a hit to verified. Secrets go to loot, never finding text.

Dual-candidate "second opinion" on generated commands — foundation (2026-08-09)

What it is. An operator-requested SECOND OPINION on a generated command: for any attack-path step, a $\rightleftharpoons$ **second opinion** control fetches, on demand, ONE AI-curated alternative candidate plus an advisory *which-is-better* verdict. The AI decides each time whether the alternative is a *different KB technique* (grounded) or a *tuned form* of the same command (ai_suggested). This does not weaken the grounding invariant — it strengthens the honest presentation of it: the primary is never touched, and every alternative is badged by kind.

The engine (`backend/alternatives.py`). Read-only, **executes nothing** (AST safety test, like `proposals.py`). `best_alternative(primary, *, goal, target, scope, by_id, search_fn)` composes with the **global /llm-config** (`llm.load_config()`) — the second opinion runs on whatever model the operator has selected — and REUSES `attack_path` 's own grounding + scope machinery one-way (`attack_path` never imports it, so no cycle): a grounded alternative resolves a real cited `entry_id` (`_resolve_entry_id + is_step_eligible`) and uses that entry's commands **verbatim**, **target-substituted** (`entry_commands + substitute_target`); a tuned alternative is the model's own command **capped + marked unverified** (`_ai_commands`); both are then scope-flagged through the same `flag_foreign_refs` a primary step gets. The verdict is shaped by a **key-whitelist** — a stray gate field (`approved / dangerous_ack`) the model emits is dropped, so the verdict can never carry a machine-actionable flag, and it never reorders or auto-selects. It **never invents an entry_id** (an unresolvable citation demotes to a tuned form, or to no alternative). On any LLM failure it **soft-fails** (alternative null + a "model unreachable" verdict) so the plan view never breaks.

Surface (attack-path). A new on-demand endpoint `POST /attack-path/alternative` (`AltStepIn` → `AlternativeOut`) wires the engine to `STATE.by_id` + the hybrid search. Because it is on-demand, the alternative comes back from its own endpoint and the frontend holds it in component state — so **AttackStep gained no fields** (the three-places schema trap is side-stepped entirely). The `Alternative` response model does declare `foreign_refs` , so that per-command scope annotation is not stripped. A self-contained `AlternativeDisclosure` component renders the alternative below the primary (never importing `PlannedCommand` , to avoid an import cycle): a `GROUNDED · kb:<id>` (accent) or `AI-SUGGESTED · VERIFY` (amber) badge, the command block, any foreign-host note, and the *which & why* verdict.

Verification. 10 new tests: engine unit tests (grounded-verbatim, tuned capped+unverified, choice-none, verdict has no gate field, LLM-unreachable soft-fail, never-invents-an-entry-id, scope-adaptation) + the AST executes-nothing test + two endpoint tests. Full backend suite **1375 pass** (the one red — `PUT`

`/cockpit/windows/profiles/{profile_id}` vs the CORS method contract — is **pre-existing and unrelated**, present at the parent commit; a windows-profiles route to convert `PUT`→`PATCH` separately). Frontend: **tsc 0, lint 11 (baseline held), next build exit 0**. Screen **LOOKED AT** live: on an exploitation step whose primary was a NoSQL attempt for a SQLi goal, the second opinion returned a **GROUNDED · kb:tool-sqlmap** alternative with the entry's commands substituted to `target.test` and a verdict that correctly flagged the mismatch — exercised end-to-end on both `claude-agent-sdk/opus` (API) and local `qwen3:8b` (browser).

Scope of this build. This began as the FOUNDATION surface (engine + attack path) and was then completed to all five structured surfaces plus the model picker in the same session:

- **Cockpit orchestrator** — `POST /cockpit/proposals/{id}/alternative` in `main.py` (cross-cutting: it reads the KB, which the cockpit package must not). It reads the queued proposal and **does not touch the Proposal model** — the "exactly one place approval is expressed" line is untouched, there is still no approval field. The `:proposals queue viewer` (`/cockpit/proposals` , in the launcher's plan band) lists the queue with each row's gate verdict and the $\Rightarrow$ second opinion; its review buttons record a decision and **run nothing** (the copy says so). Looked at live (empty state); list/review contract locked by a test.
- **AD / cloud / kill-chain graphs** — `POST /cockpit/{ad,cloud,killchain}/alternative` . The alternative is *a different way to abuse the SAME edge/seam*, grounded via a **category-filtered candidate set** (AD→active-directory/windows, cloud→cloud, kill-chain→unrestricted) — enforcing the same domain restriction each edge grounder does, while giving the model candidate diversity to pick a genuinely different technique. Wired into all three orchestrator screens.
- **Home model dropdown** — the launcher's llm status cell became an inline quick-switch: no-key providers switch in place (claude-agent-sdk opus/sonnet/haiku, local ollama models via `/ollama-models`) by writing the global `/llm-config` ; keyed providers route to the existing settings modal. The rail payload stays status-only, so `test_home_summary` is unchanged. (A `overflow:hidden` on the rail clipped the dropdown — moved corner-clipping onto the end cells; caught by looking at it.)

Also cleared a **pre-existing** unrelated red on the way: `PUT /cockpit/windows/profiles/{id}` → PATCH, to satisfy the CORS-method contract (the windows-profile route, not this feature).

The shared engine (`alternatives.py`) is written once and reused by every surface via a `fetcher` -prop disclosure on the frontend. **Chat stays excluded** (no discrete command object). Verified: **16 new tests** (engine 8 · attack-path endpoint 2 · proposal 3 · graph 3) + full suite **1382 pass**; tsc 0 · lint 11 · next build 0; the attack-path and home surfaces were looked at live. Spec: `docs/superpowers/specs/2026-08-09-dual-candidate-commands-and-model-picker-design.md` ; plan: `docs/superpowers/plans/2026-08-09-dual-candidate-foundation-engine-attackpath.md` .